

# Resource management in IaaS cloud



where to place the VMs



how much to allocate

# 1

What can a resource manager  
do for you

# Models of VM scheduler



# 3 Implementing VM schedulers

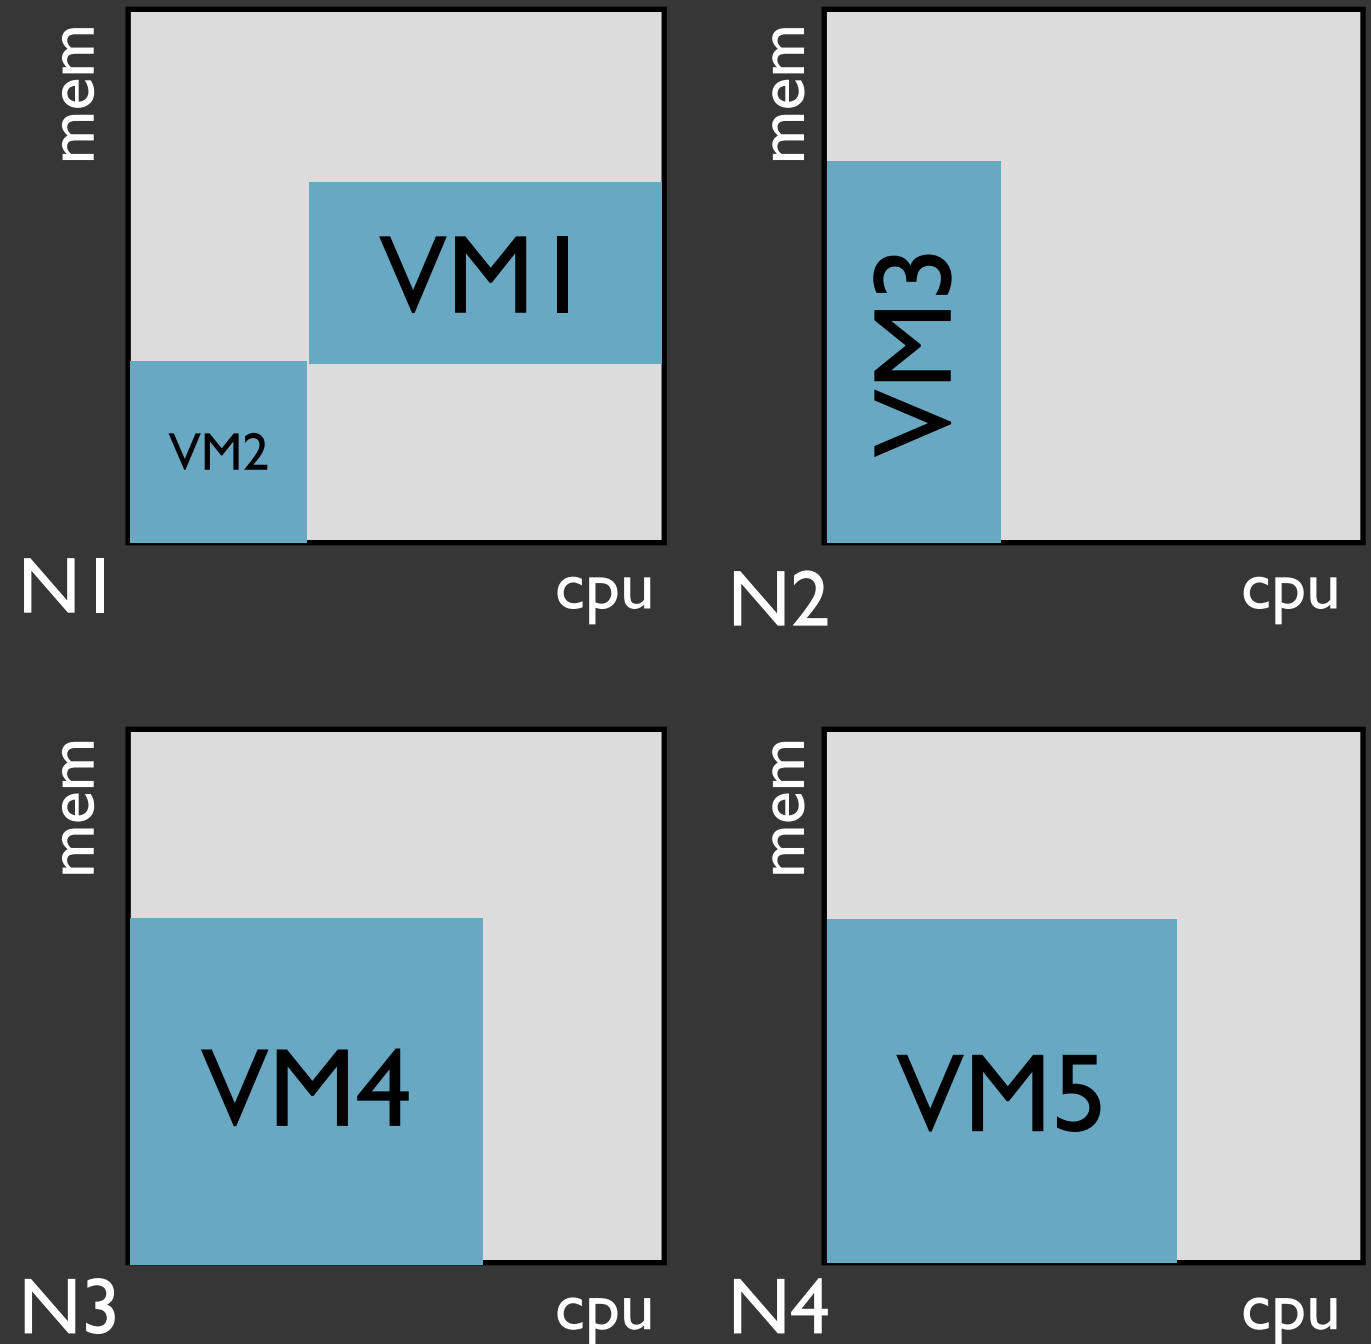
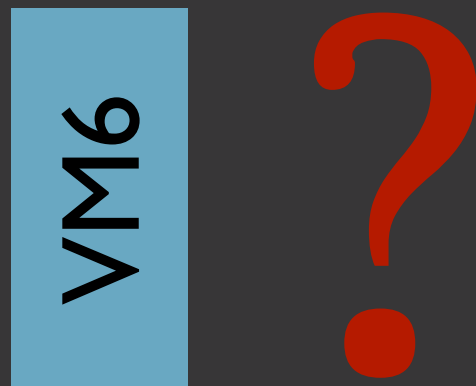
# 1

What can a resource manager  
do for you



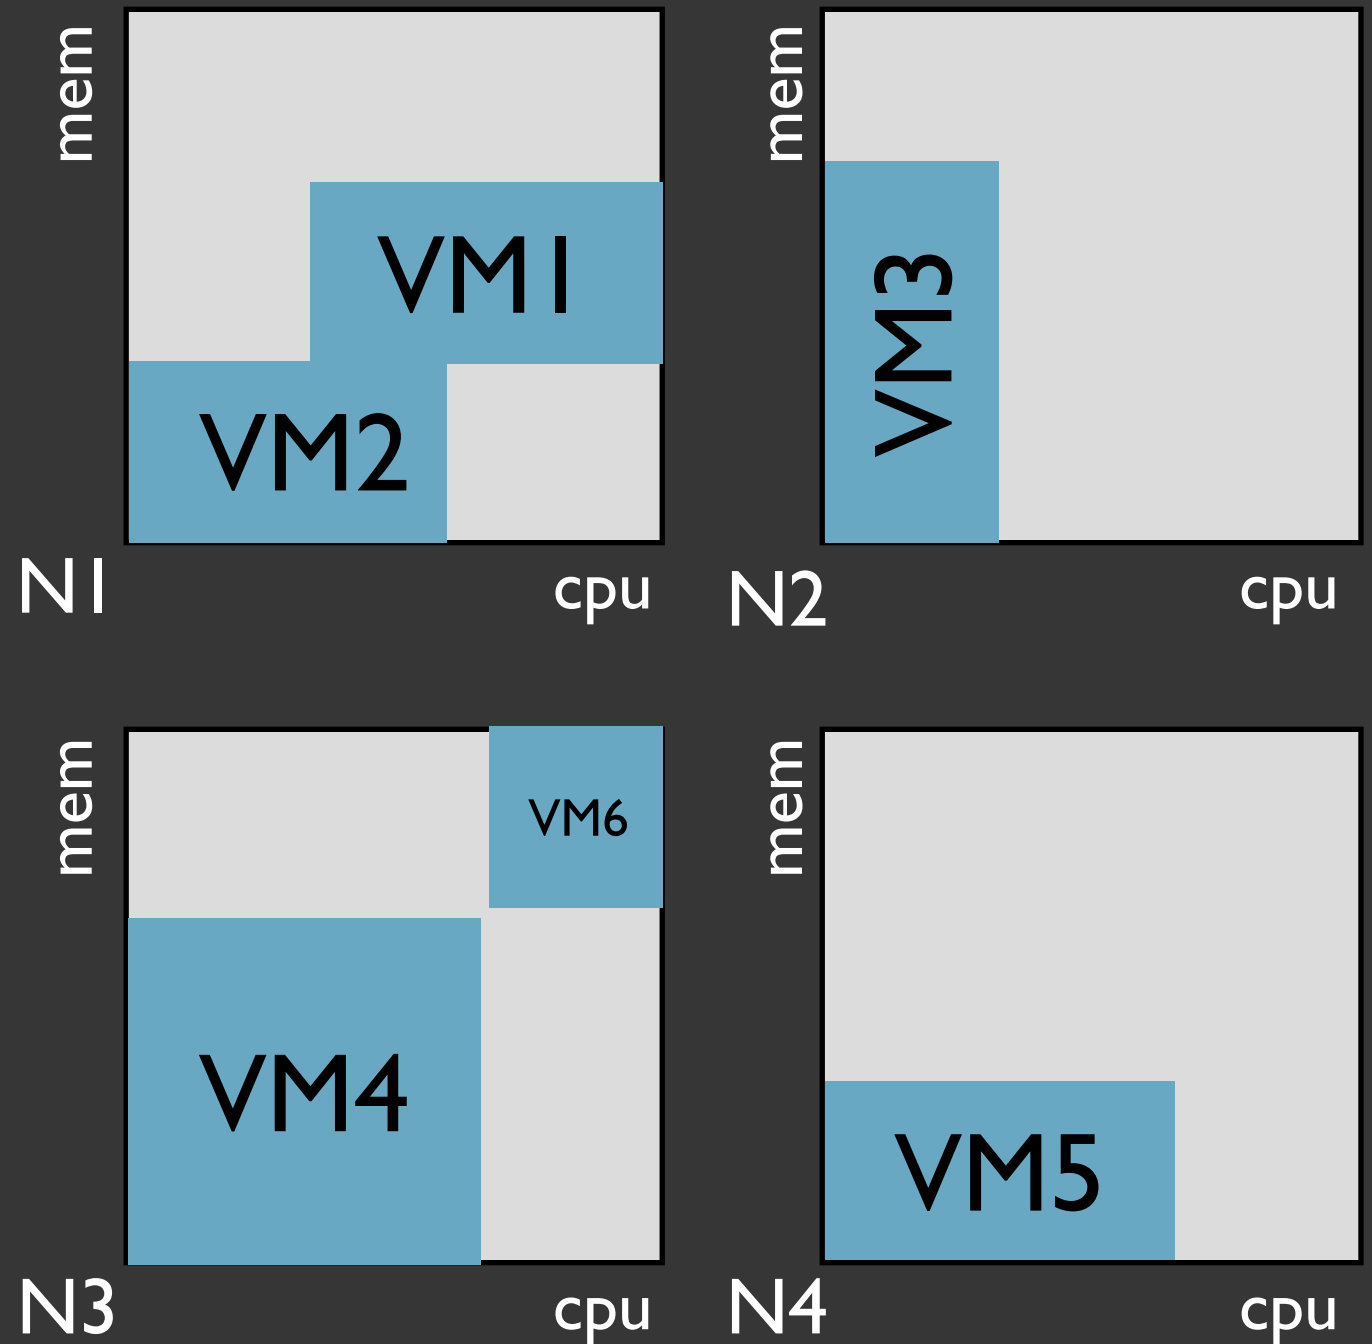
provide the awaited Quality of Service  
address the management objectives

the RM should solve that





the RM should prevent that



# Service

# Level

# Agreement

a contract where a service is formally defined

performance indicator

MTTR

MTBF

availability  $(\text{MTBF} / (\text{MTBF} + \text{MTTR}))$

pricing

penalties

...



# Google Cloud Platform

Monthly uptime percentage	Percentage of monthly bill for the respective Covered Service which does not meet SLA that will be credited to future monthly bills of Customer
99.00% - < 99.95%	10%
95.00% - < 99.00%	25%
< 95.00%	50%

uptime as the only  
Key Performance Indicator (KPI) ?

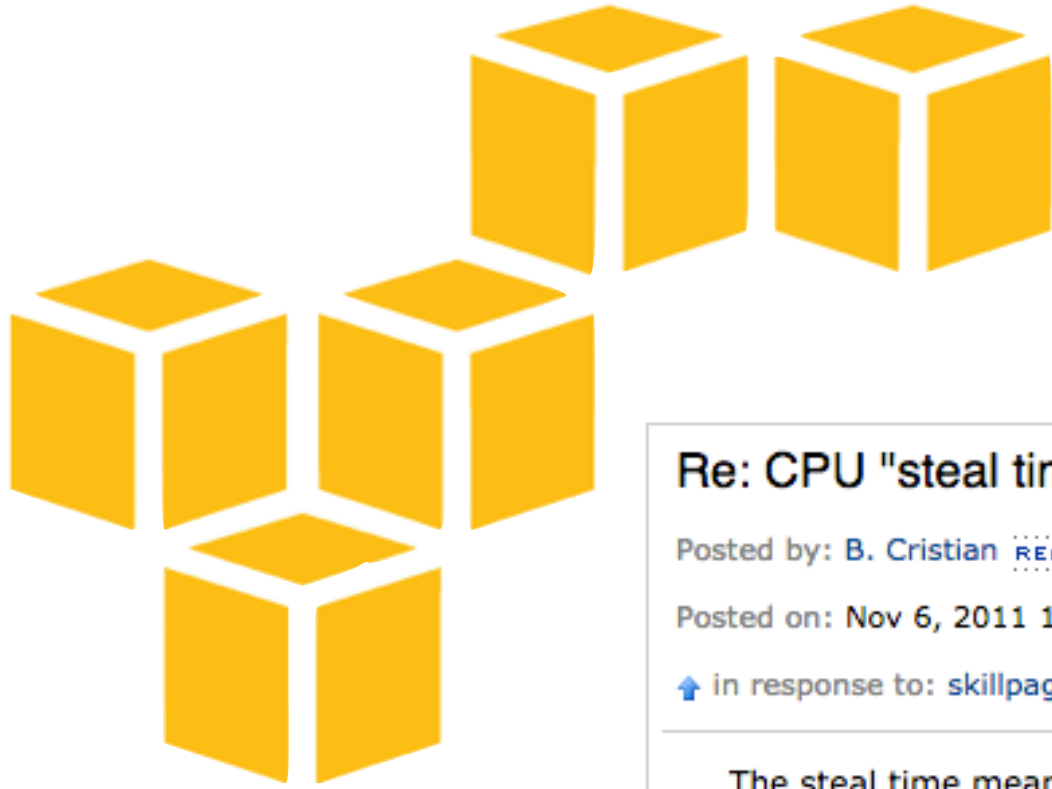
# Cloud Performance SLAs

## What customers want

guaranteed capacities (CPU, memory, bw)  
guaranteed latencies/throughput (network, load balancer, disk)

## What providers give

-



# Re: CPU "steal time" - is Netflix right?

Reply

Posted by: [B. Cristian](#) REAL NAME™

Posted on: Nov 6, 2011 12:32 AM

in response to: [skillpages](#)

The steal time means you're trying to go over your allocated resources. AWS doesn't give you a dedicated amount of resources, it gives you access to more, so that when the host machine isn't under heavy load you can use more resources, which would otherwise be wasted.

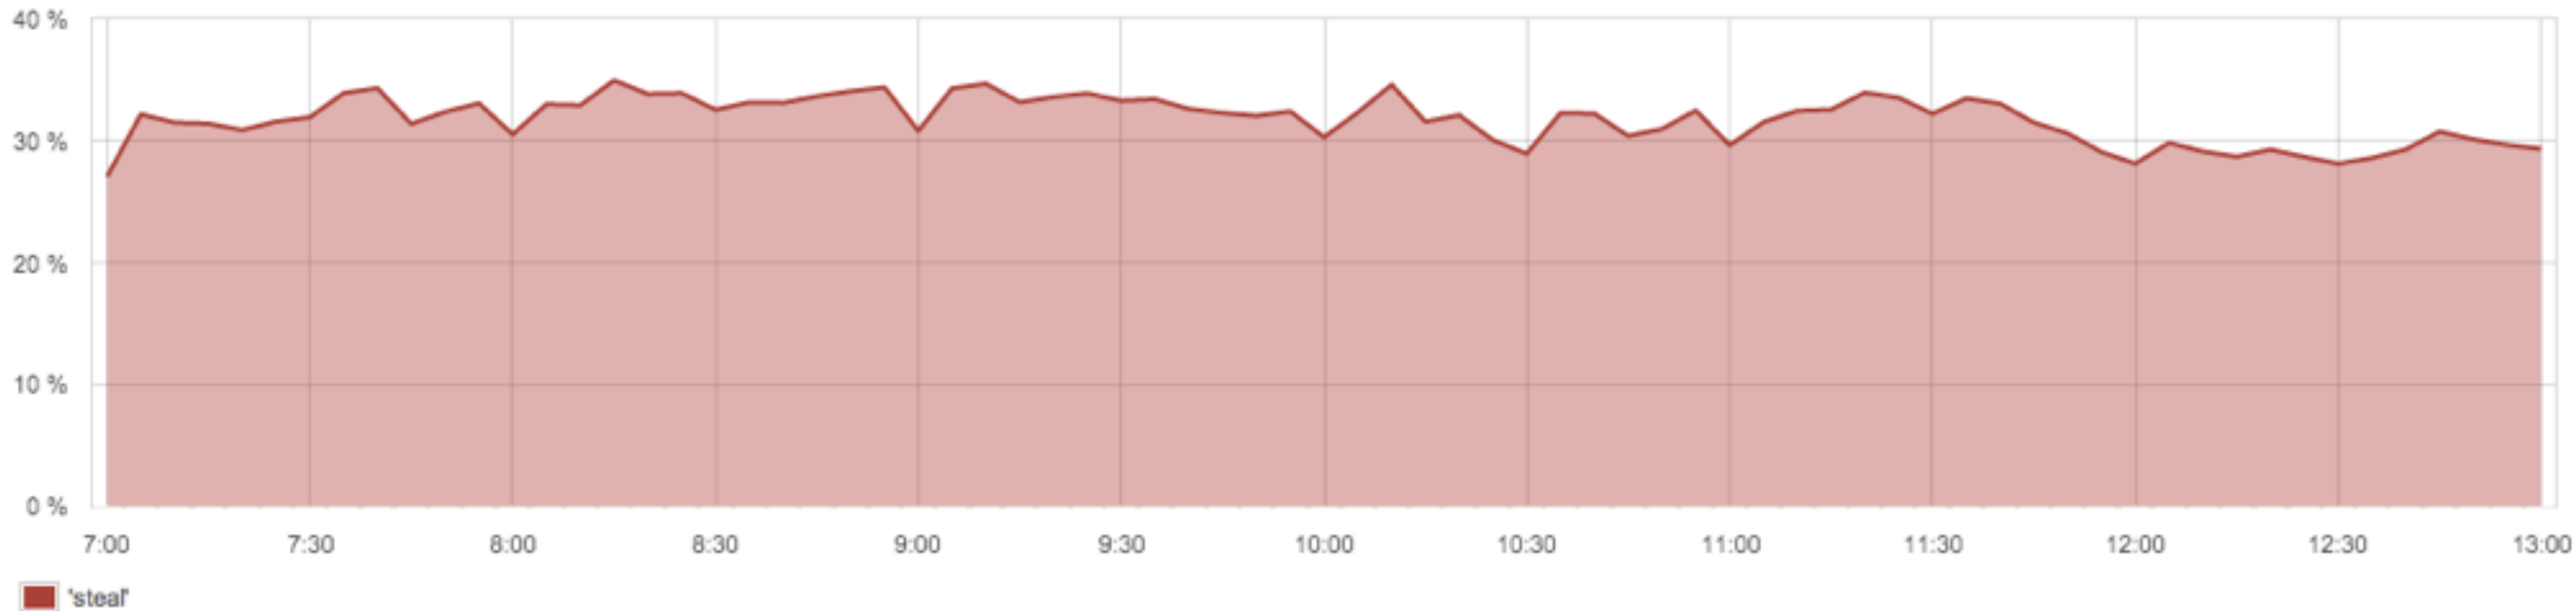
- Overview
- HBase Requests
- HBase Regions
- HBase Split
- HBase Memstore
- HBase Store
- HBase Compactions
- HBase FS
- HBase Block Cache
- CPU & Mem**
- Disk
- Network
- JVM

auto granularity From: 2013.04.19 07: To: 2013.04.19 13: (GMT-5:00) America/New York 30m 1h 6h 1d 2d 1w 30d 60d

- Memory Details
- Load
- Swap

CPU Details - 2013.04.19 07:00 to 2013.04.19 13:00 ☐ hide extended chart ☒ stacked

- ☐ 'user'
- ☐ 'system'
- ☐ 'wait'
- ☐ 'interruption'
- ☐ 'soft interruption'
- ☐ 'nice'
- ☒ 'steal'
- ☐ 'idle'



# the objective

provider side

*min(x)* or *max(x)*

# atomic objectives

*min(penalties)*

*min(Total Cost Ownership)*

*min(unbalance)*

...

# composite objectives

using weights

$$\min(\alpha x + \beta y)$$

How to estimate coefficients ?

useful to model sth. you don't understand ?

$$\min(\alpha \text{ TCO} + \beta \text{ VIOLATIONS})$$

€ as a common quantifier:

$$\max(\text{REVENUES})$$





Should a client ask for an objective



Should a client ask for an objective

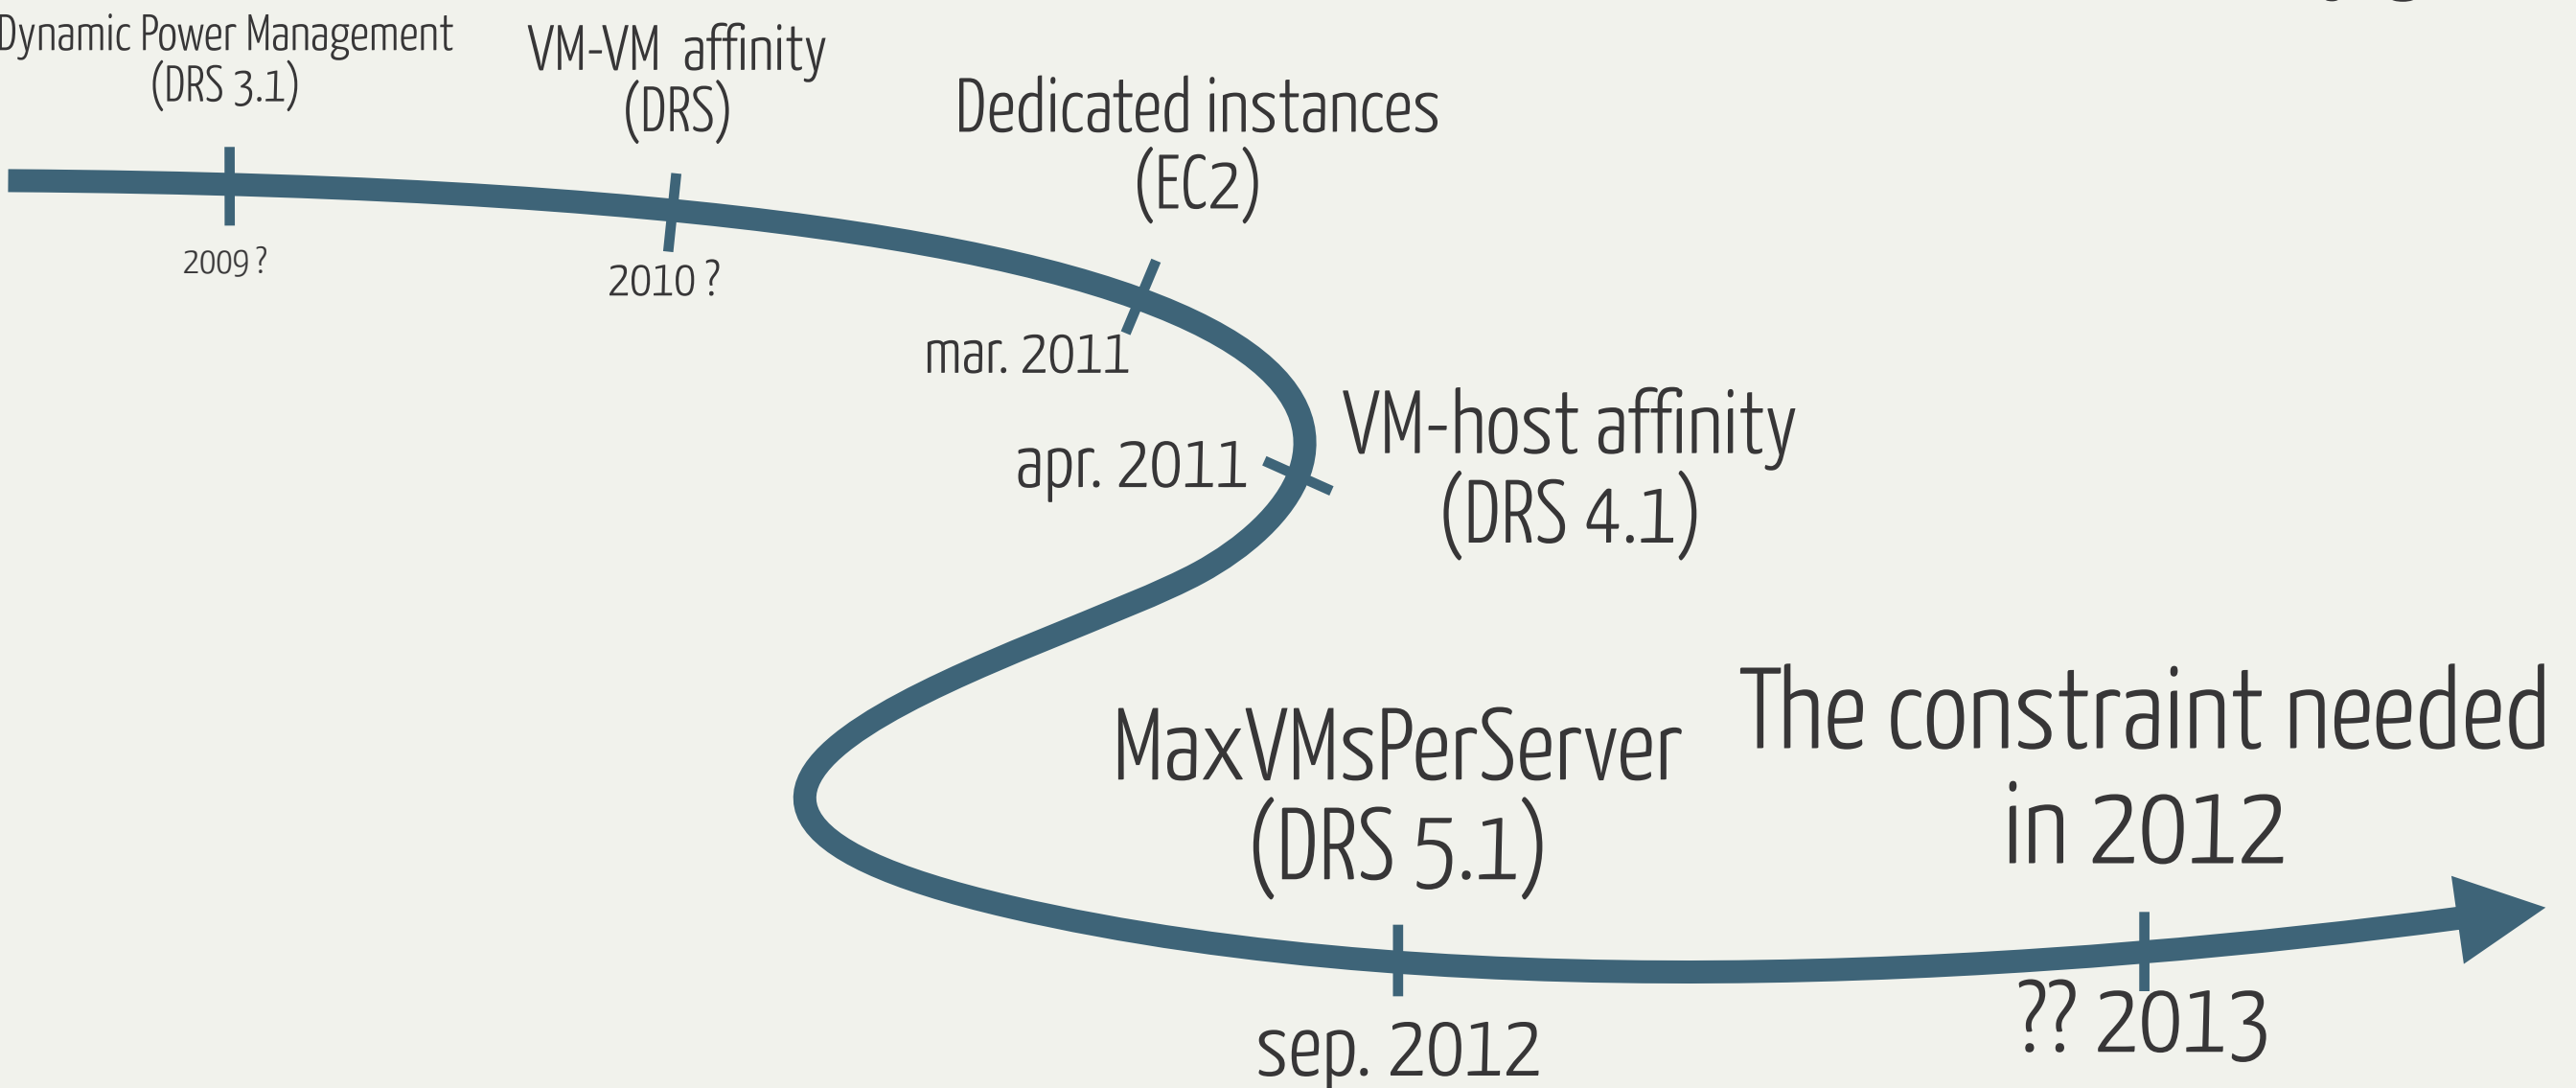
no if the provider cannot prove it is achieved

one can check the latency is below a threshold

proving a latency is minimal ...

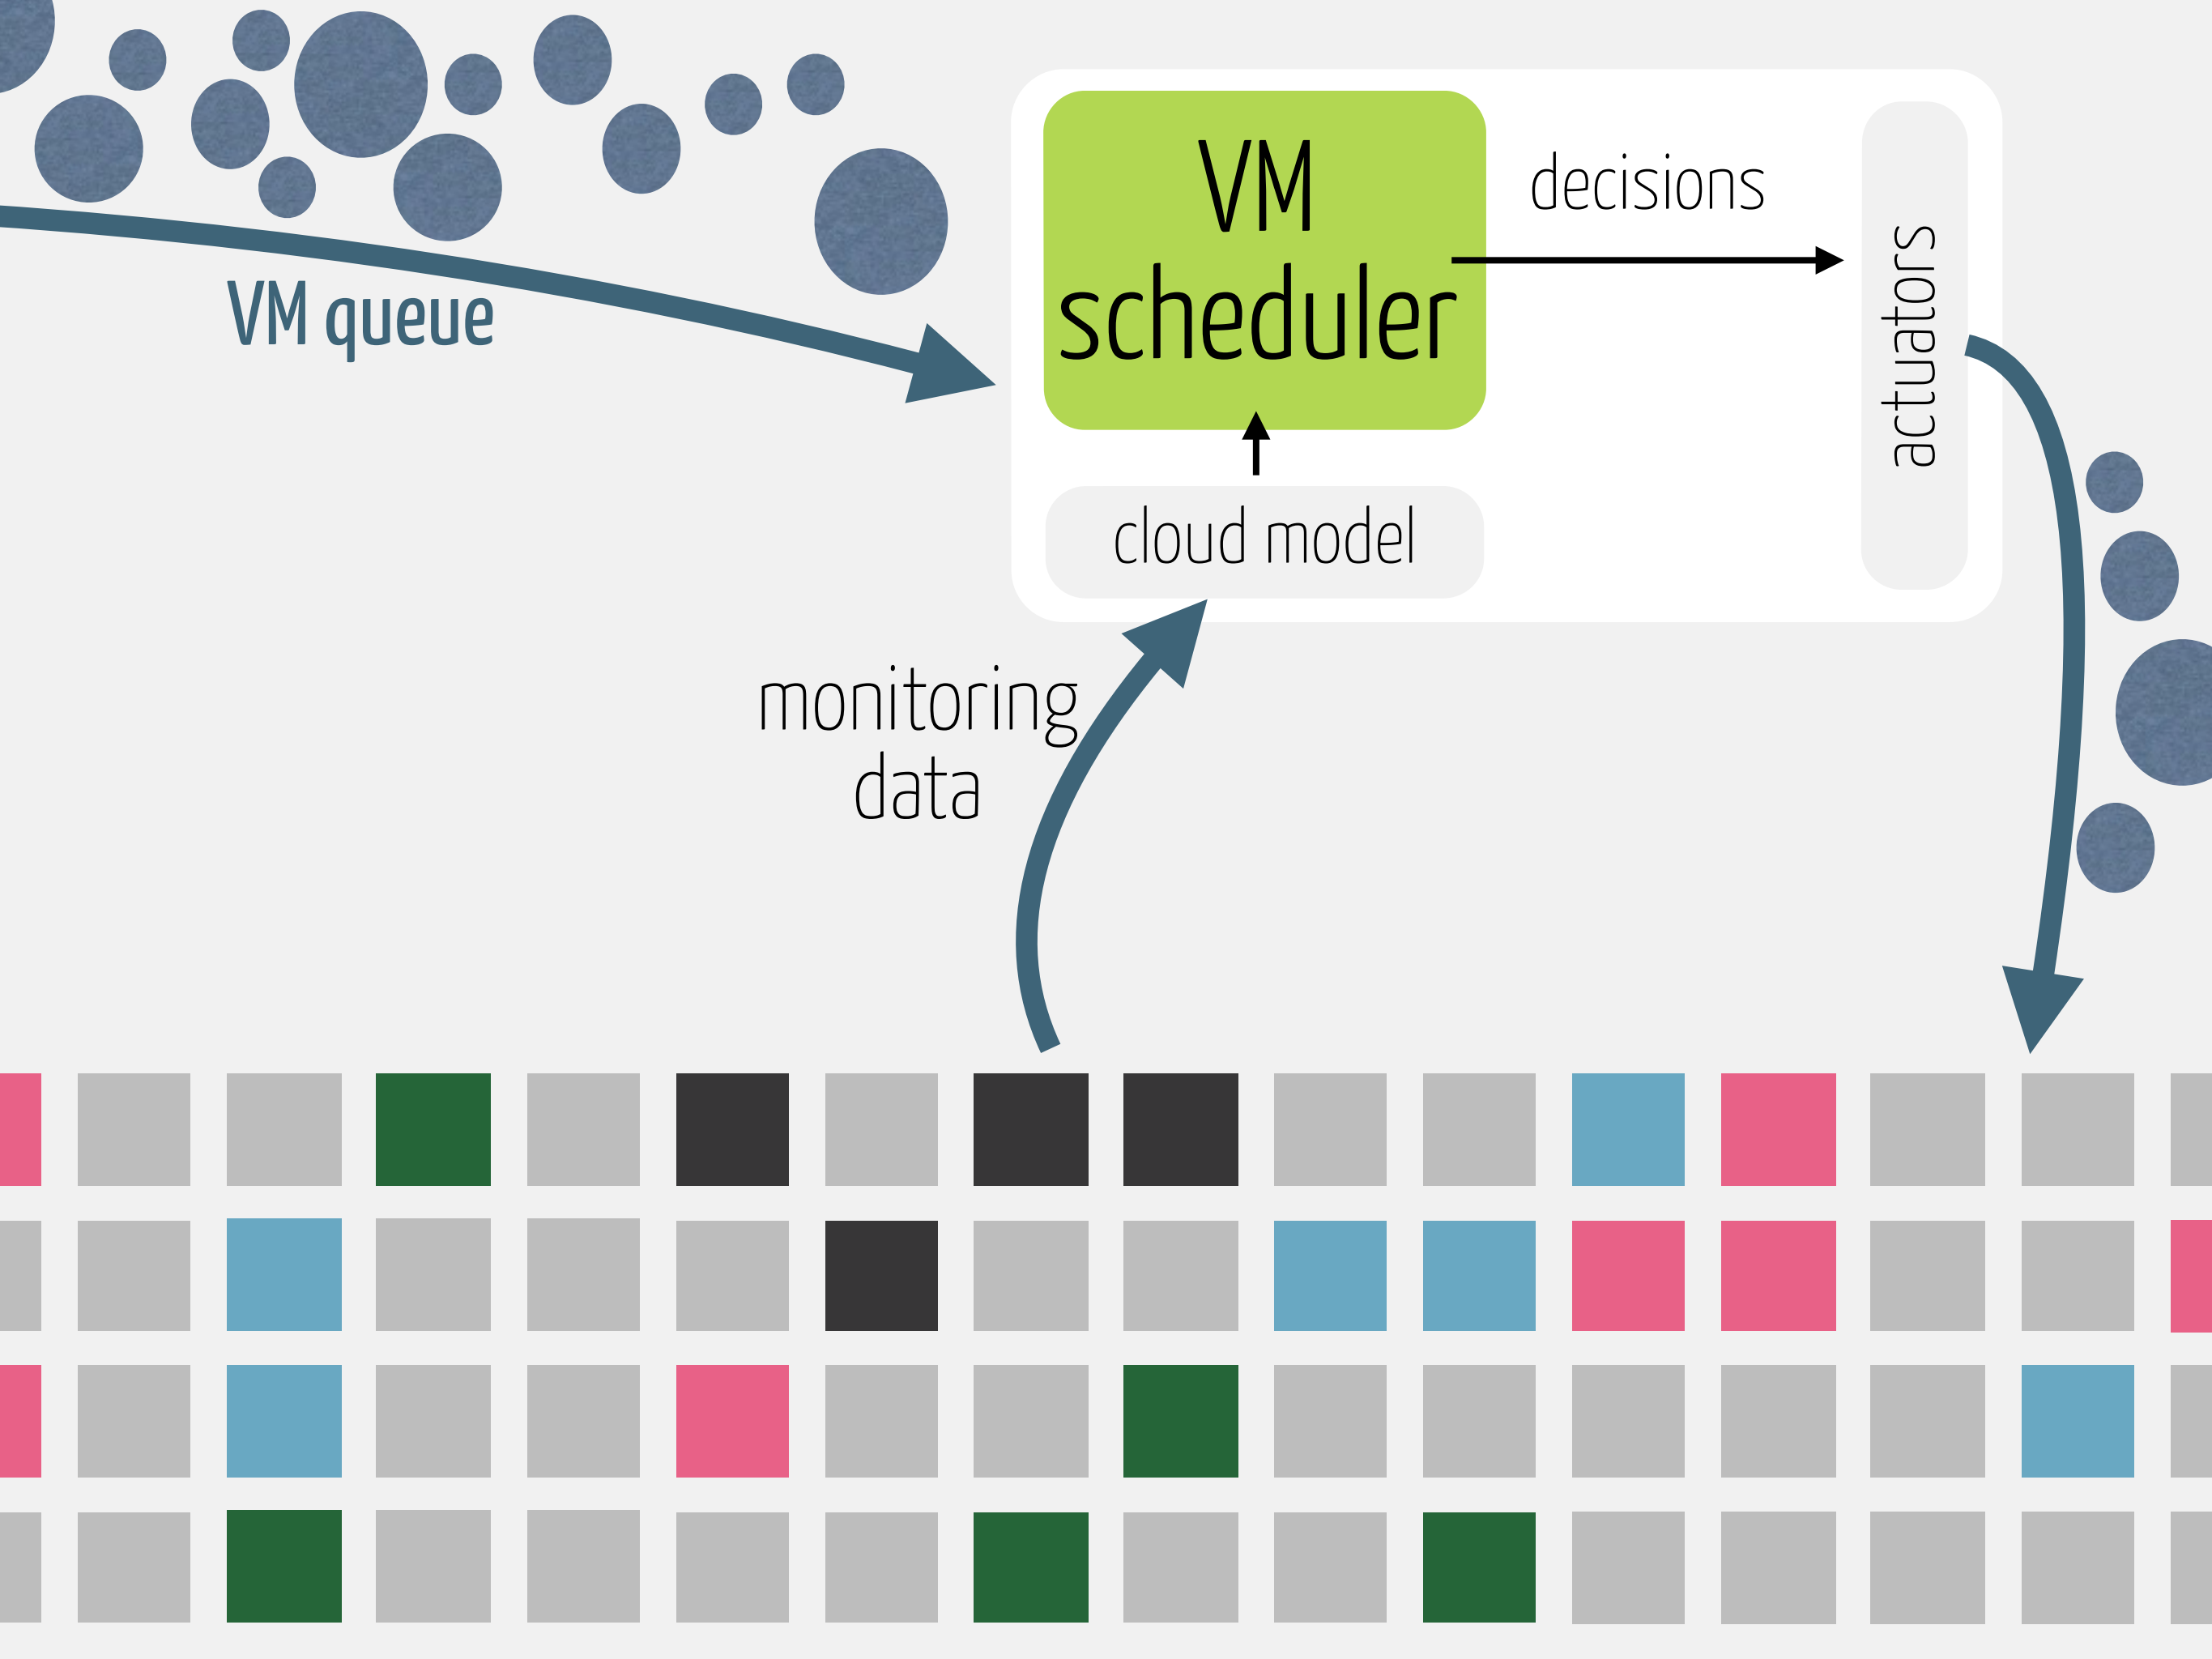
# SLAs and objectives are evolving

there is no holy grail



# Models of VM schedulers





# scheduling models

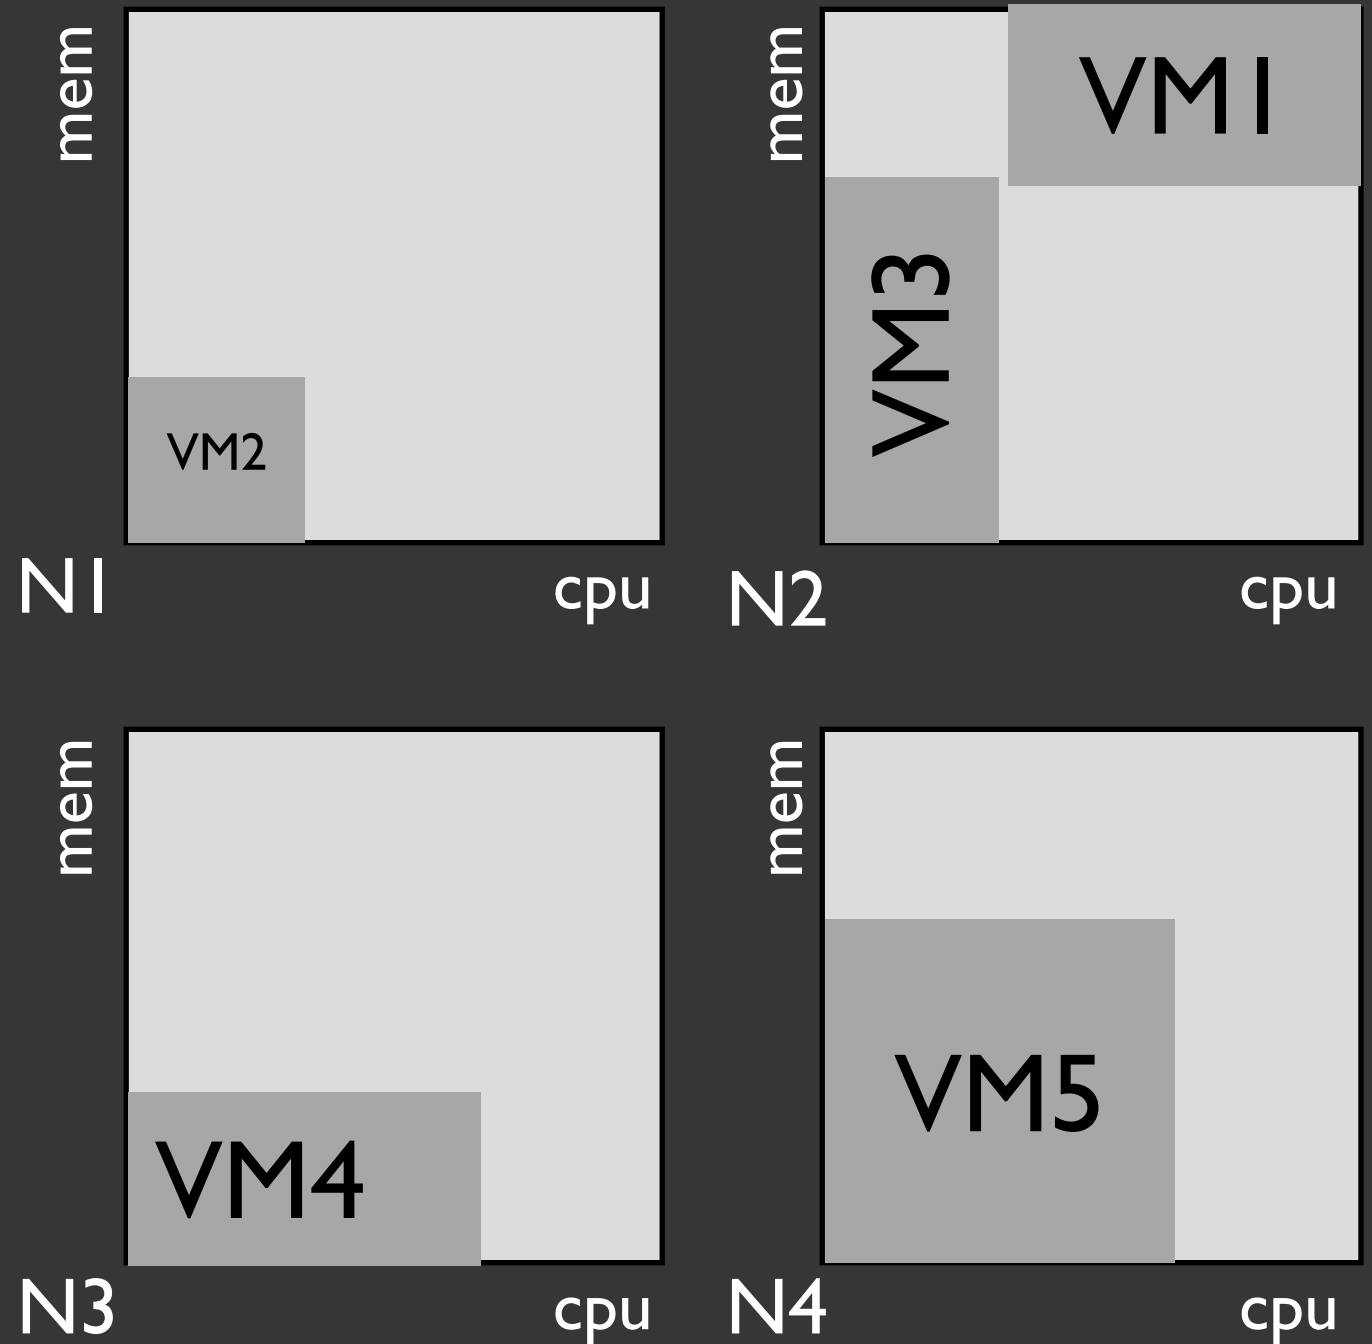
static  
dynamic schedulers

static  
dynamic resource allocation

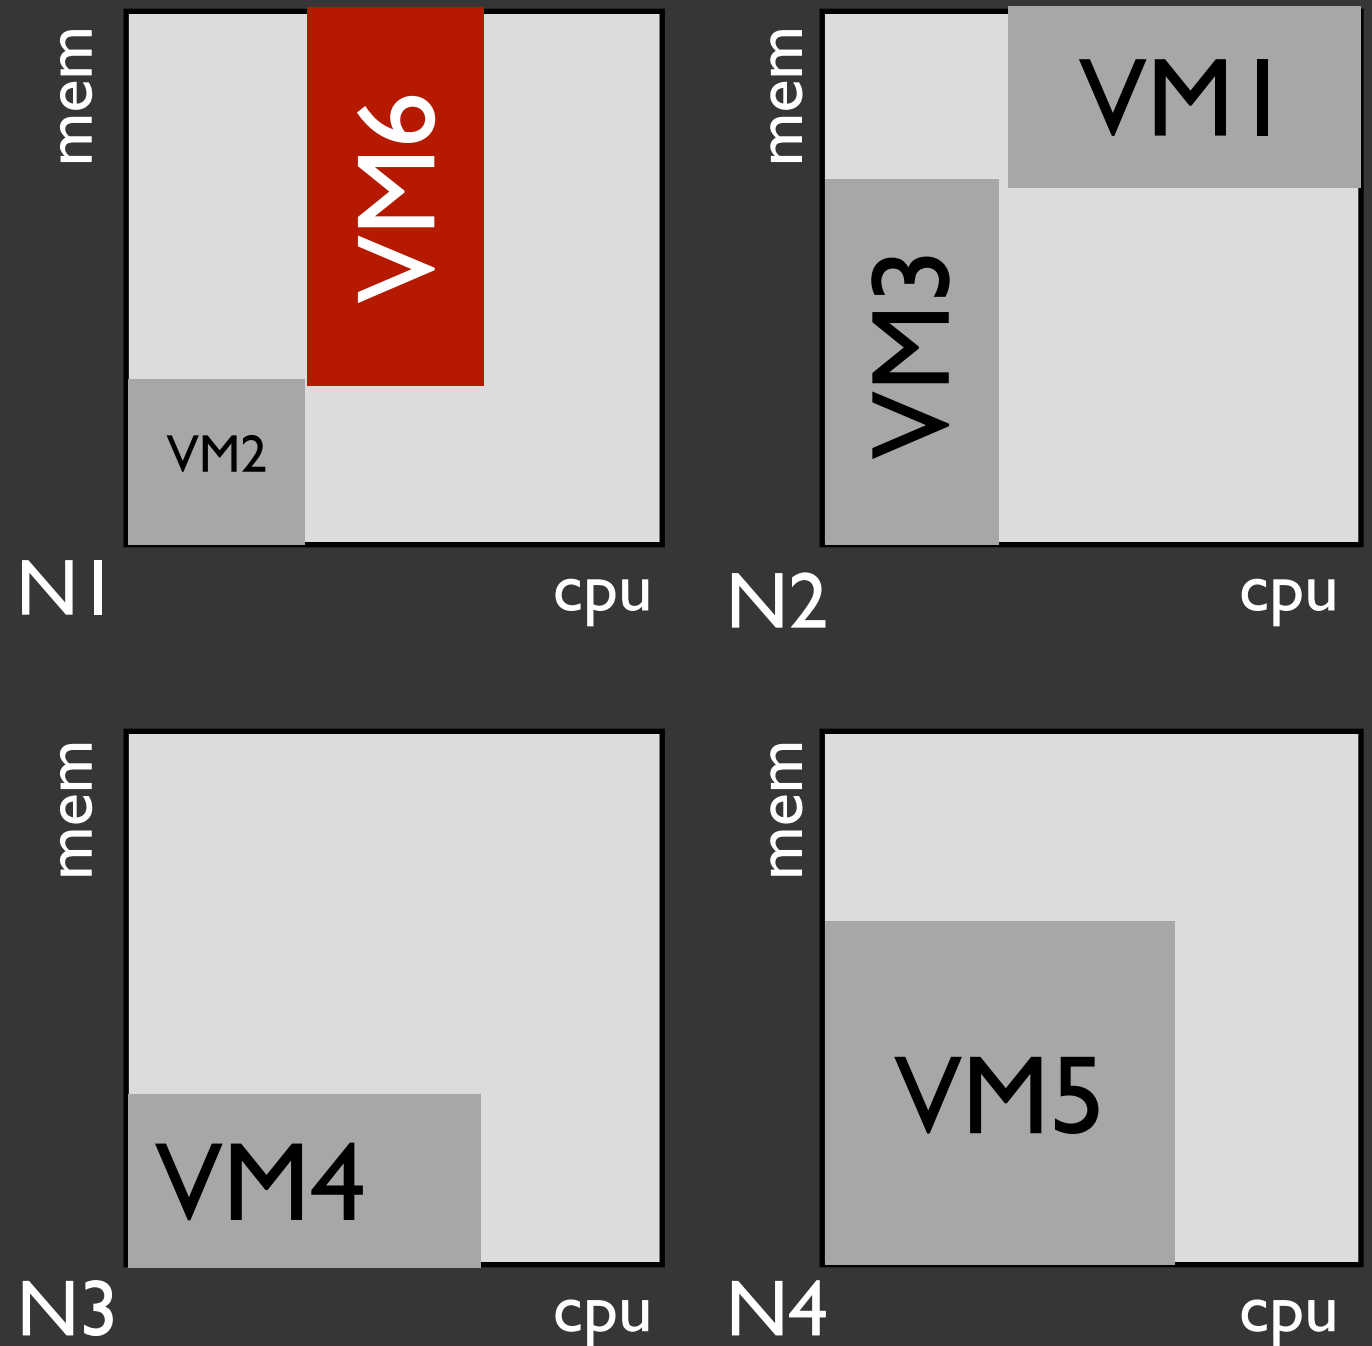
# static schedulers



VM6

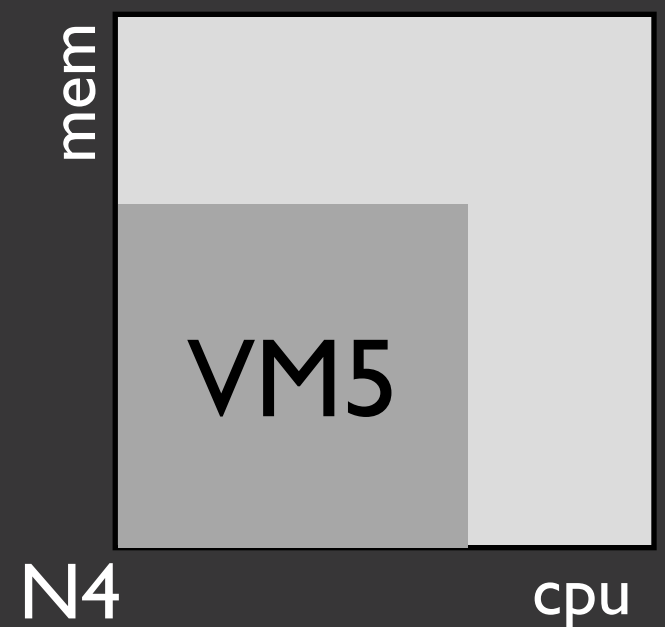
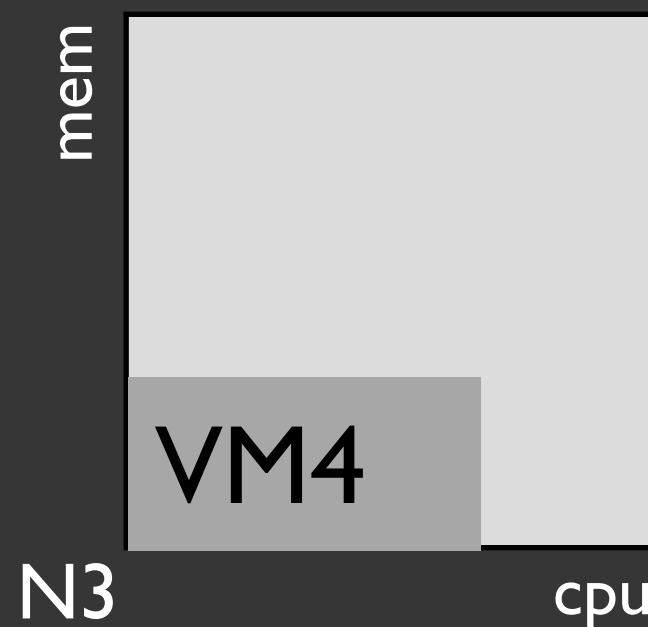
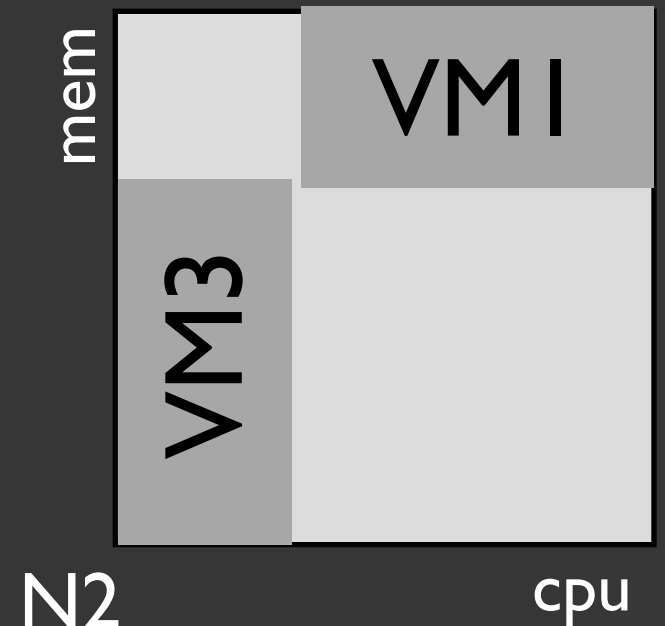
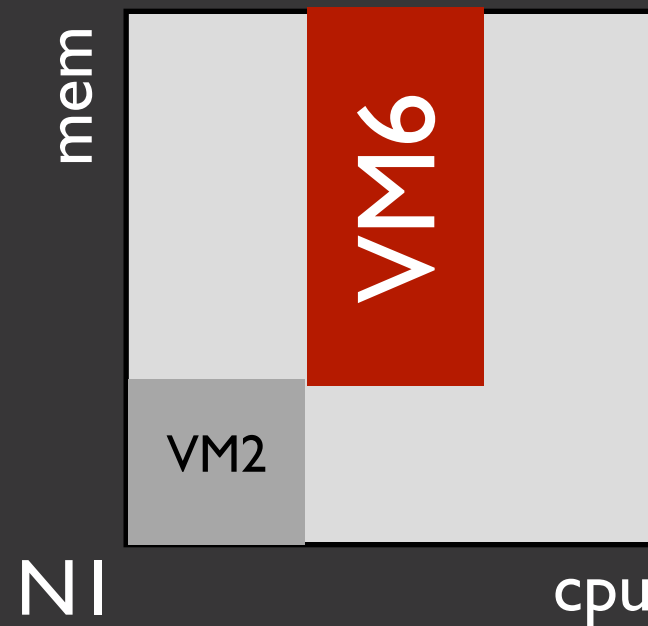


# static schedulers





# static schedulers



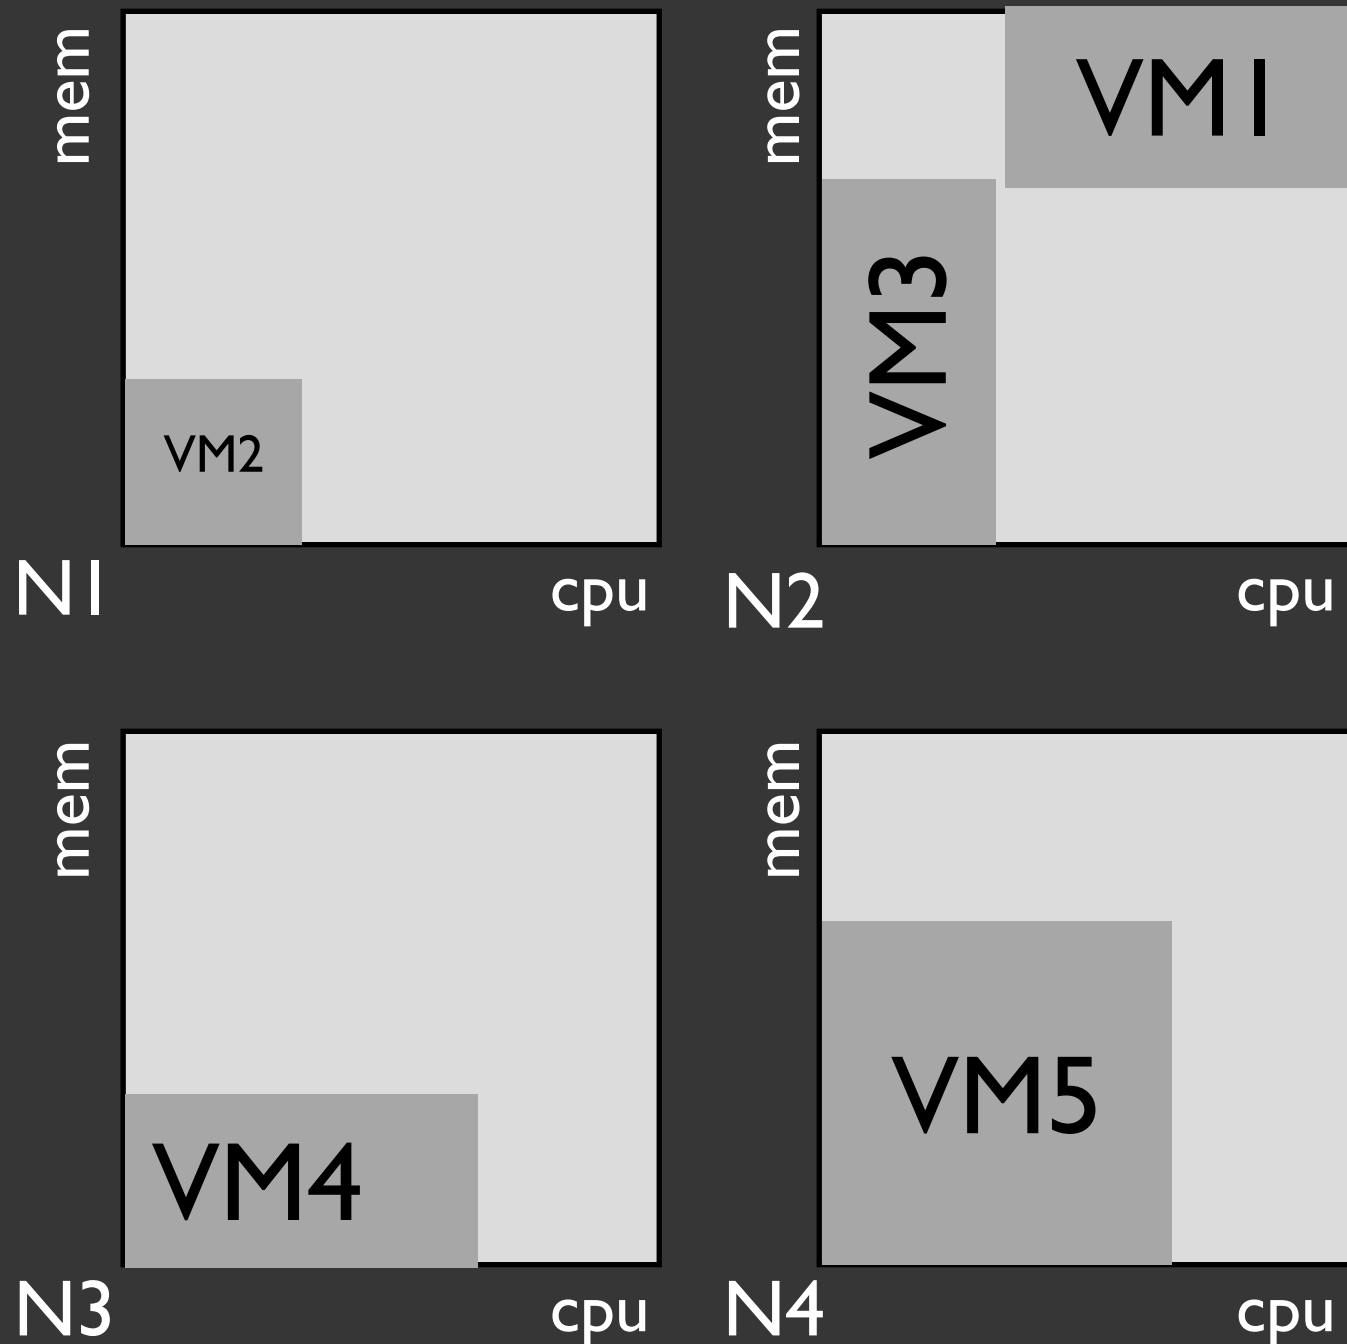
# Combinatorial problem !

past decisions impact the future ones

# static schedulers

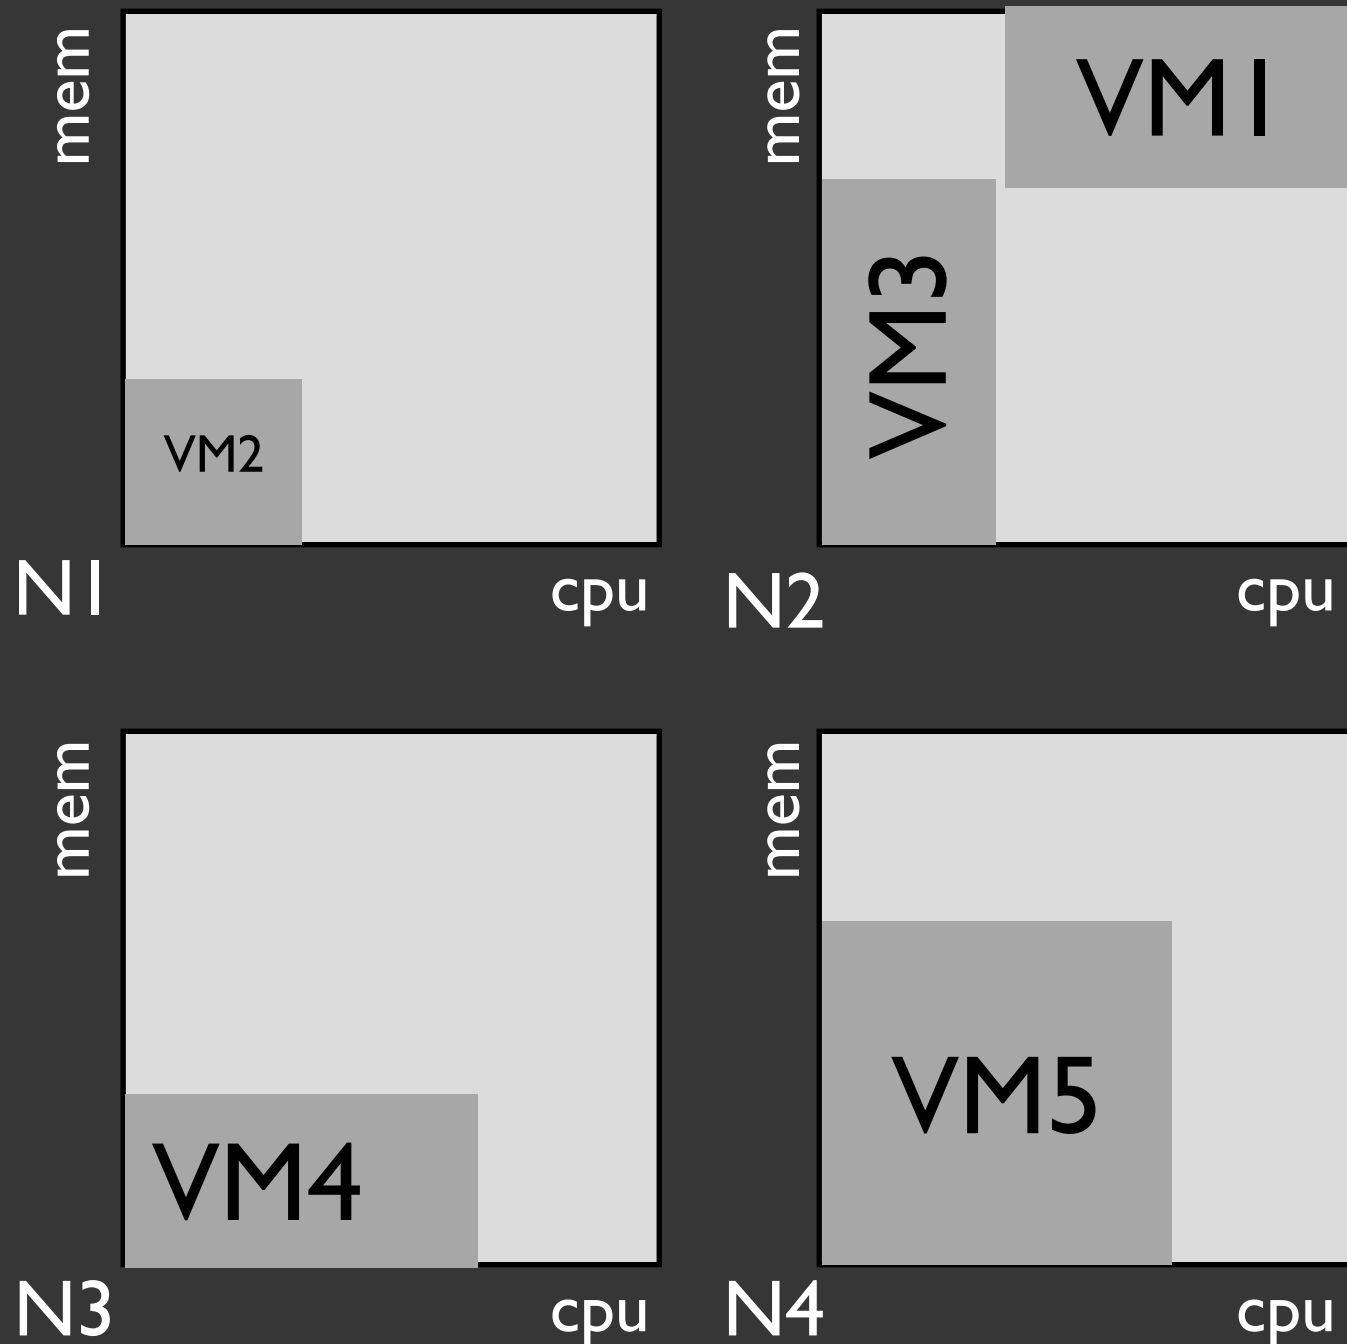
buffering might help

VM6



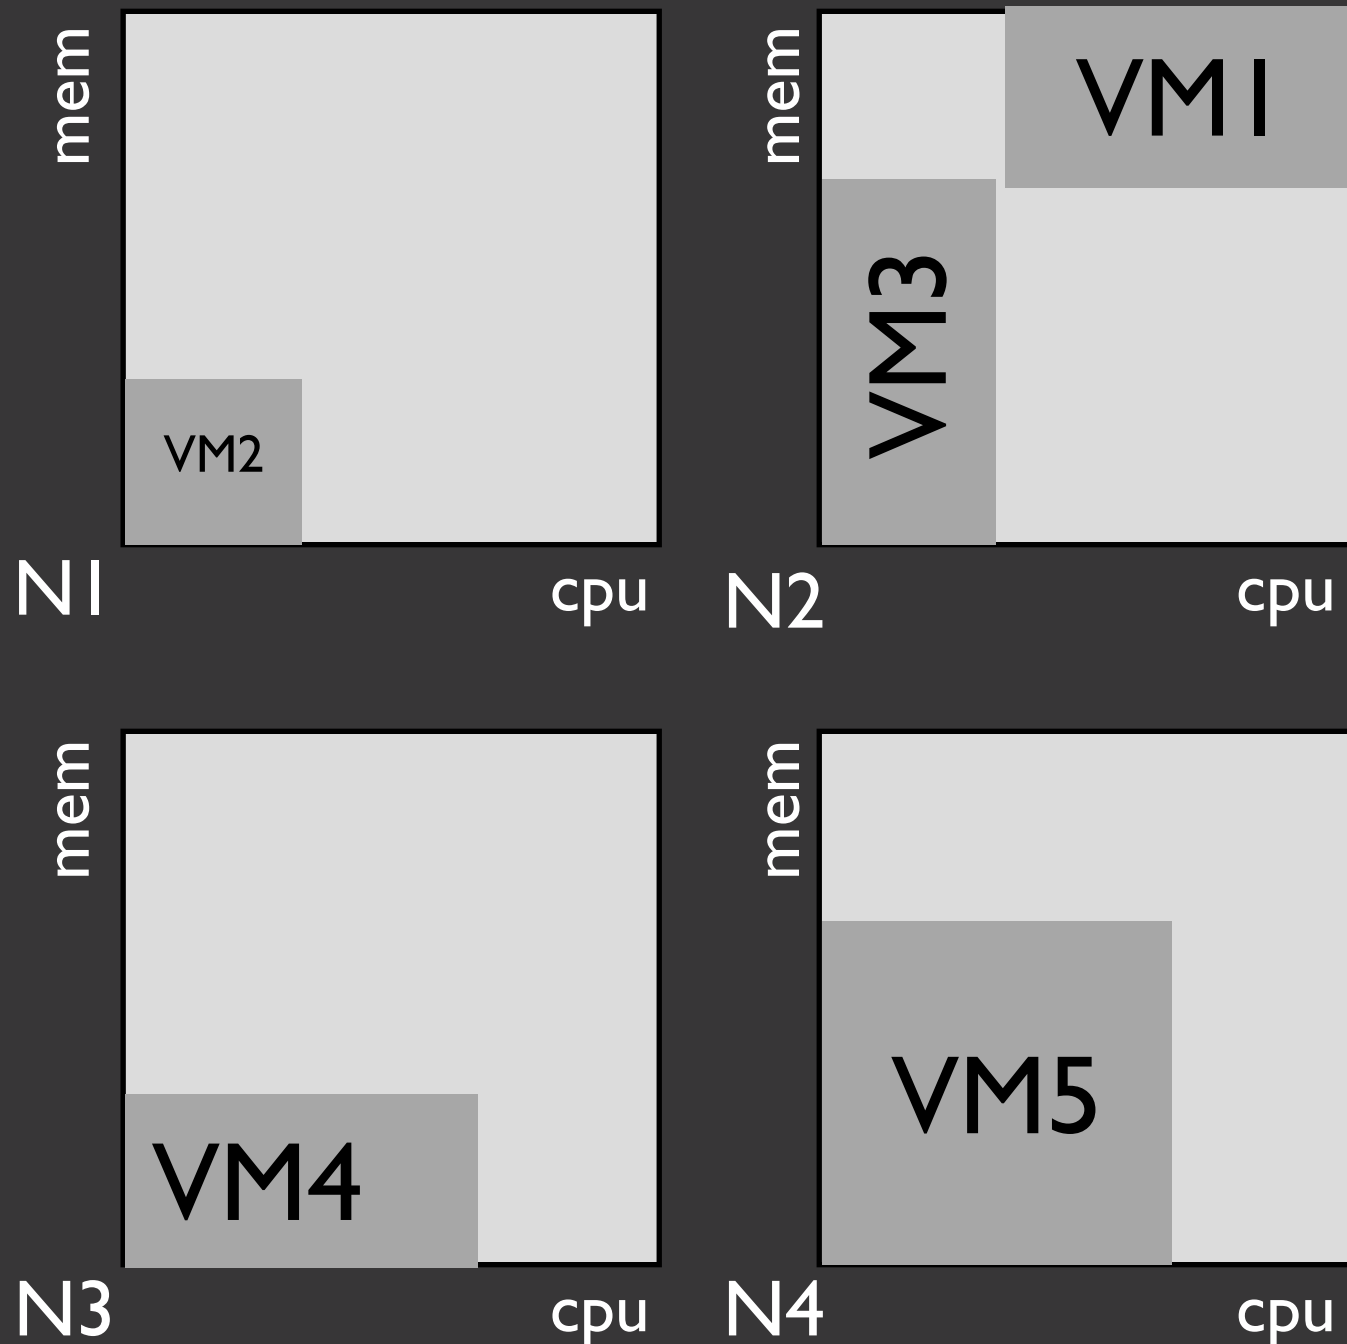
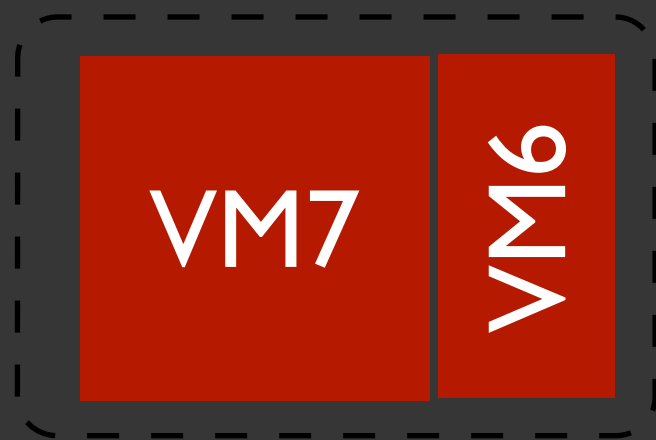
# static schedulers

buffering might help



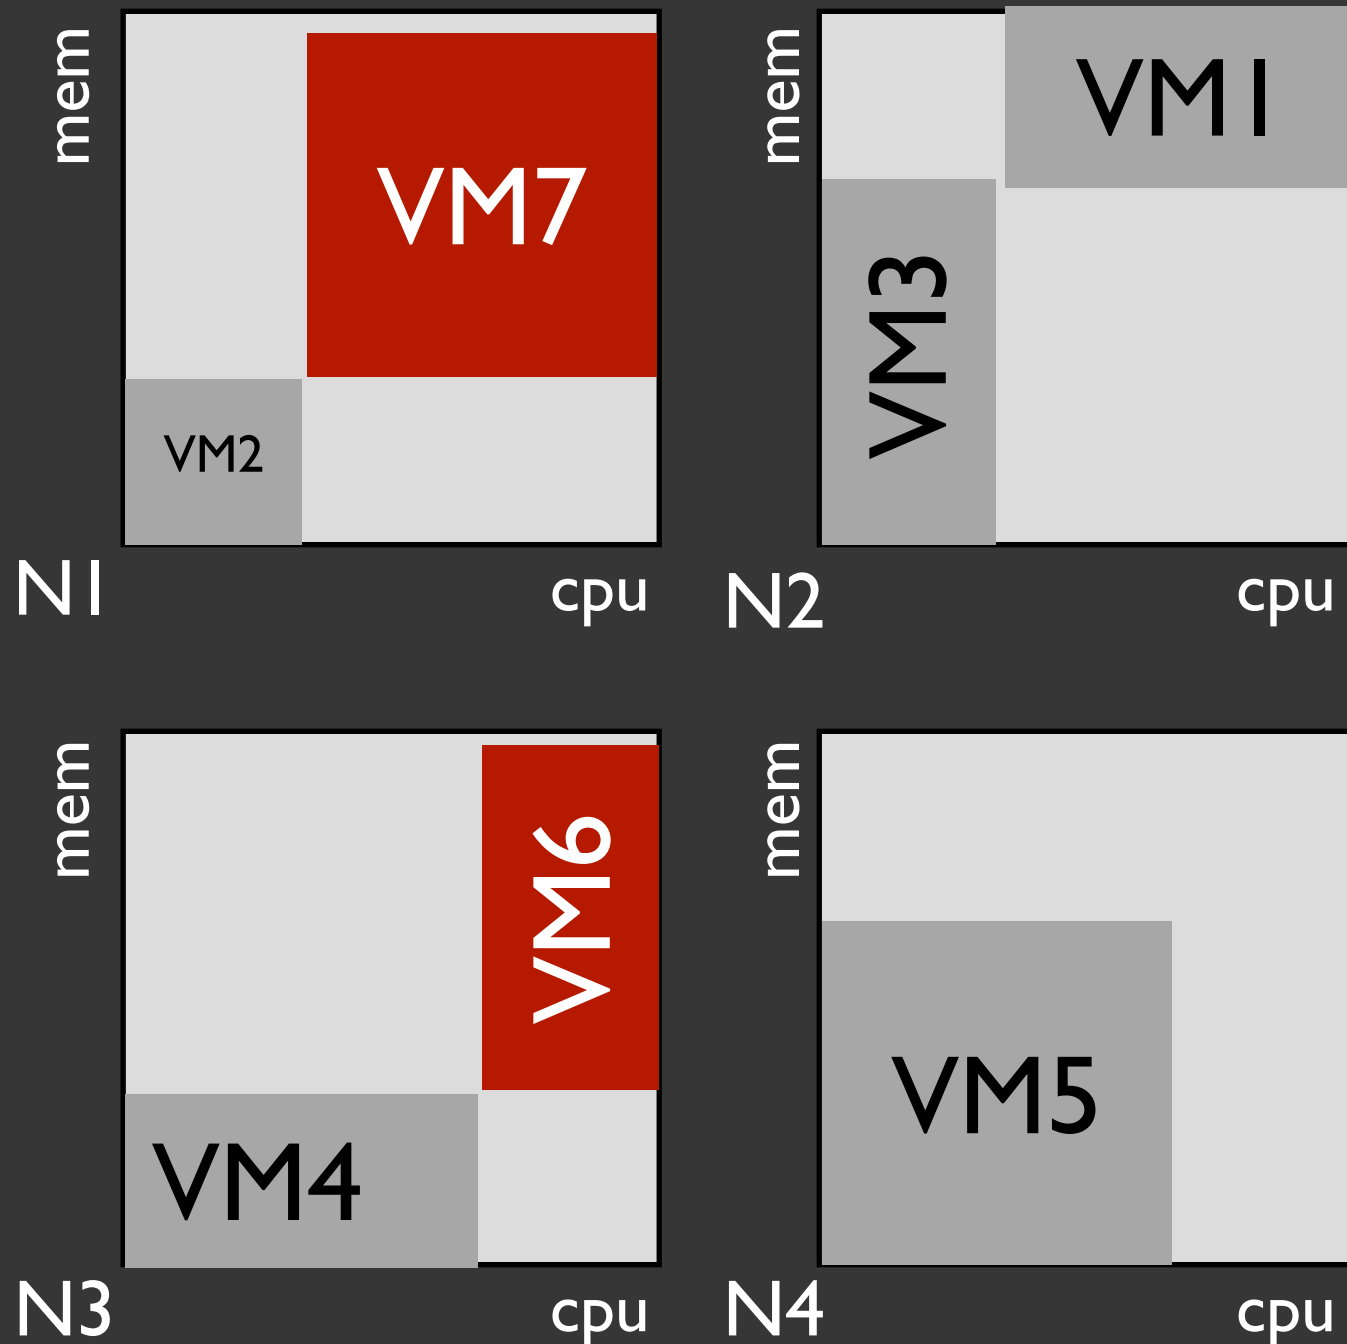
# static schedulers

buffering might help



# static schedulers

buffering might help

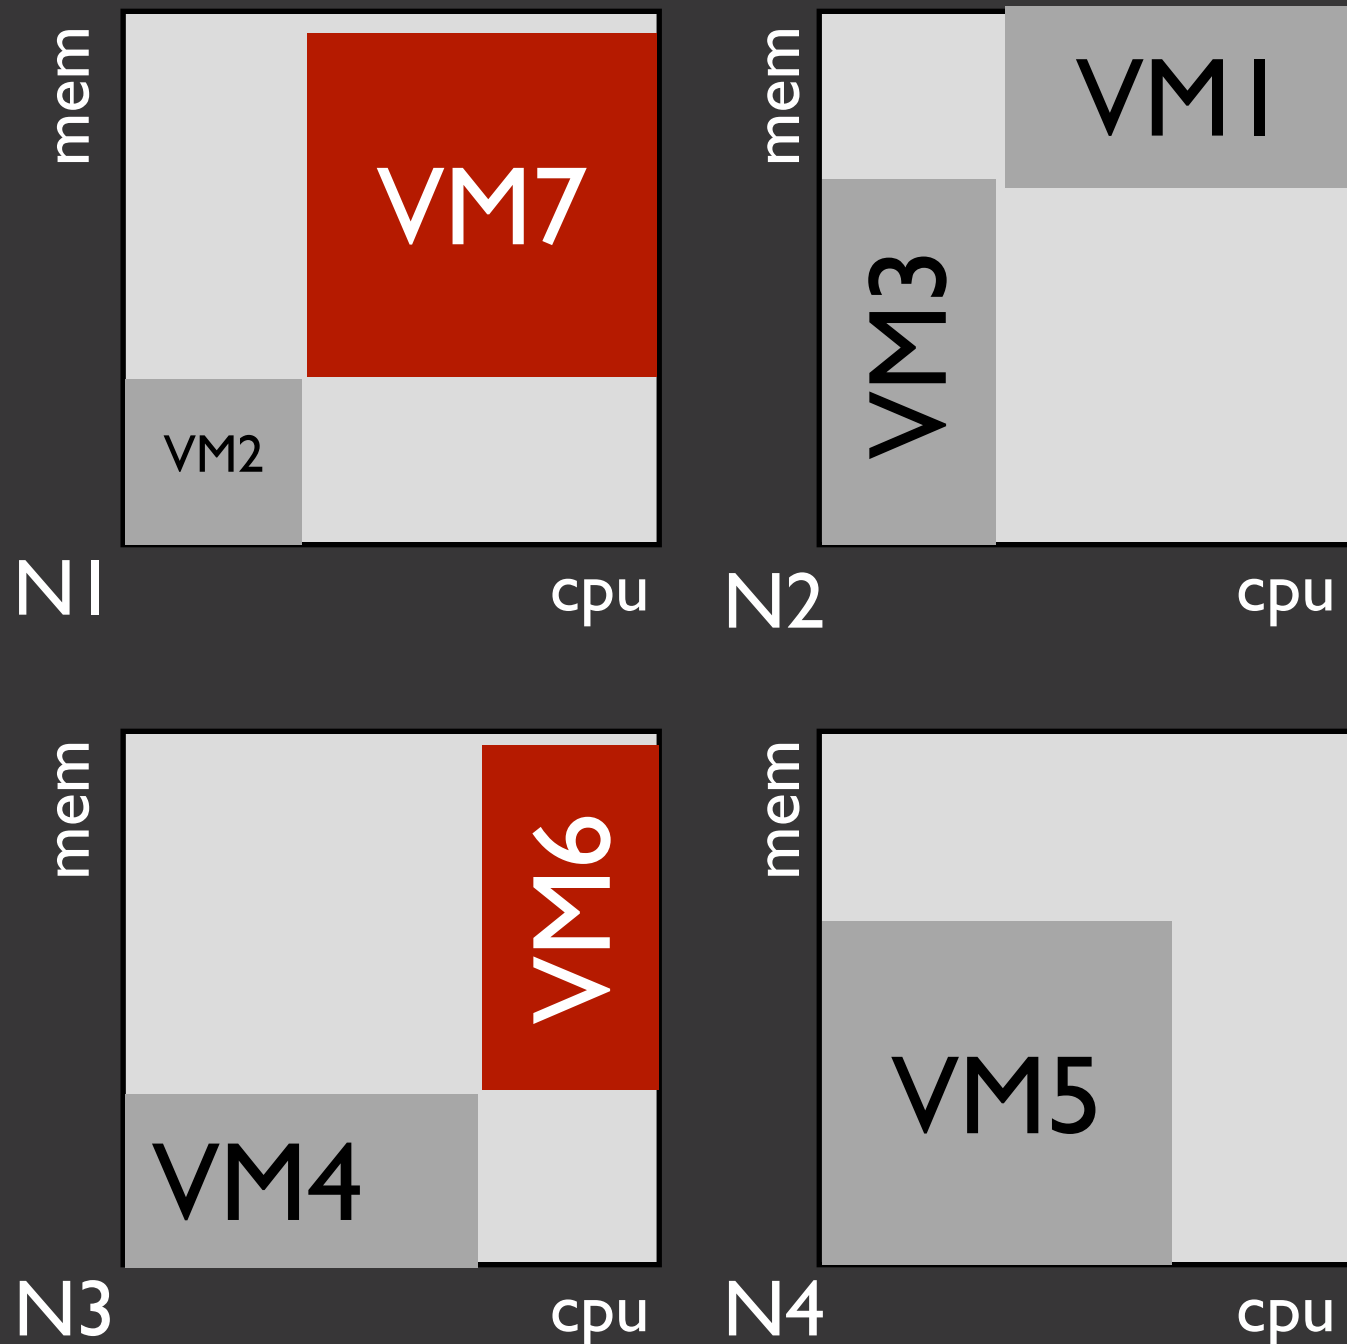


# static schedulers

buffering might help



incoming rate  
starvation



# static schedulers

manipulate the VM queue only  
react on VM arrival



# static schedulers

## pros

- easy to model
- simple standard actions
- manage a few elements

# static schedulers

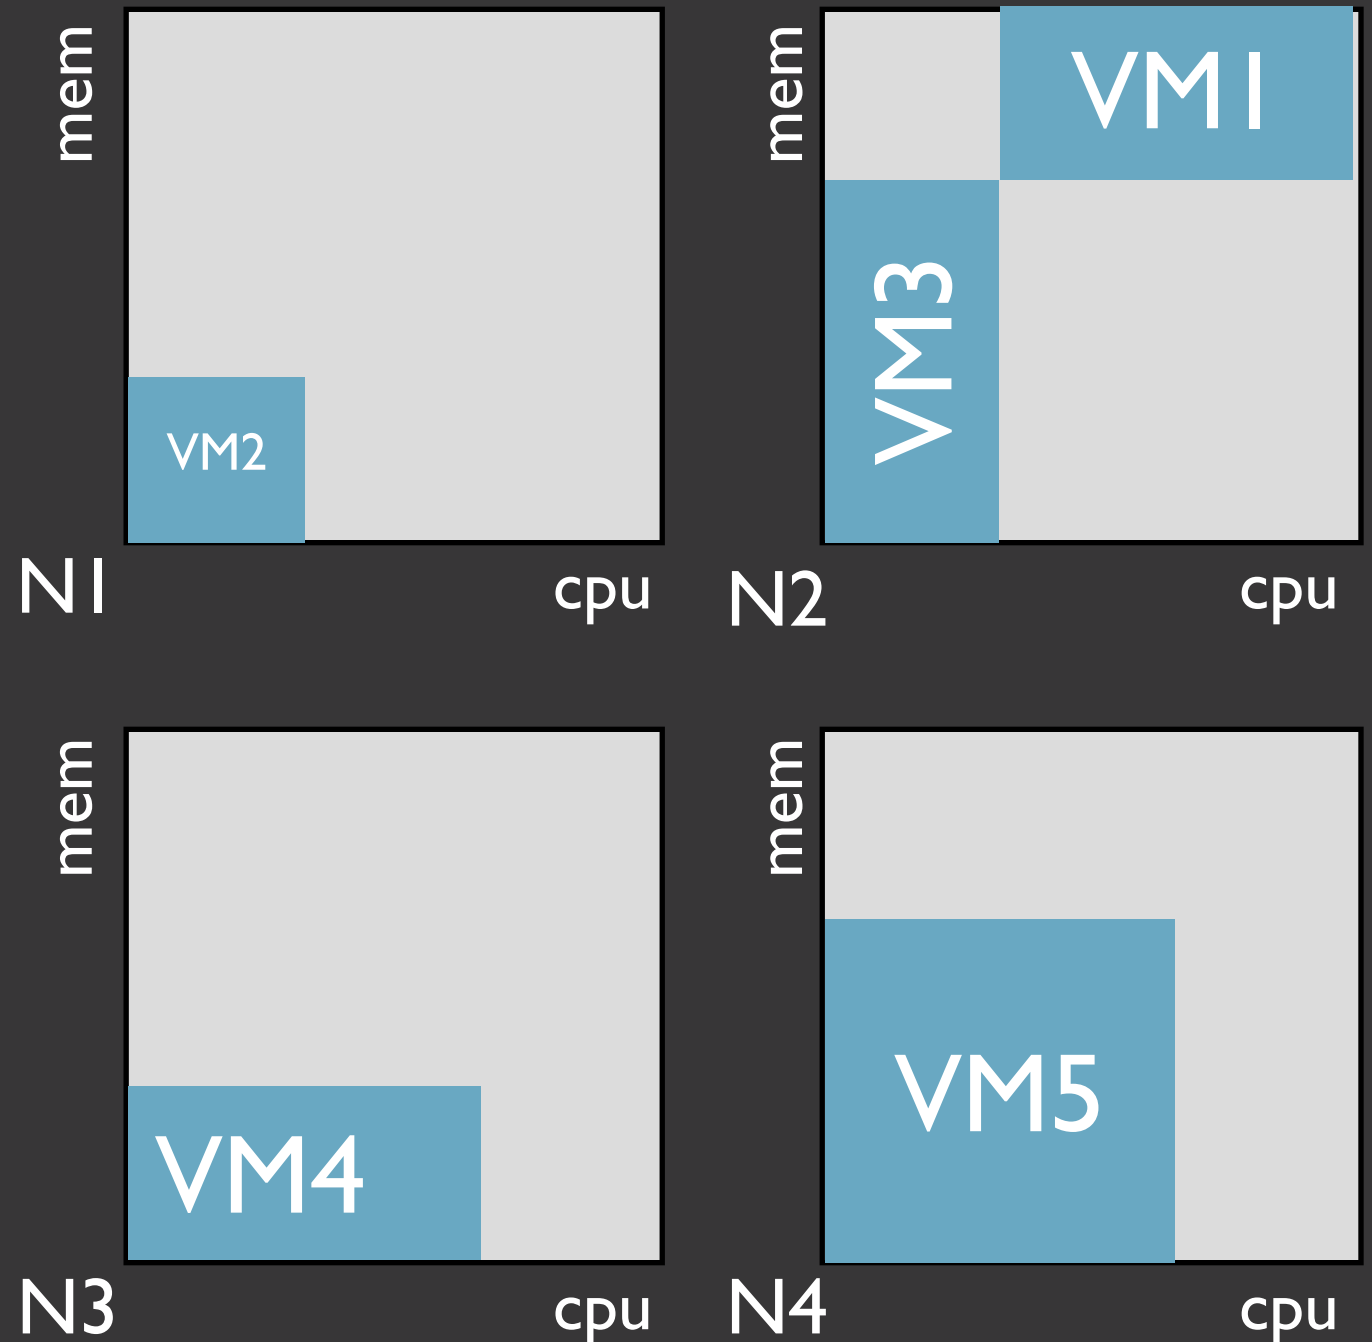
cons

hard to perform fine  
grain optimisation

# dynamic schedulers

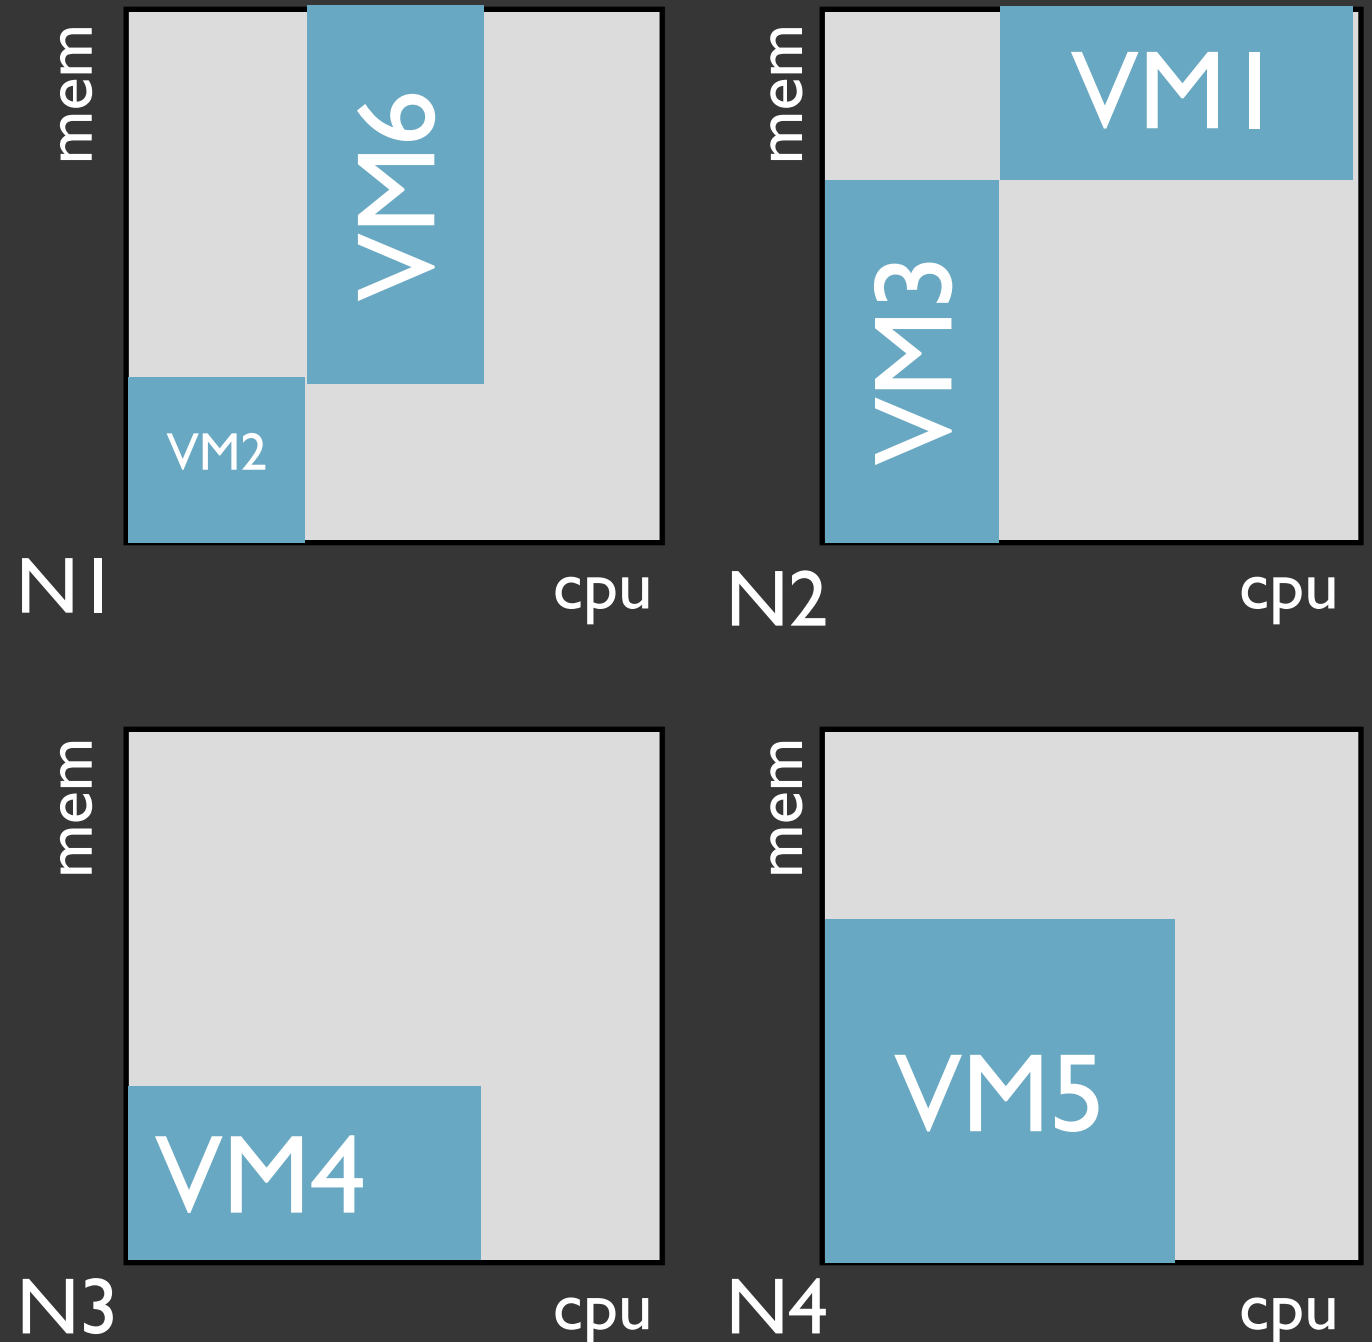
consider every VMs

VM6



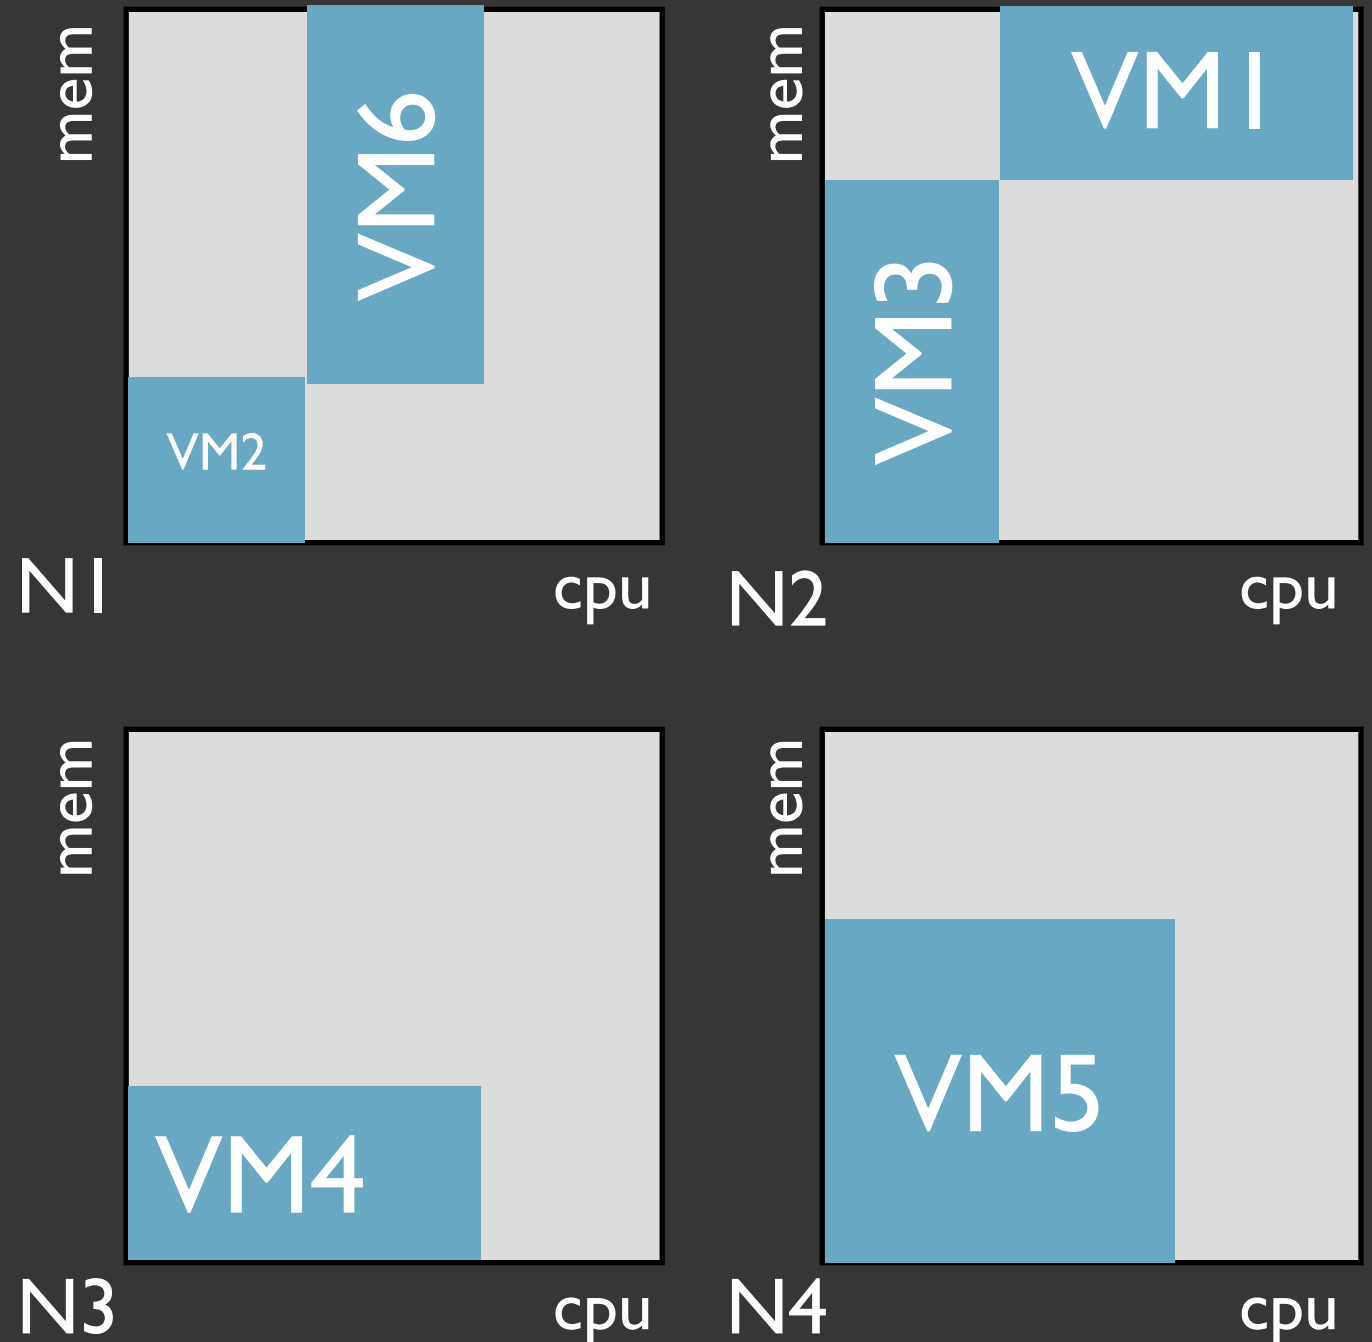
# dynamic schedulers

consider every VMs



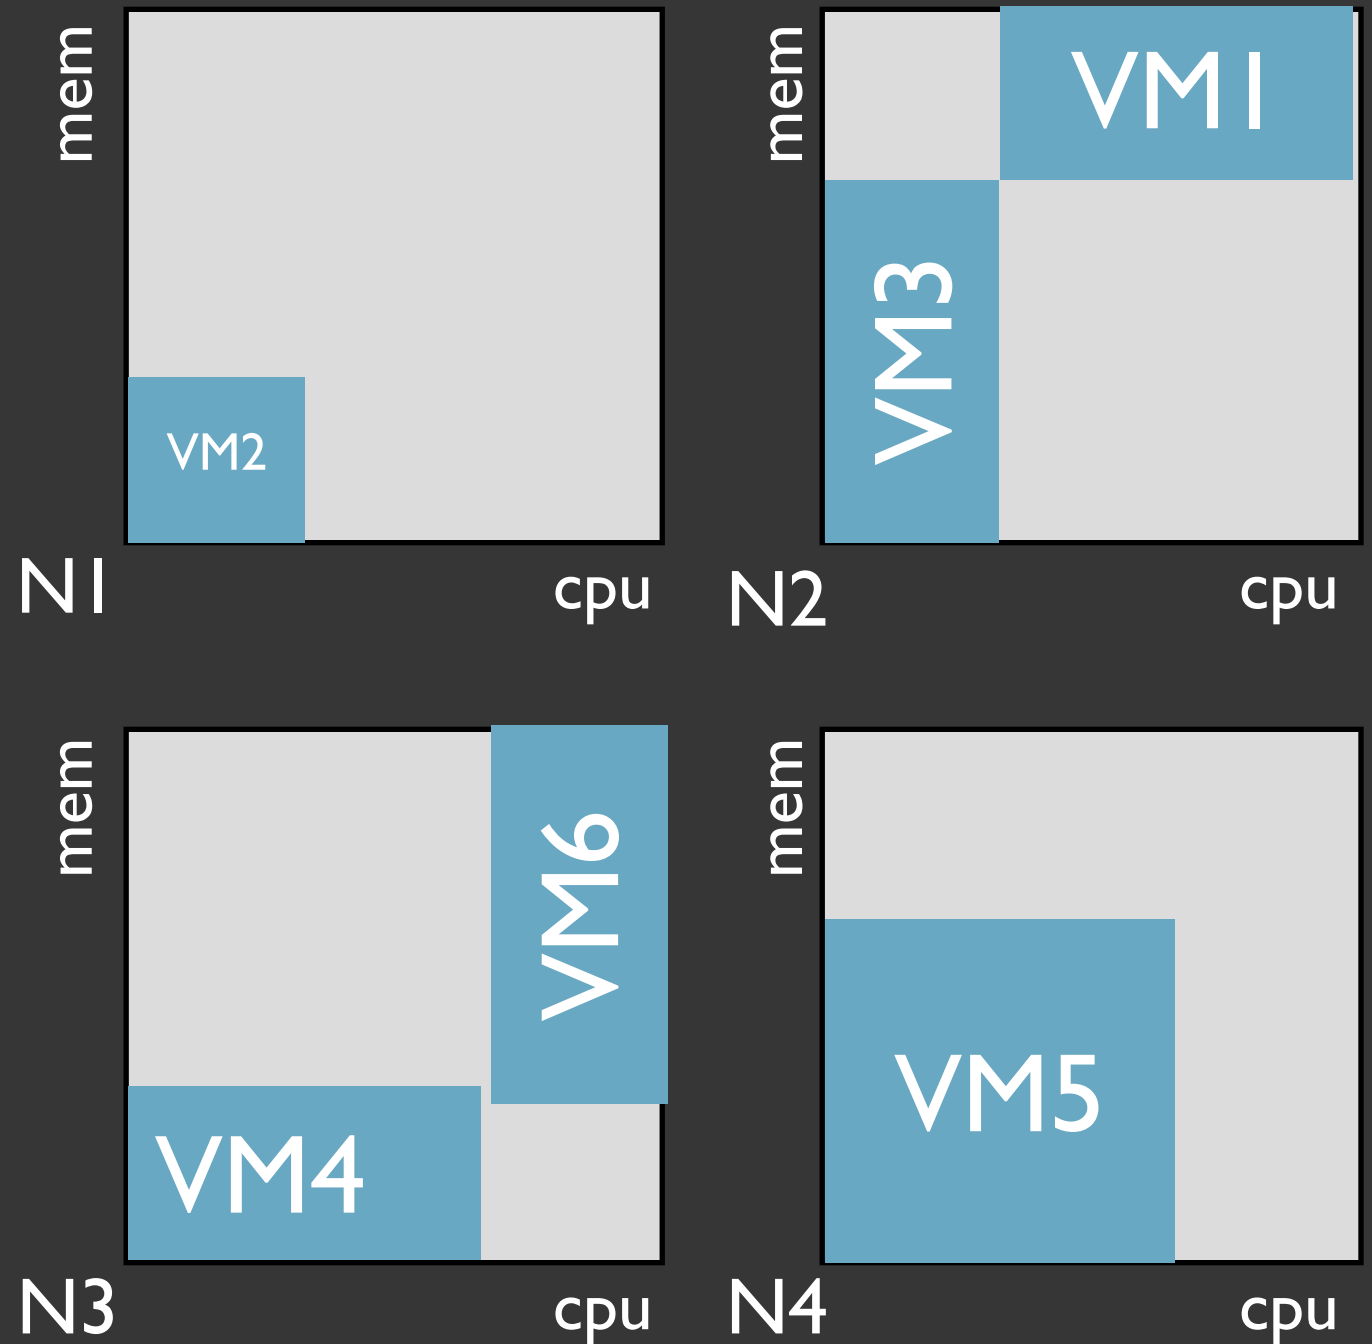
# dynamic schedulers

consider *every* VMs



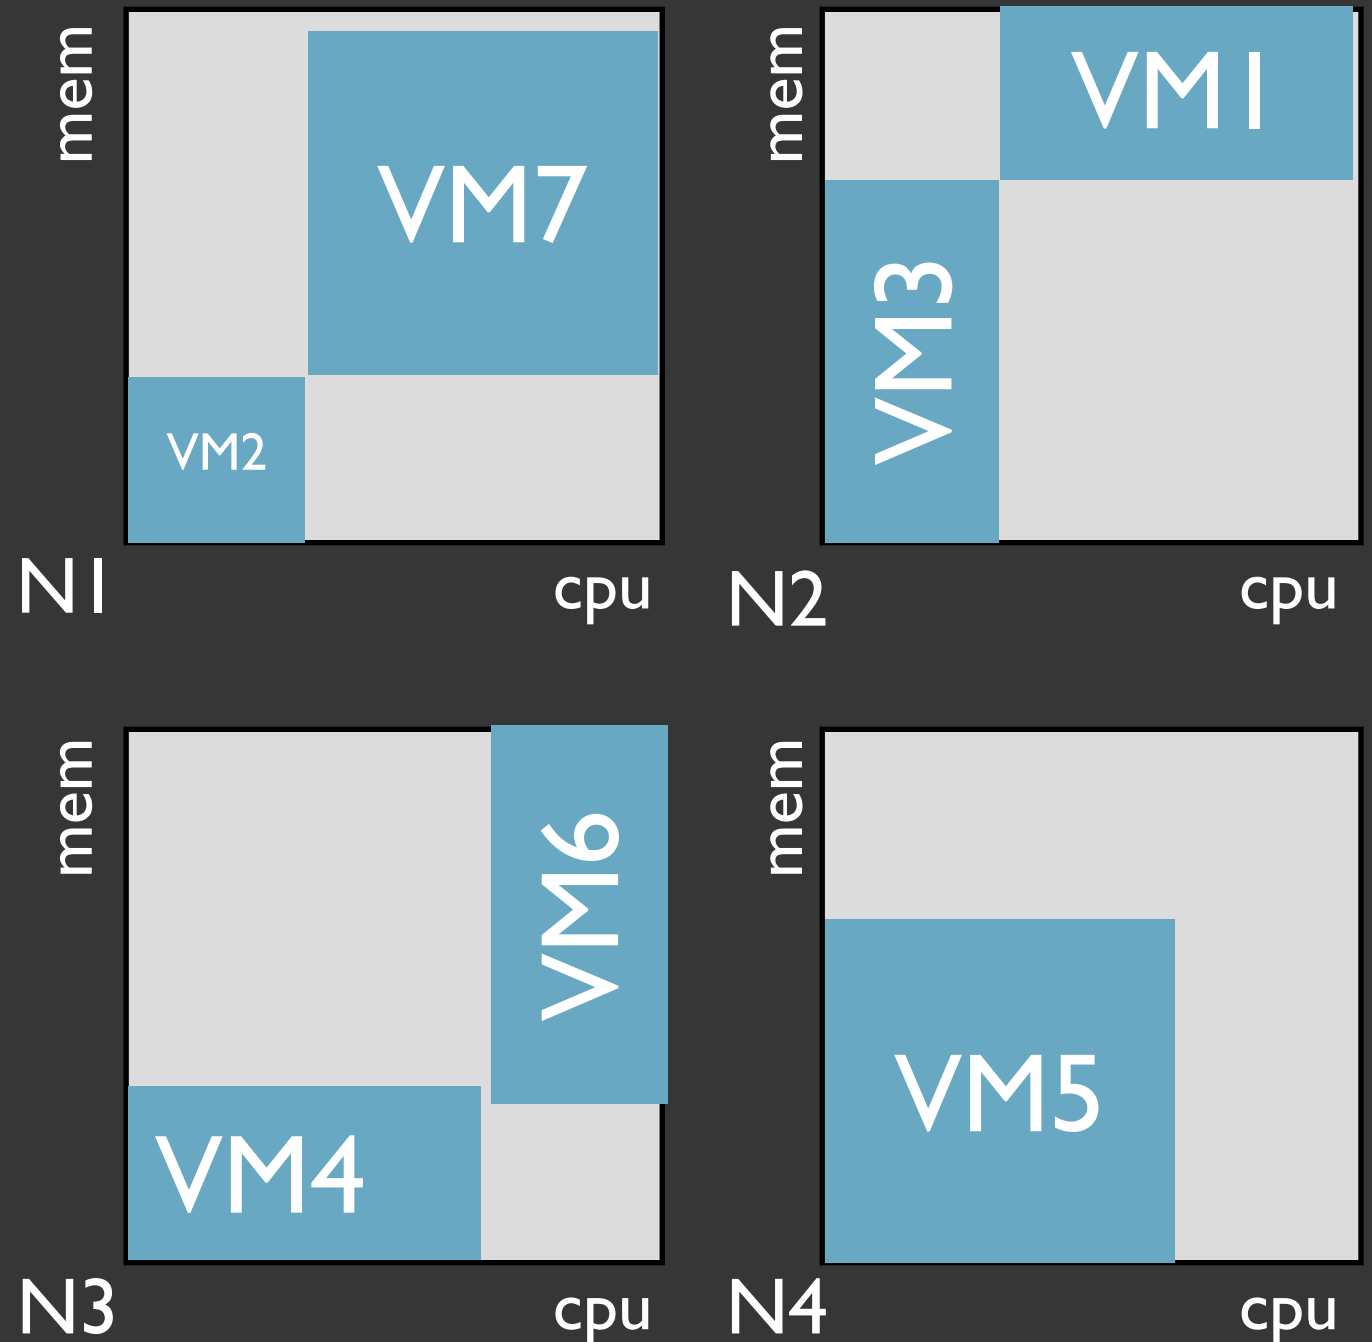
# dynamic schedulers

consider *every* VMs



# dynamic schedulers

consider *every* VMs



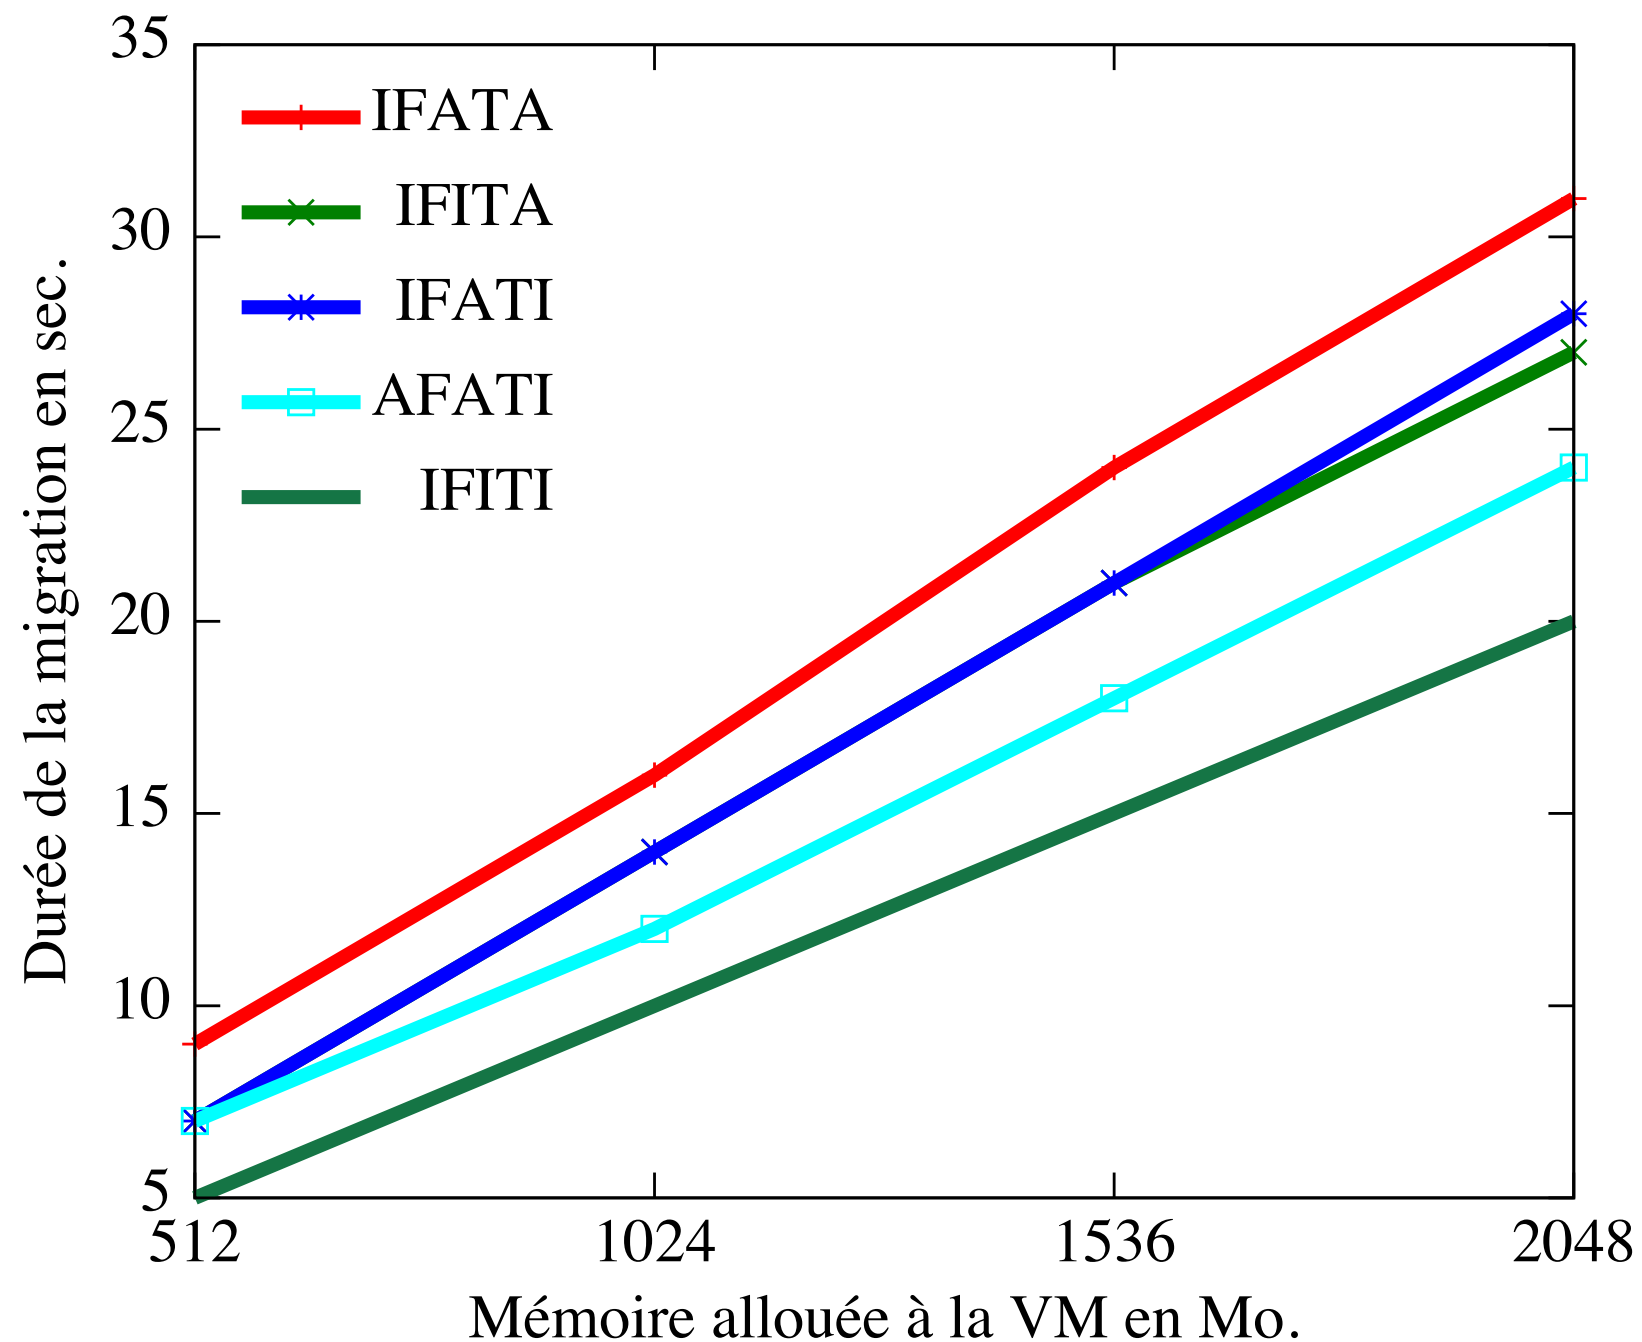
# dynamic schedulers

manipulate the VM queue  
reconfigure the schedule with migrations



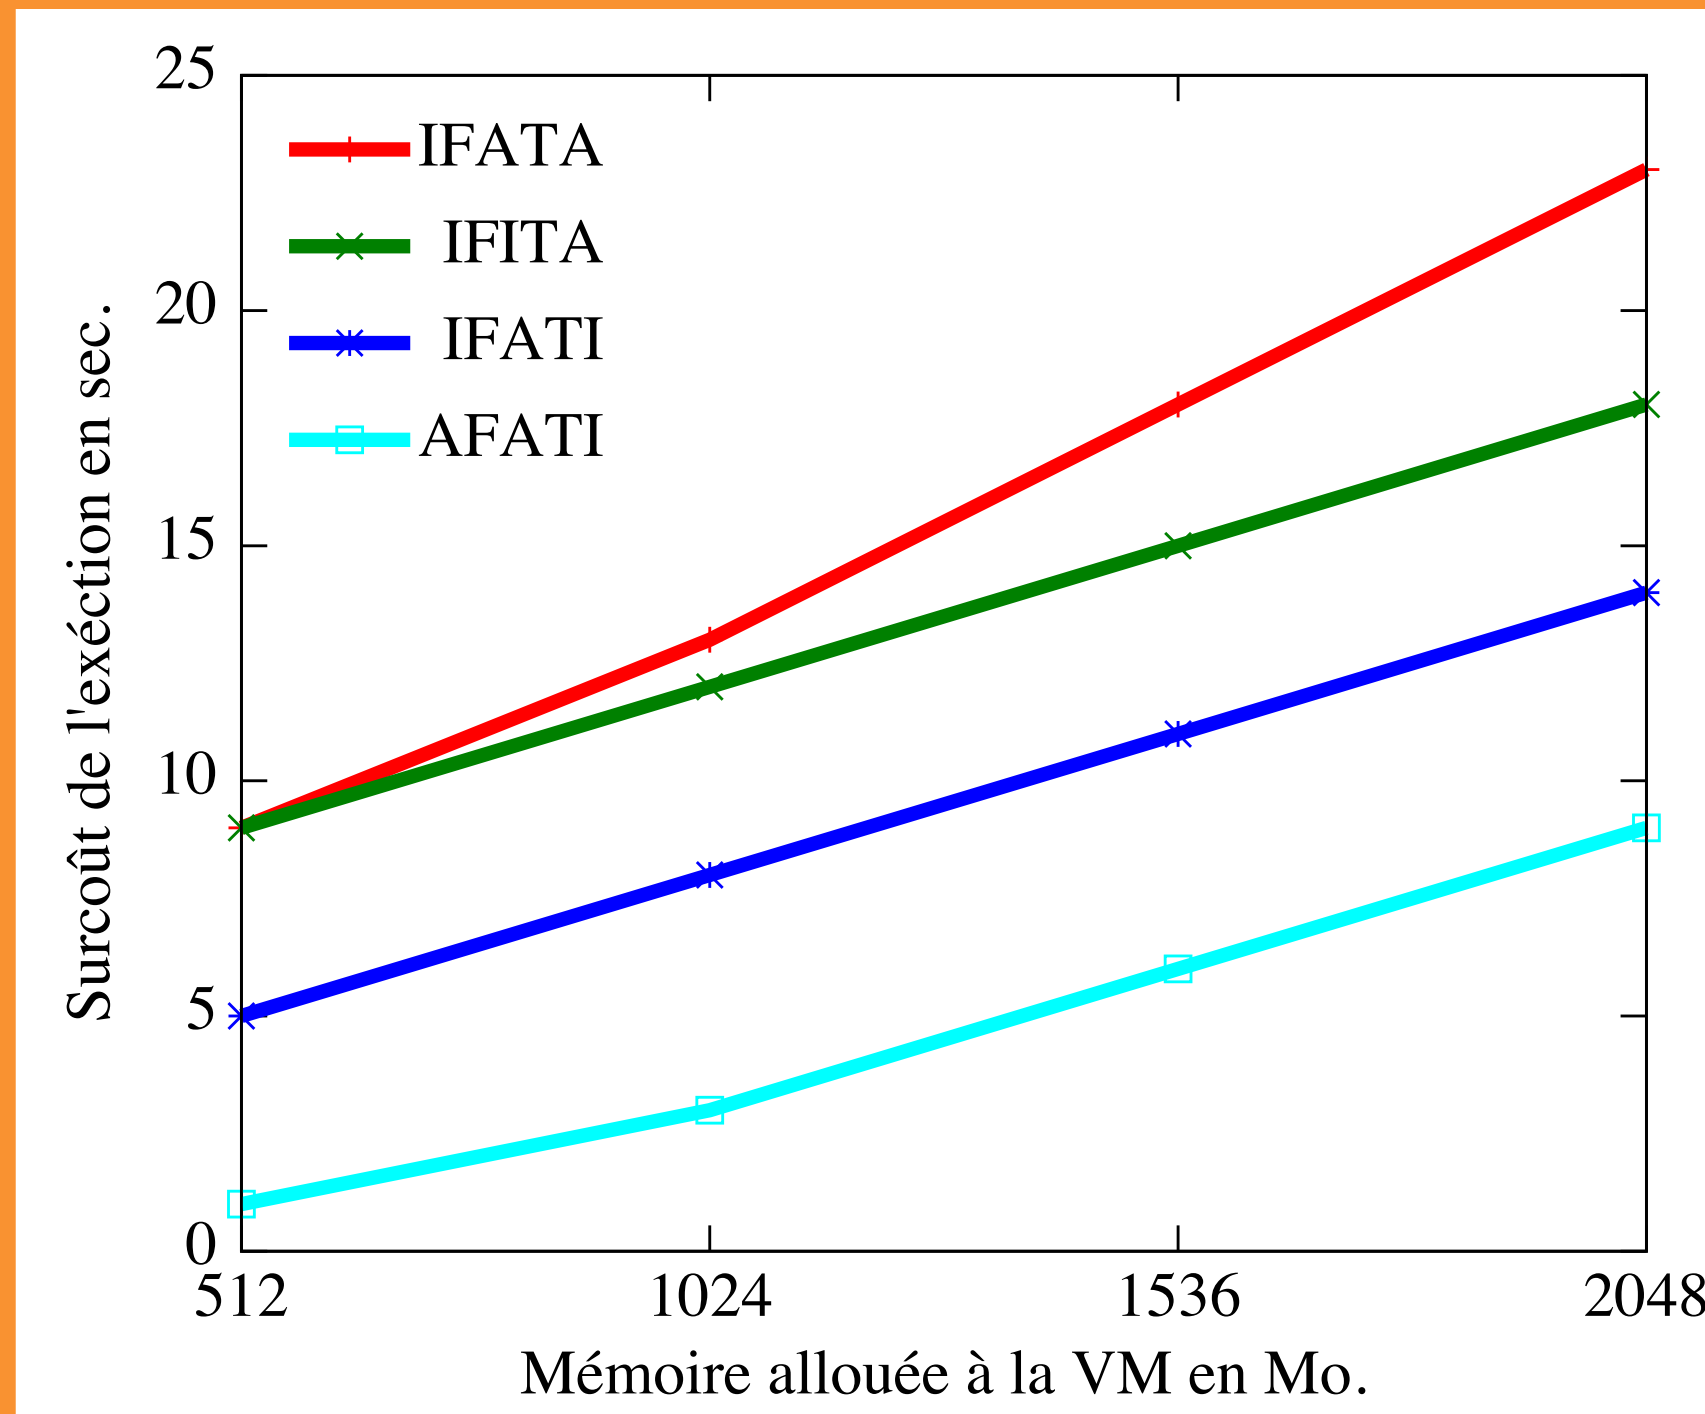
# dynamic schedulers

a migration takes time



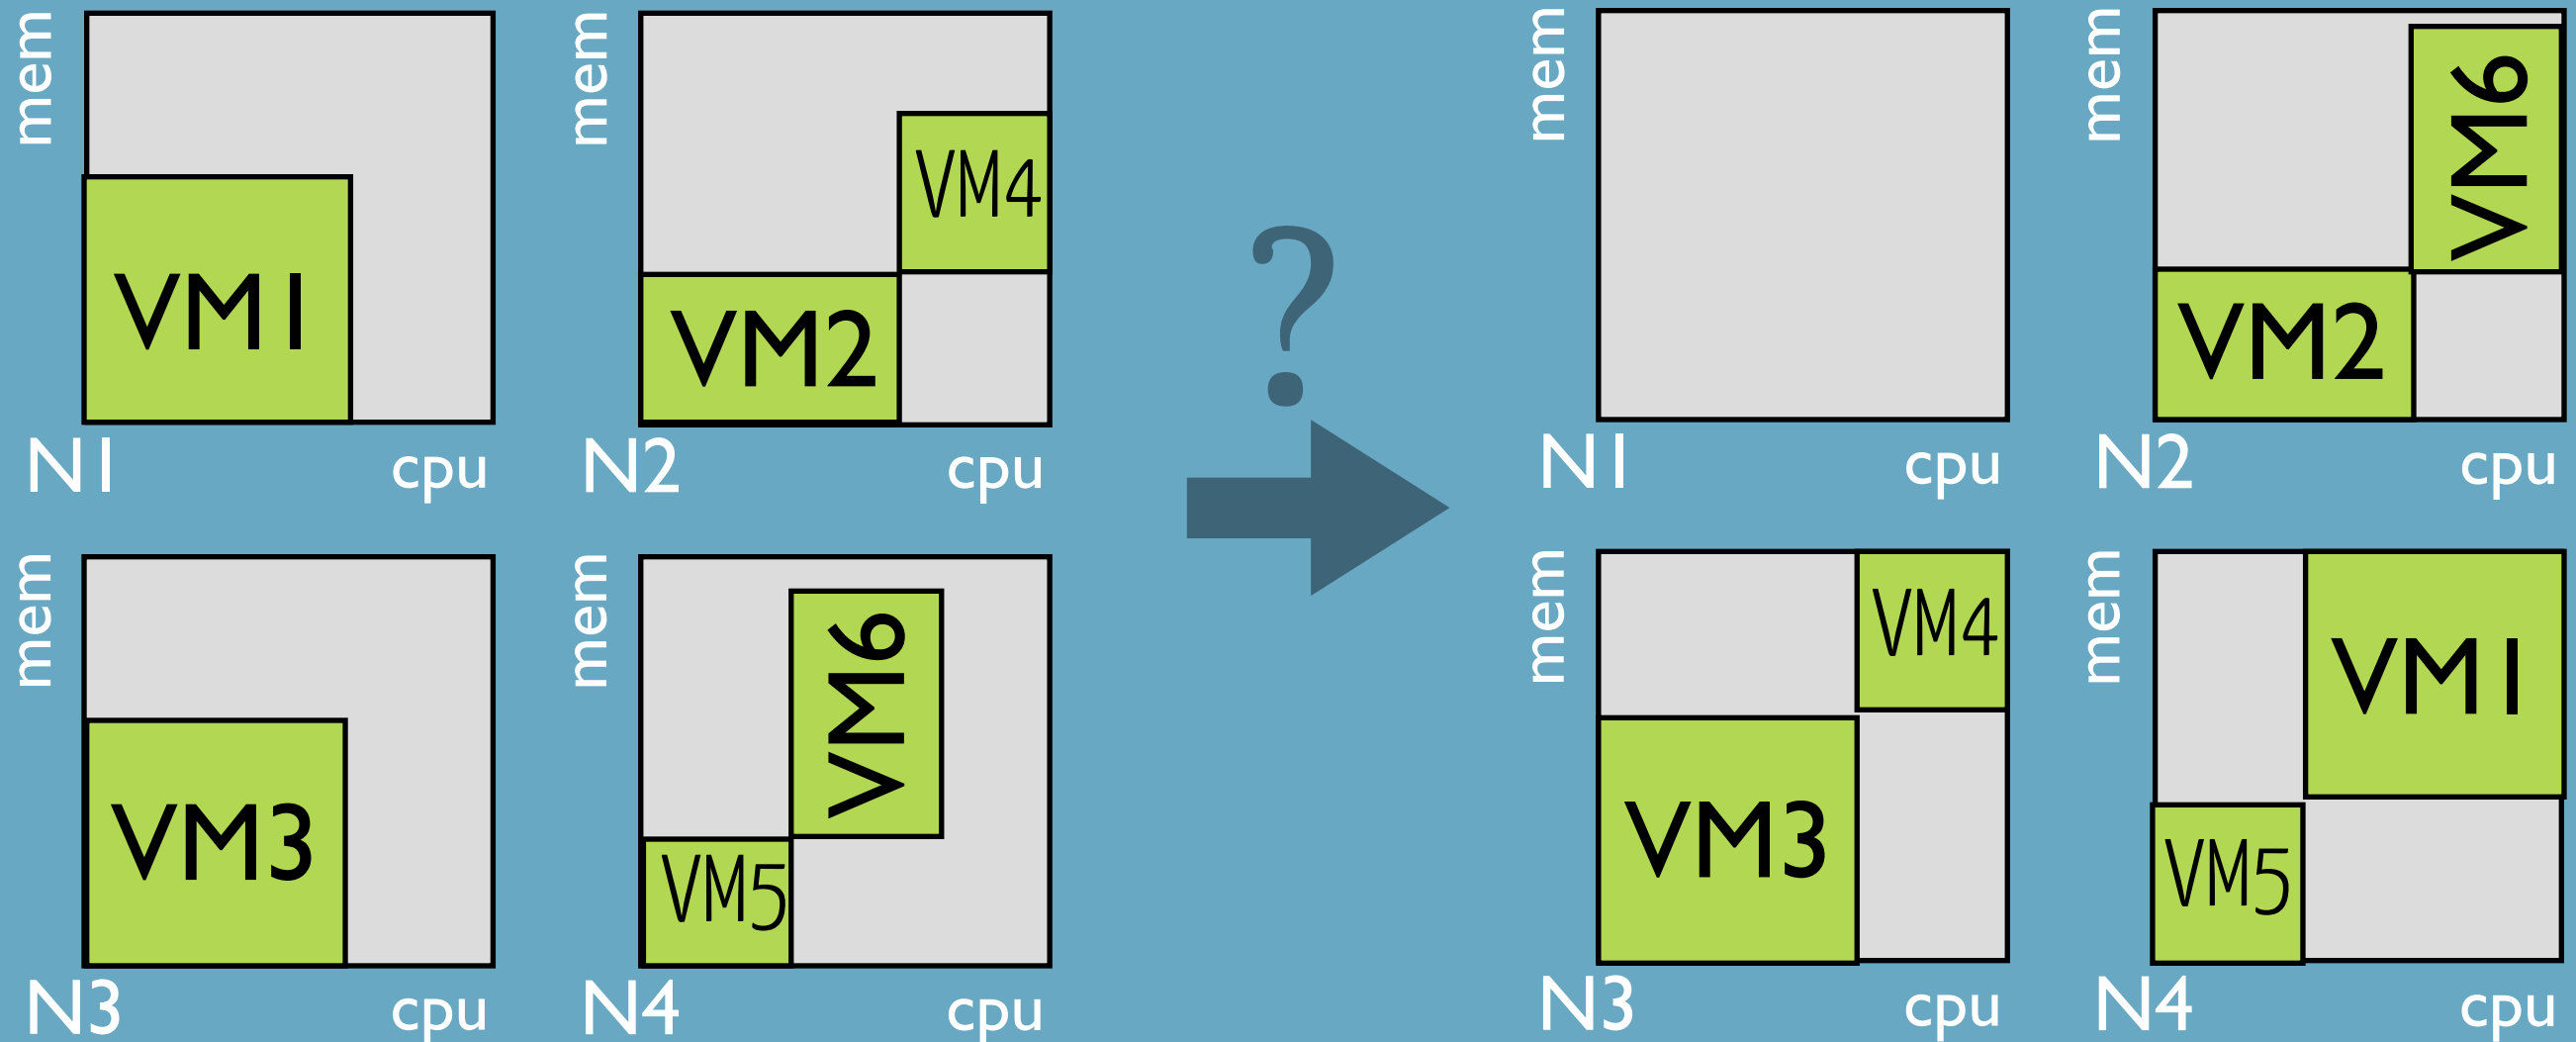
# dynamic schedulers

a migration impacts performance



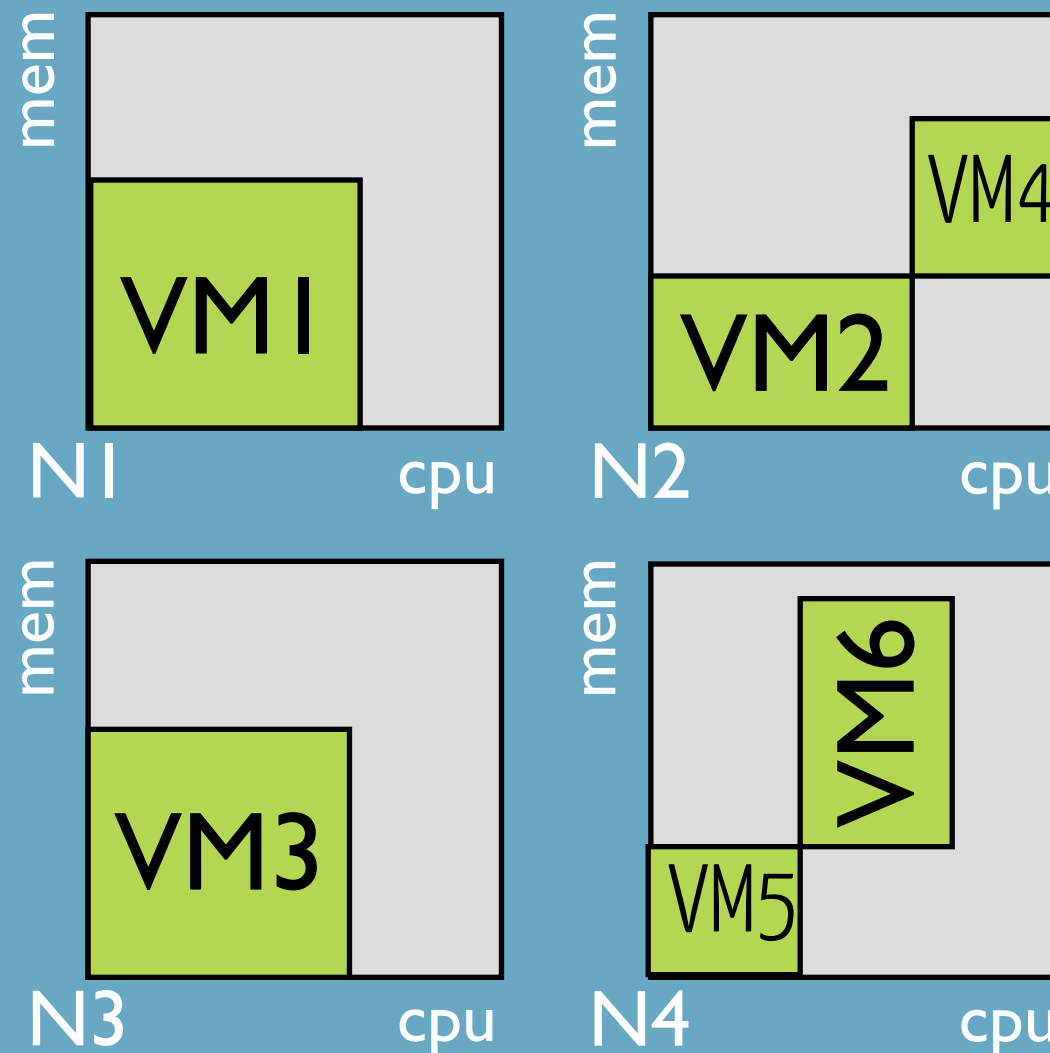
# dynamic schedulers

dependency management



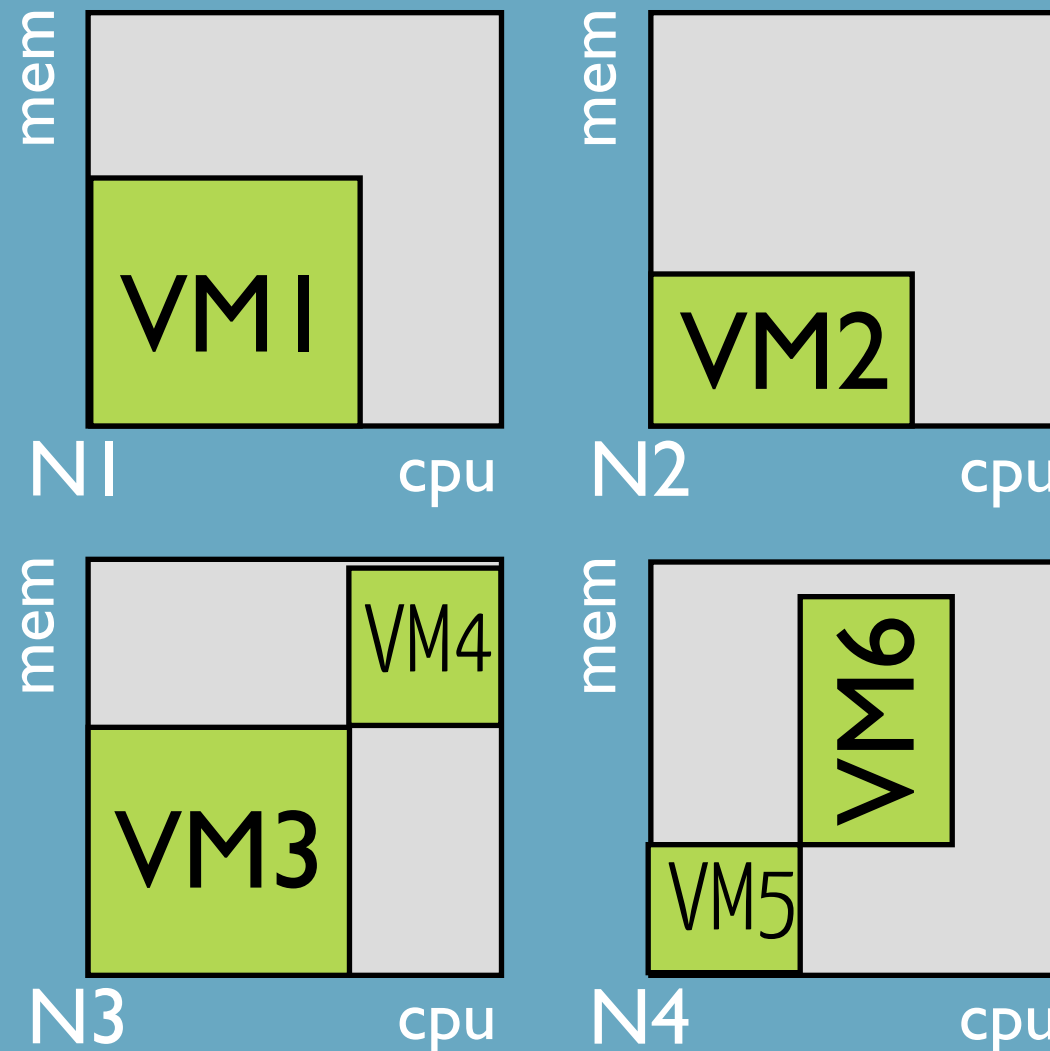
# dynamic schedulers

dependency management



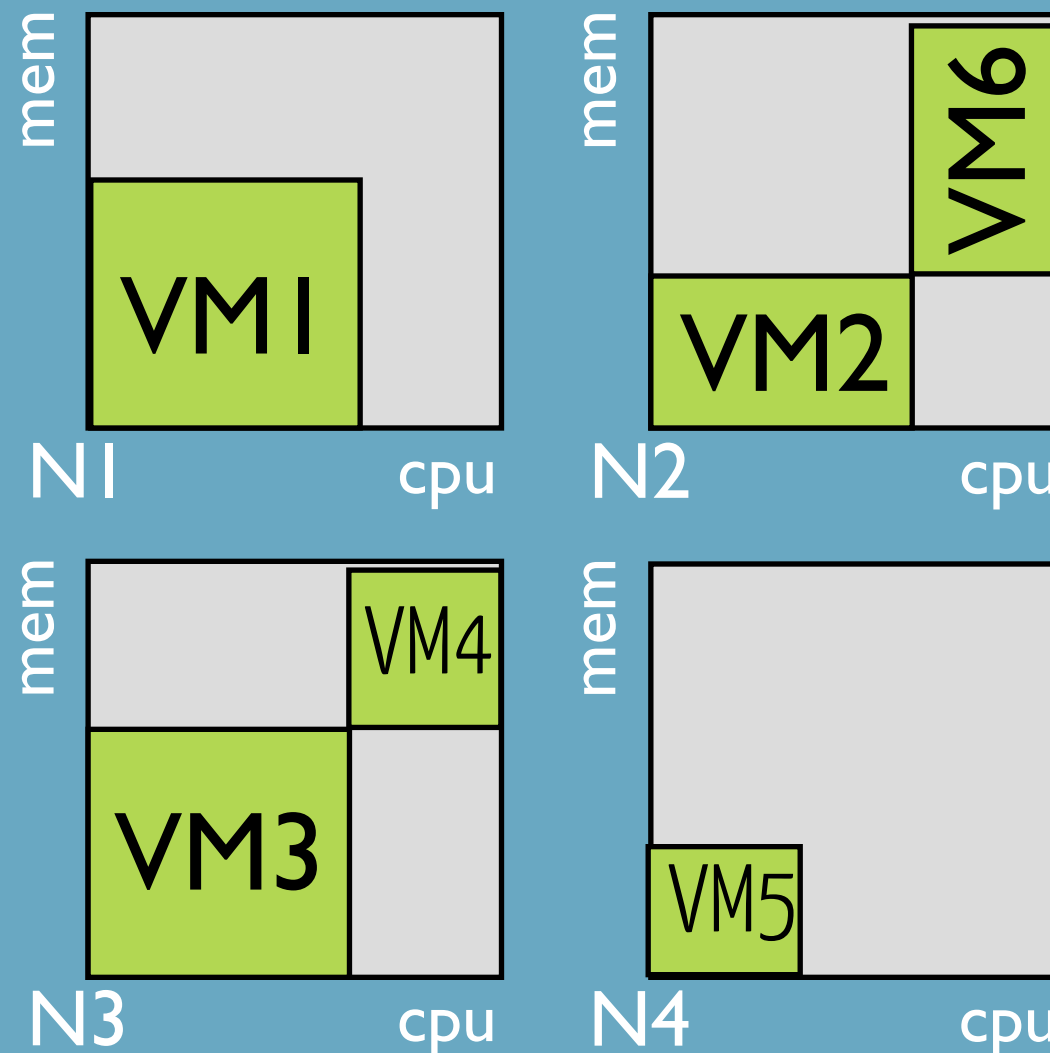
# dynamic schedulers

dependency management



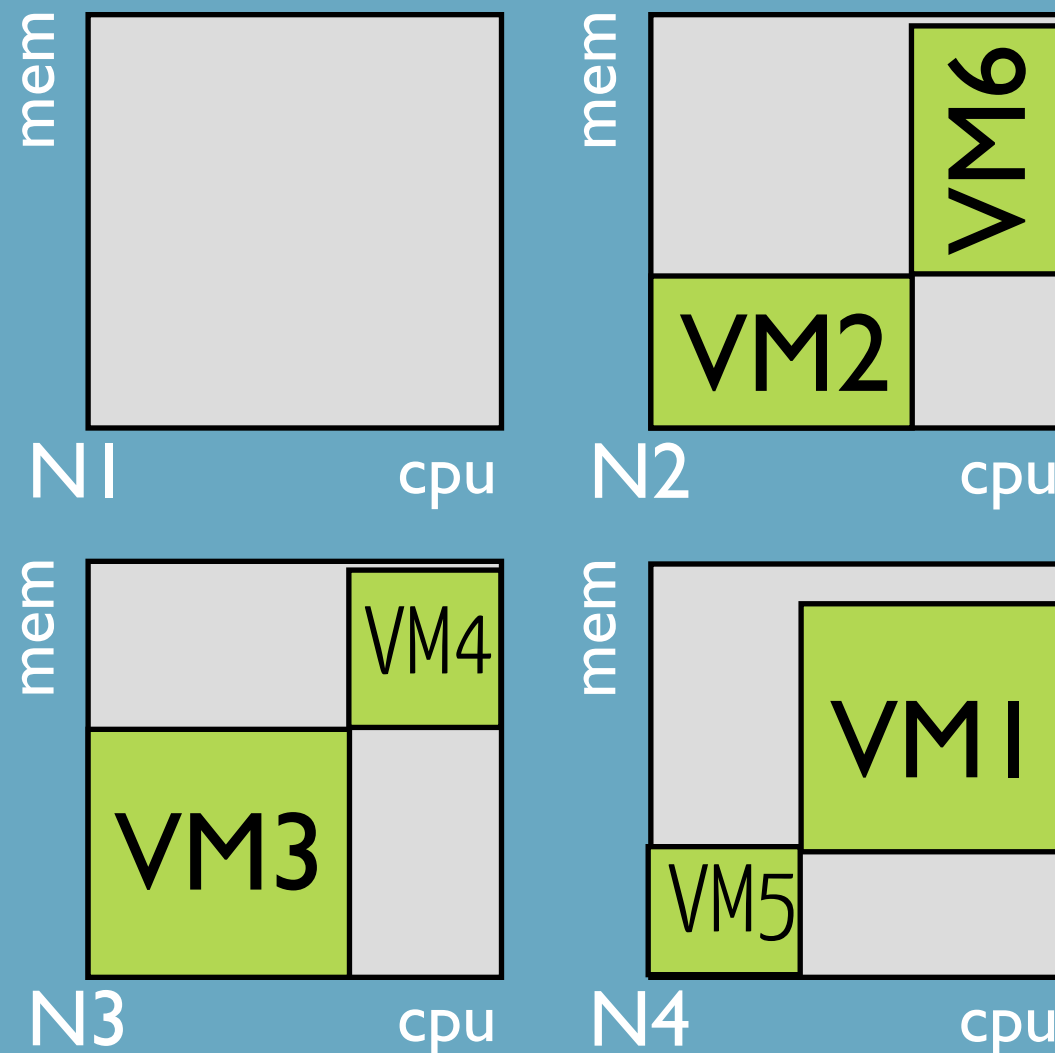
# dynamic schedulers

dependency management



# dynamic schedulers

dependency management

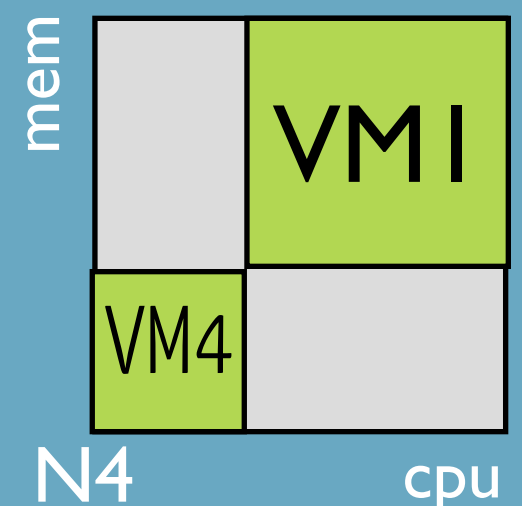
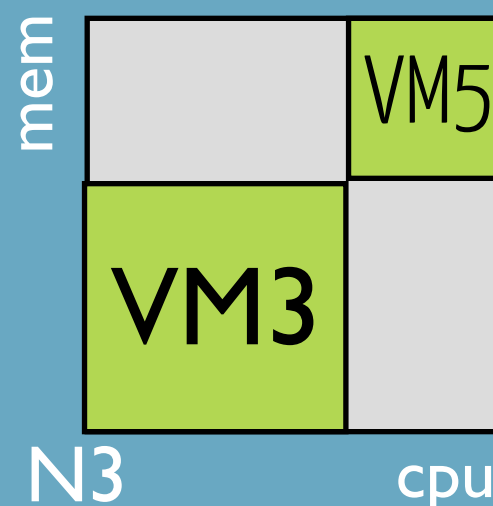
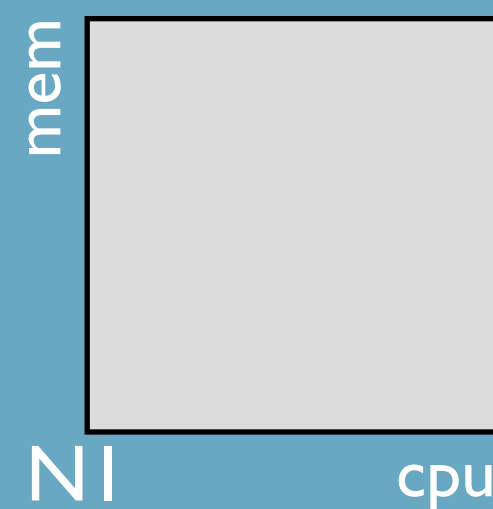
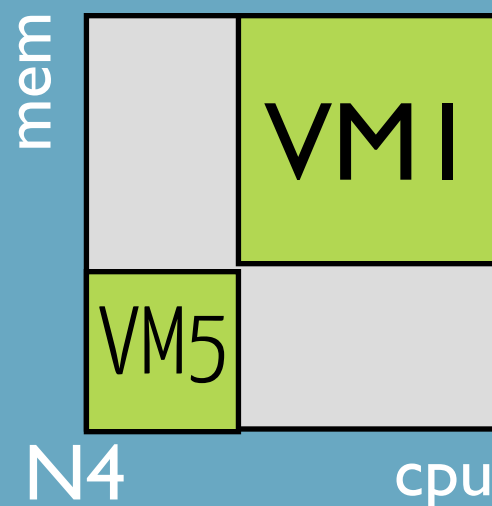
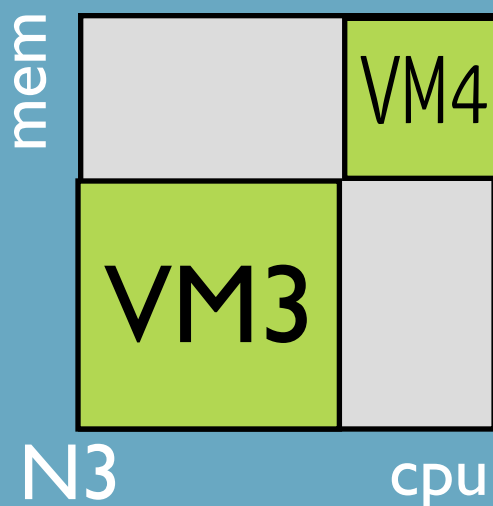
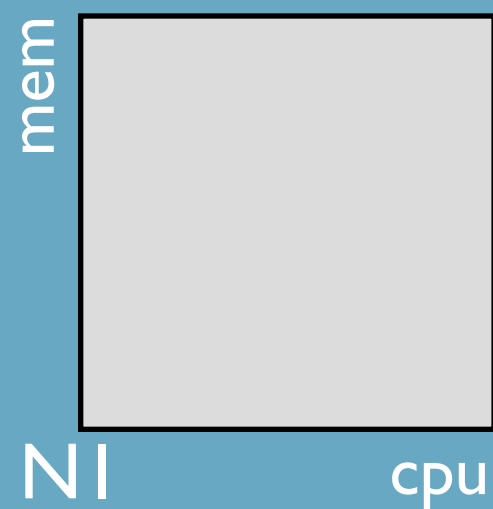


# dynamic schedulers

cyclic dependencies

**anti-affinity(VM3,VM4)**  
min(#onlineNodes)

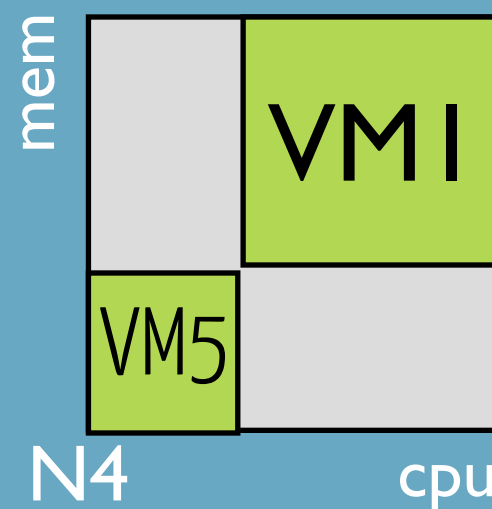
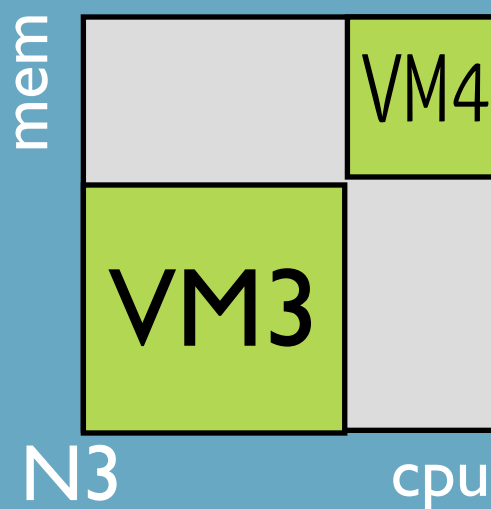
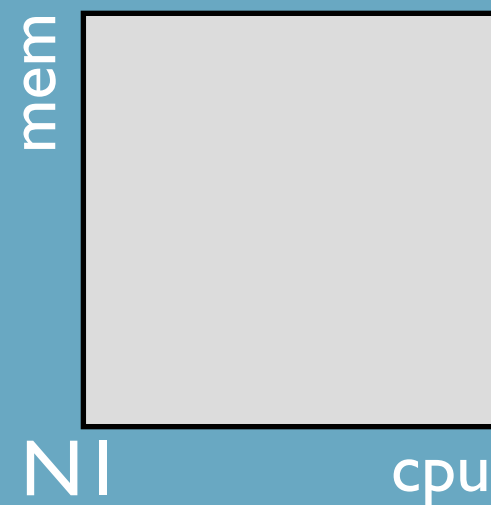
anti-affinity(VM3,VM4)  
min(#onlineNodes)





# dynamic schedulers

cyclic dependencies

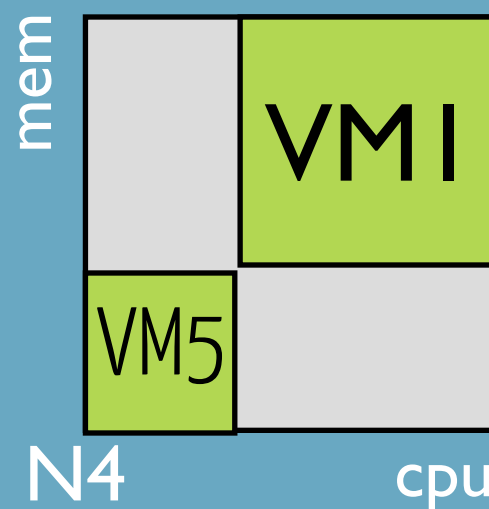
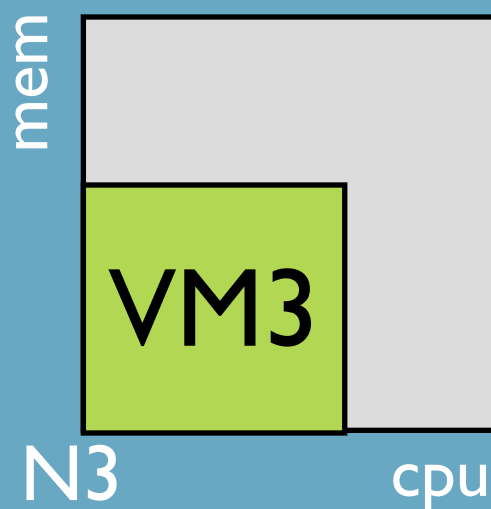
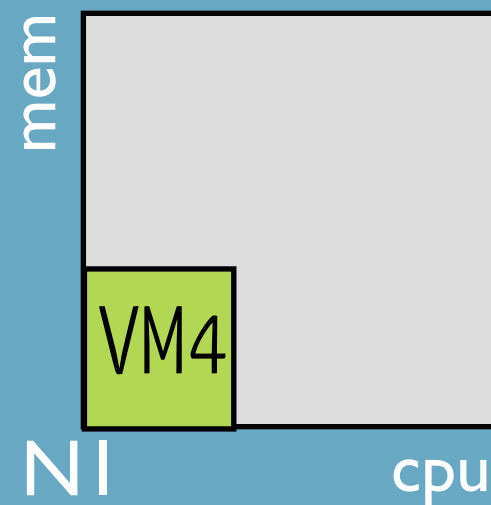


a pivot to break the cycle

fix or prevent the situation?

# dynamic schedulers

cyclic dependencies

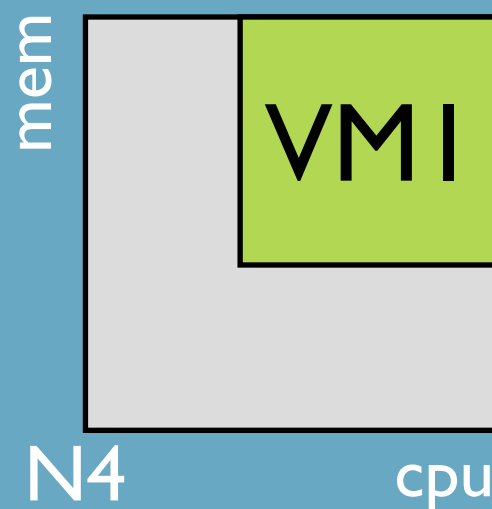
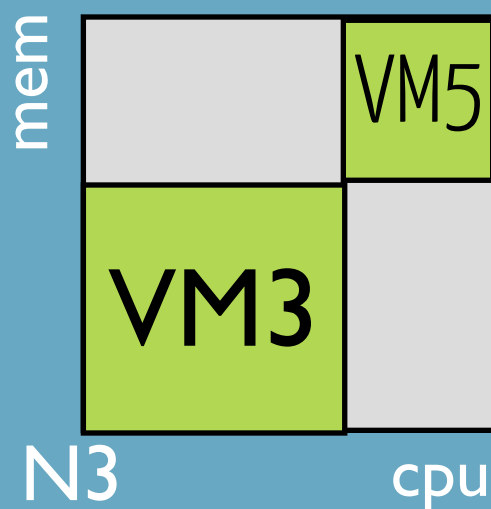
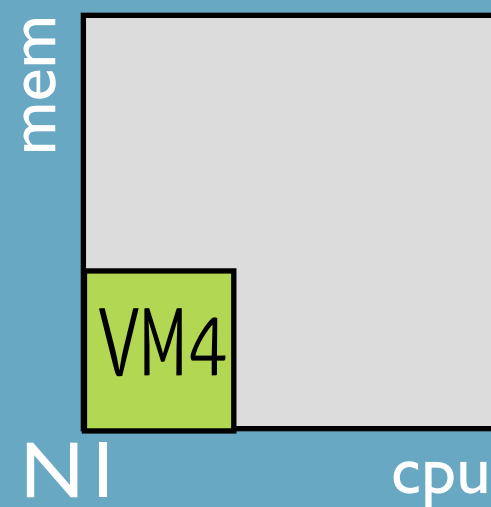


a pivot to break the cycle

fix or prevent the situation?

# dynamic schedulers

cyclic dependencies

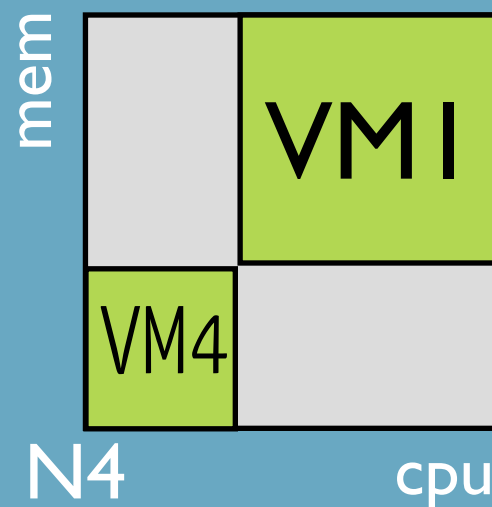
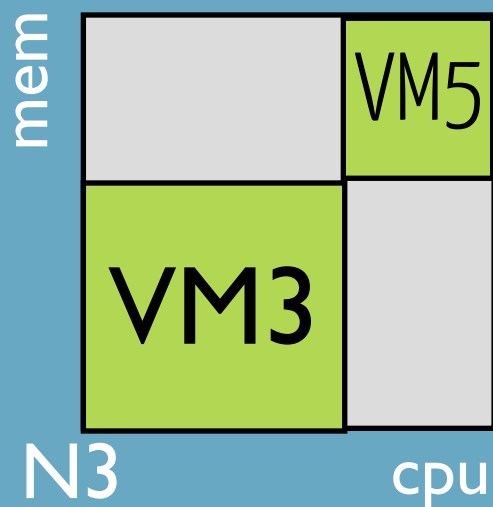
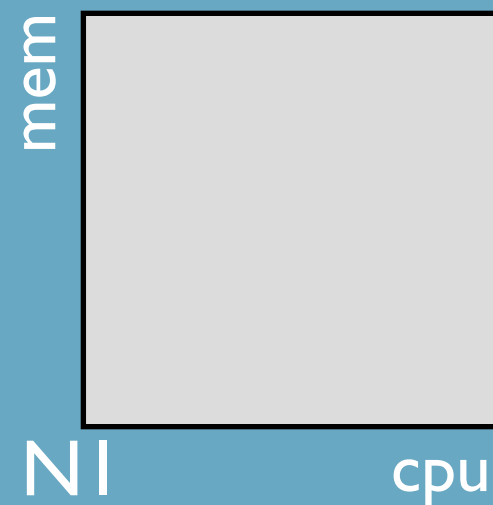


a pivot to break the cycle

fix or prevent the situation?

# dynamic schedulers

cyclic dependencies

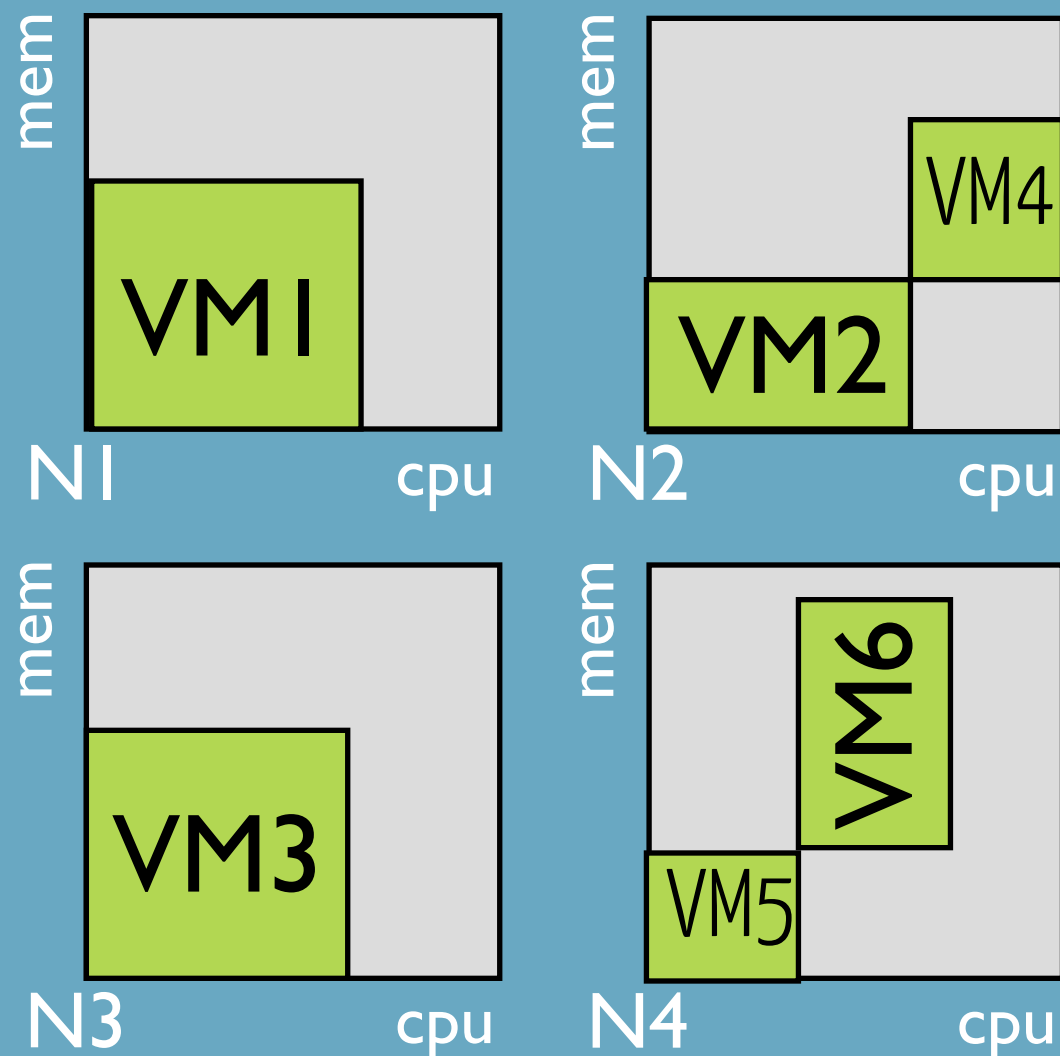


a pivot to break the cycle

fix or prevent the situation?

# dynamic schedulers

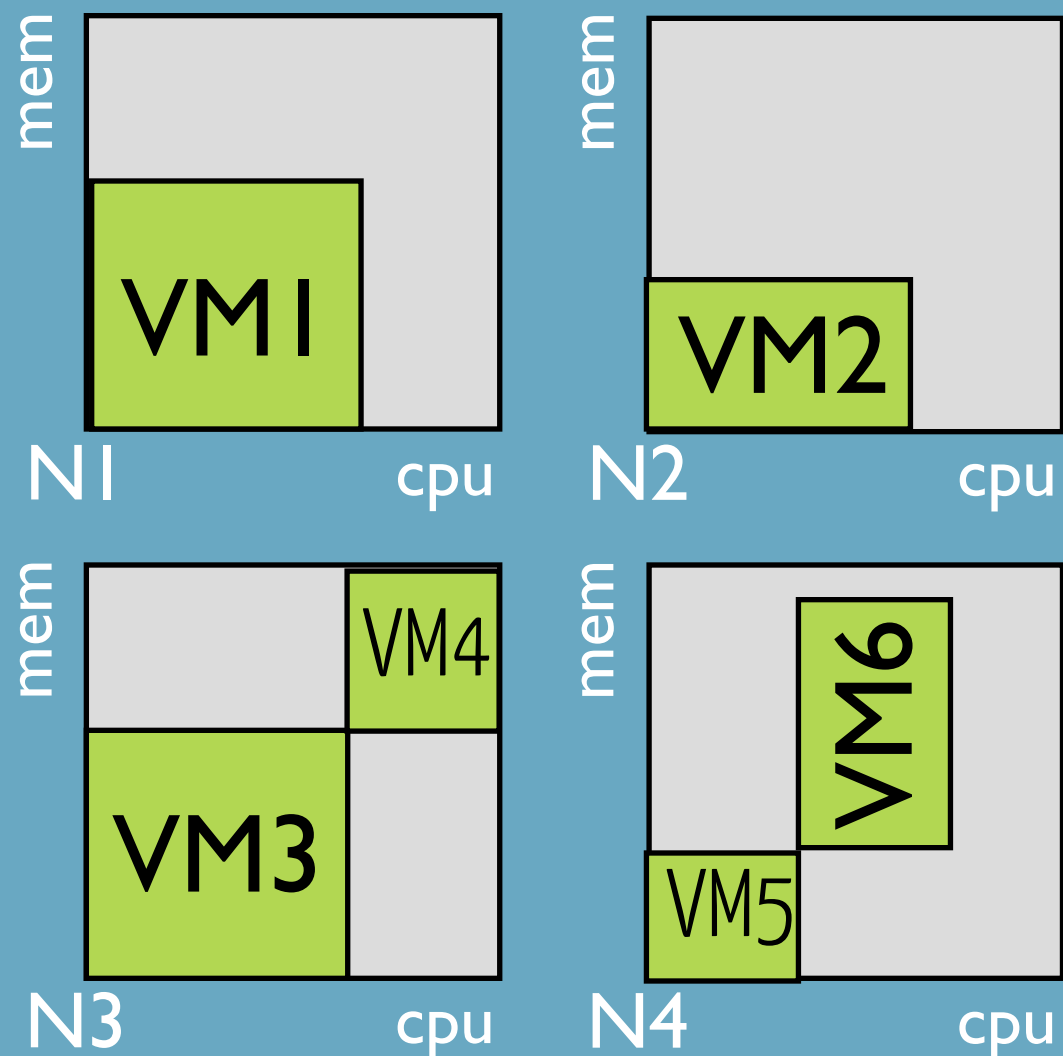
quality at a price



$$\min(\#onlineNodes) = 3$$

# dynamic schedulers

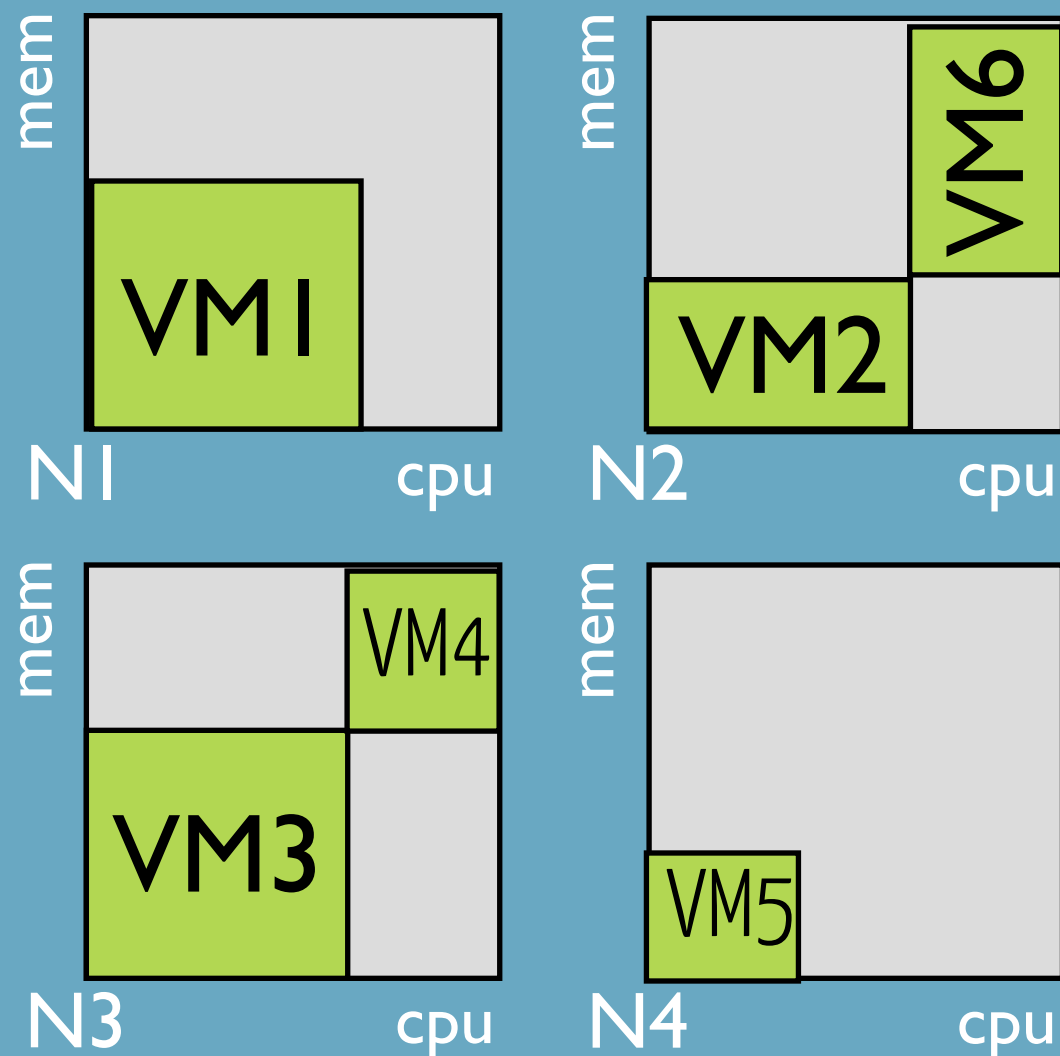
quality at a price



$$\min(\#onlineNodes) = 3$$

# dynamic schedulers

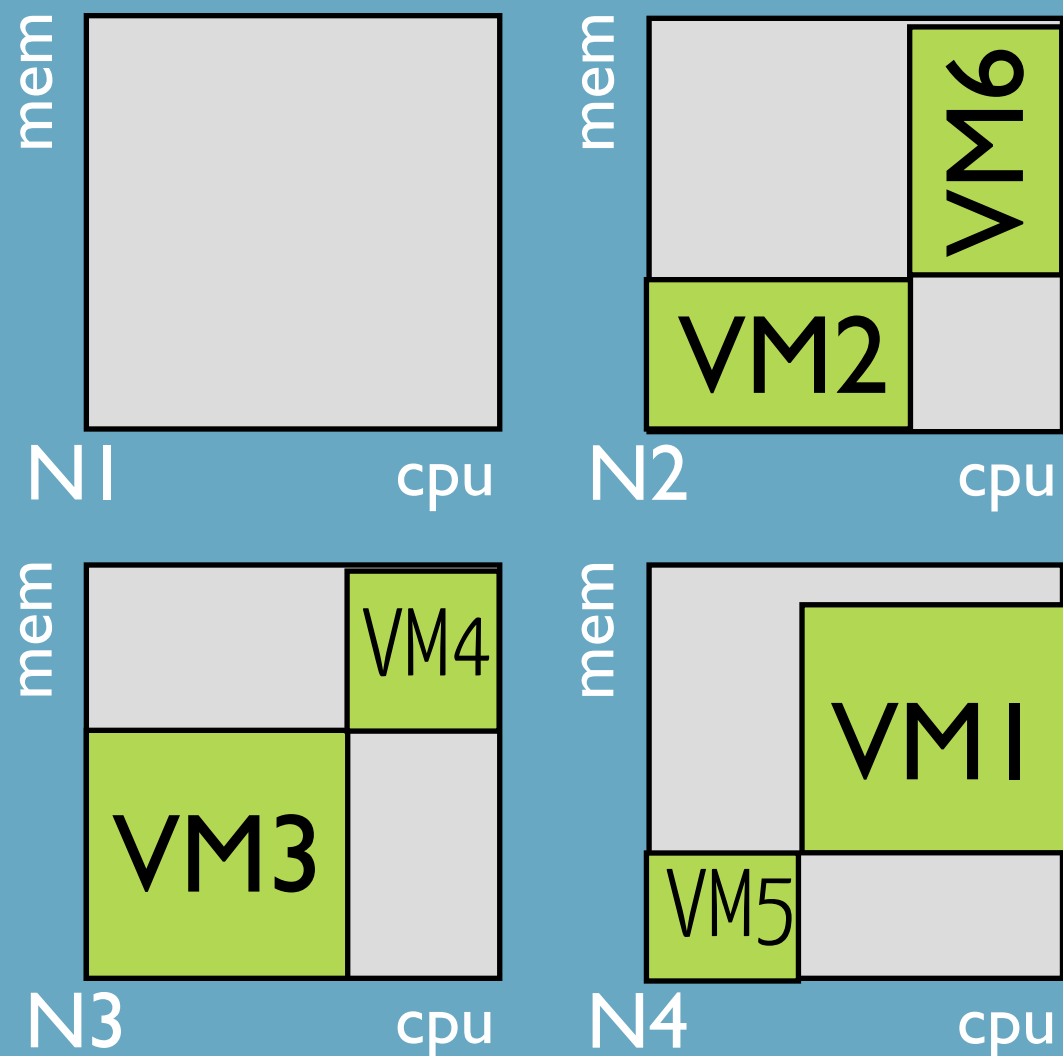
quality at a price



$\min(\#onlineNodes) = 3$

# dynamic schedulers

quality at a price

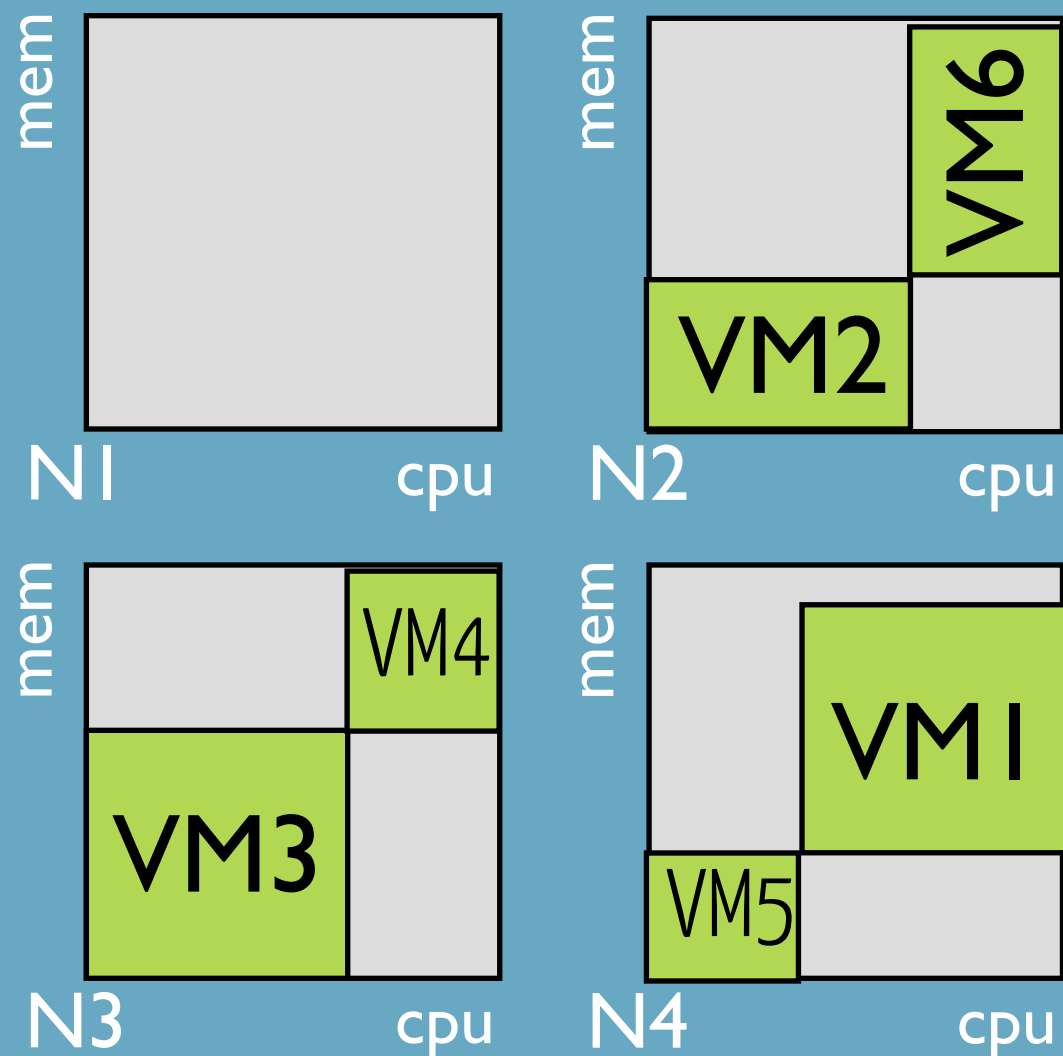


$\min(\#onlineNodes) = 3$



# dynamic schedulers

quality at a price

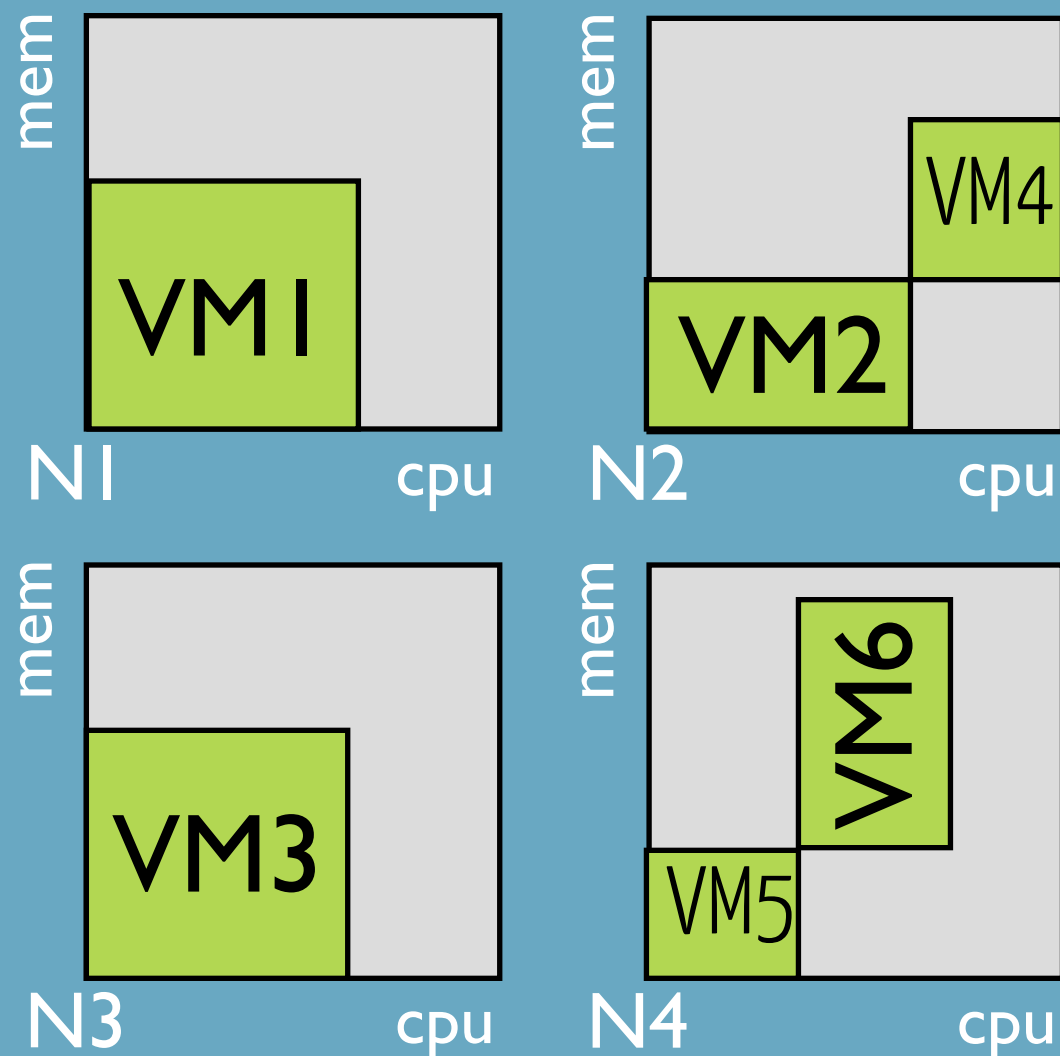


sol #1: 1m,1m,2m

$\min(\#onlineNodes) = 3$

# dynamic schedulers

quality at a price

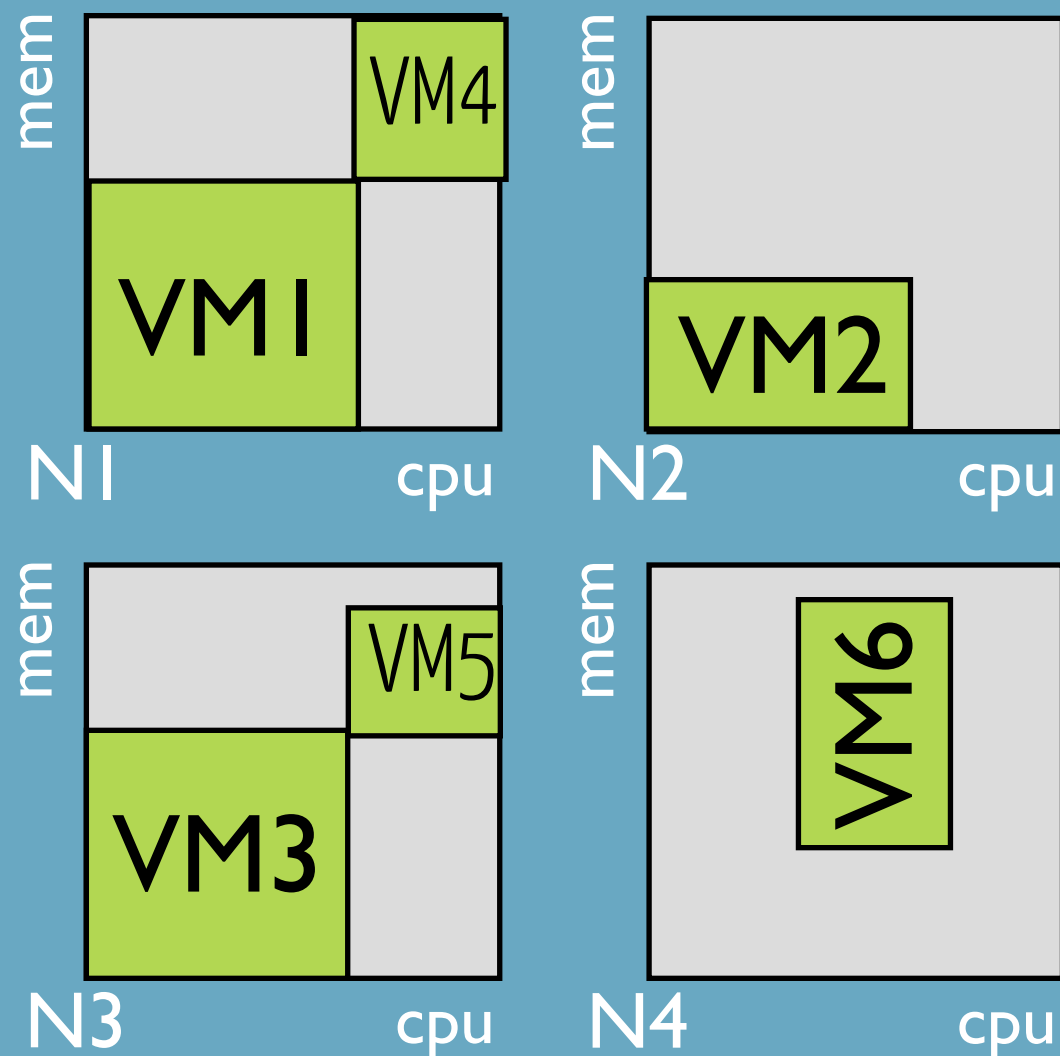


sol #1: 1m,1m,2m

$\min(\#onlineNodes) = 3$

# dynamic schedulers

quality at a price

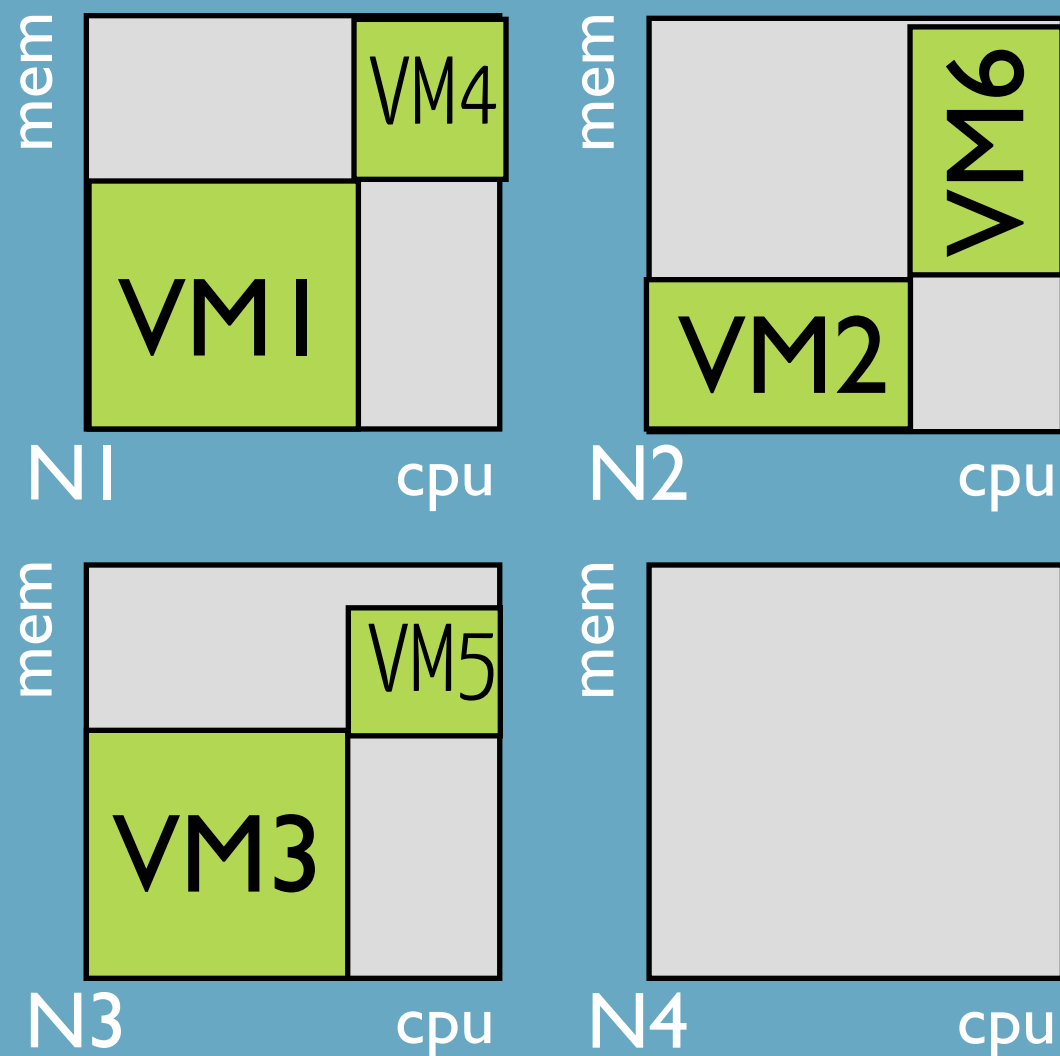


sol #1: 1m,1m,2m

$\min(\#onlineNodes) = 3$

# dynamic schedulers

quality at a price

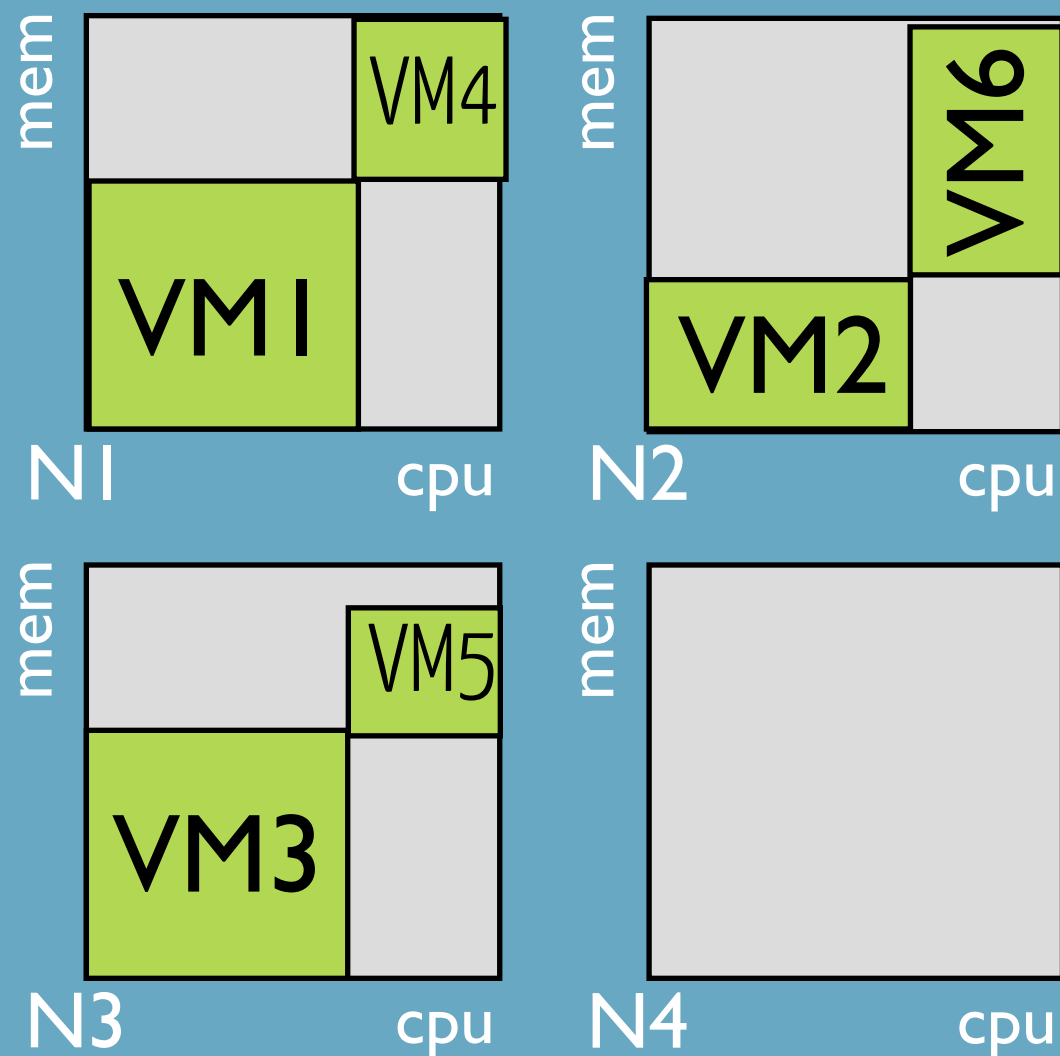


sol #1: 1m,1m,2m

$\min(\#onlineNodes) = 3$

# dynamic schedulers

quality at a price



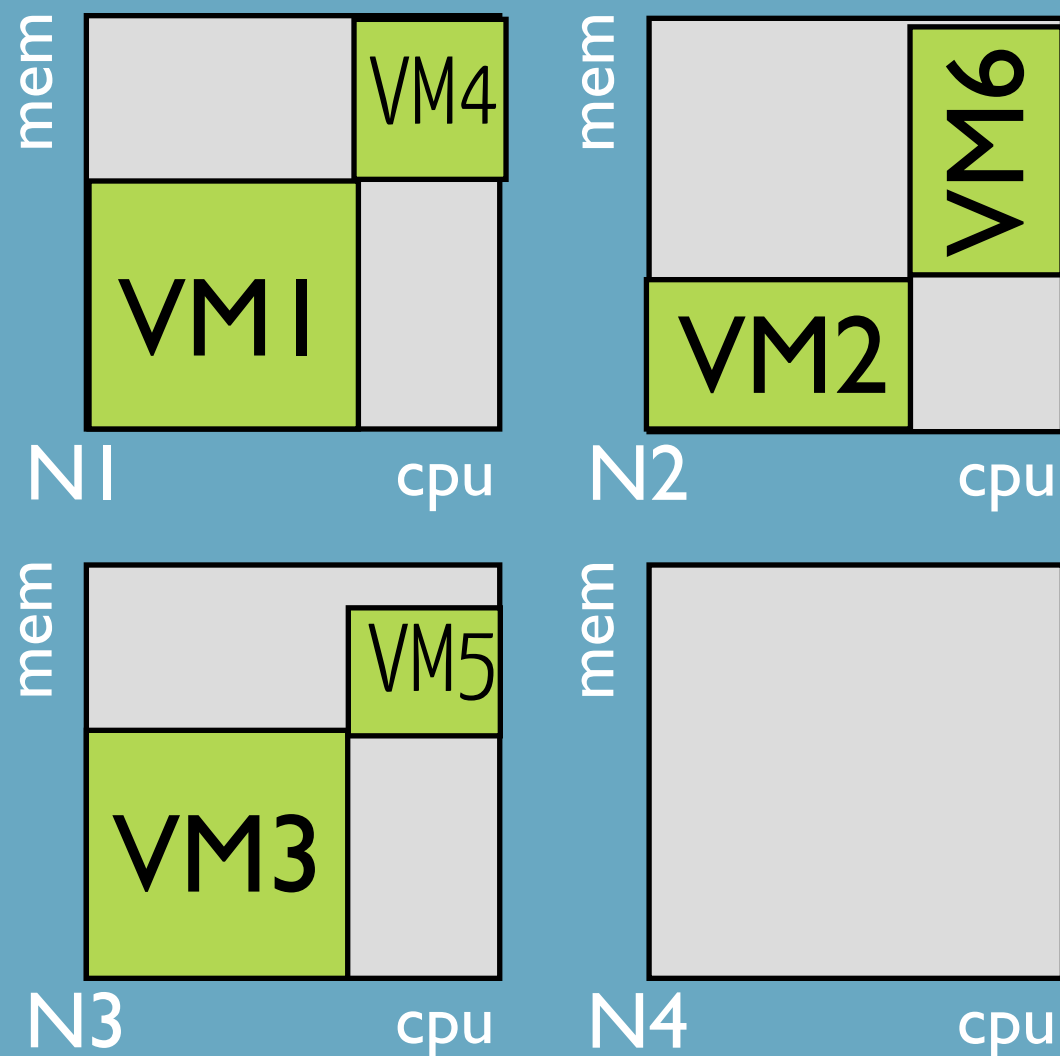
sol #1: 1m, 1m, 2m

sol #2: 1m, 1m  
1m

$\min(\#onlineNodes) = 3$

# dynamic schedulers

quality at a price



sol #1: 1m, 1m, 2m

sol #2: 1m, 1m  
1m

lower MTTR  
(faster)

$\min(\#onlineNodes) = 3$

# dynamic schedulers

quality at a price

the objective should reflect  
reconfiguration costs

*$\min(MTTR), \min(\#migrations), \dots$*

# dynamic schedulers

## pros

continuous optimisation  
through reconfiguration



# dynamic schedulers

## cons

- harder to model
- harder to scale
- technically expensive
- costly reconfiguration

# dynamic scheduling for the win?

in theory yes but that's theory

benefits depends on

- the workload
- the objective/ SLAs
- the infrastructure

dynamic scheduling is rare  
in public or large clouds

scheduling models

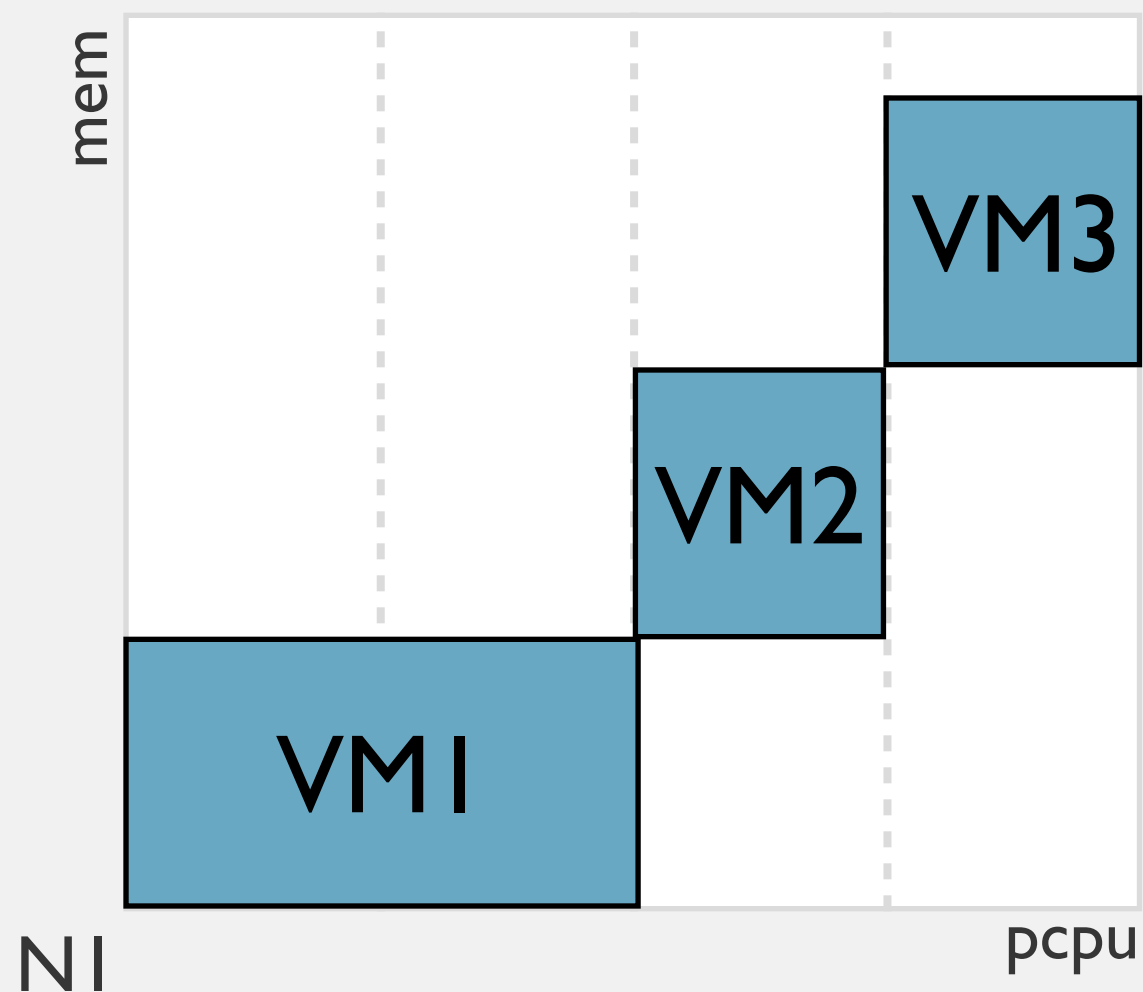
static  
dynamic

resource allocation

static  
dynamic

schedulers

# static resource allocation



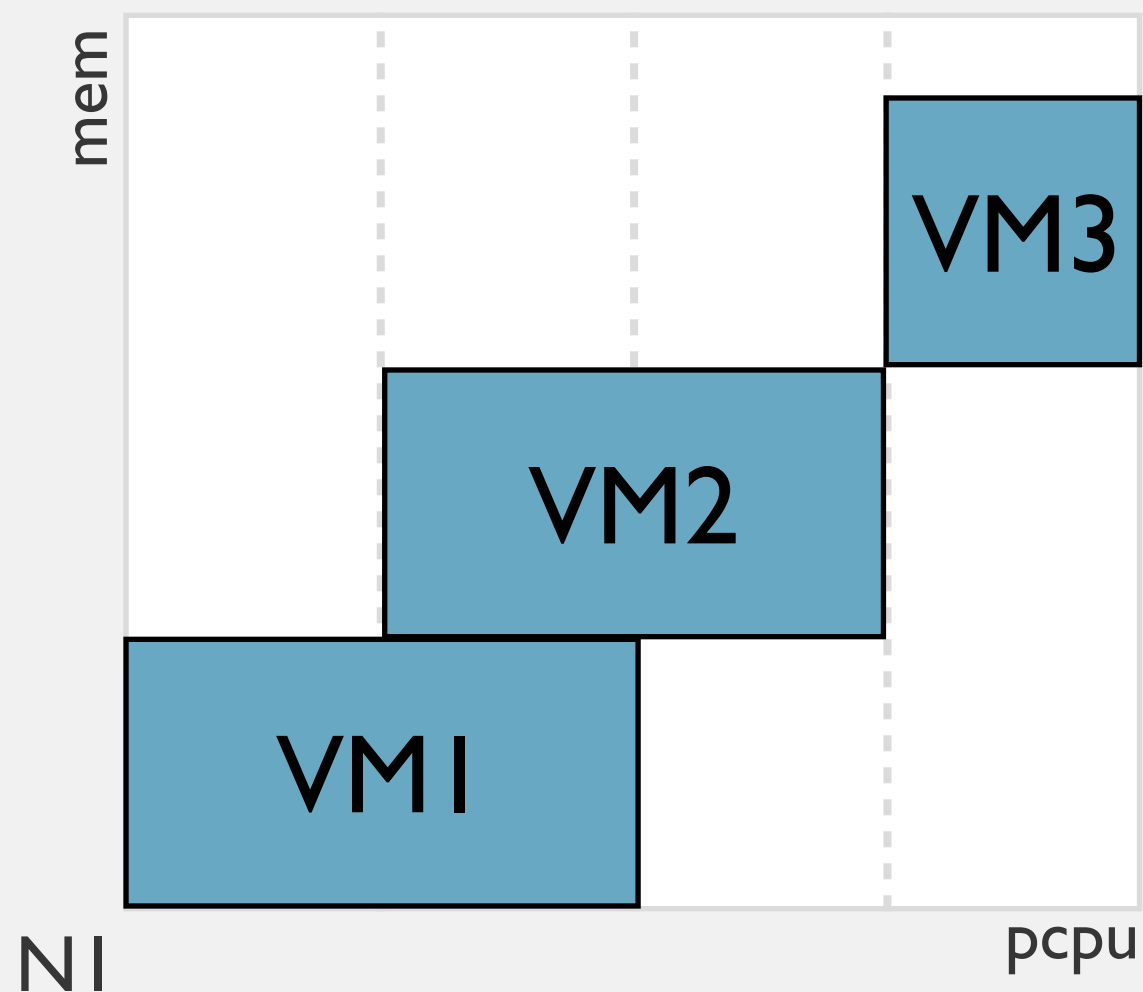
cpu, ram, i/o, bandwidth  
allocated once for all

allocation != utilisation

**no sharing**  
**conservative allocation**

# static resource allocation

sharing (overbooking)

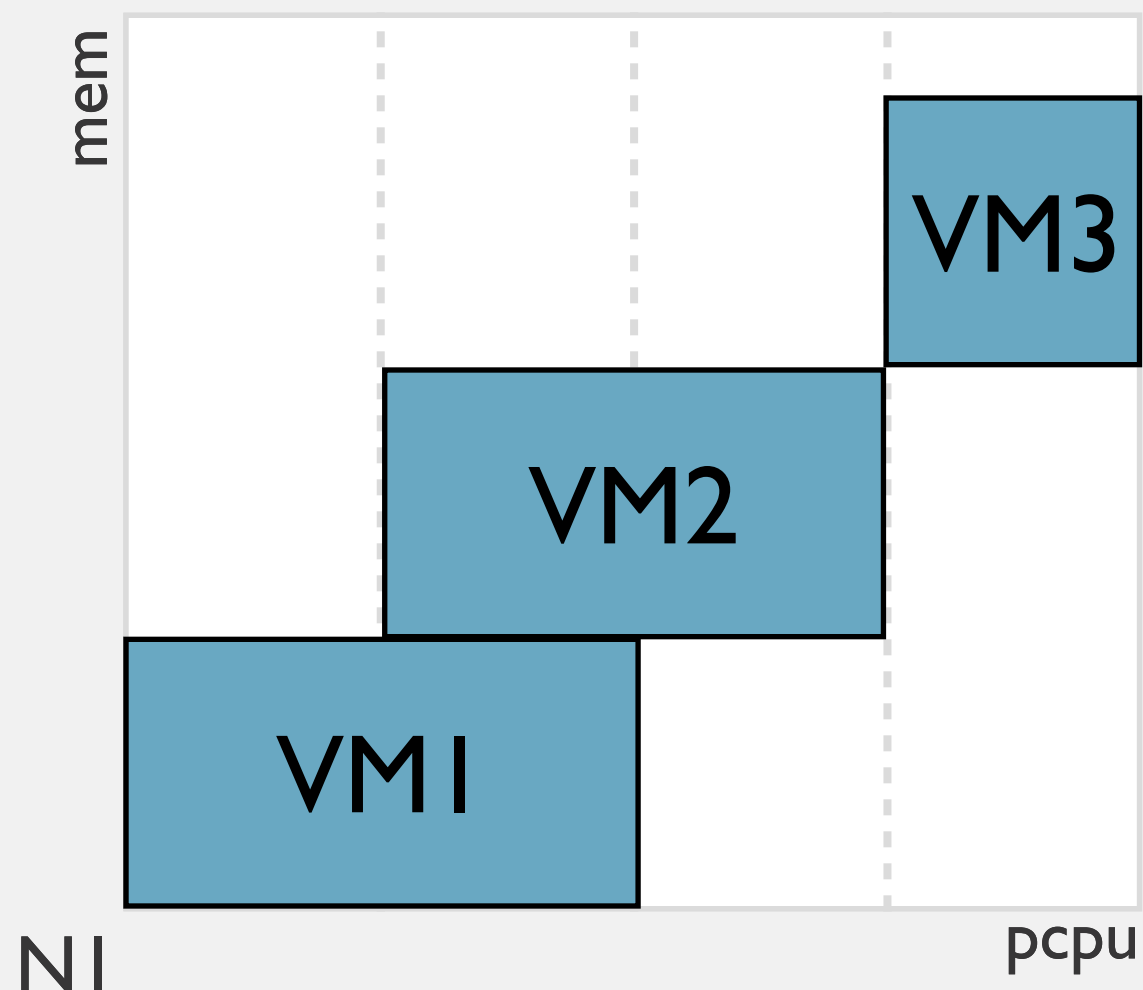


performance loss  
with concurrent accesses

acceptable if stated  
in the SLA

# dynamic resource allocation

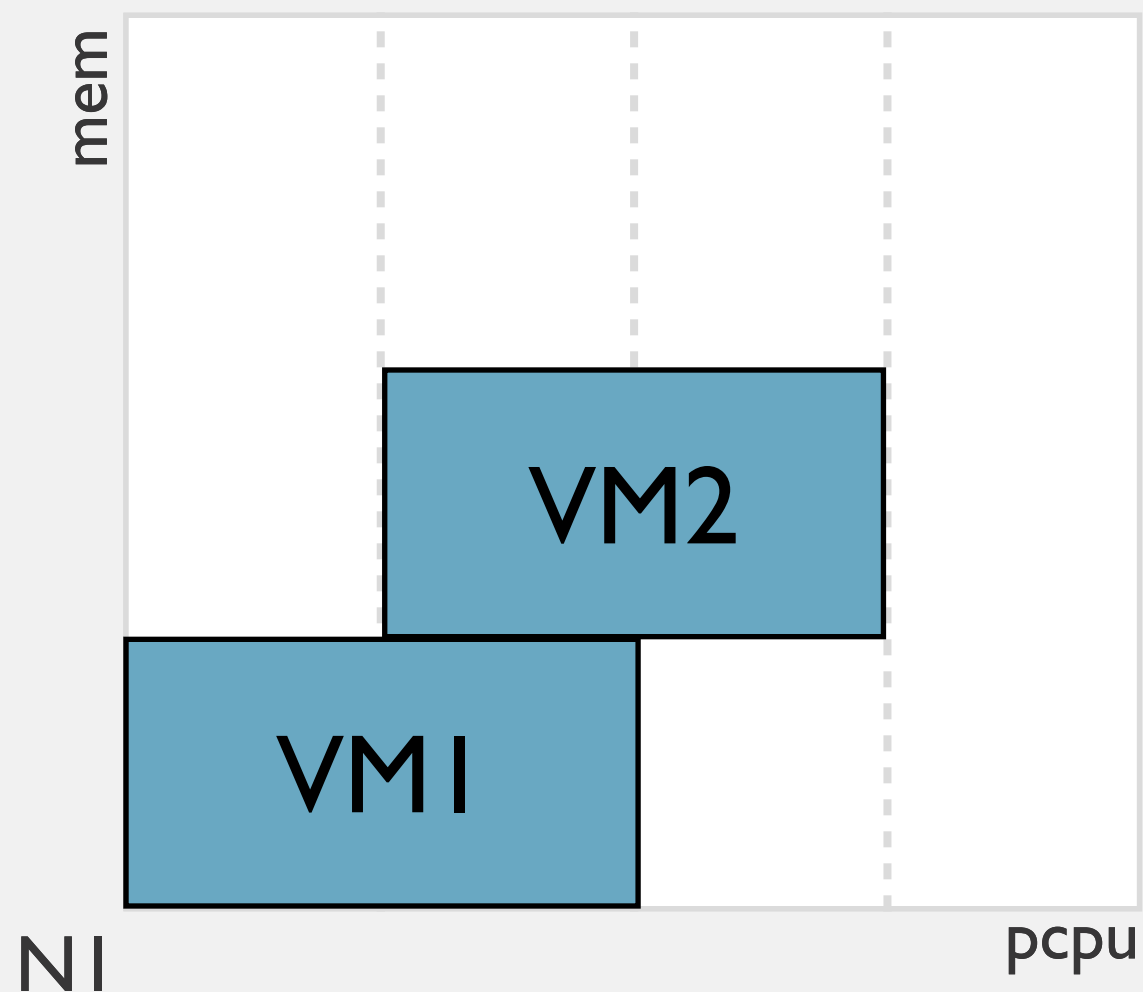
to fix violations



cpu, ram, i/o, bandwidth  
allocated can be revised

# dynamic resource allocation

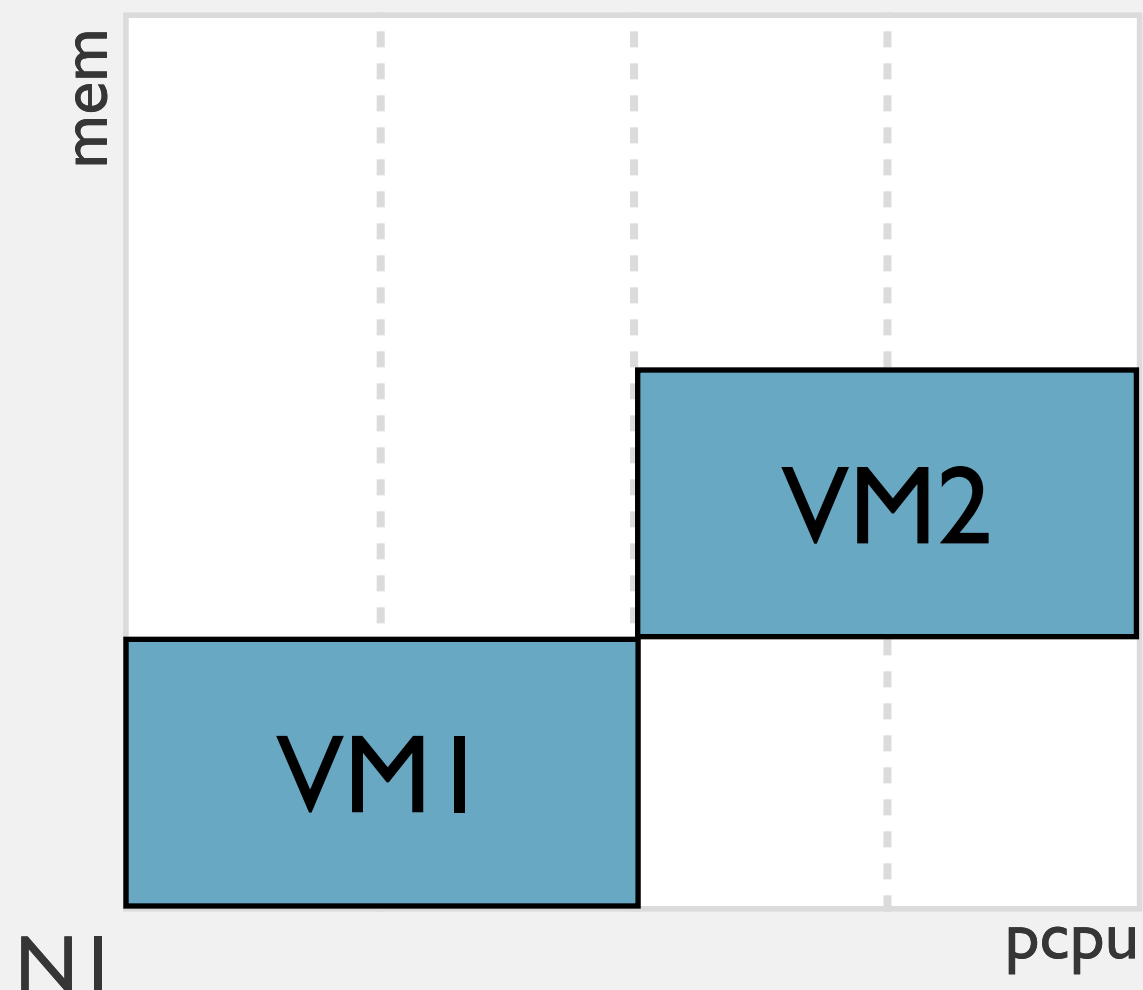
to fix violations



cpu, ram, i/o, bandwidth  
allocated can be revised

# dynamic resource allocation

to fix violations

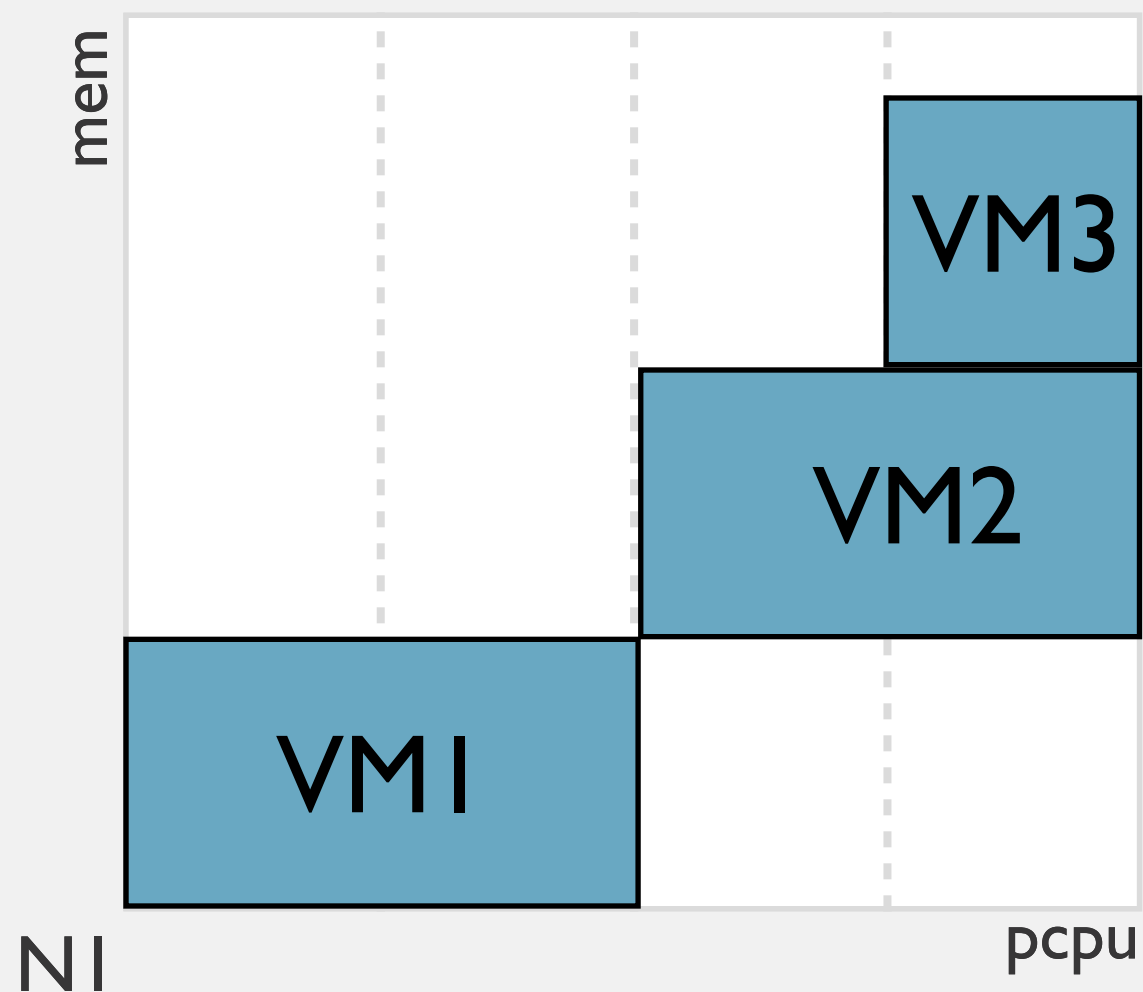


cpu, ram, i/o, bandwidth  
allocated can be revised



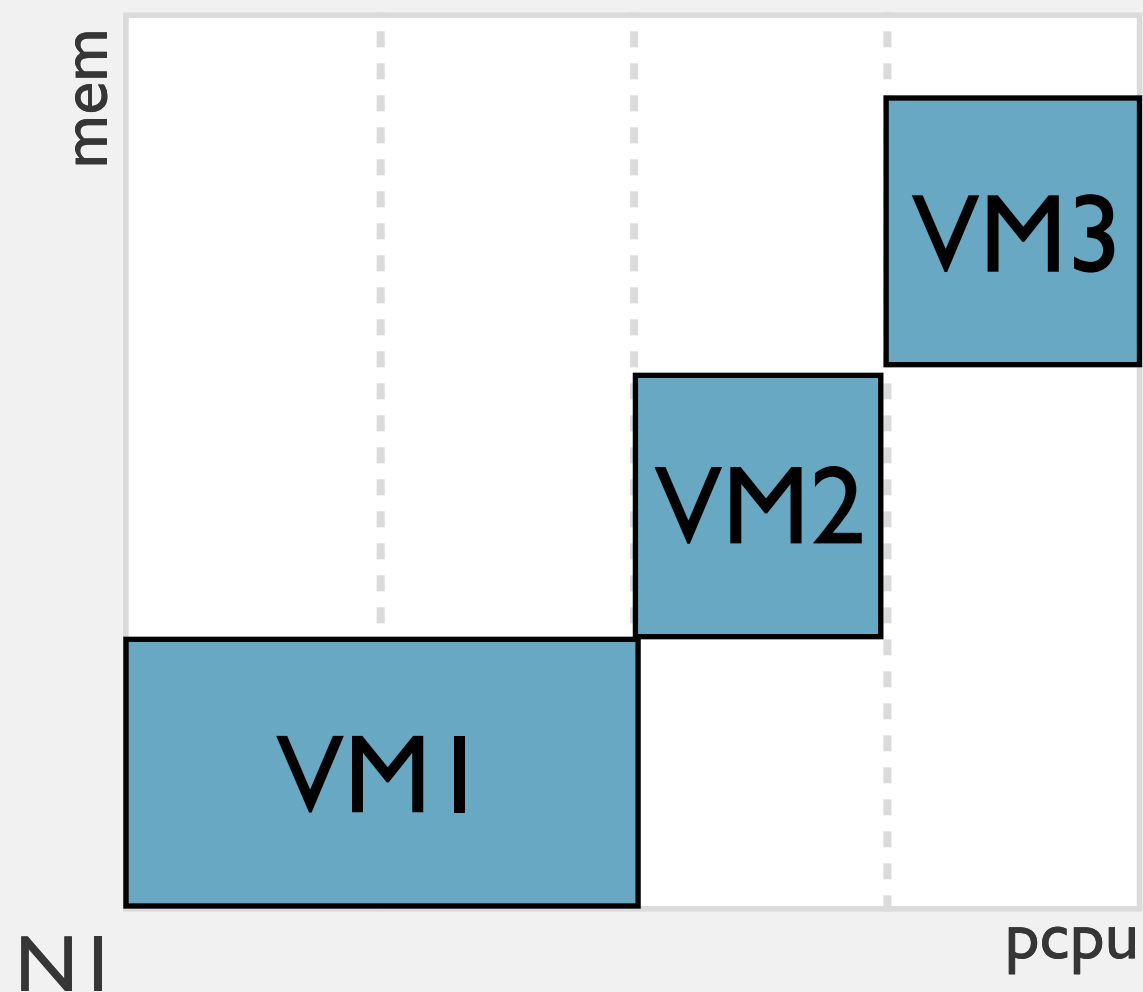
# dynamic resource allocation

to support vertical elasticity



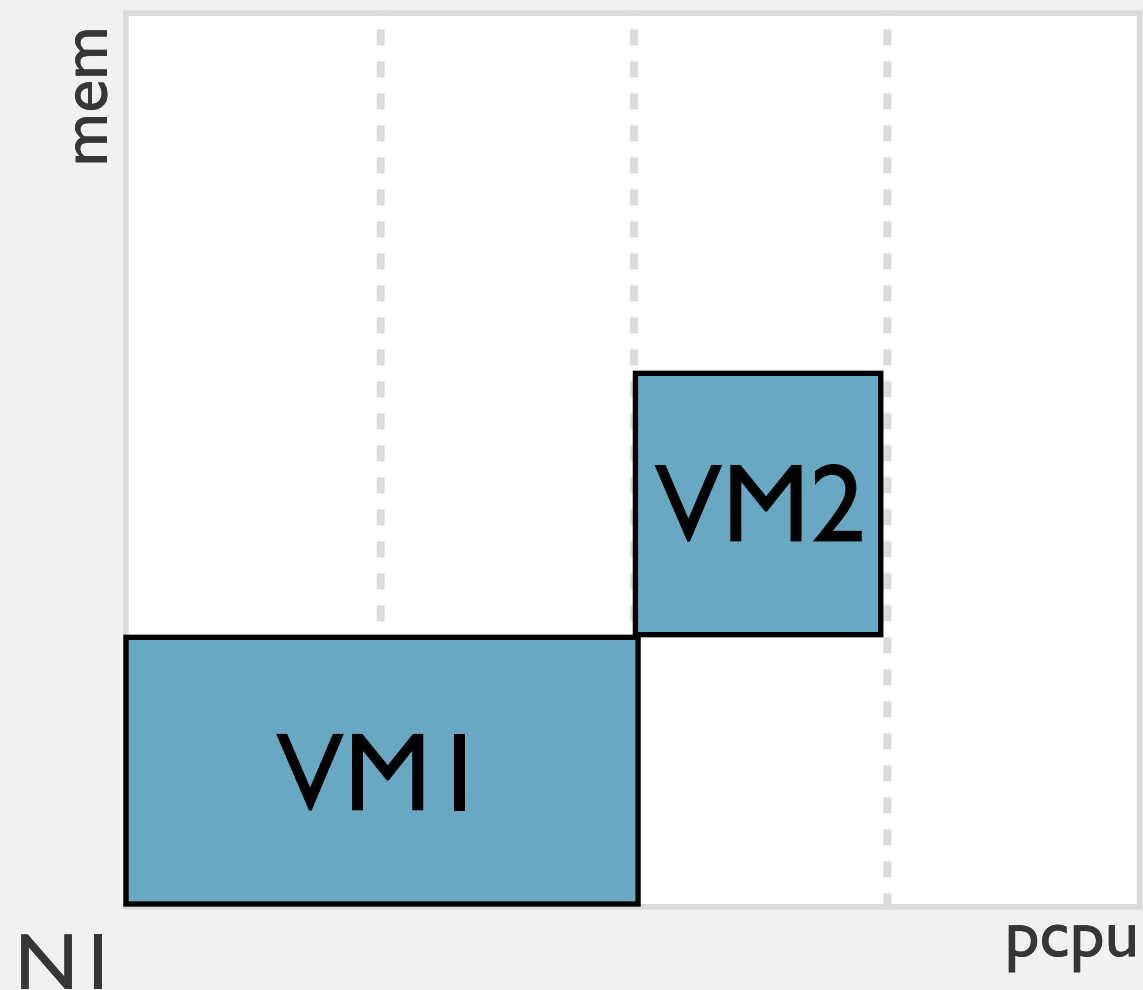
# dynamic resource allocation

to support vertical elasticity



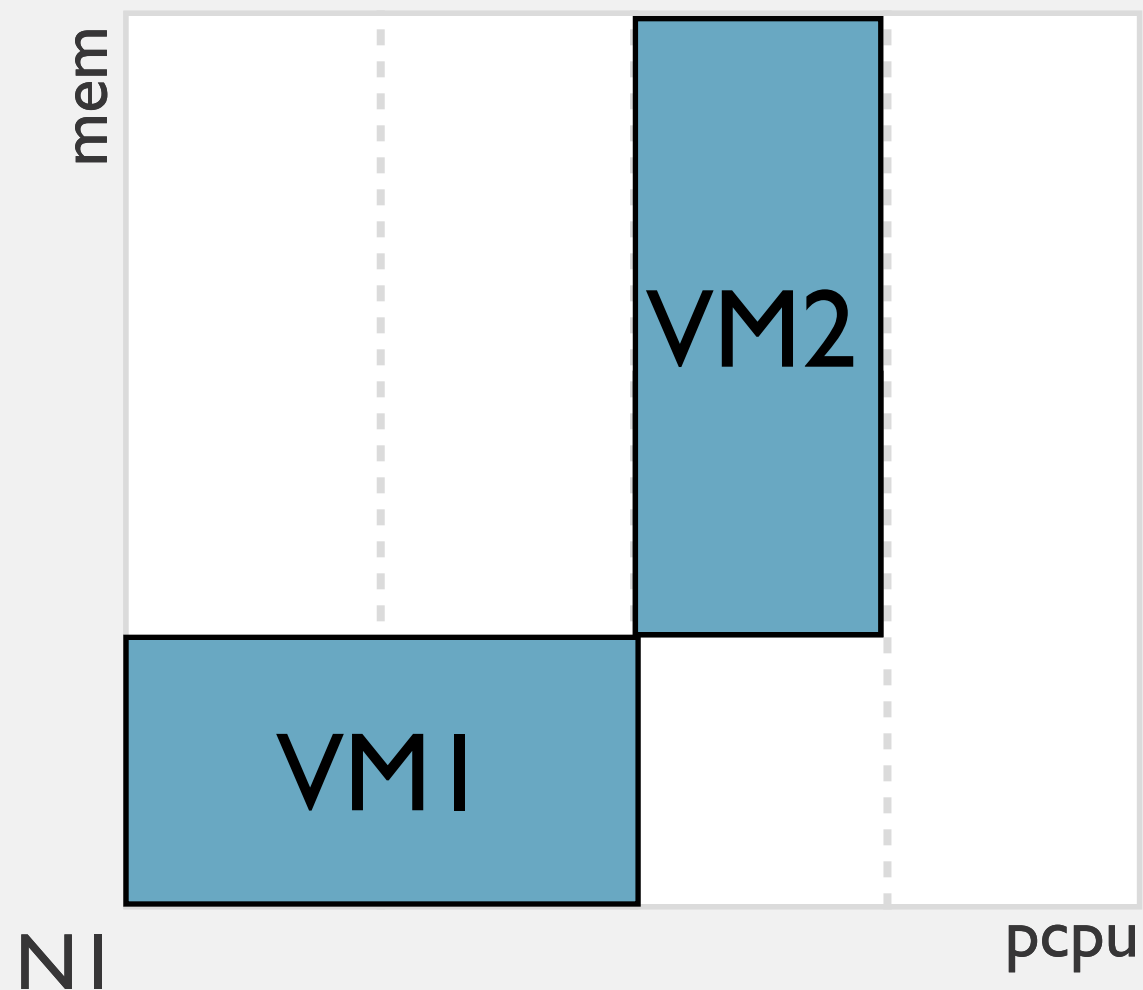
# dynamic resource allocation

to support vertical elasticity



# dynamic resource allocation

to support vertical elasticity



# dynamic resource allocation for the win ?

common on CPU + overbooking  
(inc. hosting capacity)

exceptional for memory  
(huge performance loss)

benefits depends on the workload  
the objective/ SLAs  
the infrastructure

# RECAP

The VM scheduler  
makes cloud  
benefits *real*

no holy grail



think about  
what is costly

static  
scheduling for a  
peaceful life

dynamic  
scheduling to  
cease the day

with great power  
comes great  
responsibility



# 3 Coding a VM scheduler



# VM scheduling is hard

static or dynamic scheduler ?  
allocation ?

does the workload/SLA/objective ?  
requires migration ?

maximum duration to schedule?

# VM scheduling is NP-Hard

issues with large infrastructures  
or hard problems

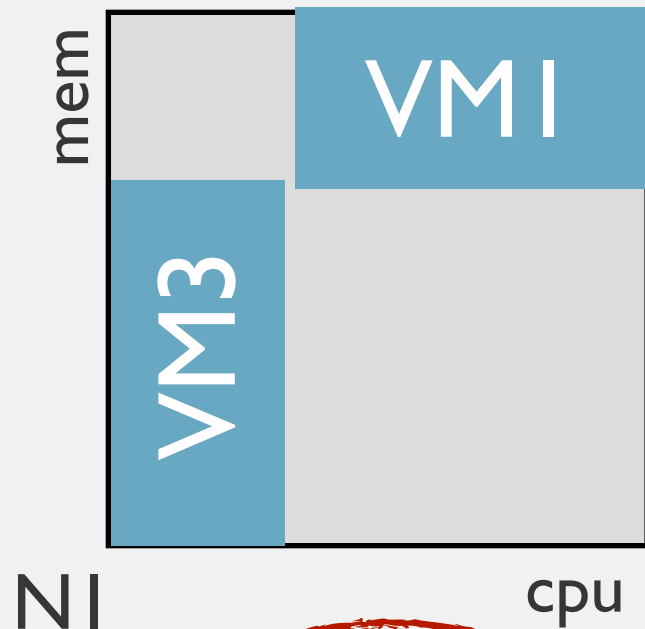
fast adhoc heuristics  
despite corner cases

some use biased complete approaches  
(linear programming, constraint programming)

*like him* ←

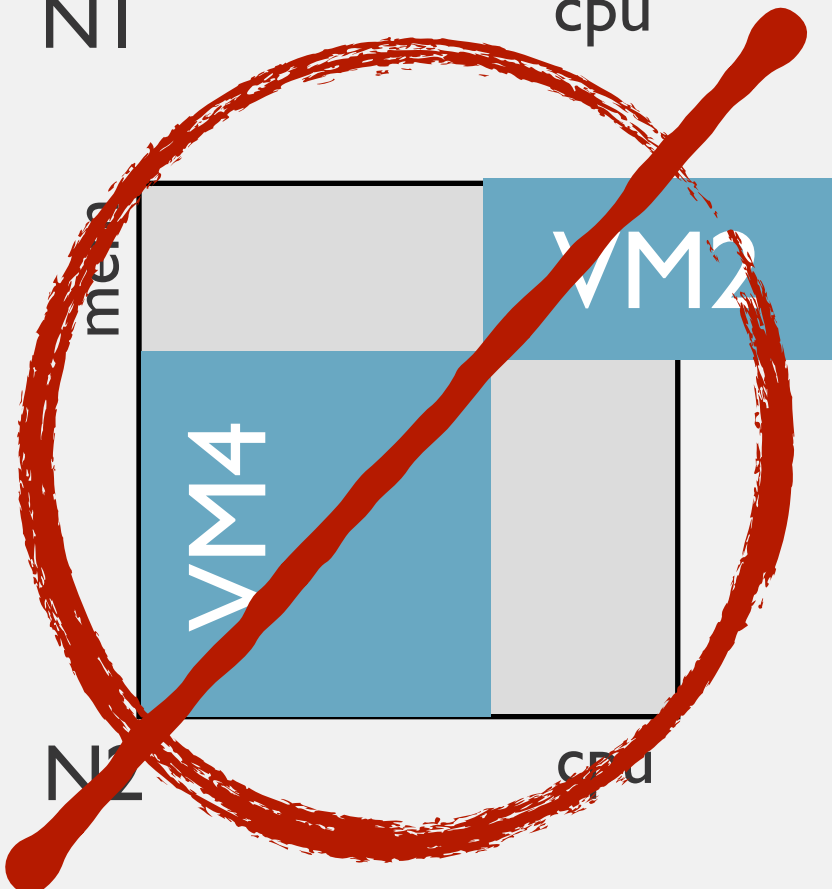


# vector packing problem



items with a finite volume to  
place inside finite bins

a generalisation of the bin  
packing problem



the basic to model the infra.  
1 dimension = 1 resource



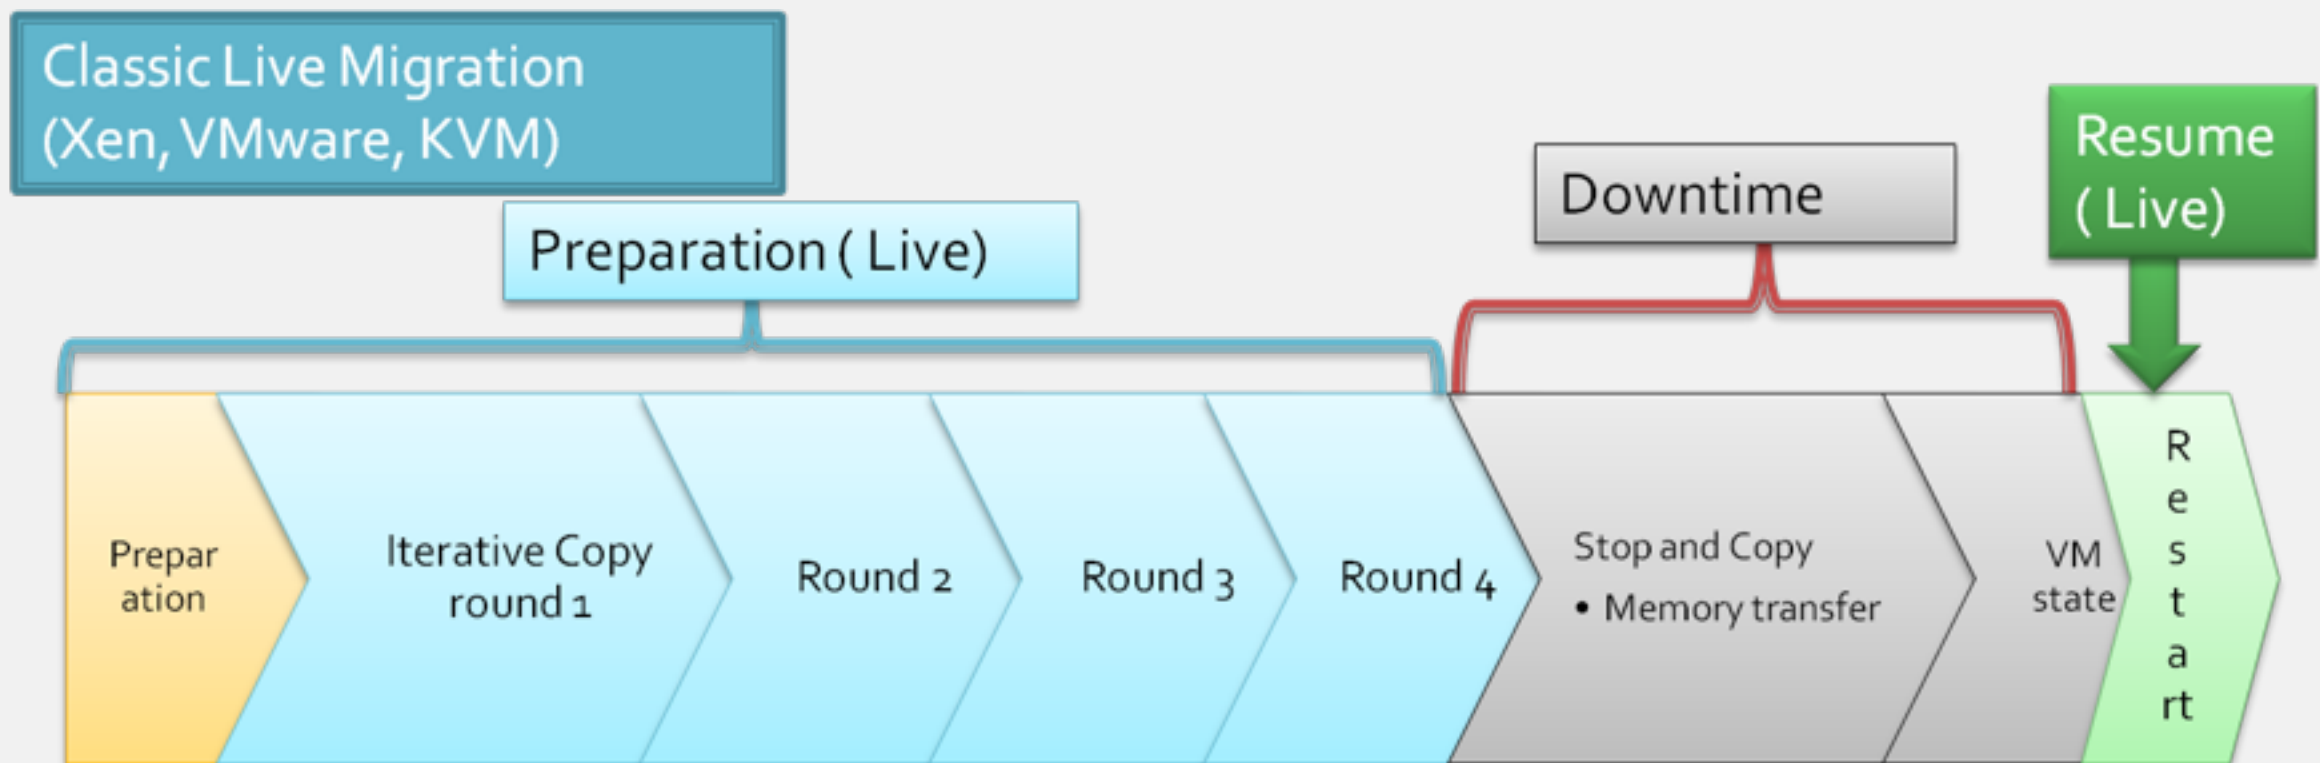
Which resource can be modeled as a  
packing dimension



Which resource can be modeled as a  
packing dimension

CPU, memory, disk IO, licences cardinality, network boundaries, ...  
~~end to end network~~

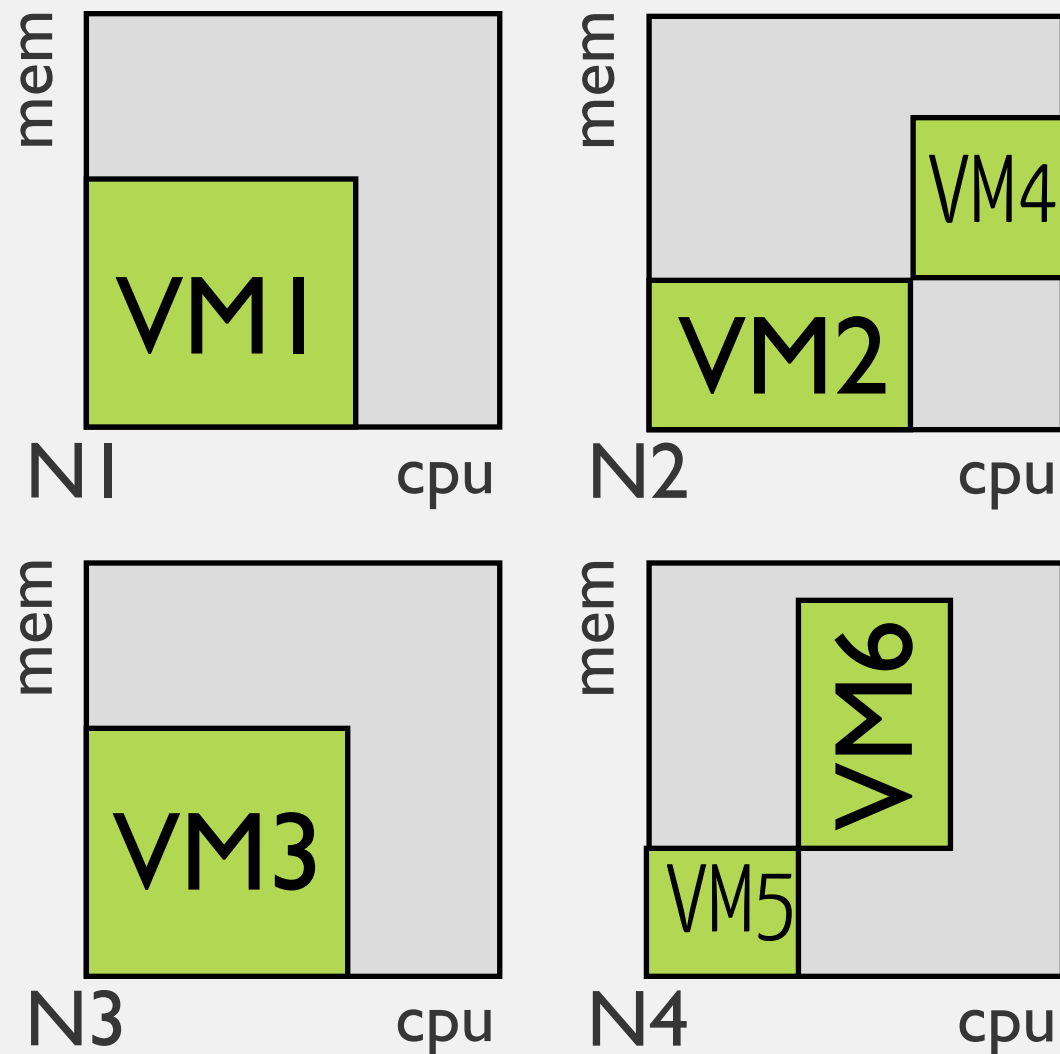
# how to support migrations



temporary,  
resources are used on the source and the destination nodes

# how to support migrations

a simple way



n-phases vector packing

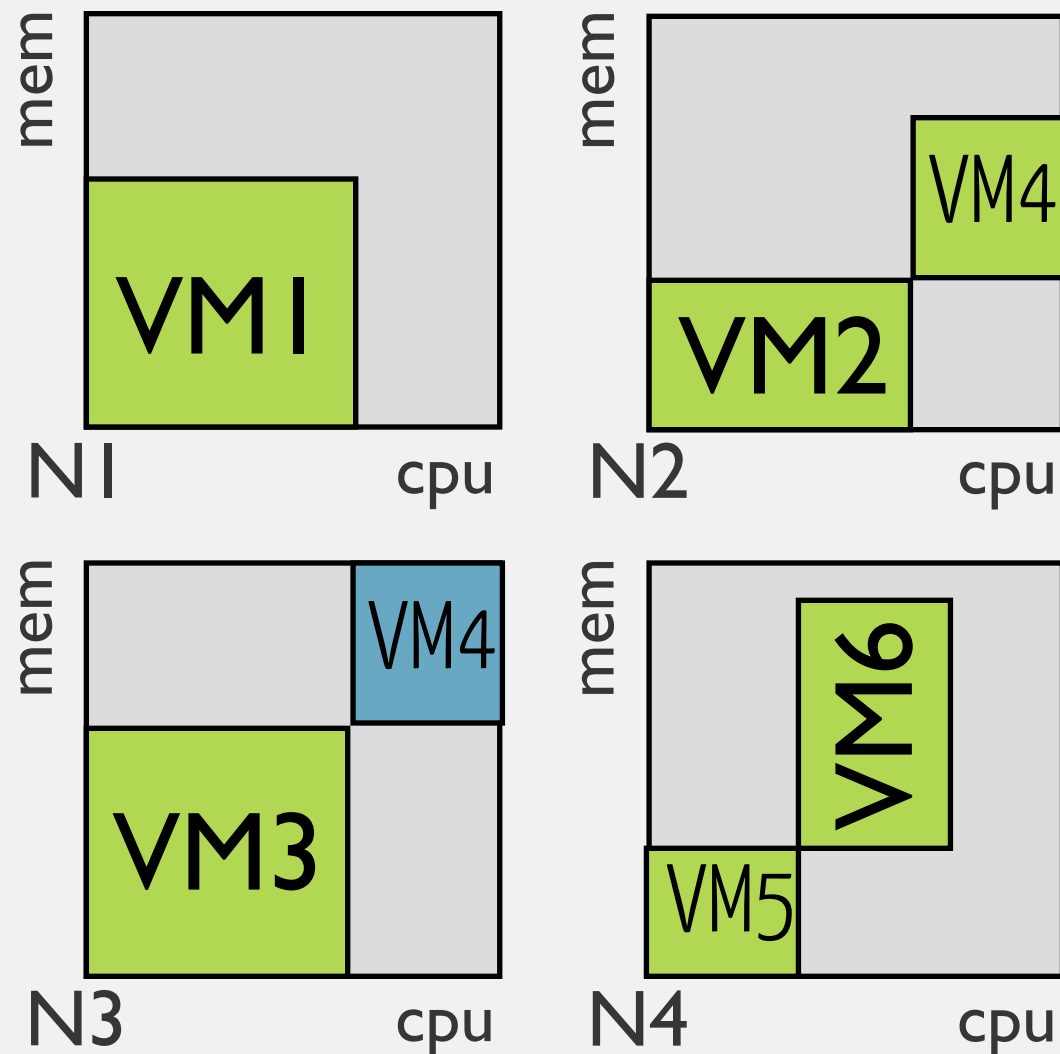
VM duplication between 2 phases  
to simulate the migration

easy to implement

hard to manipulate to  
compute long term previsions

# how to support migrations

a simple way



n-phases vector packing

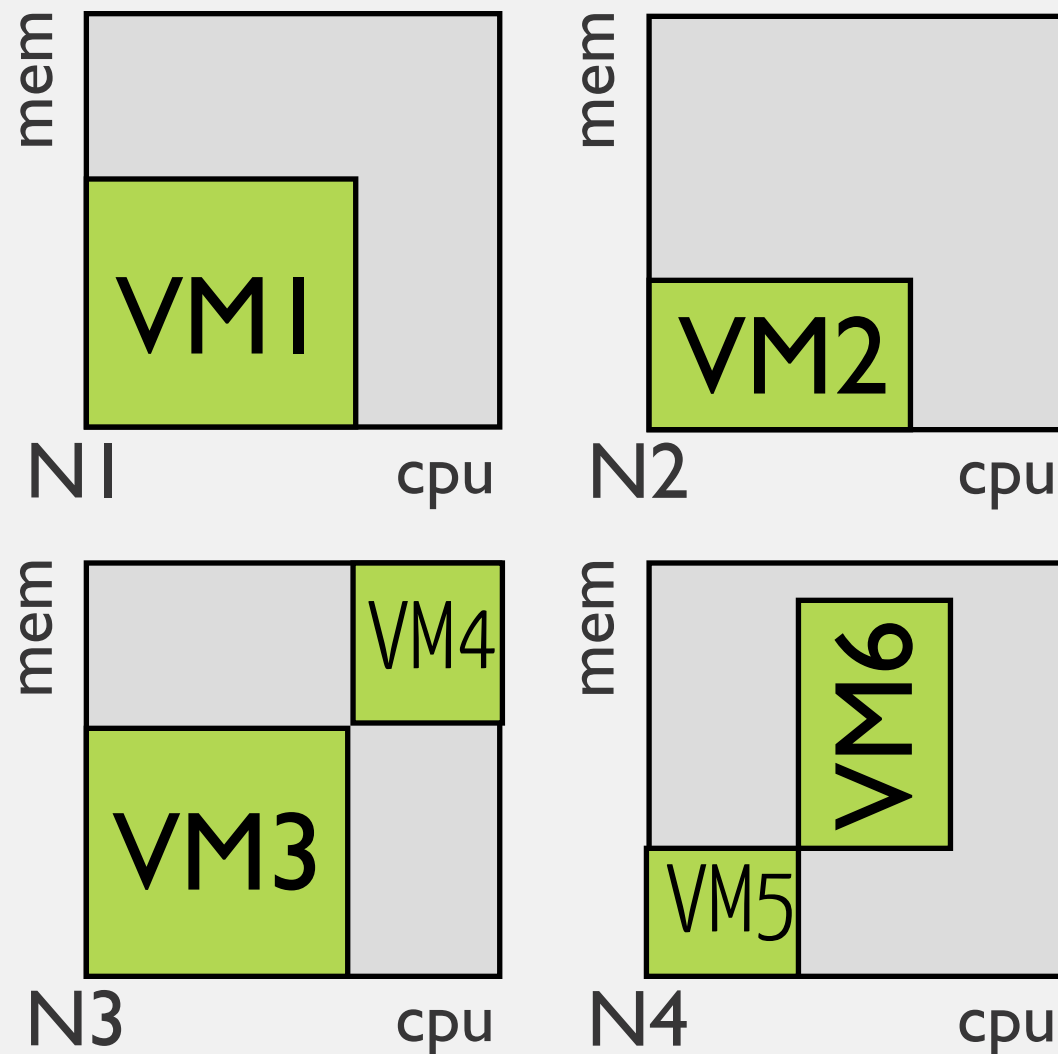
VM duplication between 2 phases  
to simulate the migration

easy to implement

hard to manipulate to  
compute long term previsions

# how to support migrations

a simple way



n-phases vector packing

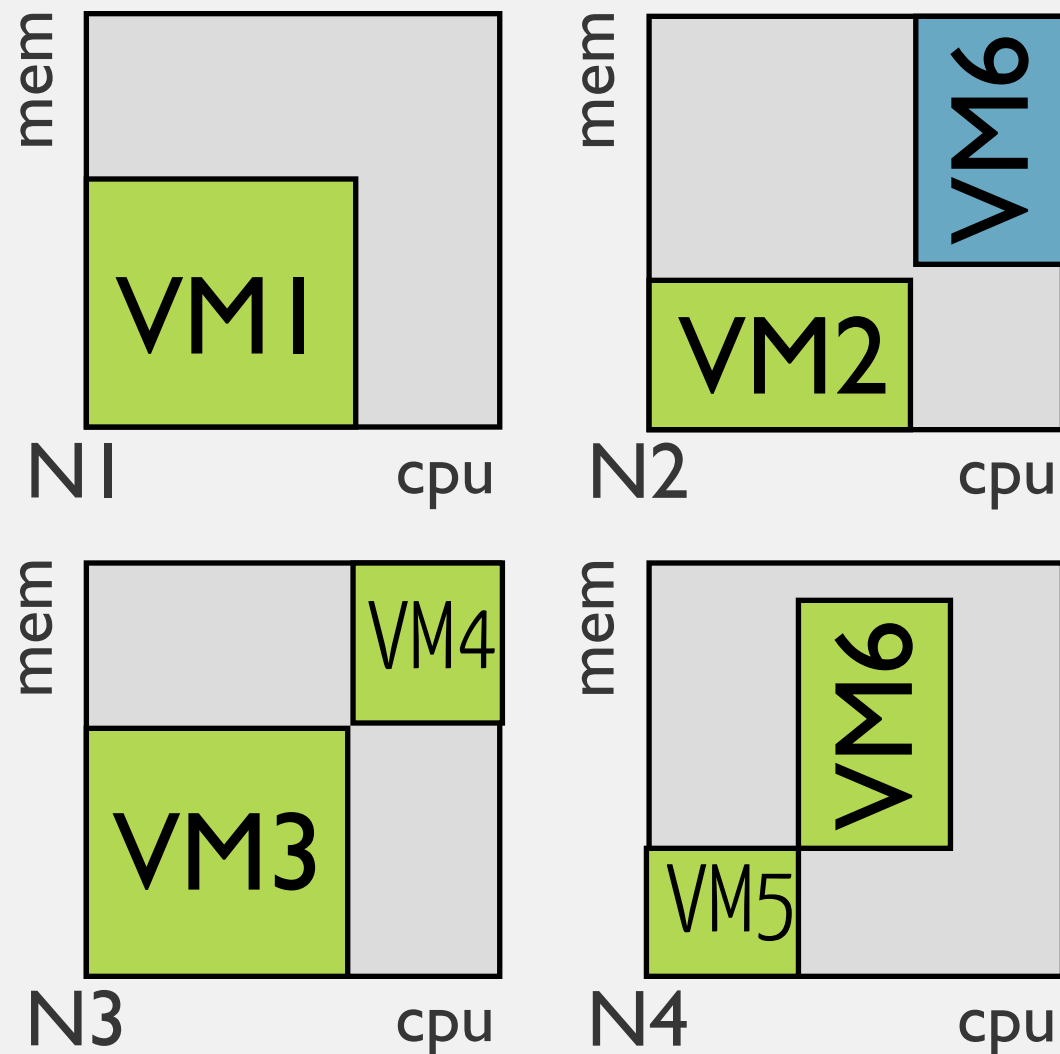
VM duplication between 2 phases  
to simulate the migration

easy to implement

hard to manipulate to  
compute long term previsions

# how to support migrations

a simple way



n-phases vector packing

VM duplication between 2 phases  
to simulate the migration

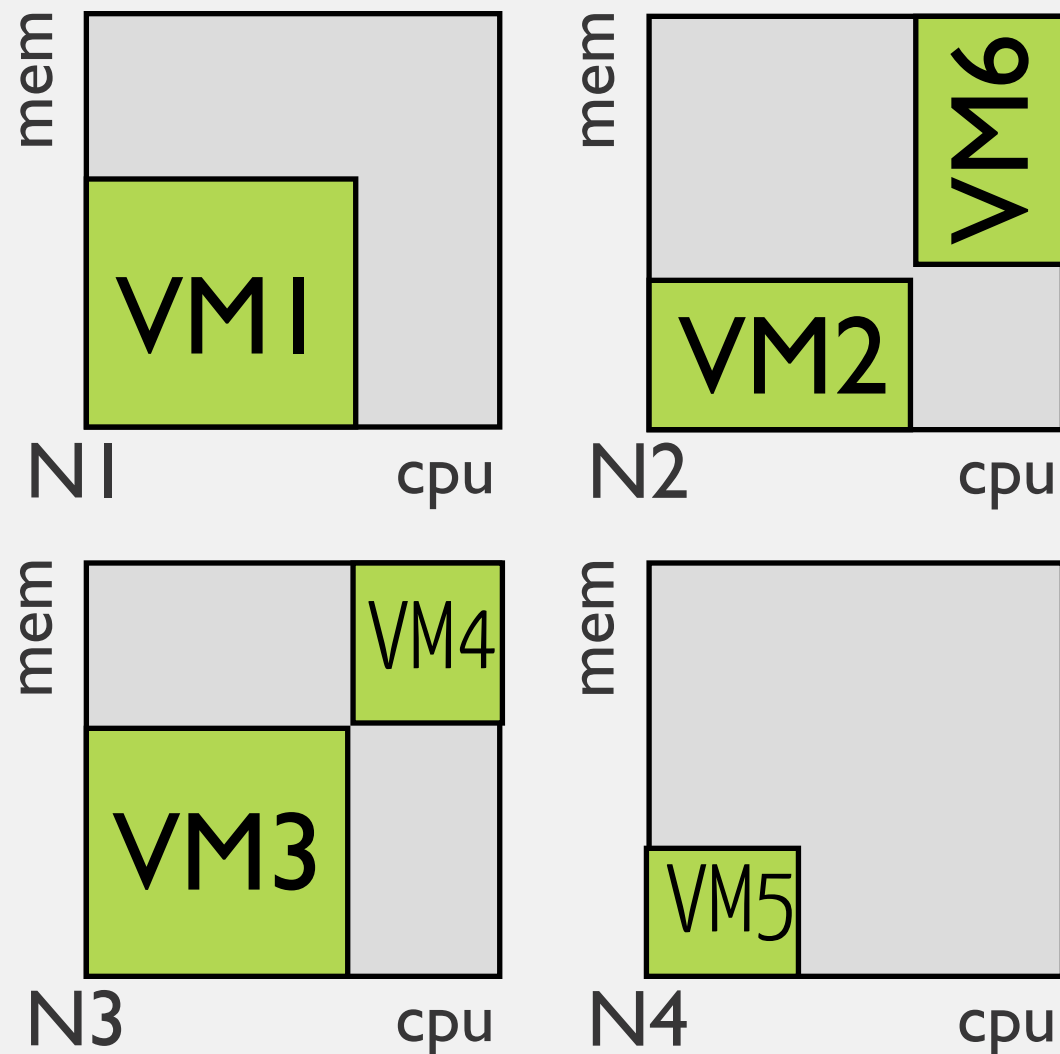
easy to implement

hard to manipulate to  
compute long term previsions



# how to support migrations

a simple way



n-phases vector packing

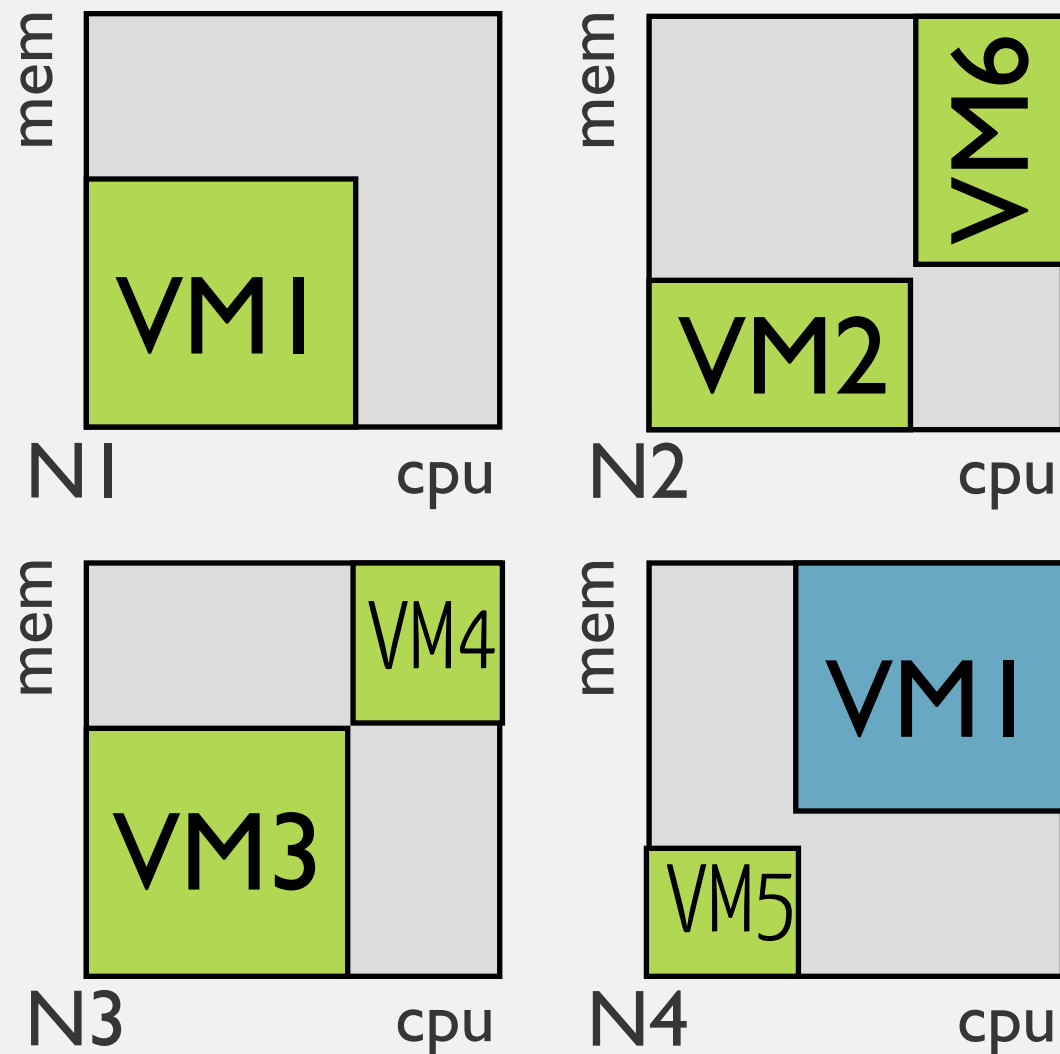
VM duplication between 2 phases  
to simulate the migration

easy to implement

hard to manipulate to  
compute long term previsions

# how to support migrations

a simple way



n-phases vector packing

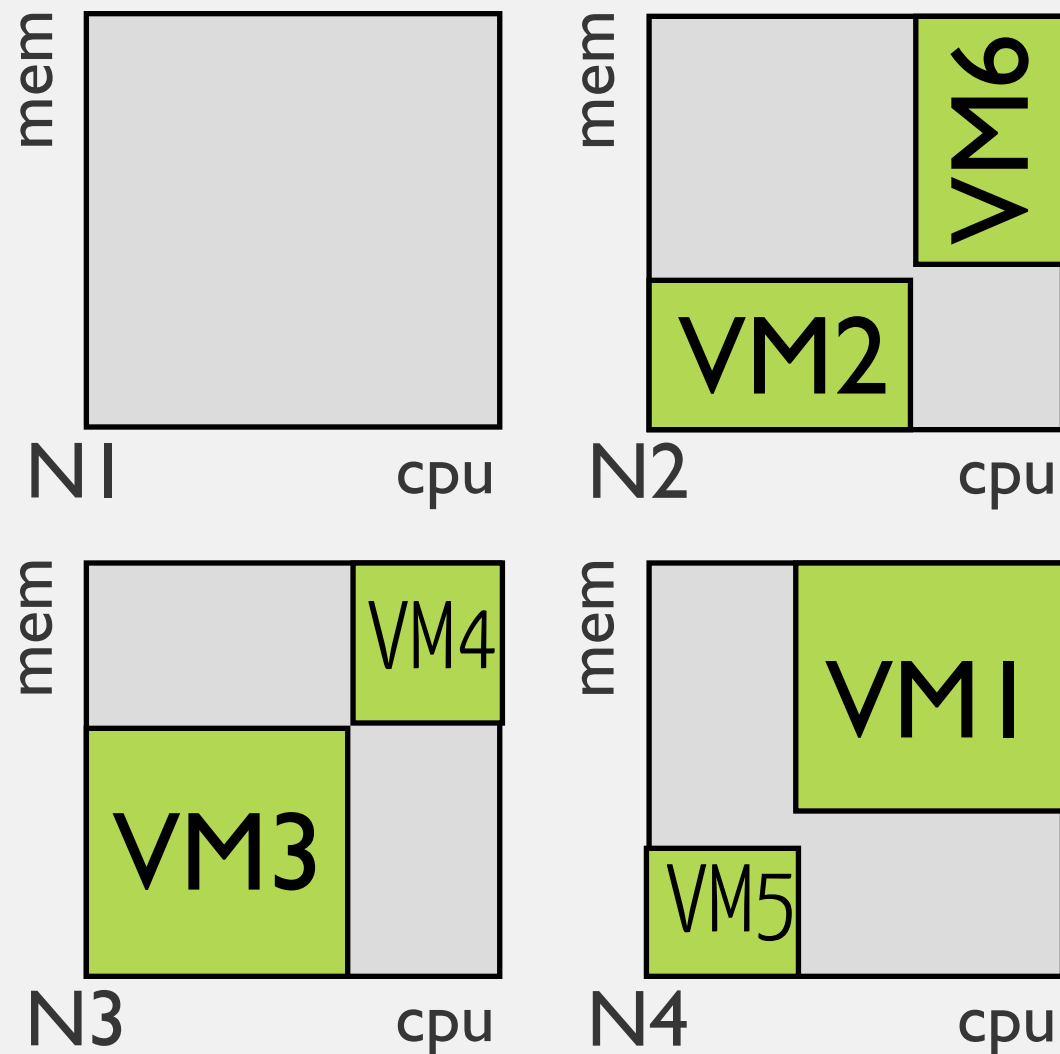
VM duplication between 2 phases  
to simulate the migration

easy to implement

hard to manipulate to  
compute long term previsions

# how to support migrations

a simple way



n-phases vector packing

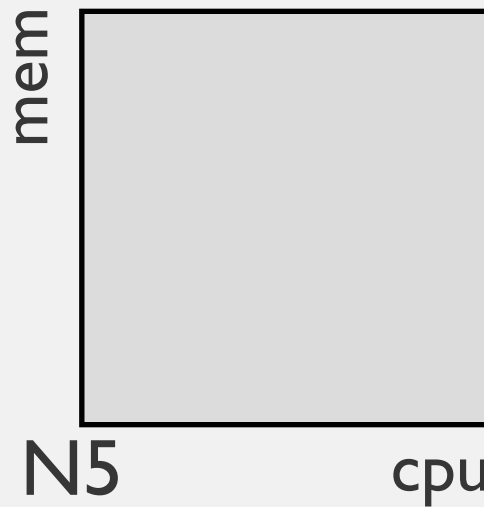
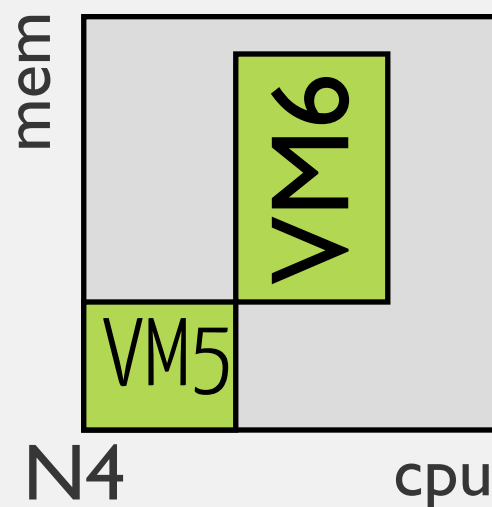
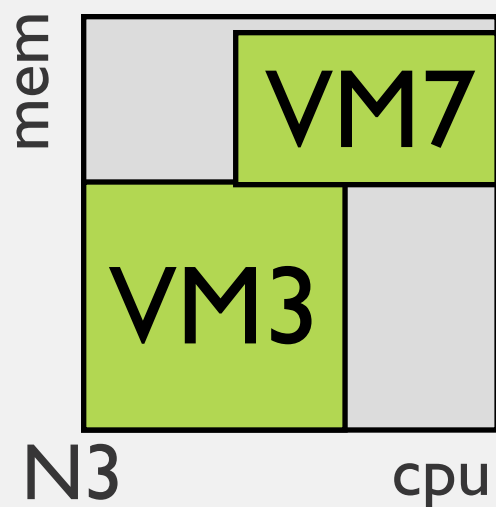
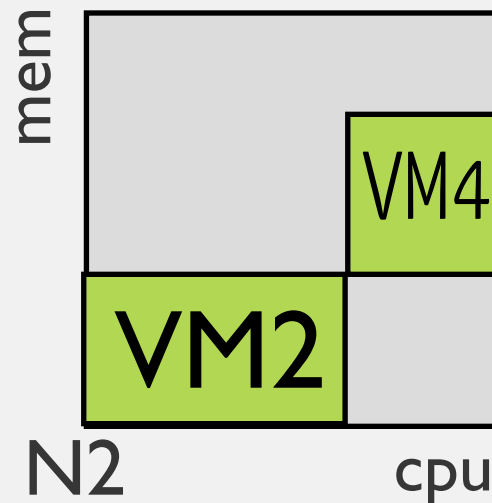
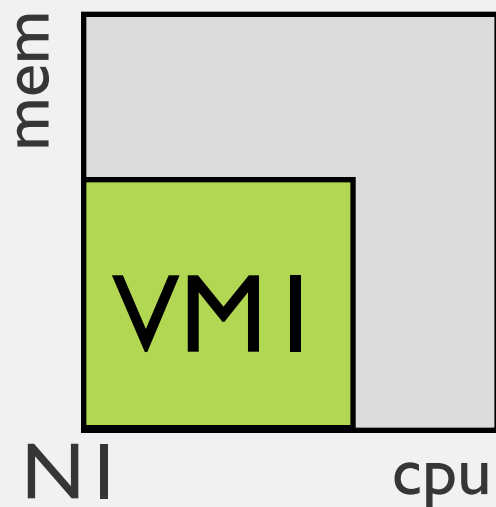
VM duplication between 2 phases  
to simulate the migration

easy to implement

hard to manipulate to  
compute long term previsions

# how to support migrations

the alternative way



1-phase  
+ compute dependencies  
step by steps

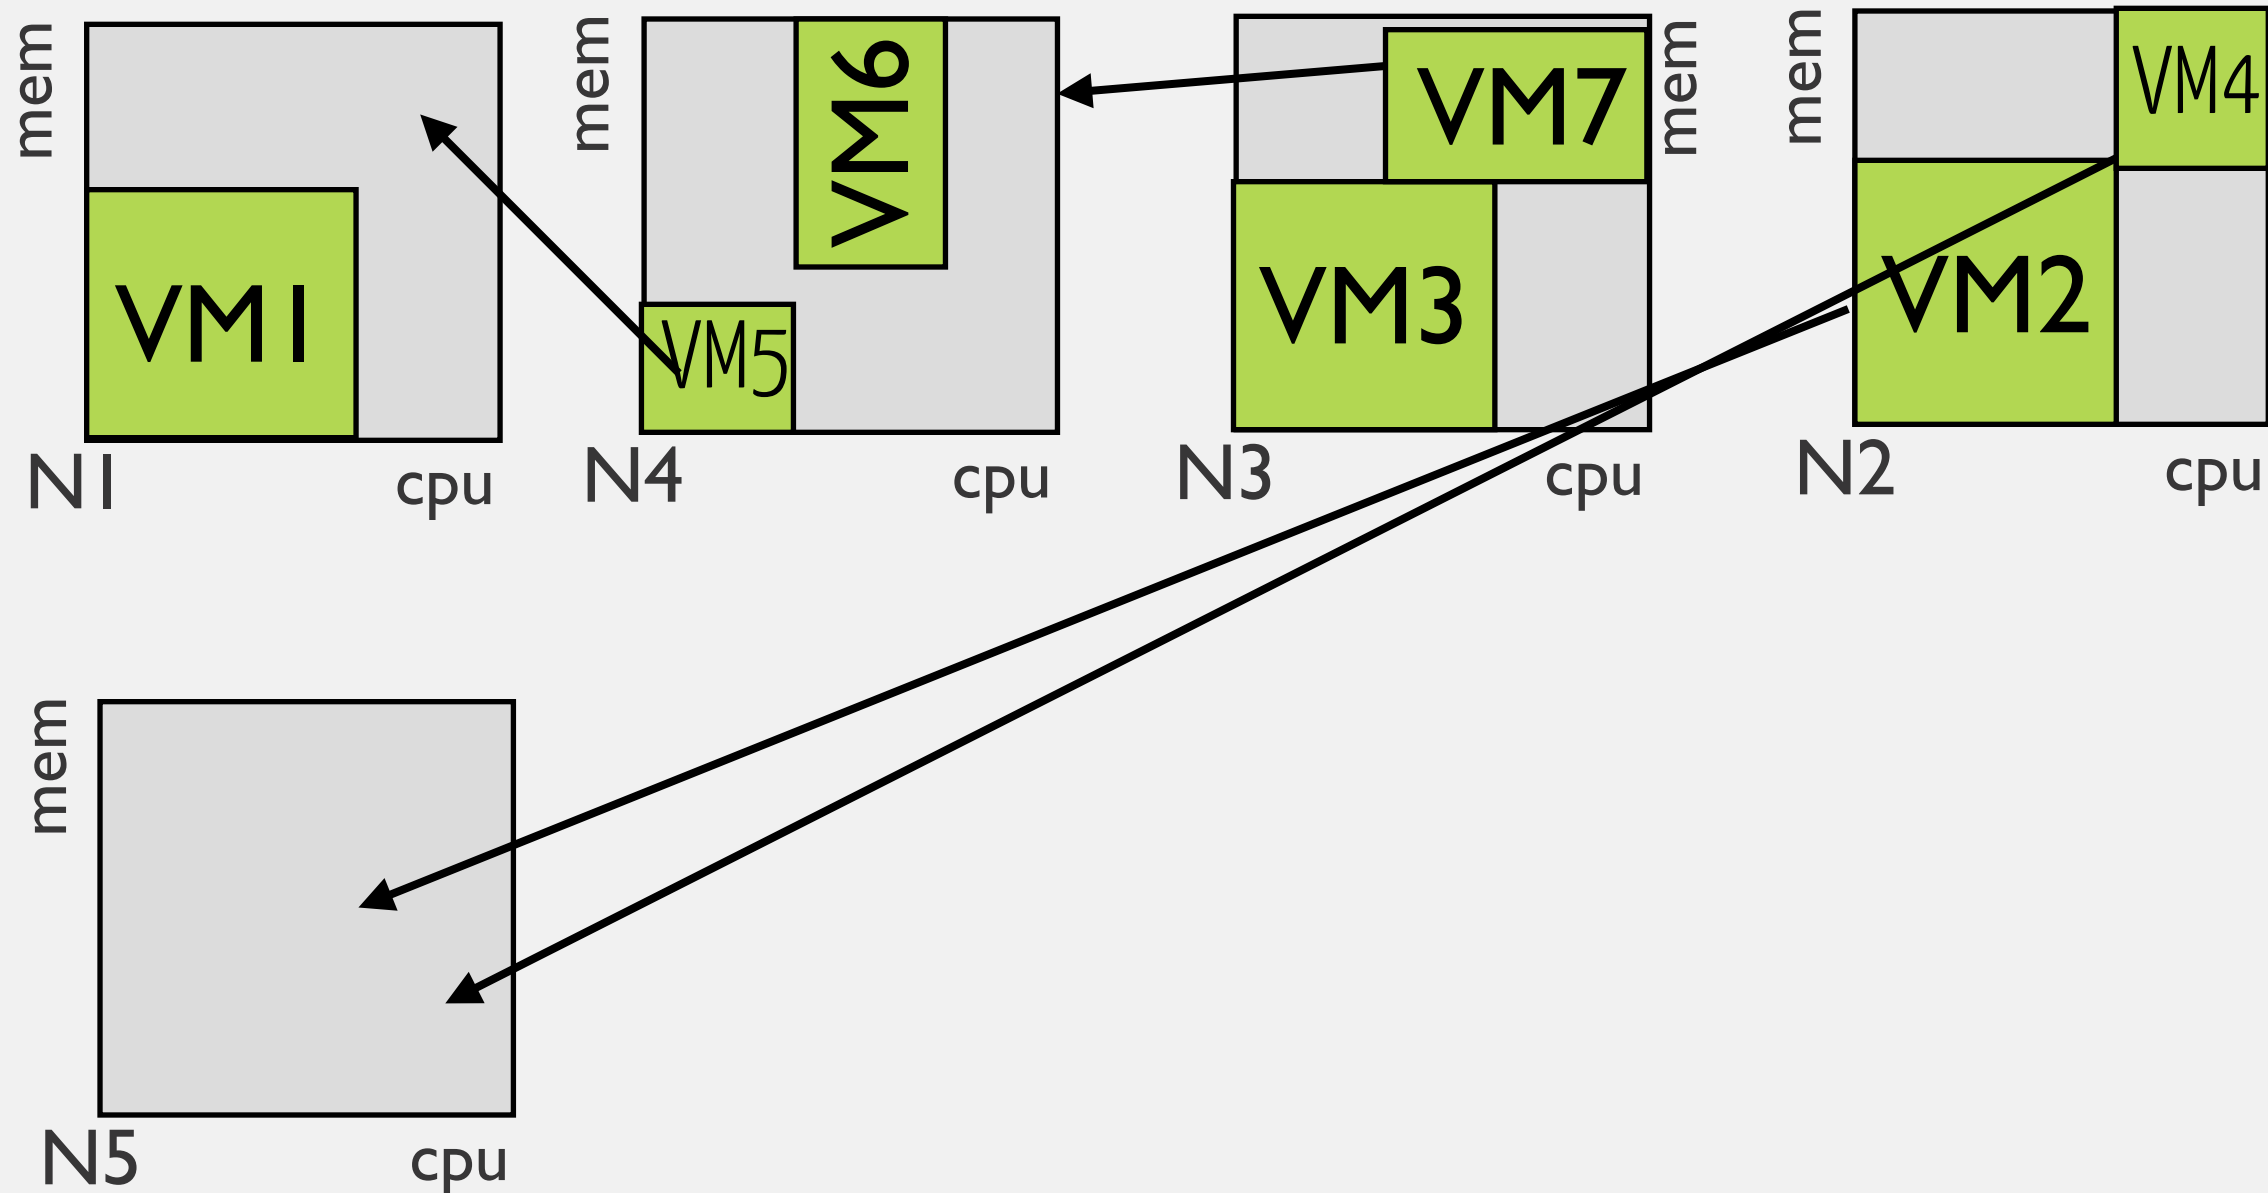
ok to implement

*offline(N2) + no CPU sharing*

! cycles  
#migrations

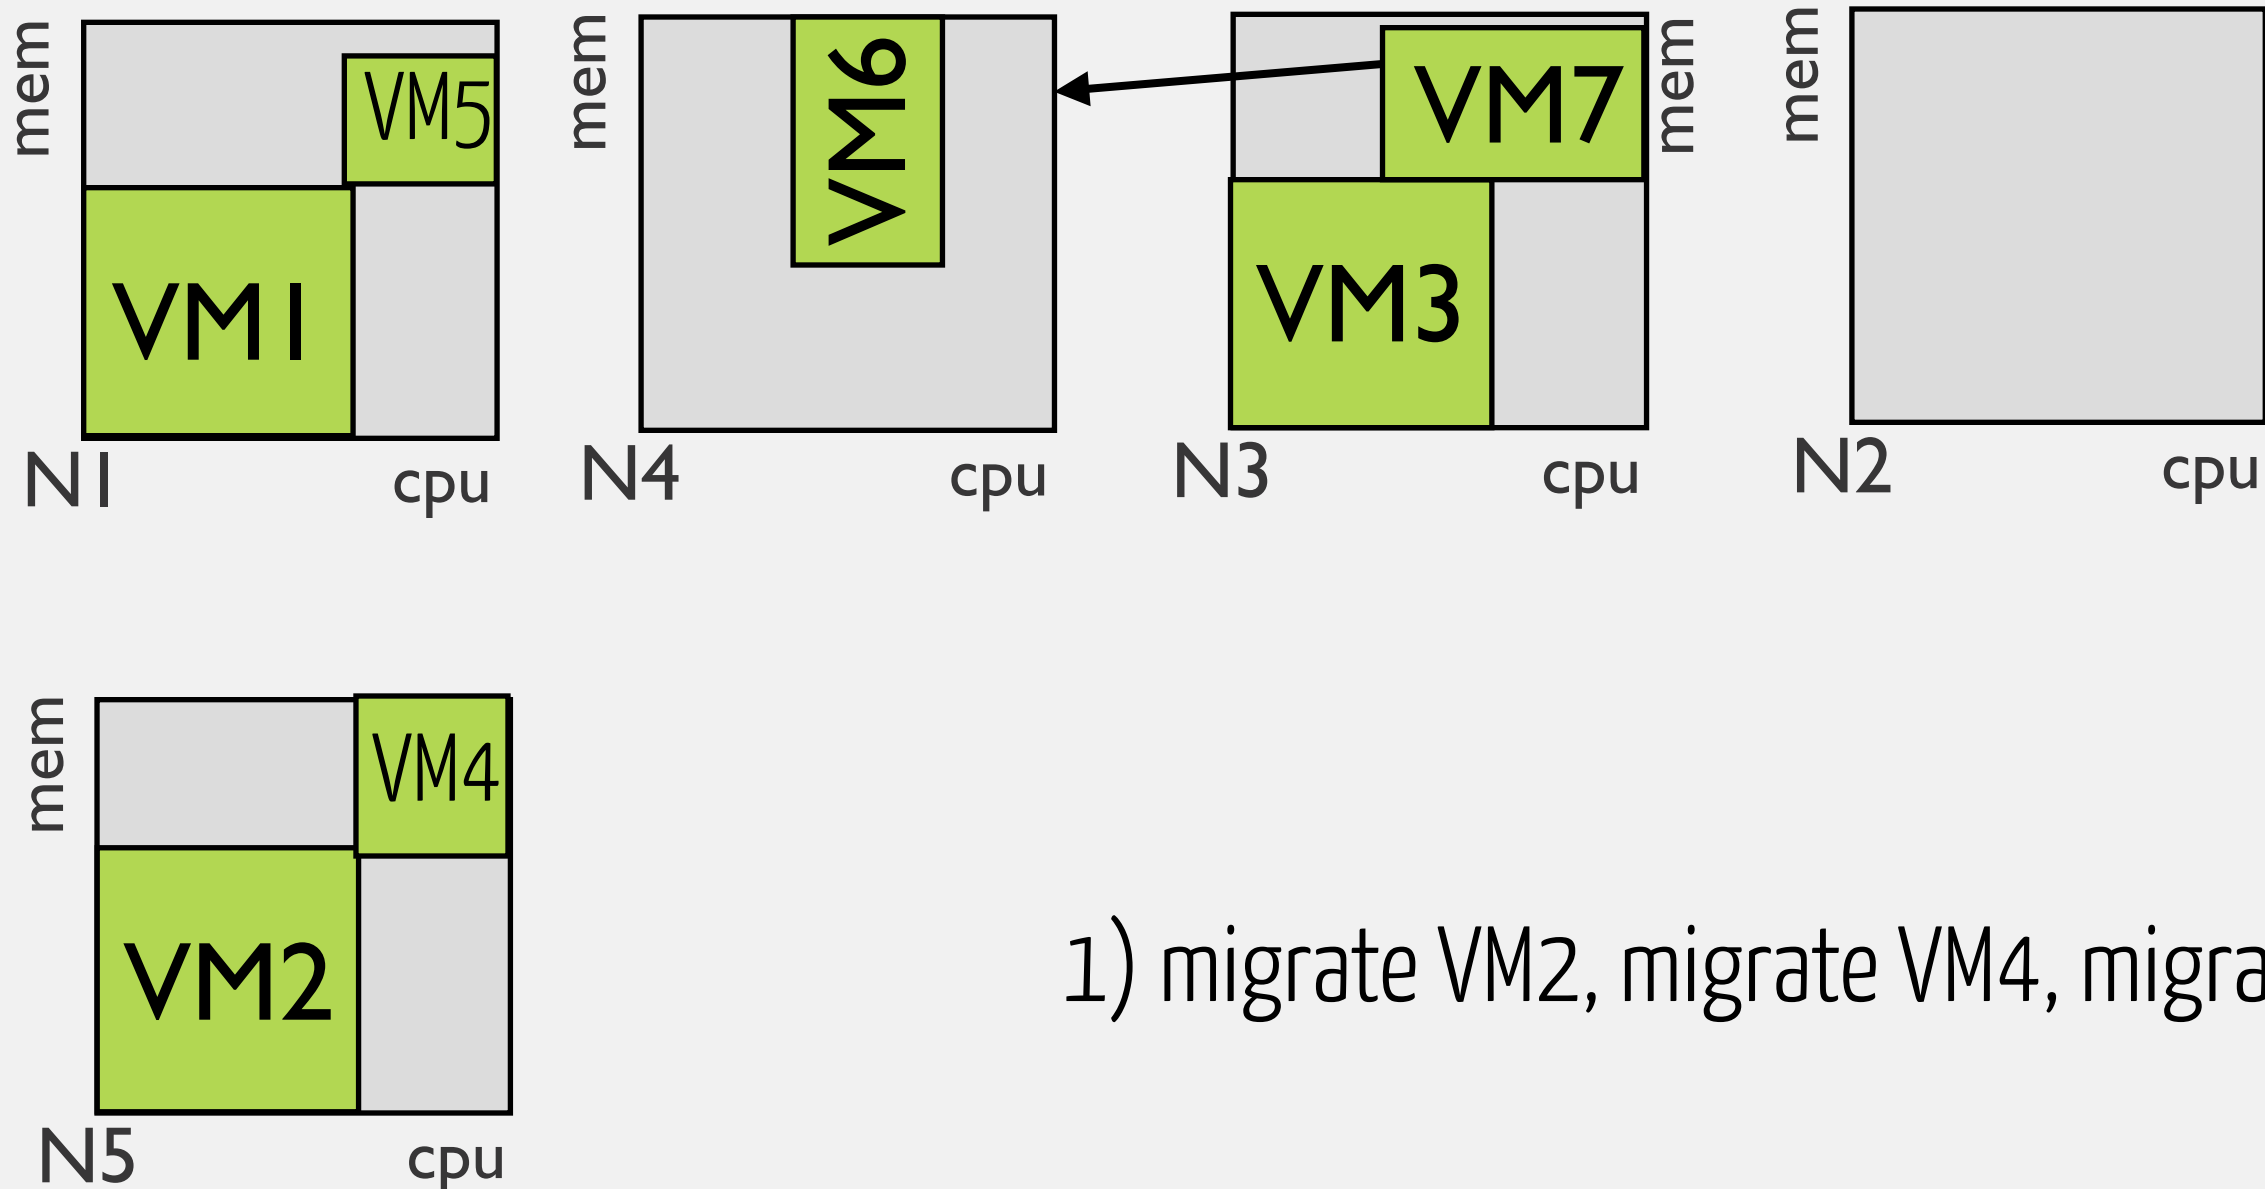
# how to support migrations

the alternative way



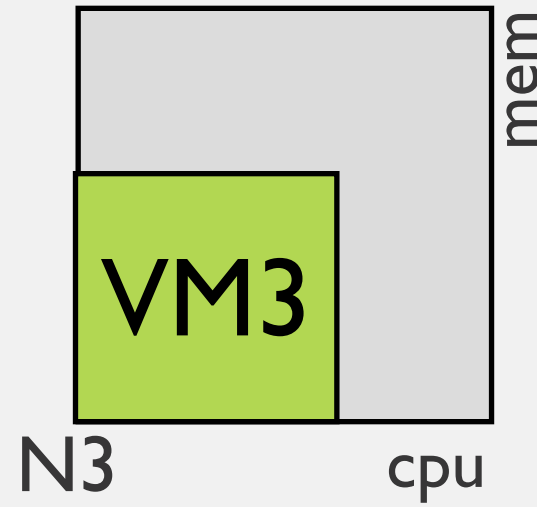
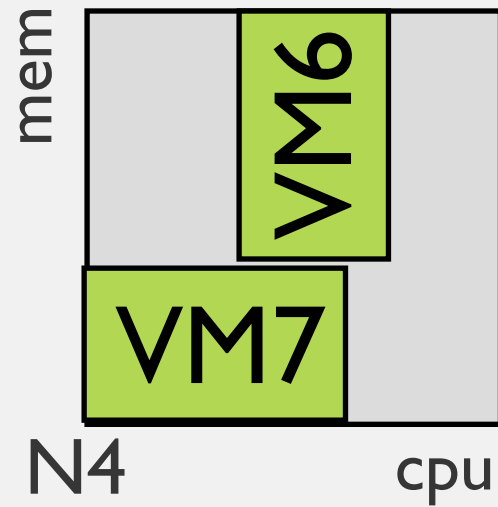
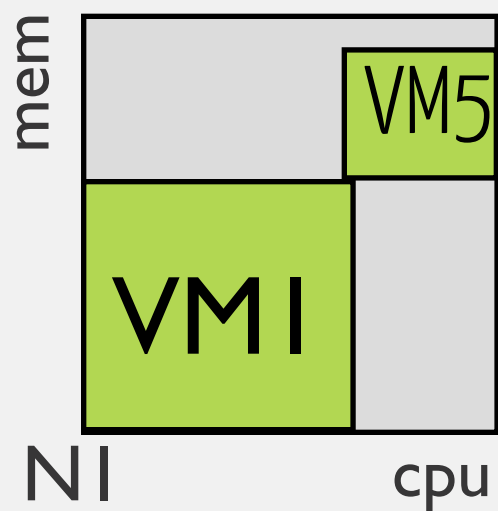
# how to support migrations

the alternative way

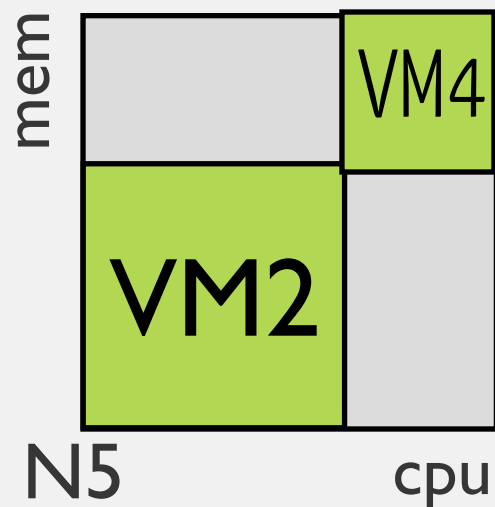


# how to support migrations

the alternative way

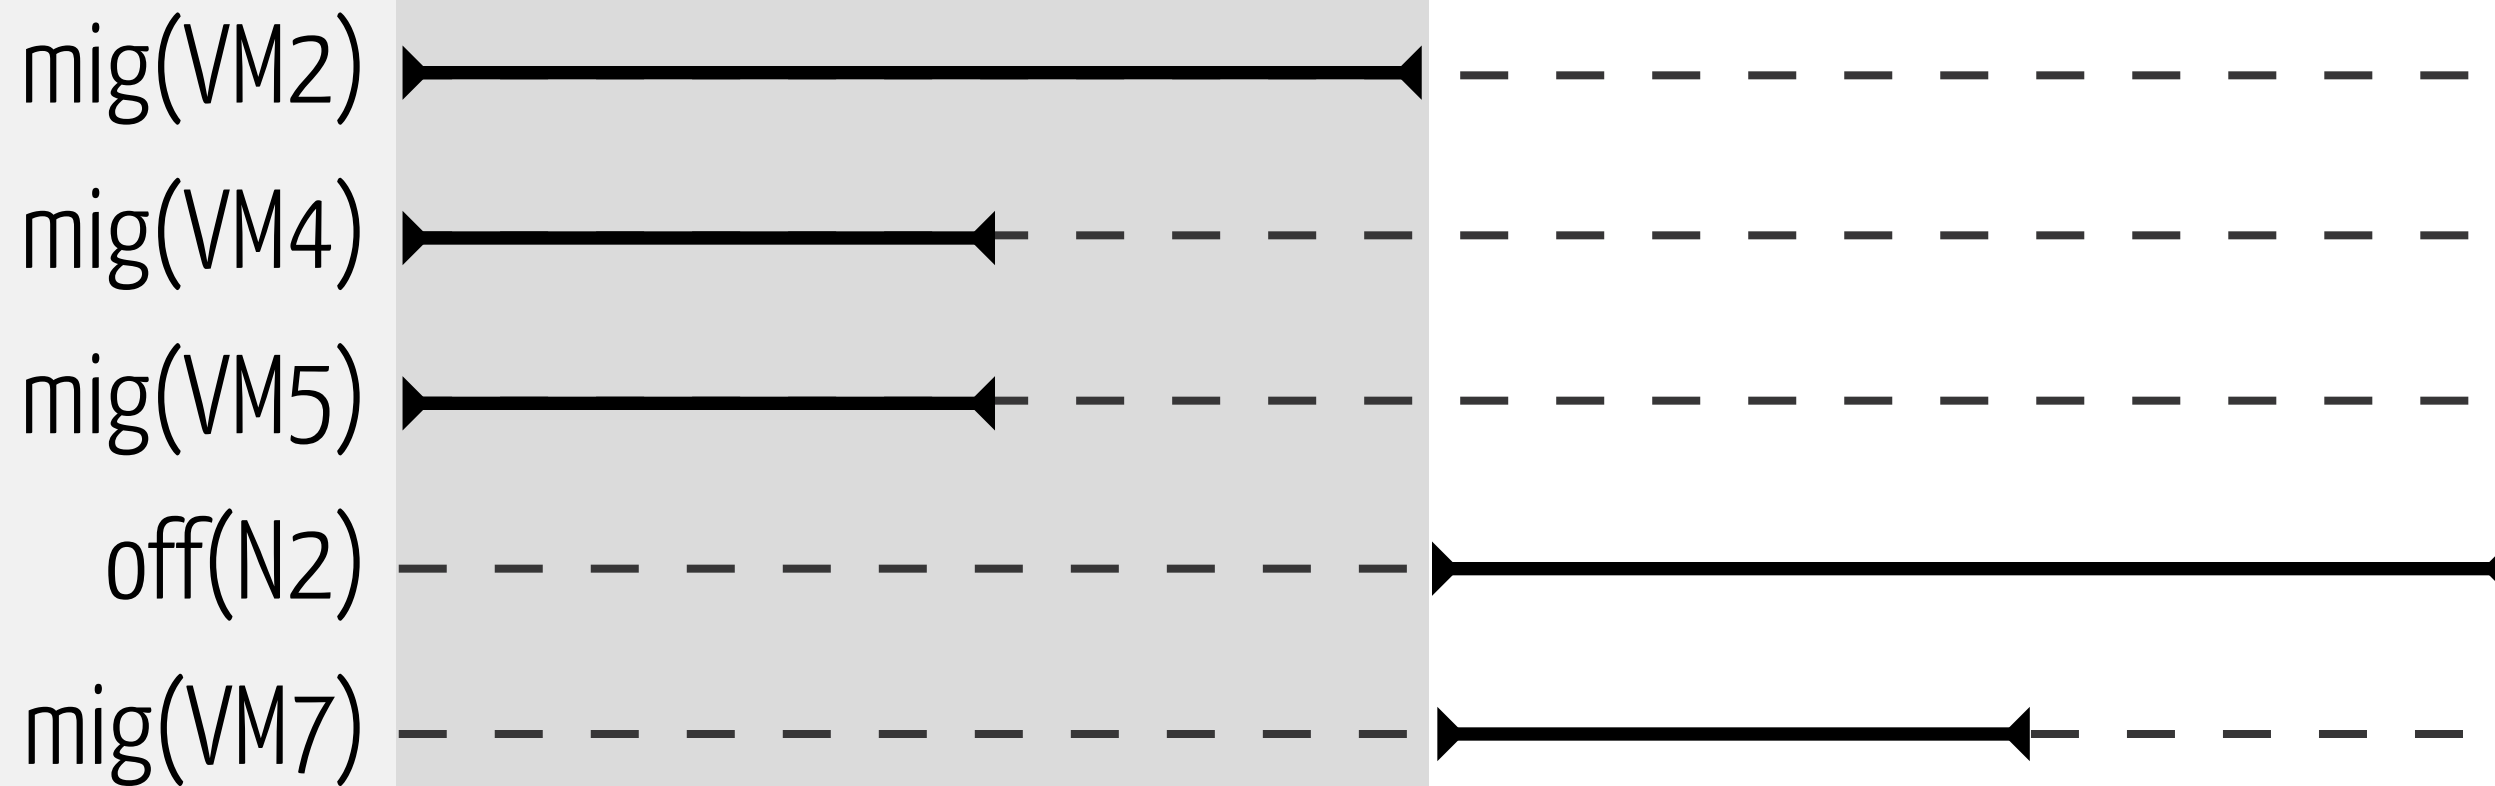


N2



- 1) migrate VM2, migrate VM4, migrate VM5
- 2) shutdown(N2), migrate VM7

# coarse grain staging delay actions



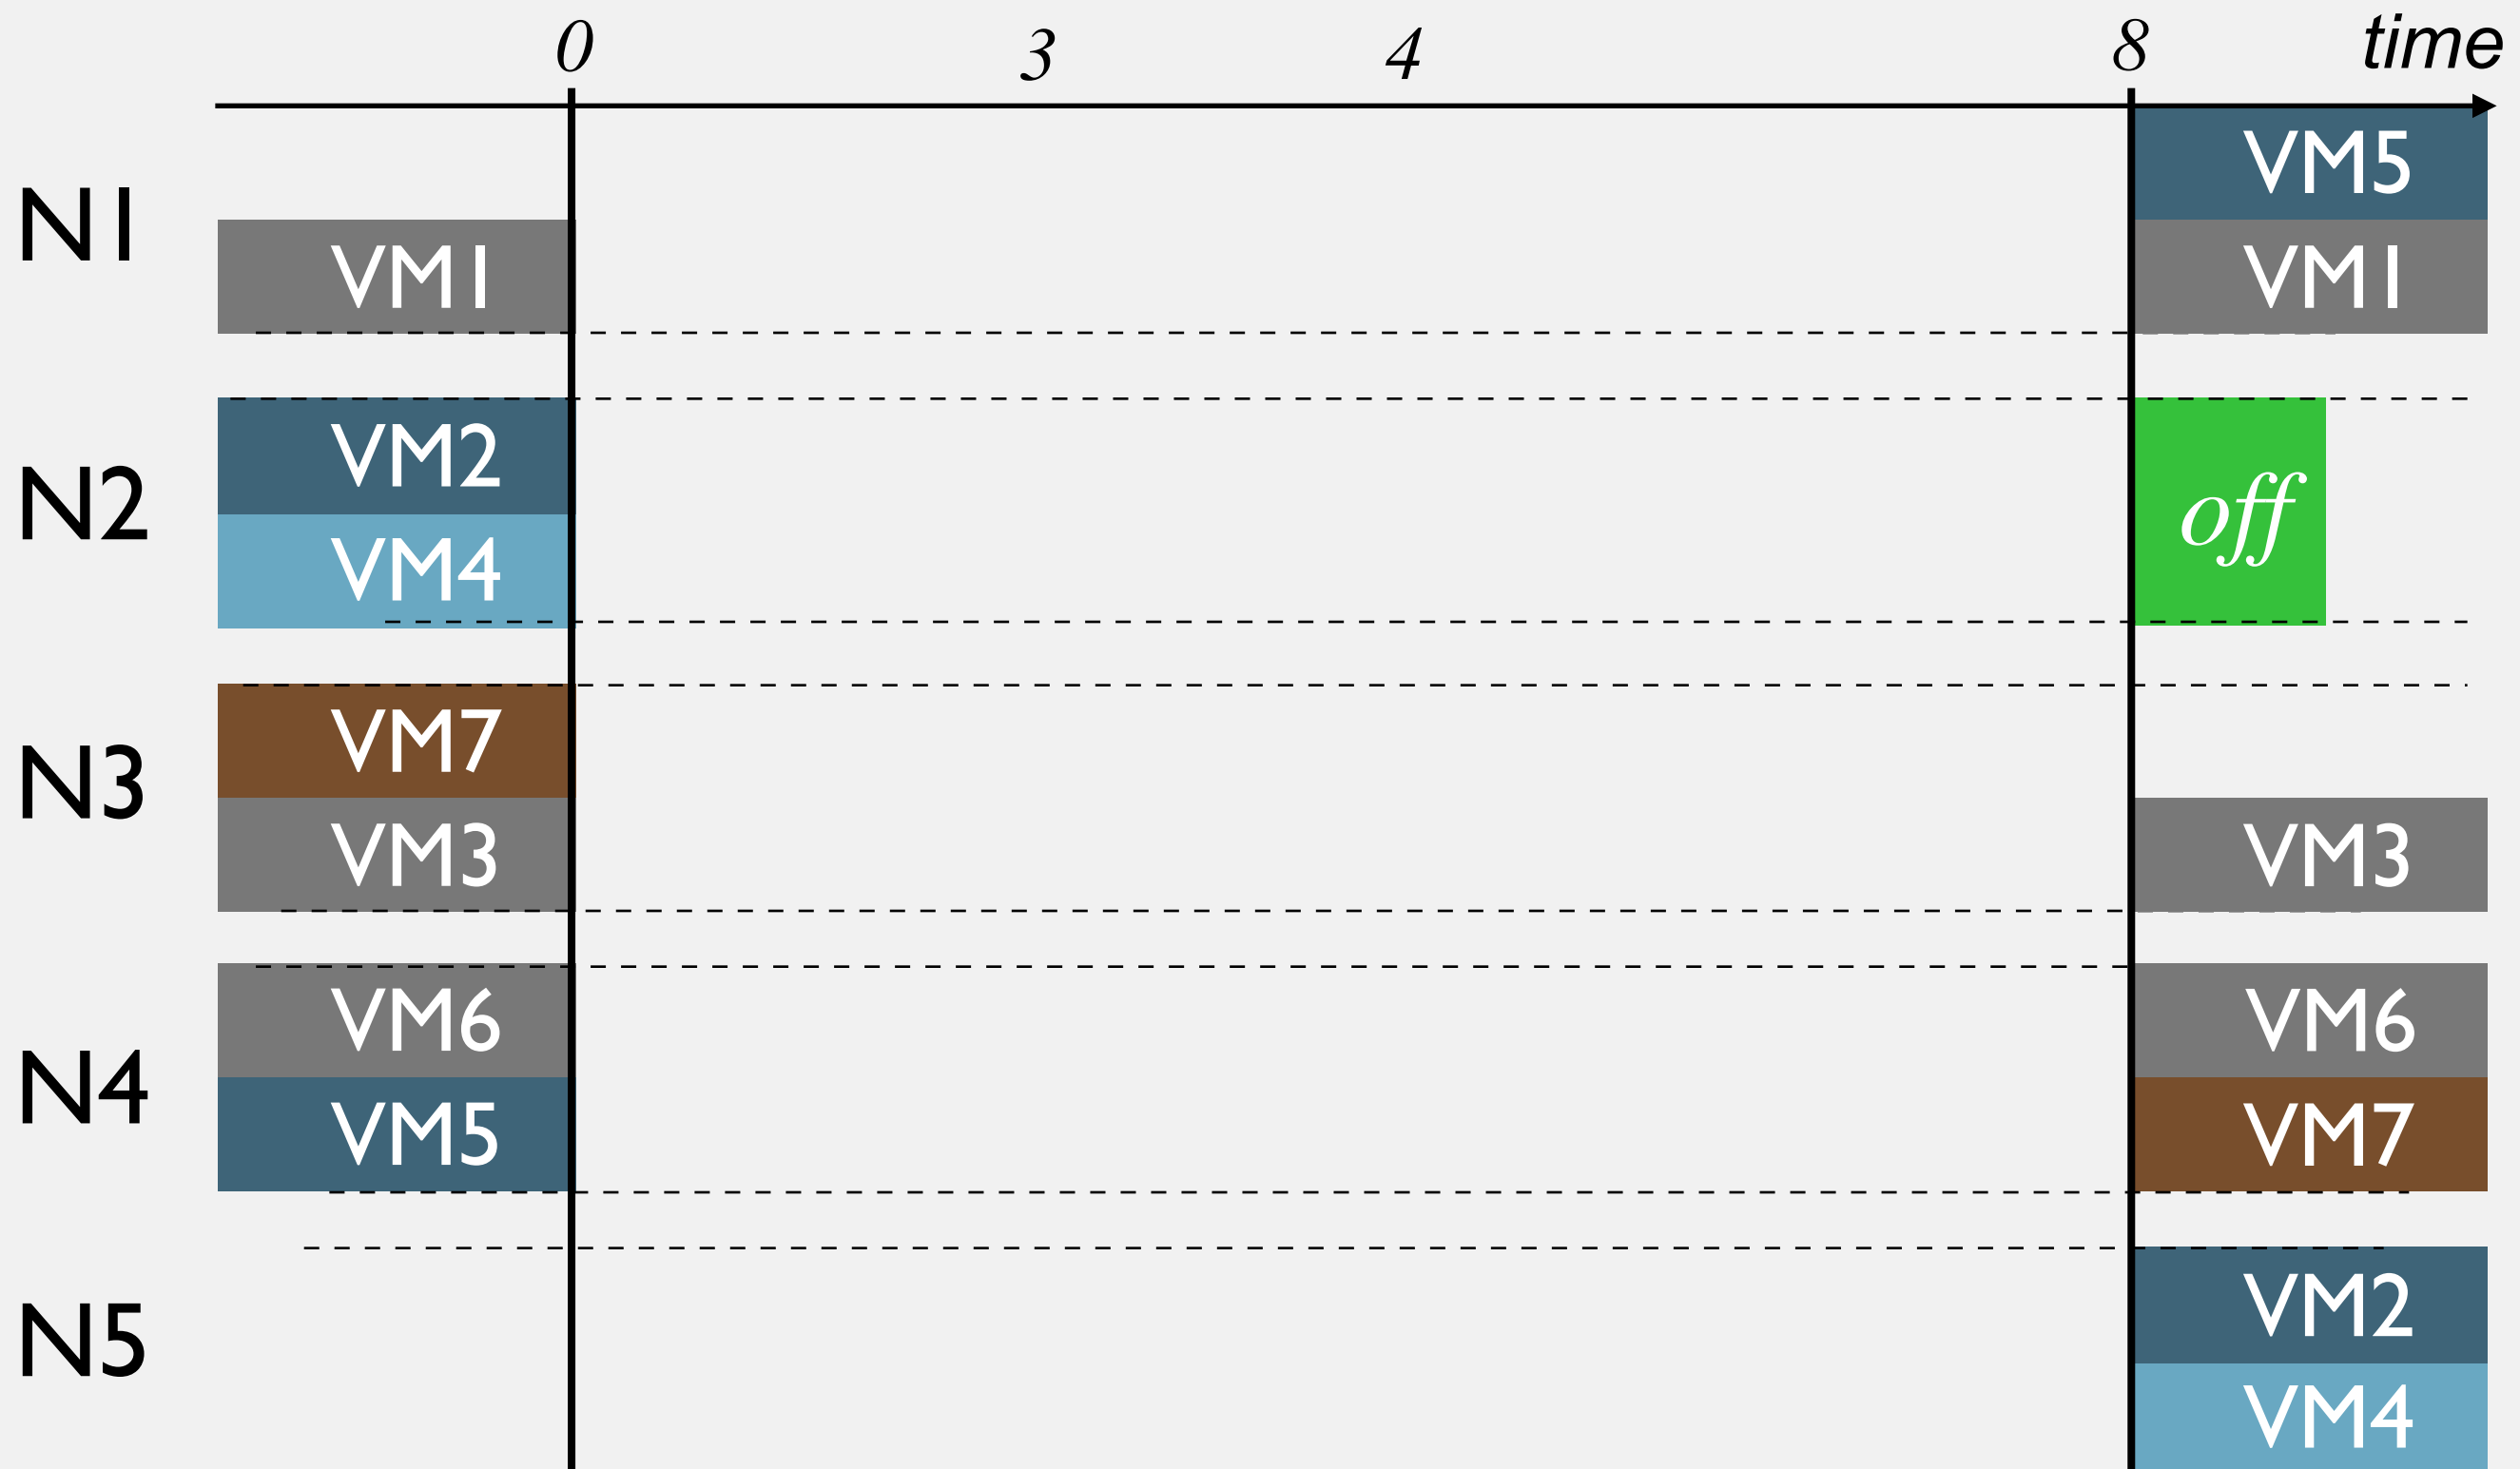
stage1

stage2 *time*



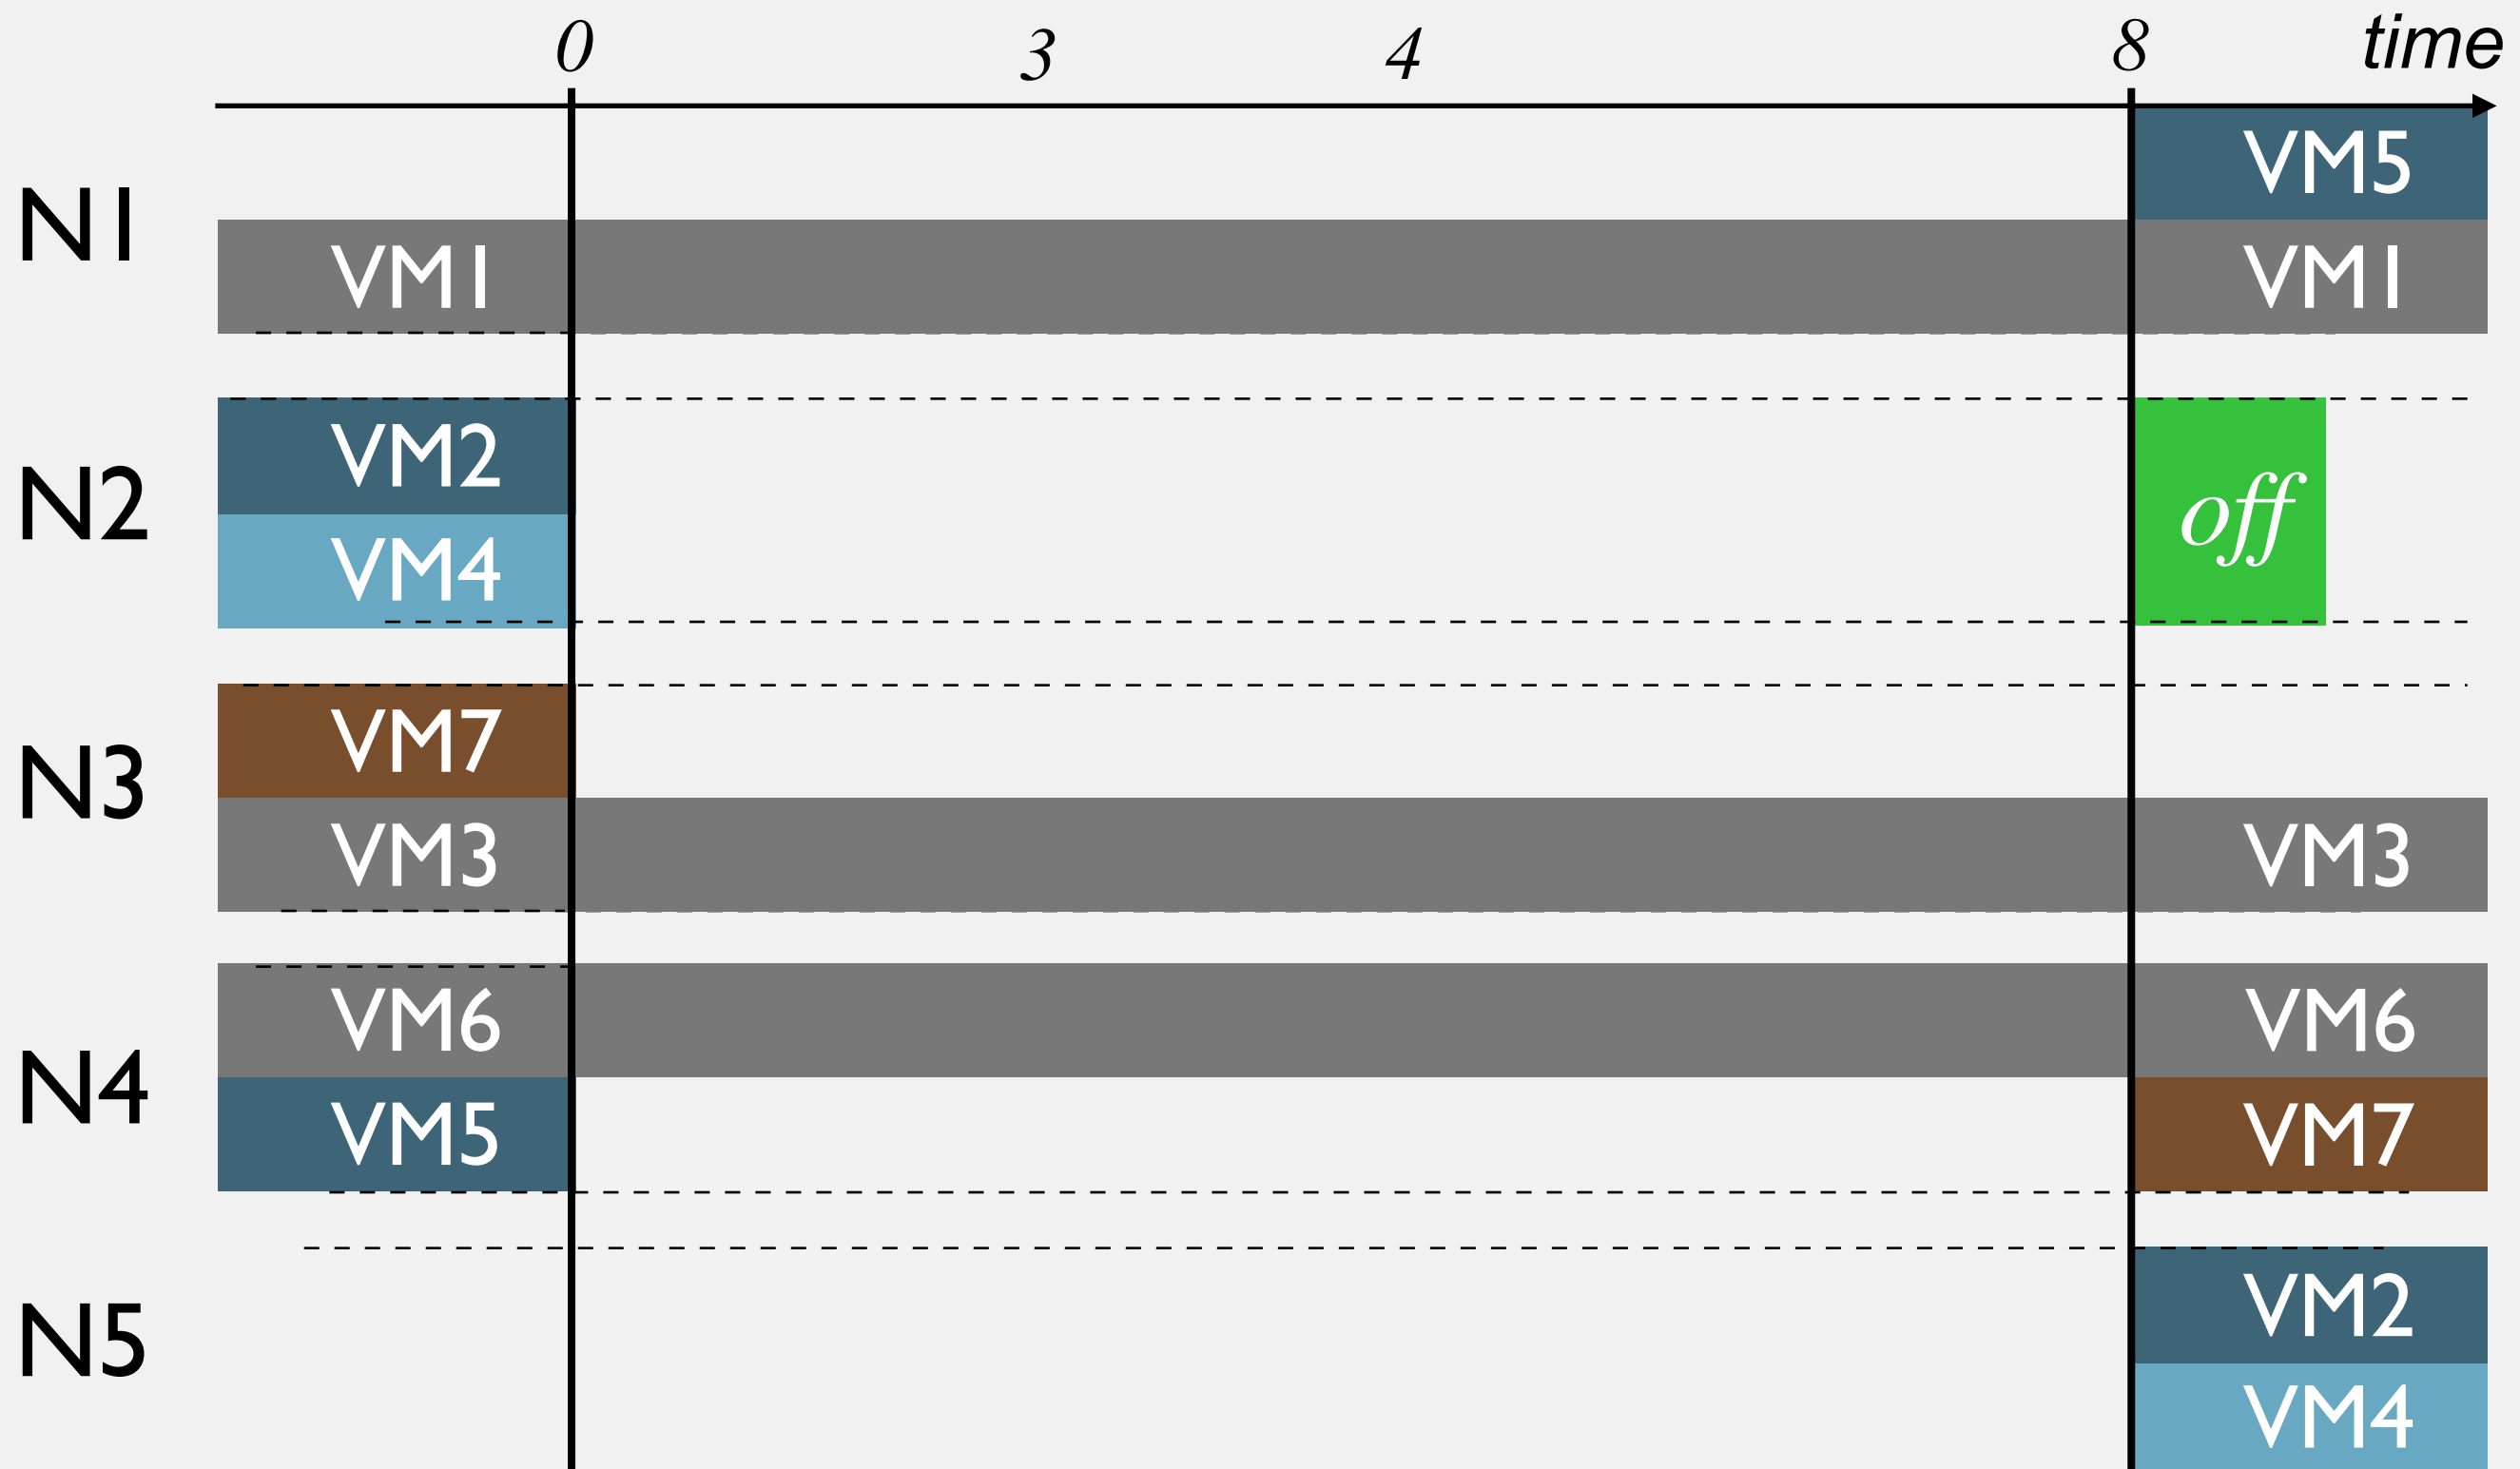
# Resource-Constrained Project Scheduling Problem

the clean way to model space and time



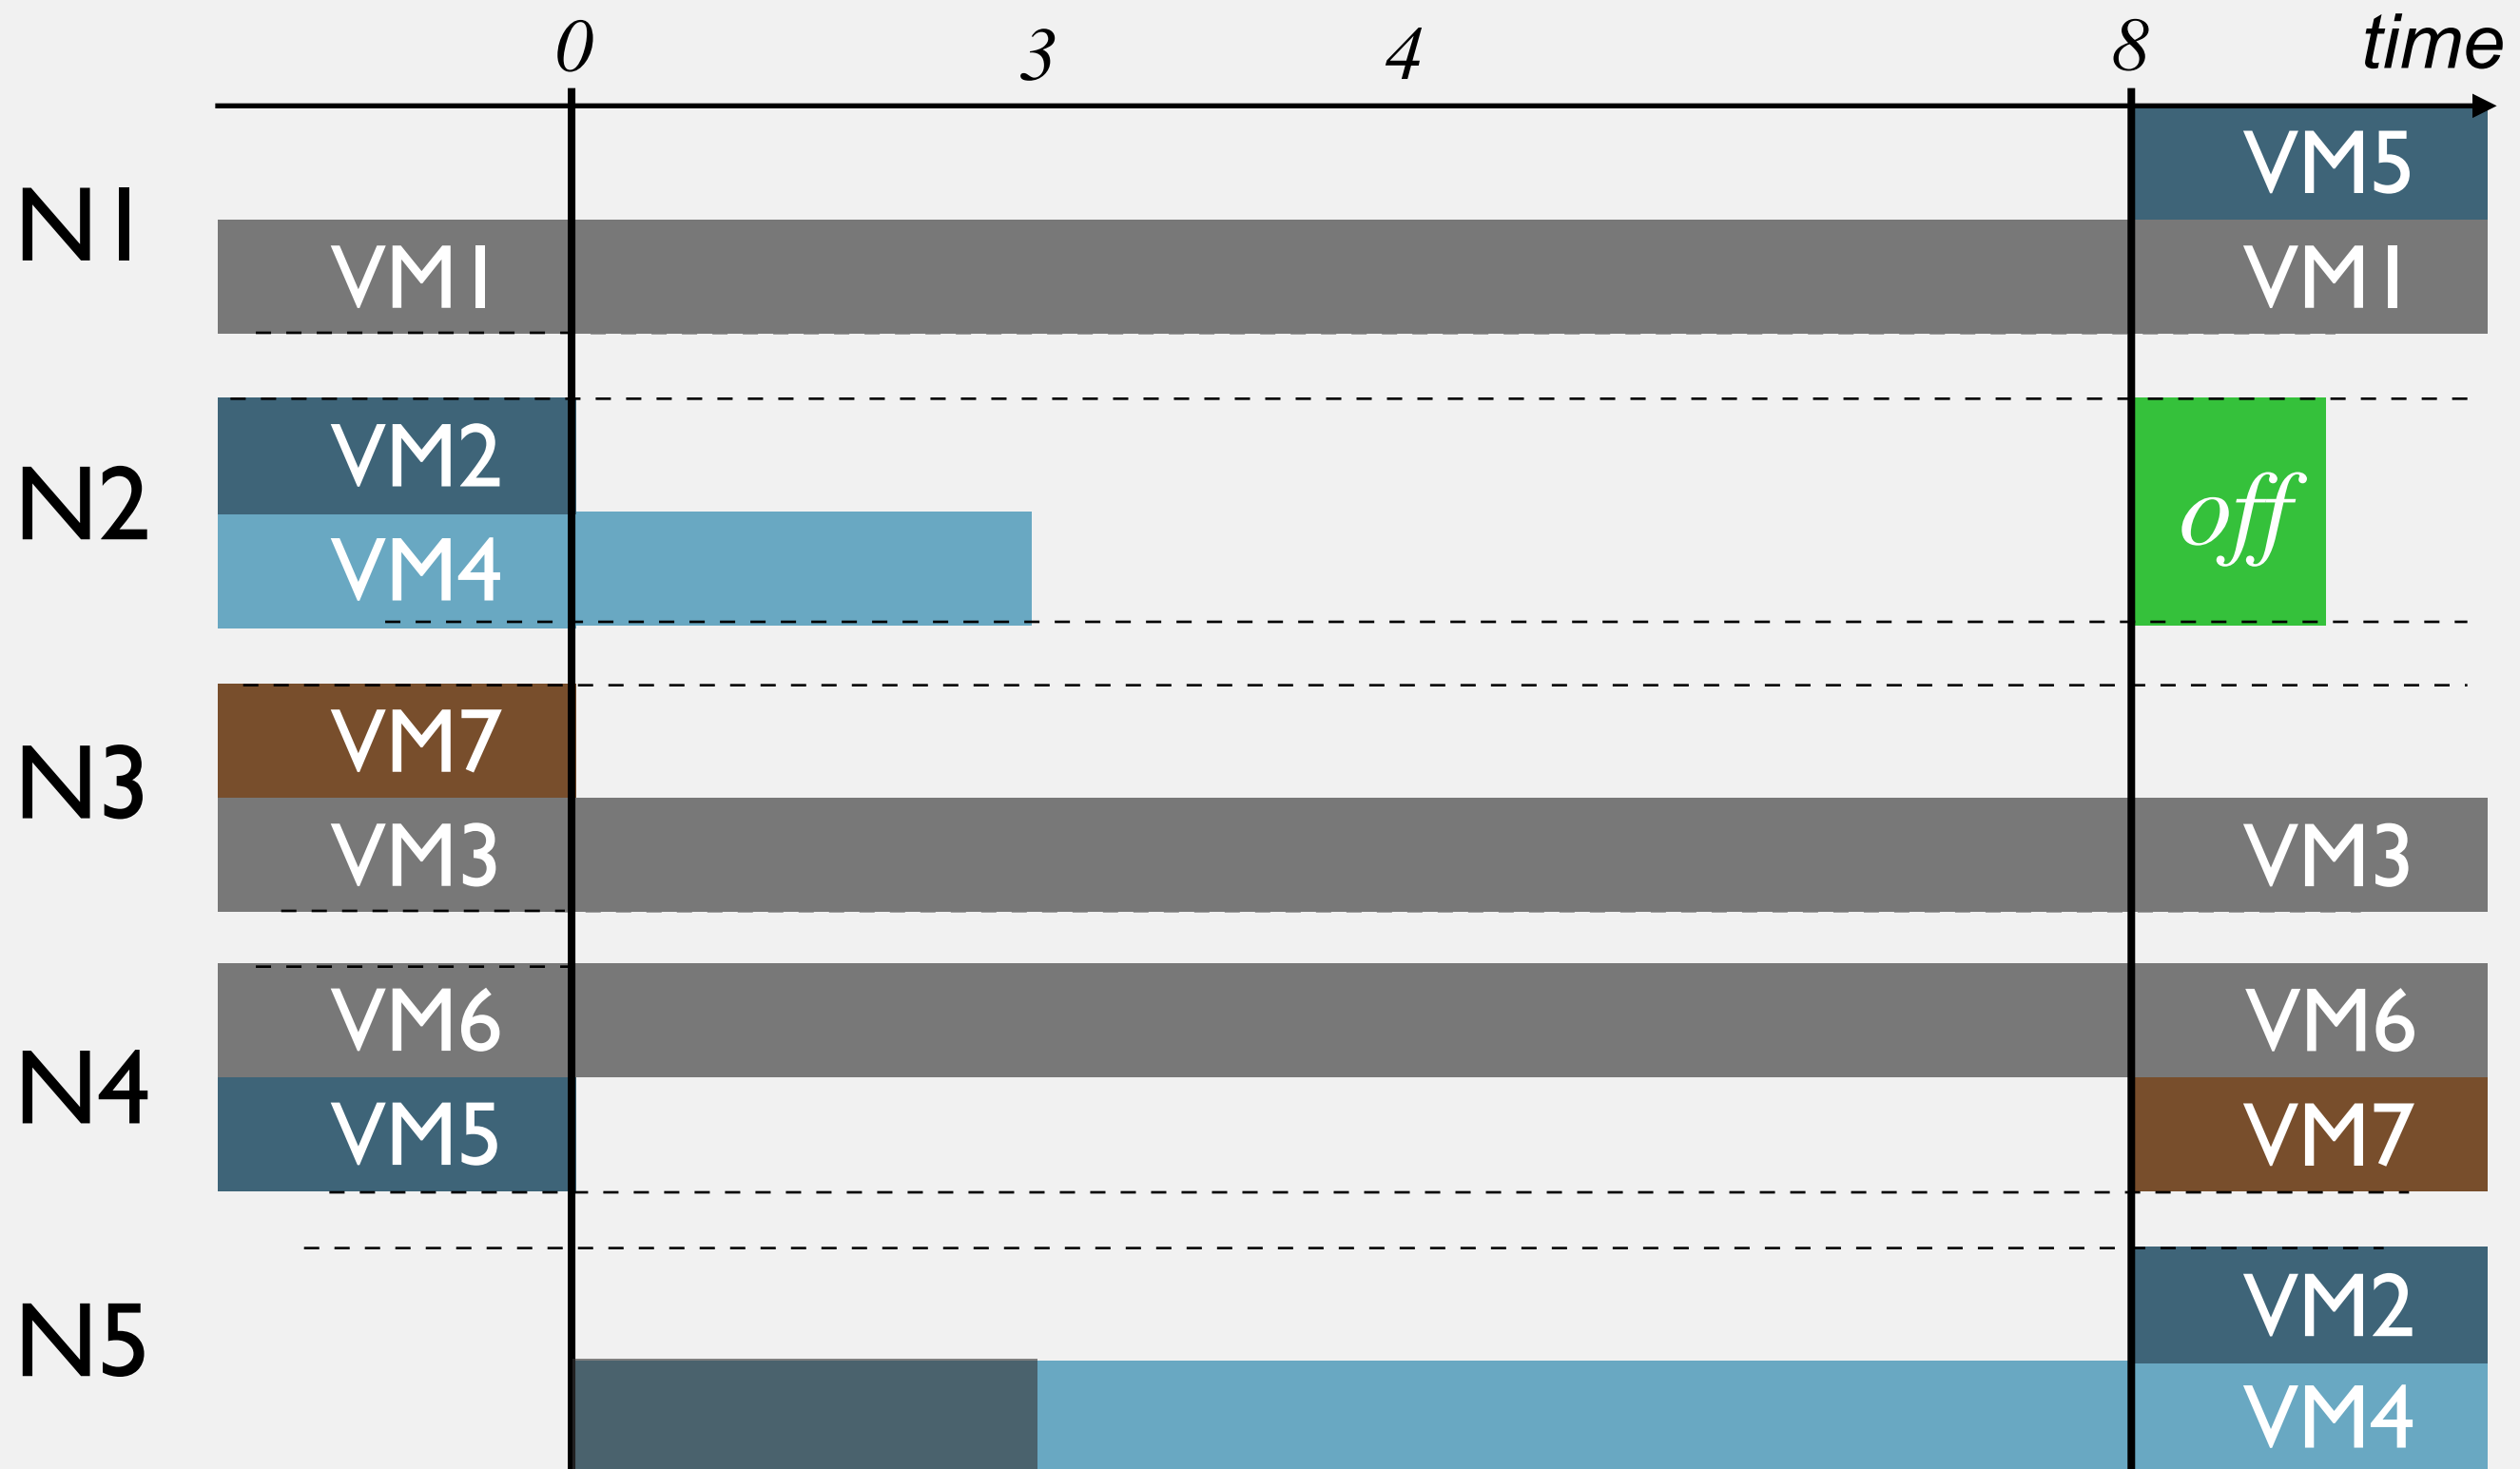
# Resource-Constrained Project Scheduling Problem

the clean way to model space and time



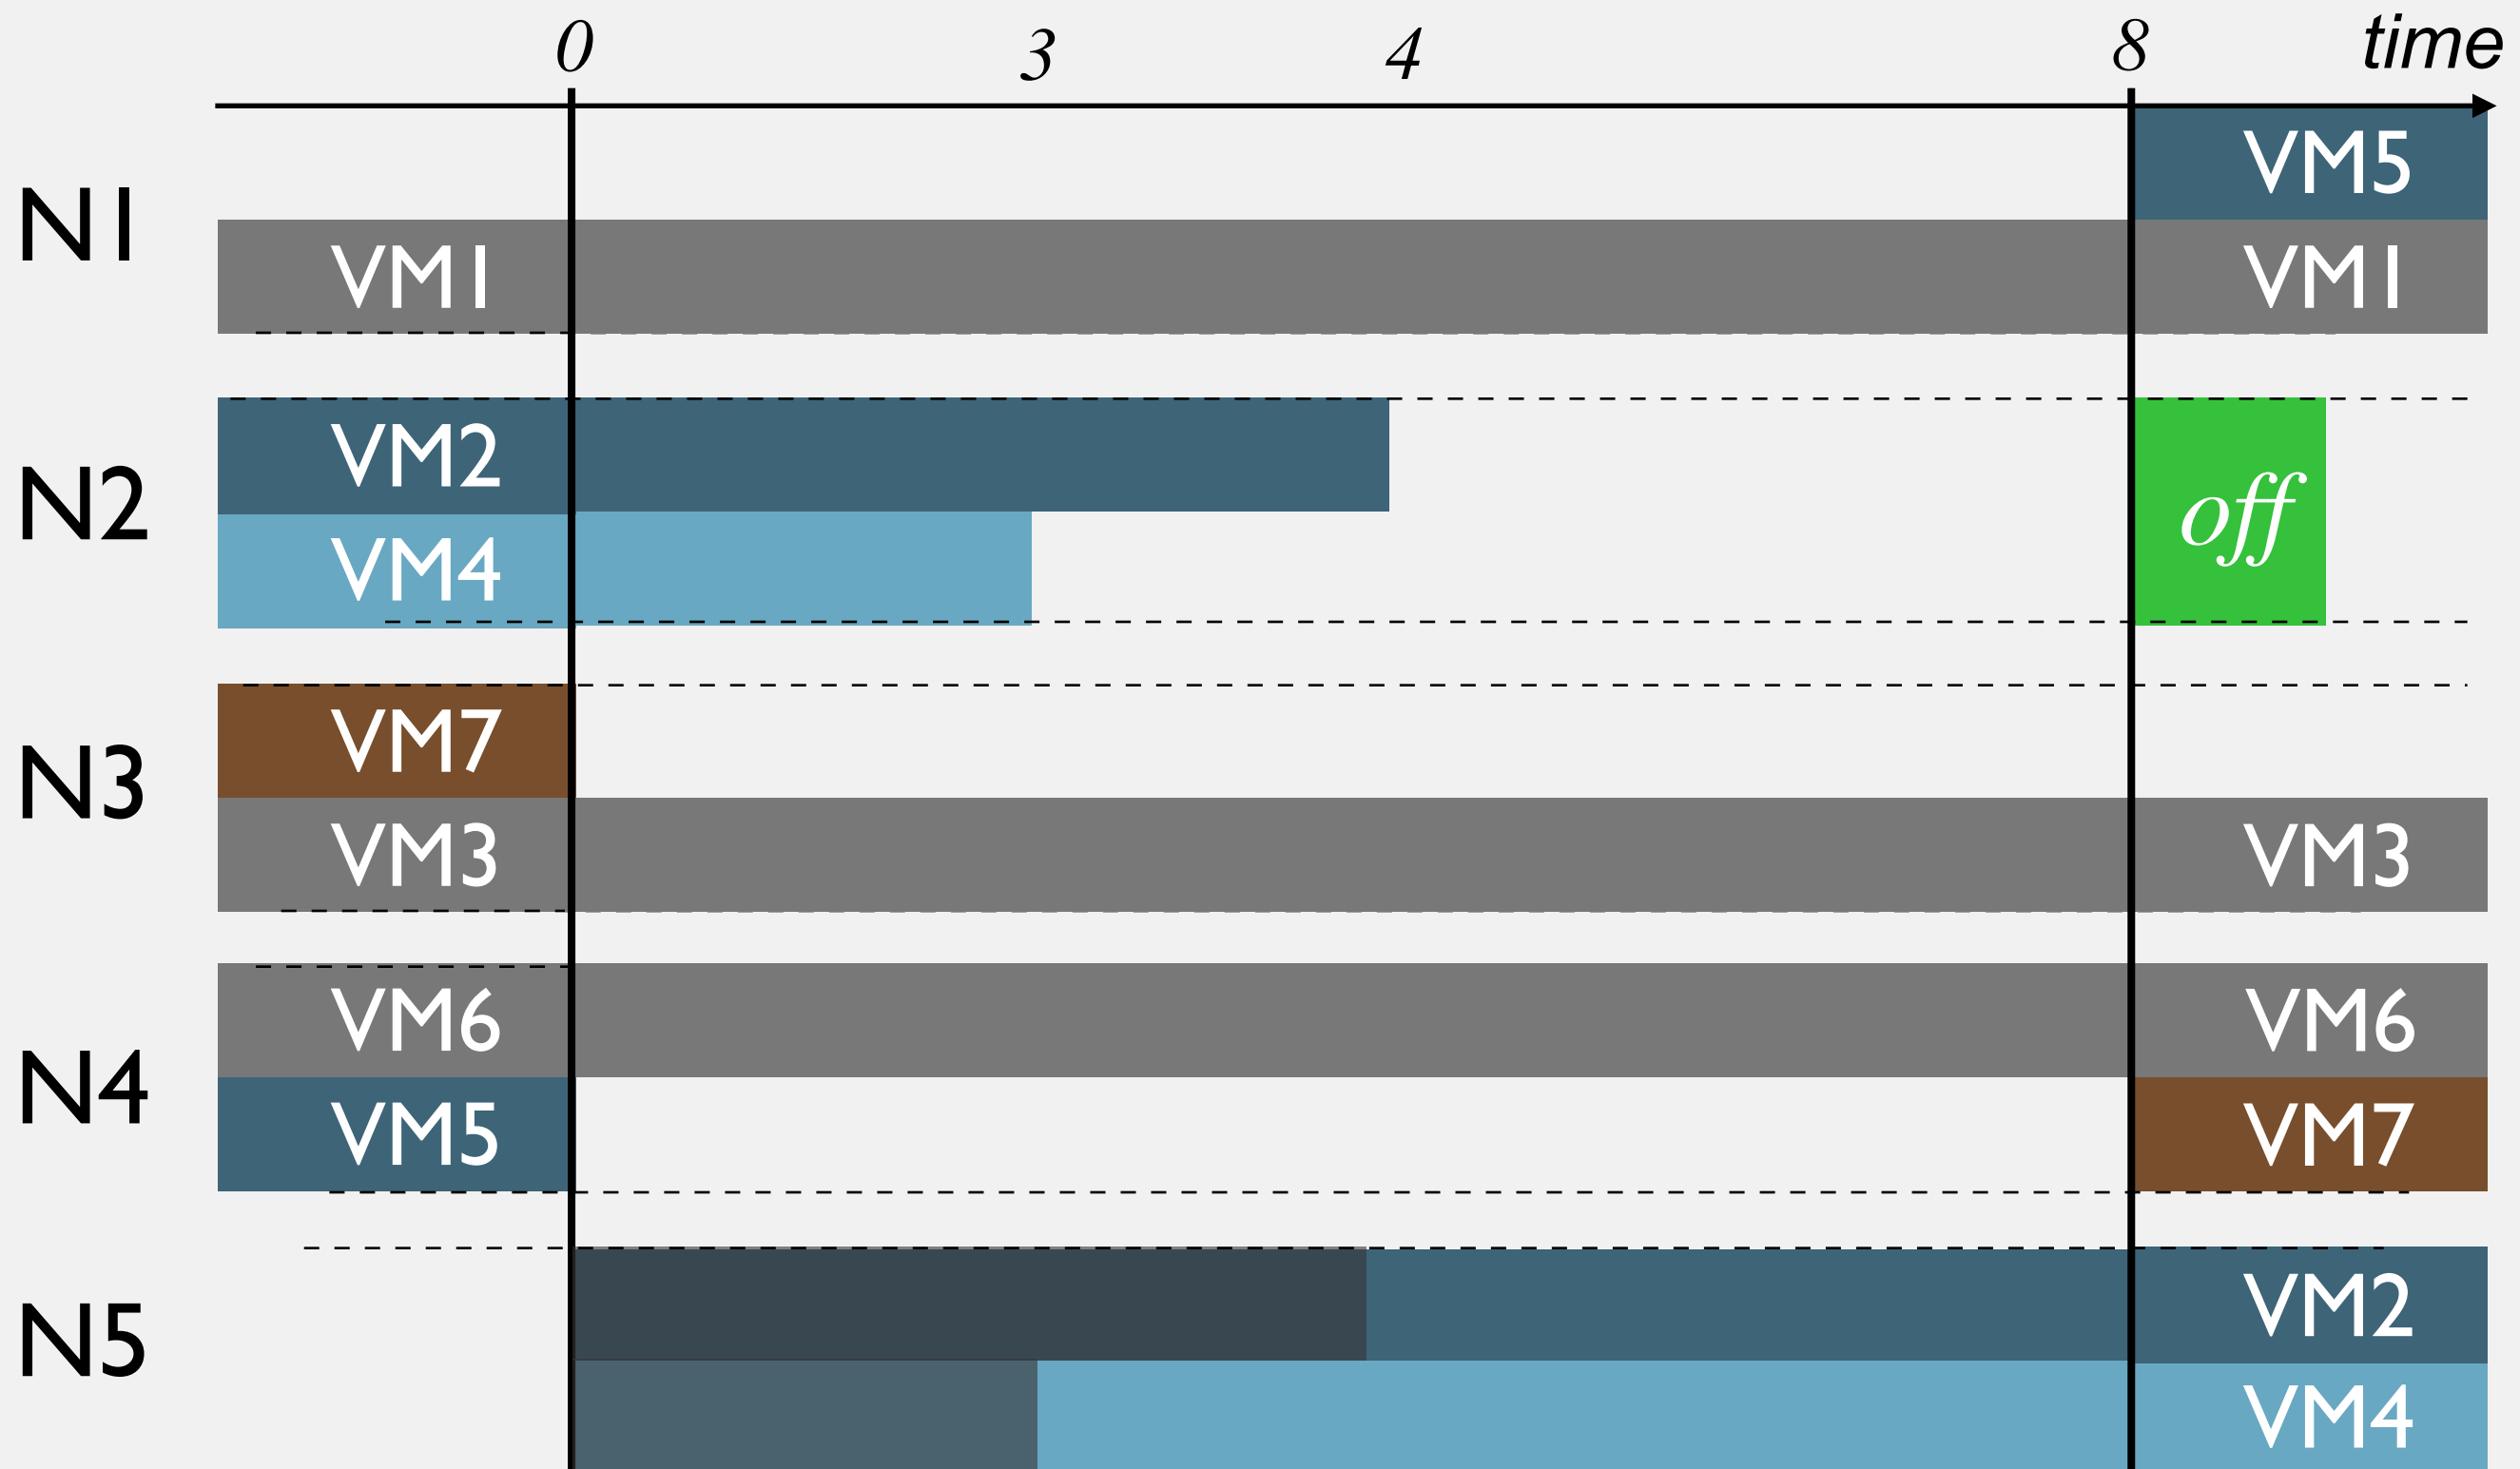
# Resource-Constrained Project Scheduling Problem

the clean way to model space and time



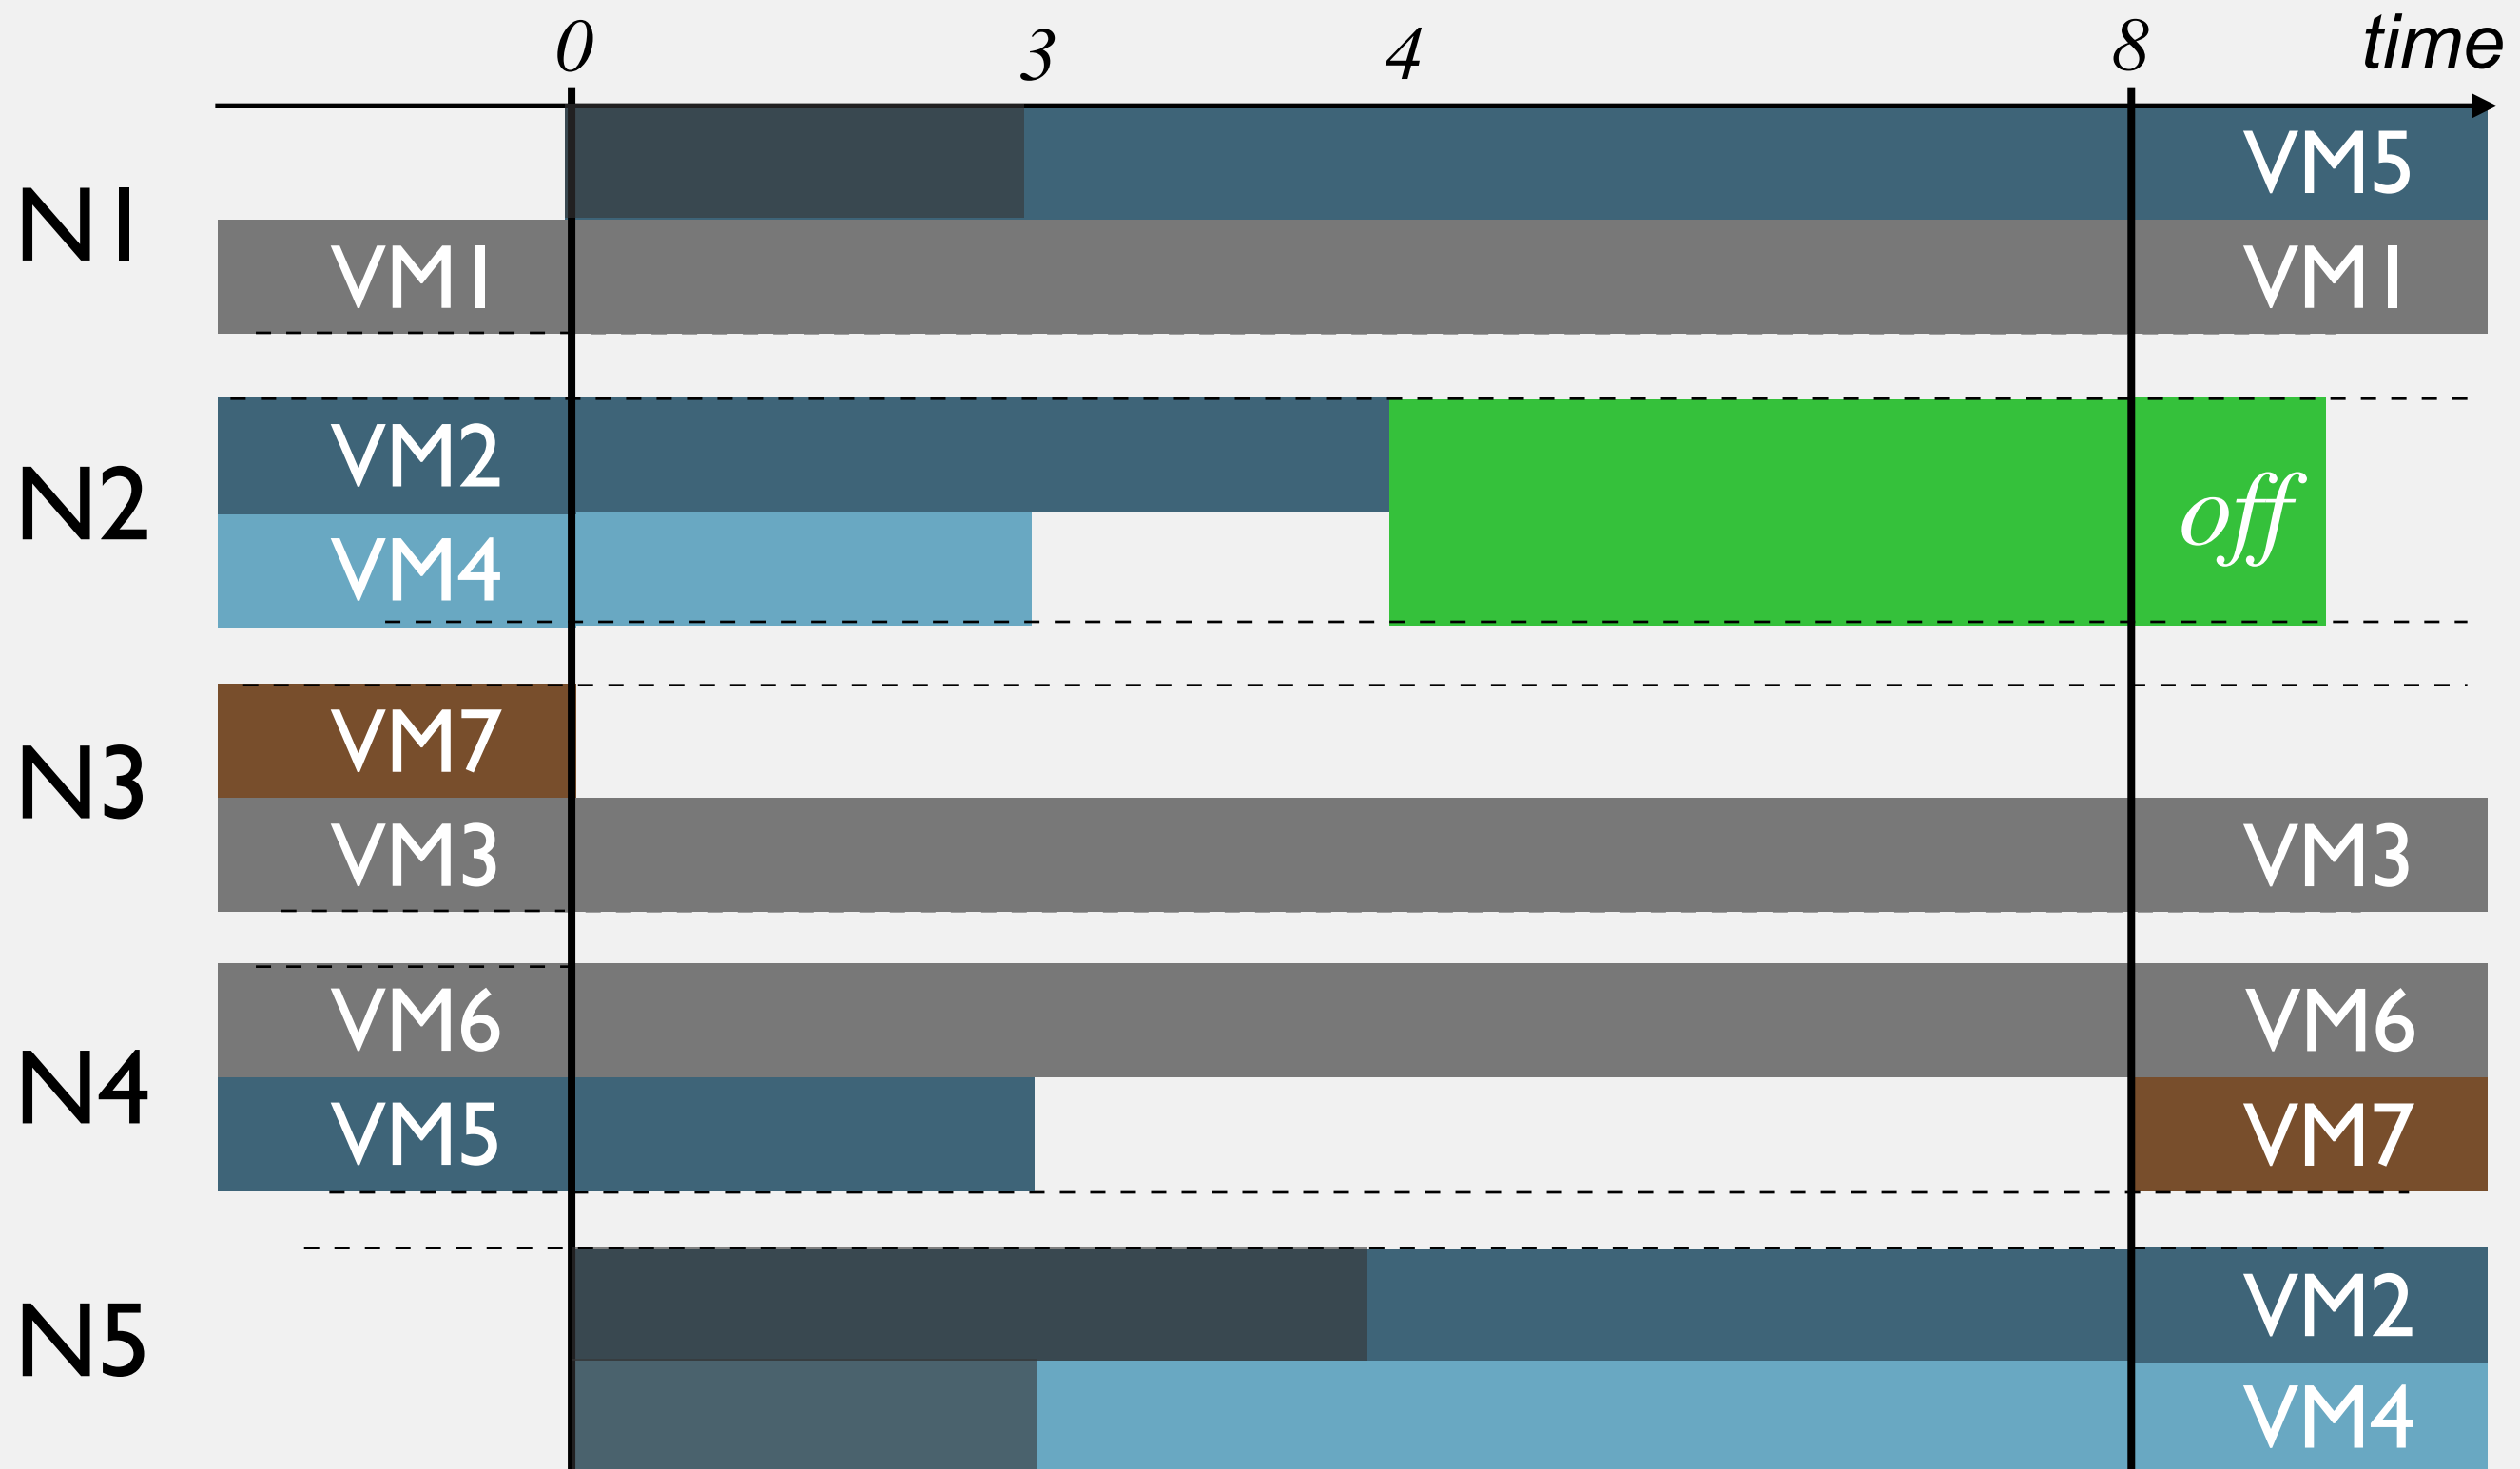
# Resource-Constrained Project Scheduling Problem

the clean way to model space and time



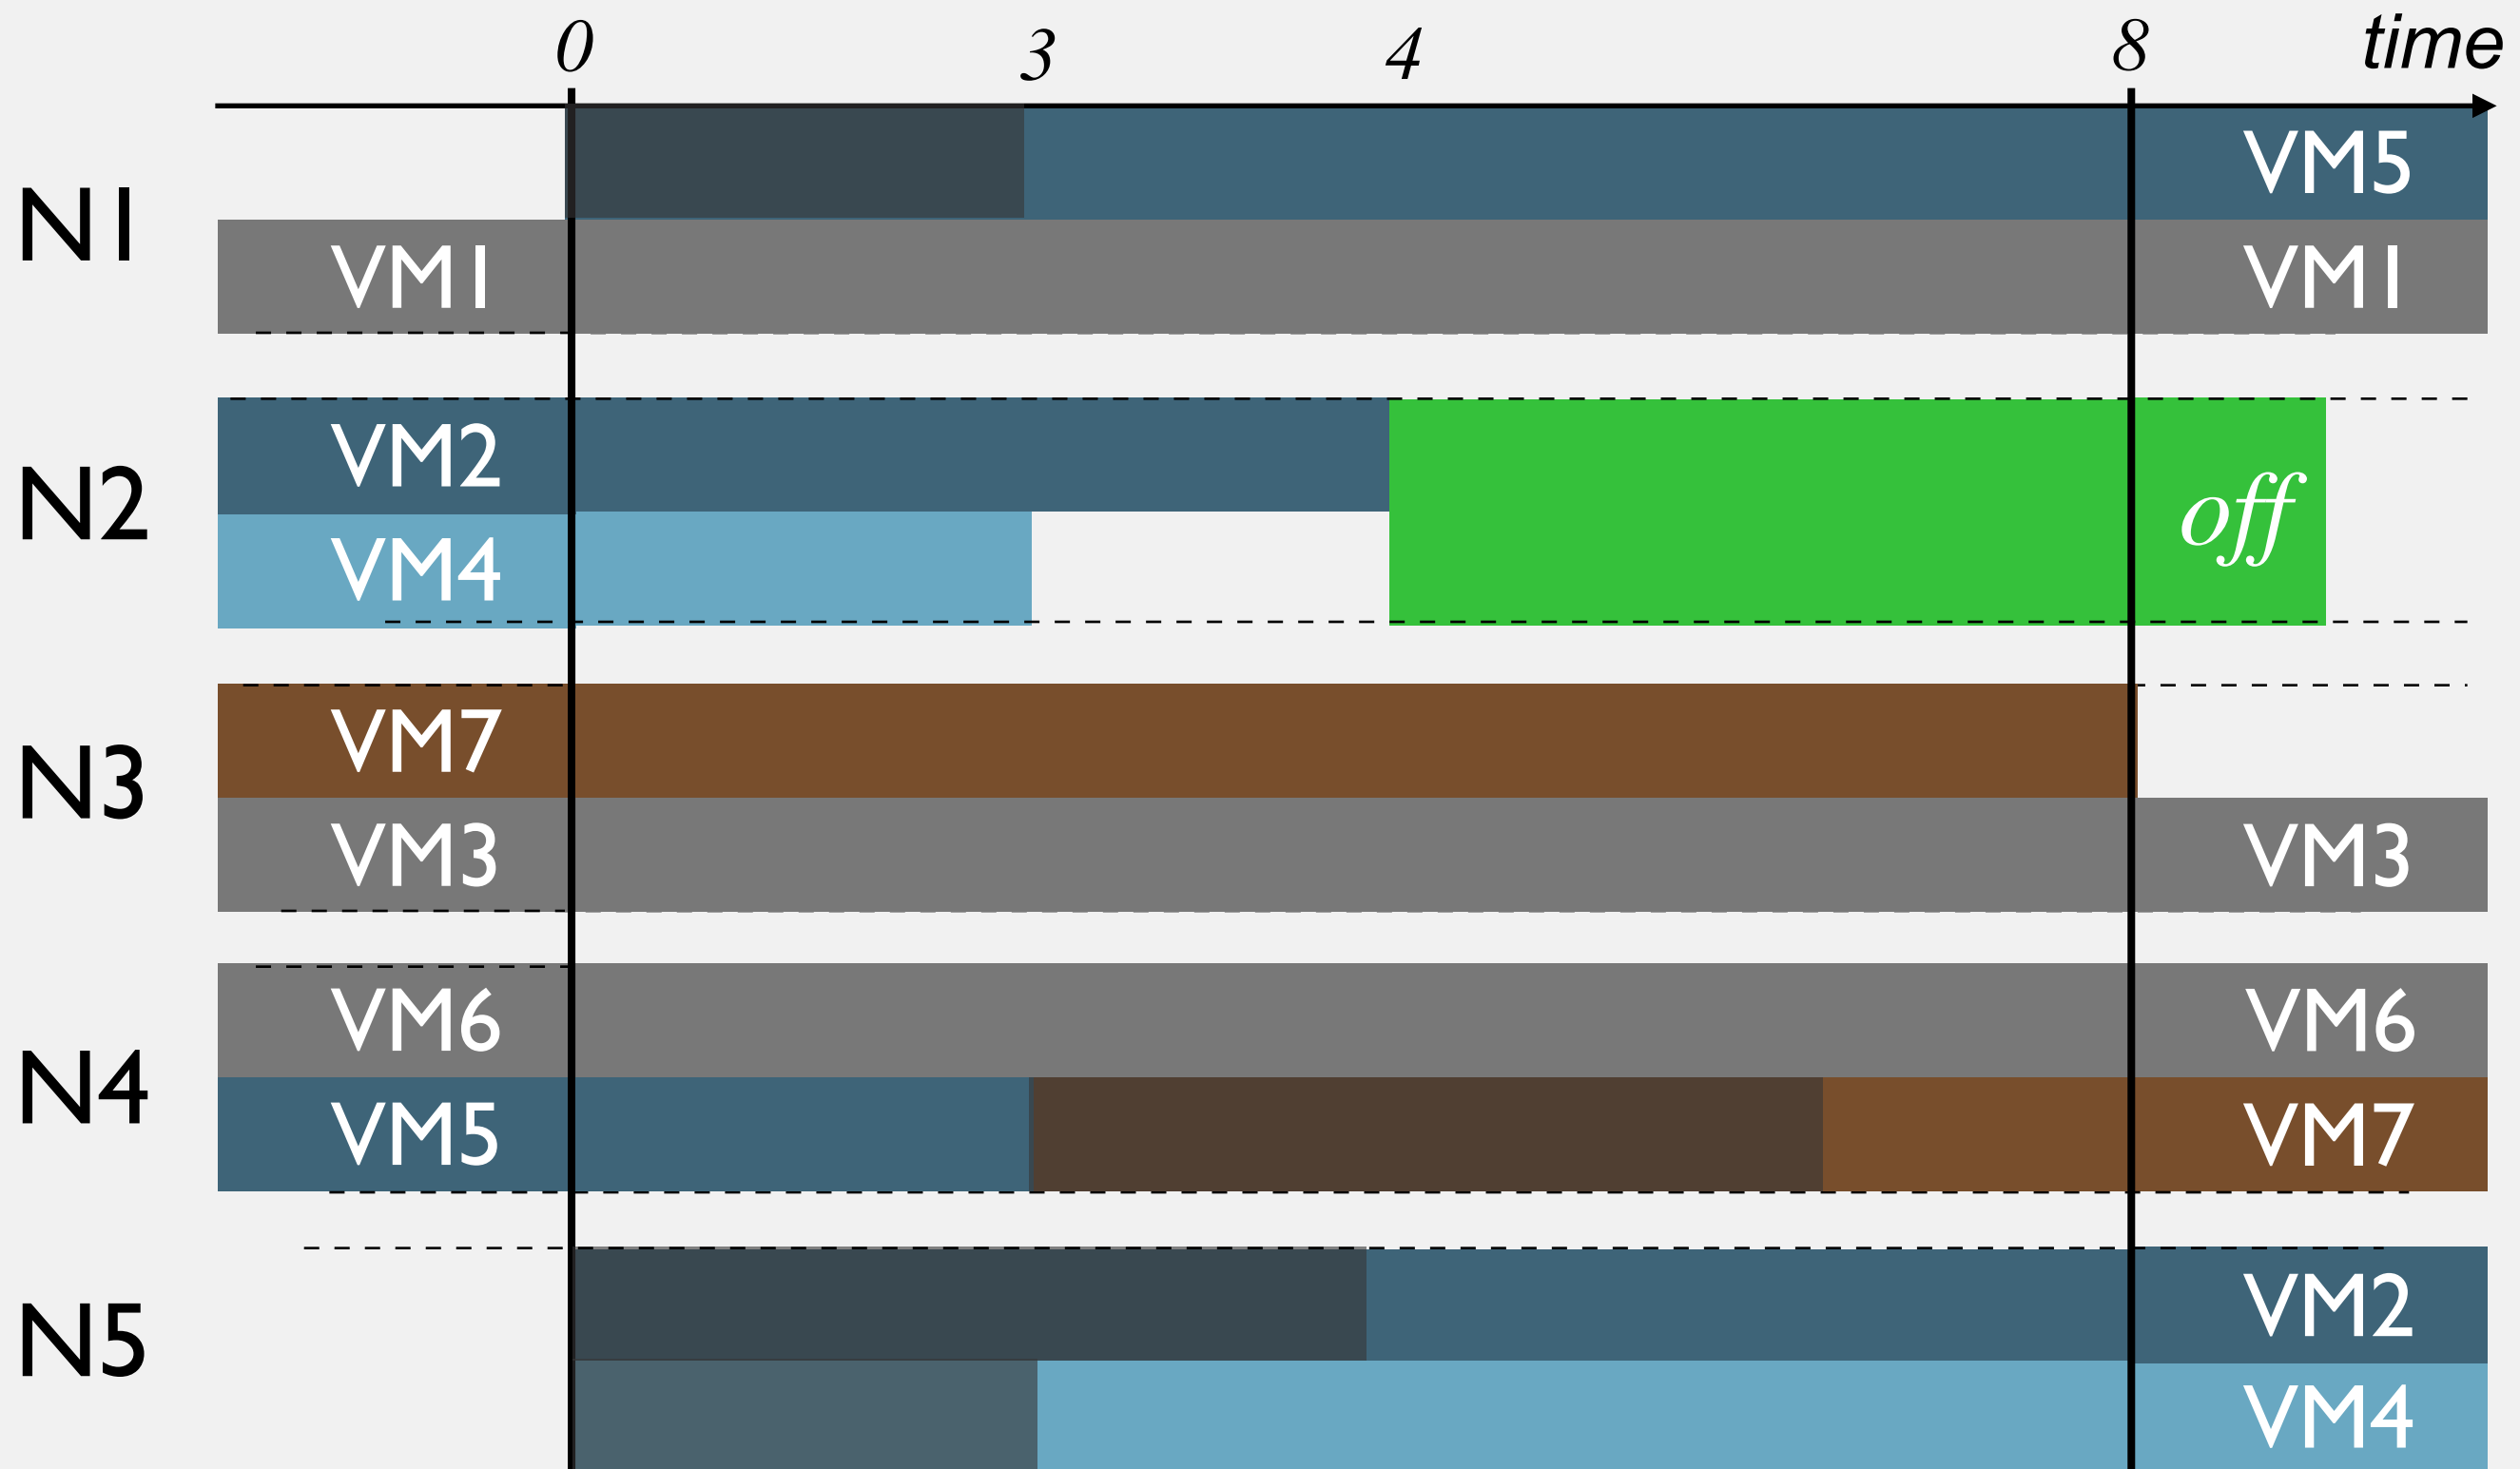
# Resource-Constrained Project Scheduling Problem

the clean way to model space and time



# Resource-Constrained Project Scheduling Problem

the clean way to model space and time



# Resource-Constrained Project Scheduling Problem

the pure way to support migrations

1 **resource** per (node x dimension), bounded capacity

**tasks** to model the VM lifecycle.

height to model a consumption

width to model a duration

at any moment, the cumulative task consumption on a resource cannot exceed its capacity

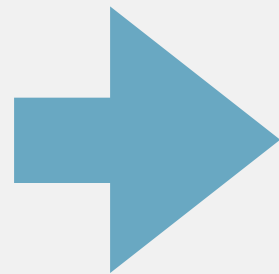
comfortable to express continuous optimisation

very hard to implement (properly)

# From a **theoretical schedule** to a **practical one**

duration may be longer  
convert to an event based schedule

0:3 - migrate VM4  
0:3 - migrate VM5  
0:4 - migrate VM2  
3:8 - migrate VM7  
4:8 - shutdown(N2)



- : migrate VM4  
- : migrate VM5  
- : migrate VM2  
!migrate(VM2) & !migrate(VM4): shutdown(N2)  
!migrate(VM5): migrate VM7



Back to  
vector packing  
based approaches

# Fit Fit Decrease (FFD)

basic VM scheduling

sort VMs in desc order

for each VM

pick the first suitable node

# Fit Fit Decrease (FFD)

basic VM scheduling

**sort** VMs in desc order

for each VM

pick the first suitable node



difficult VMs first!

# Fit Fit Decrease (FFD)

basic VM scheduling

**sort** VMs in desc order

for each VM

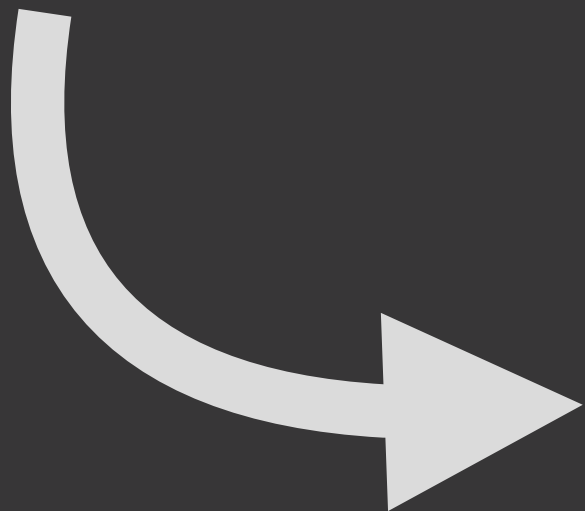
pick the first **suitable** node

difficult VMs first!

enough free resources

# Fit Fit Decrease (FFD)

example



N1

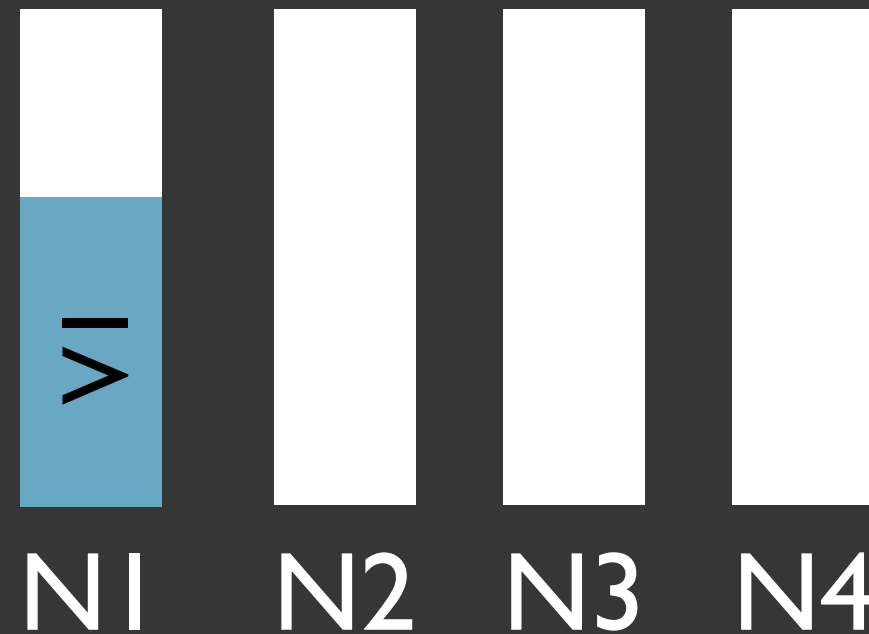
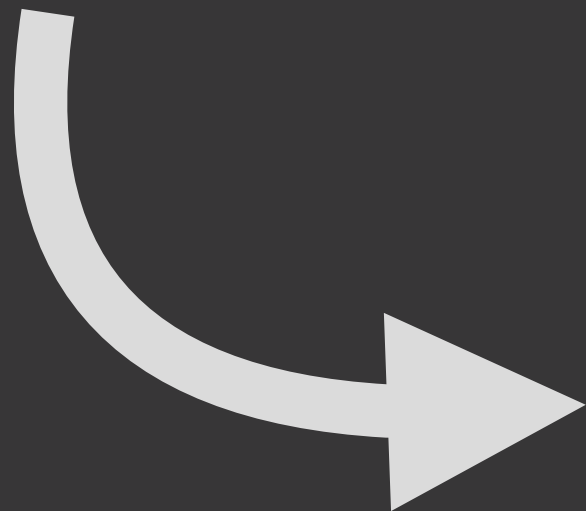
N2

N3

N4

# Fit Fit Decrease (FFD)

example

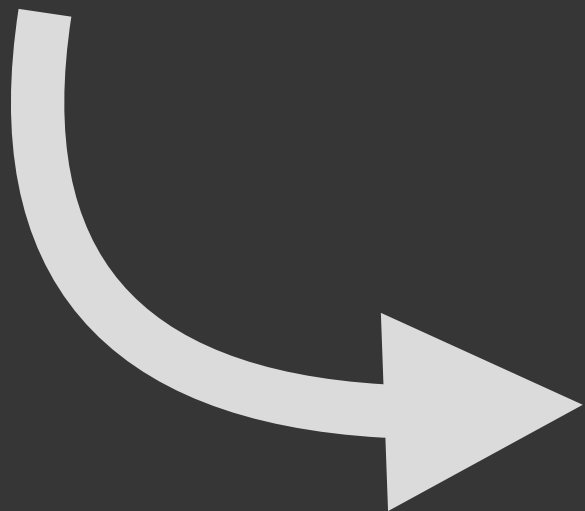


# Fit Fit Decrease (FFD)

example

V4

V3



N1



N2



N3

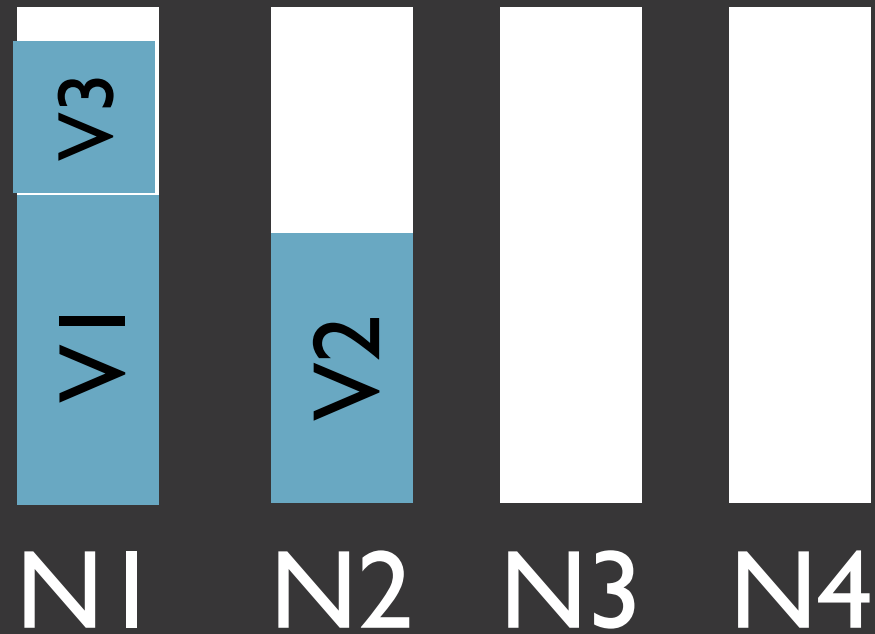
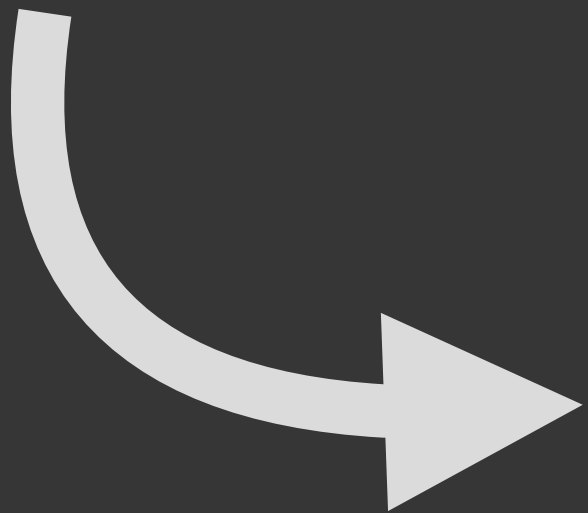


N4

# Fit Fit Decrease (FFD)

example

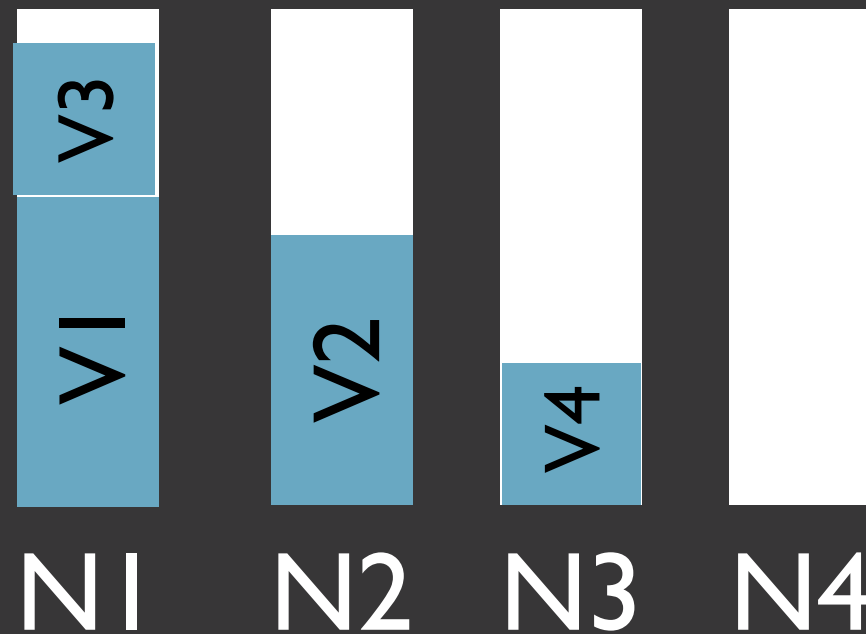
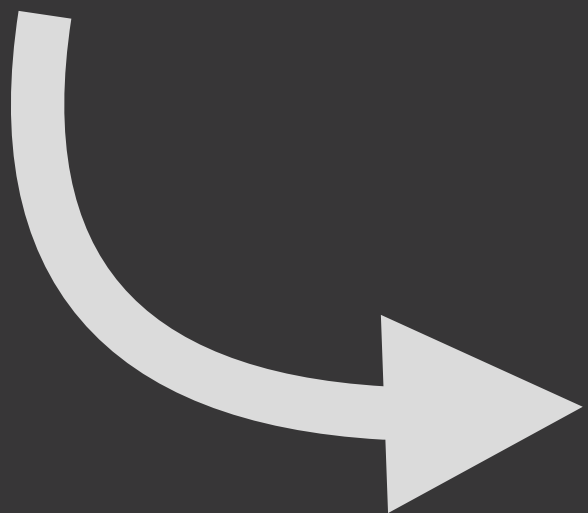
V4





# Fit Fit Decrease (FFD)

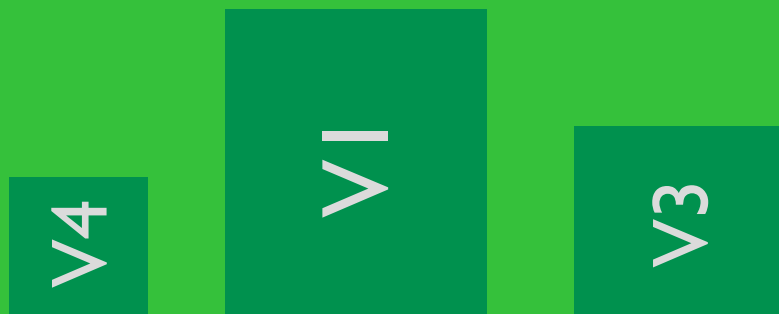
example



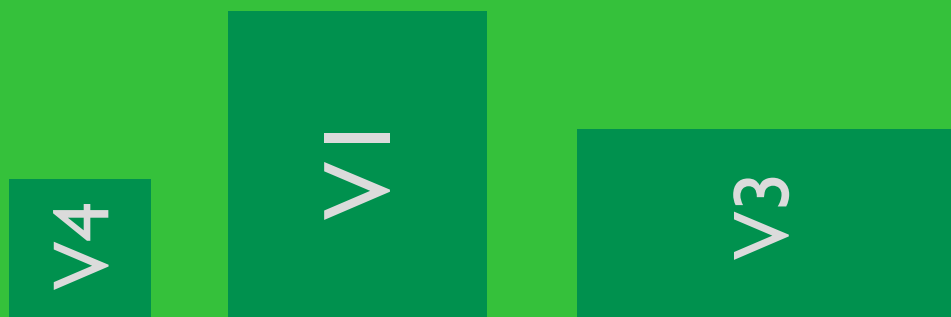
# multi-dimension sorting ?



easy, 1 dimension is varying



easy, uniform variation

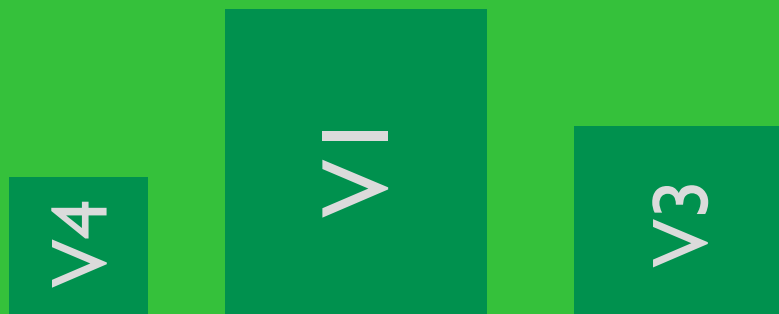


?

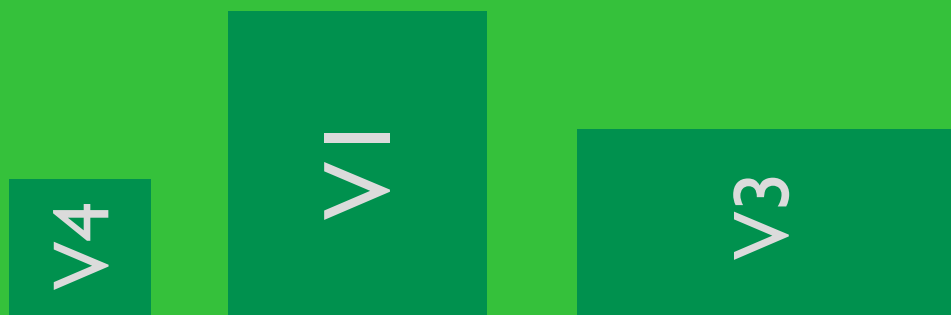
# multi-dimension sorting ?



easy, 1 dimension is varying



easy, uniform variation



sort by dimension  
aggregate dimensions  
find the most critical one

...

# Balancing VMs



Why?



# Why?

- to reduce loss in terms of failures
- to reduce hotspots
- to absorb load spikes

# balancing

in theory

$$\min(\text{stddev}([load(n_1), \dots, load(n_i)]))$$

# Worst Fit Decrease (WFD)

balancing, a practice

**sort** VMs in desc order

for each VM

pick the **suitable node**

with the **highest remaining**

space

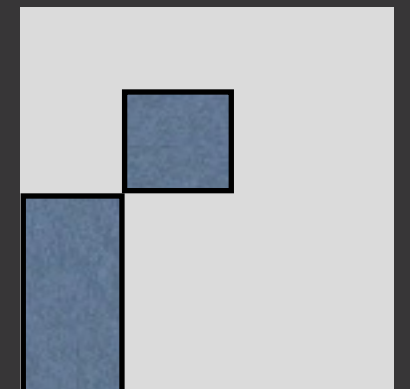
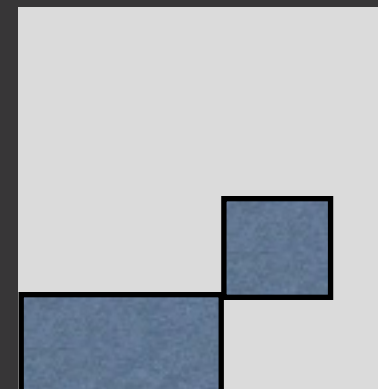


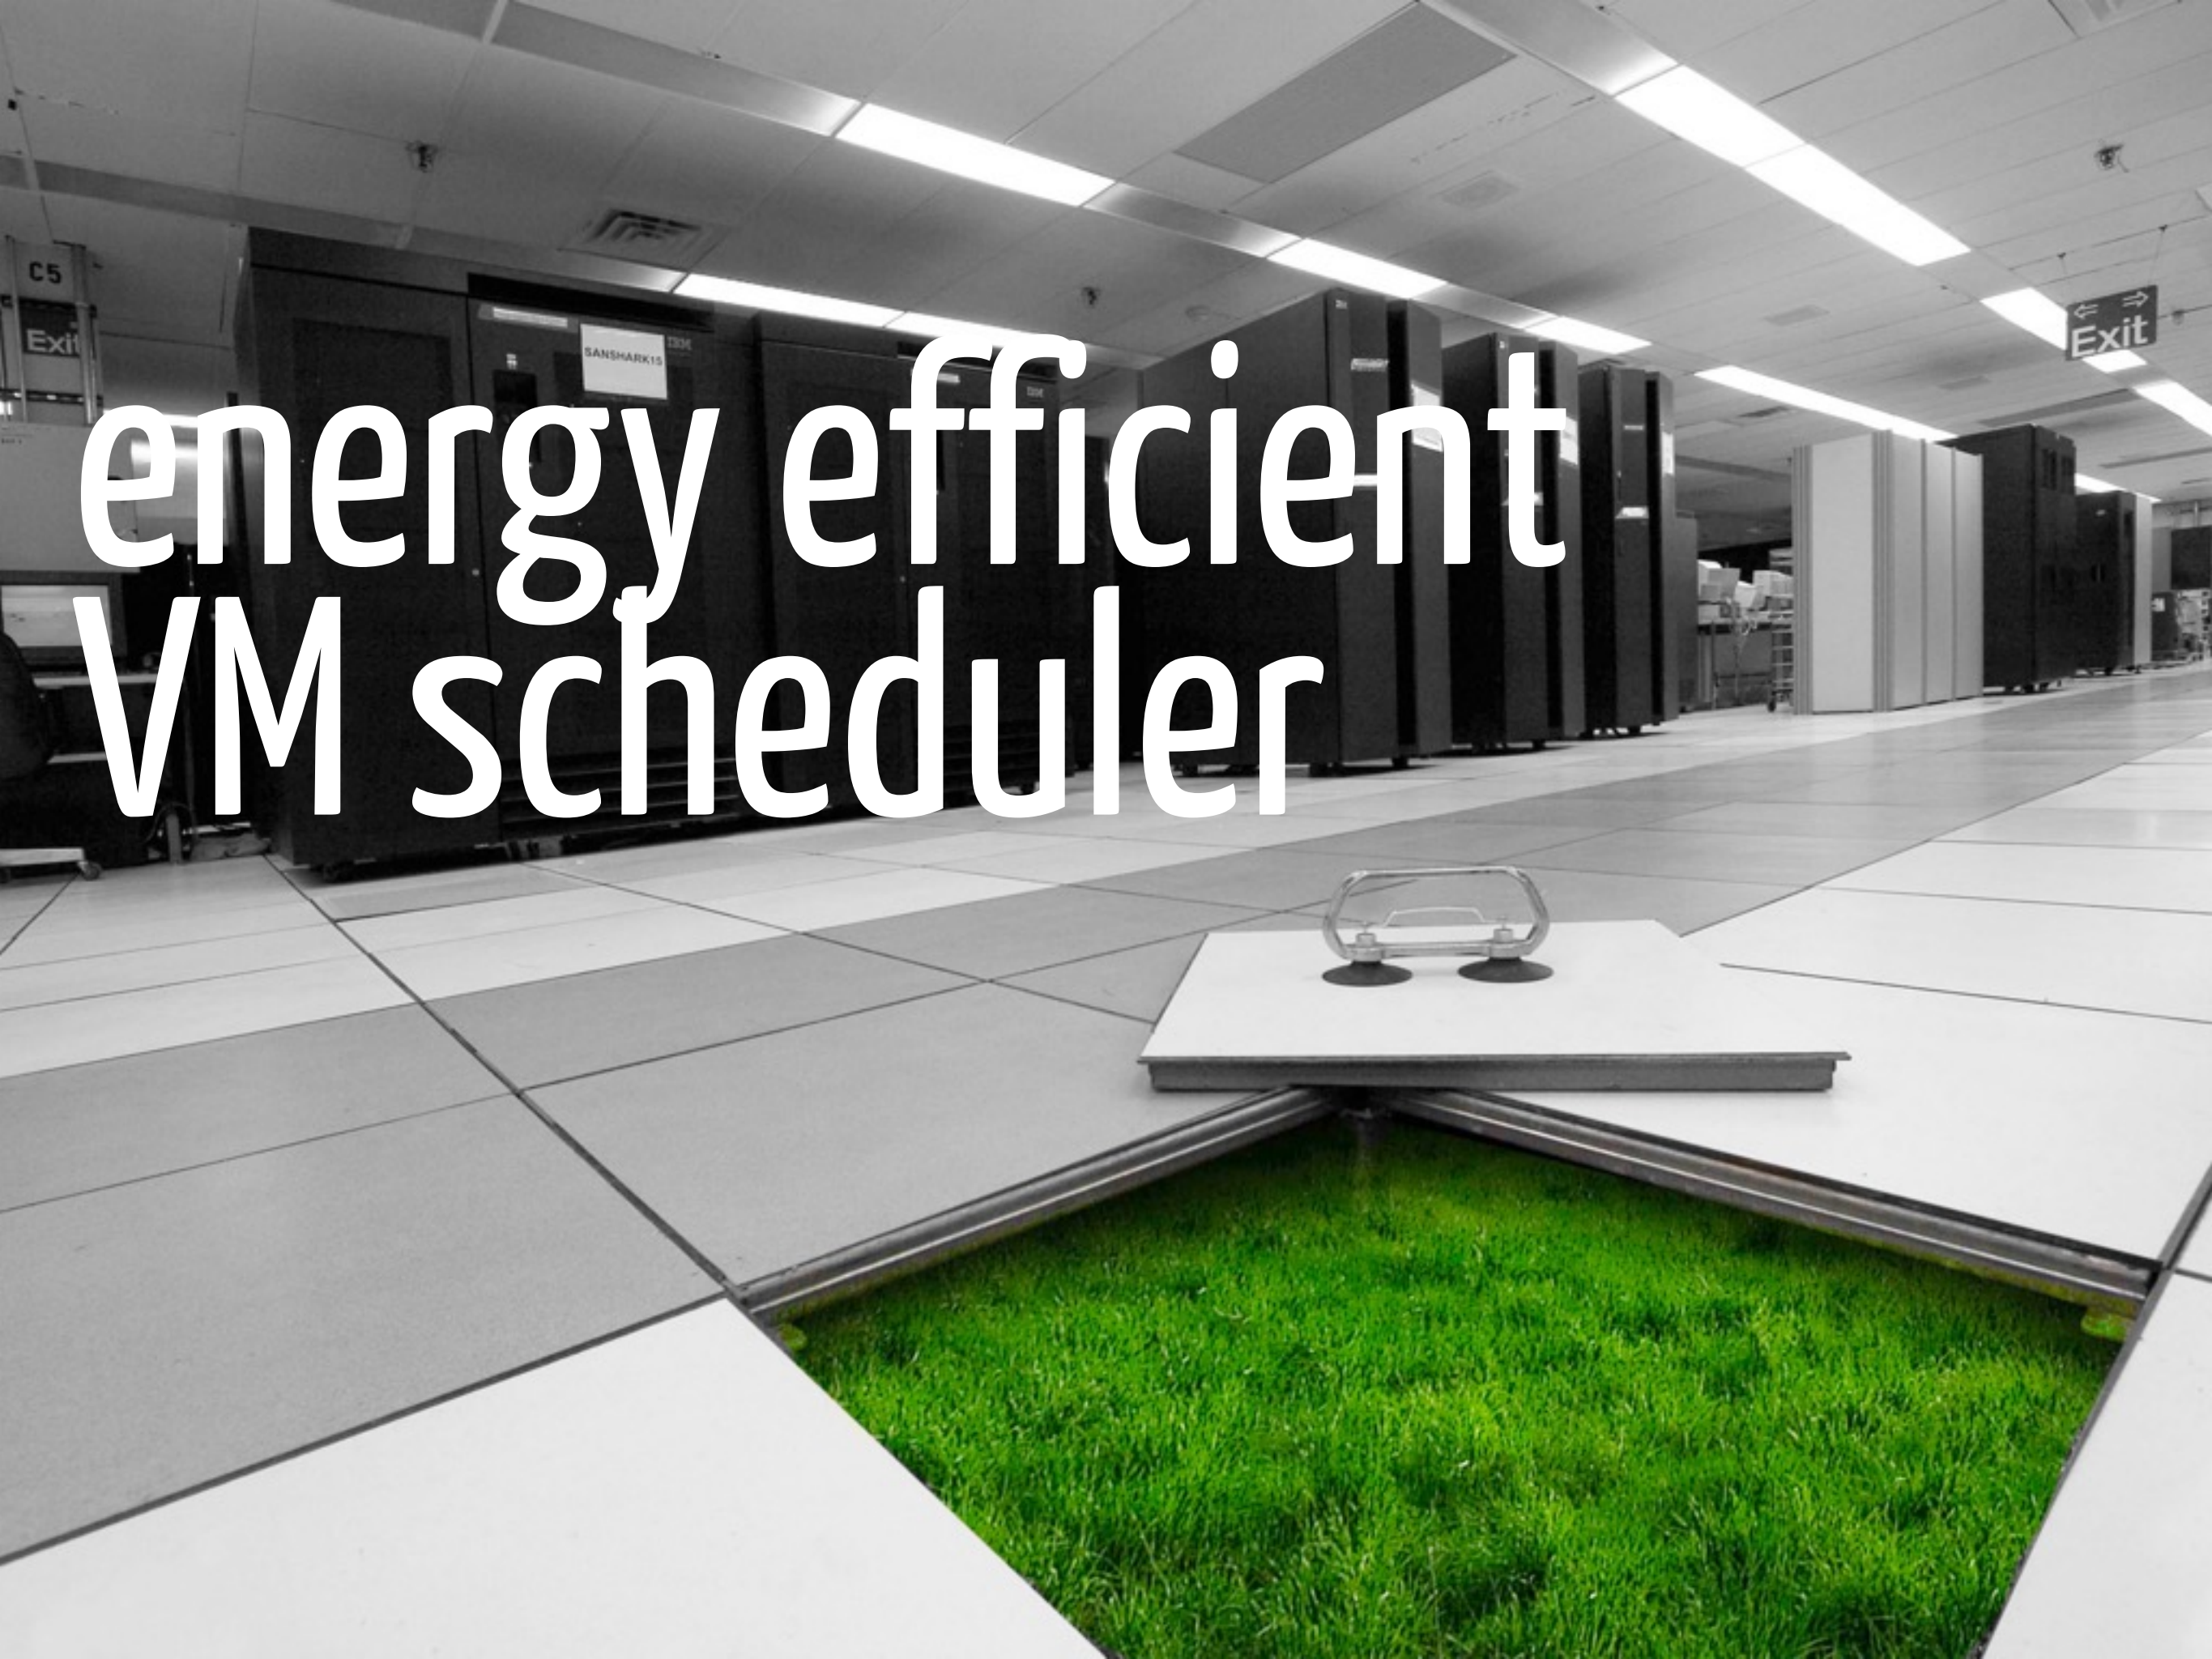


balancing over  
multiple dimensions

# balancing over multiple dimensions

dimension normalisation  
worst dimension  
dimension aggregation





energy efficient  
VM scheduler



# USA

**1.5%**

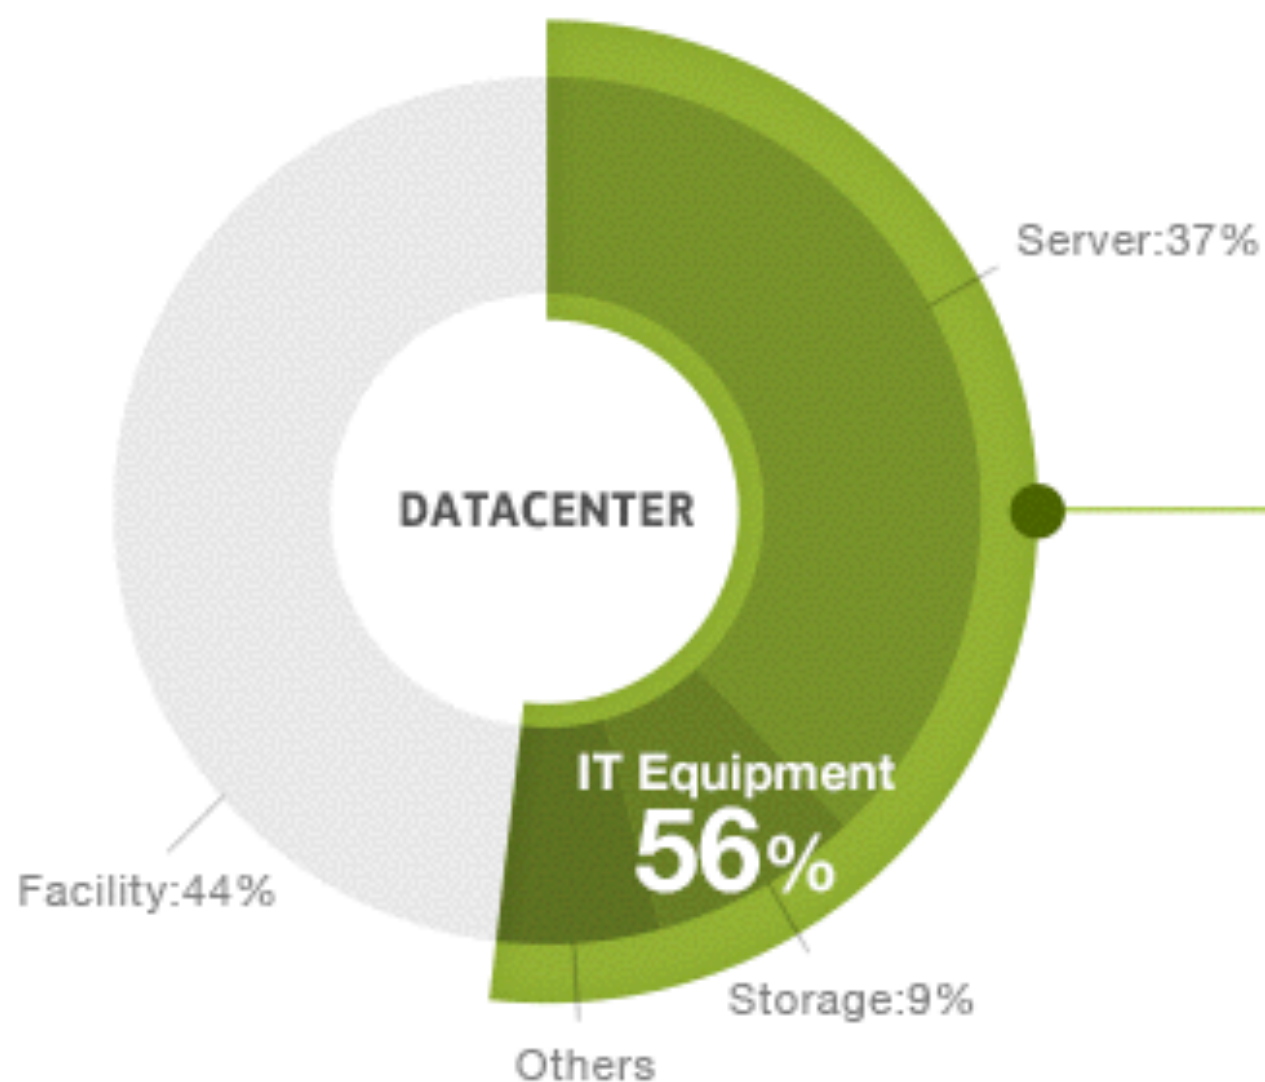
of the 2005 budget

**3%**

in 2010 ?

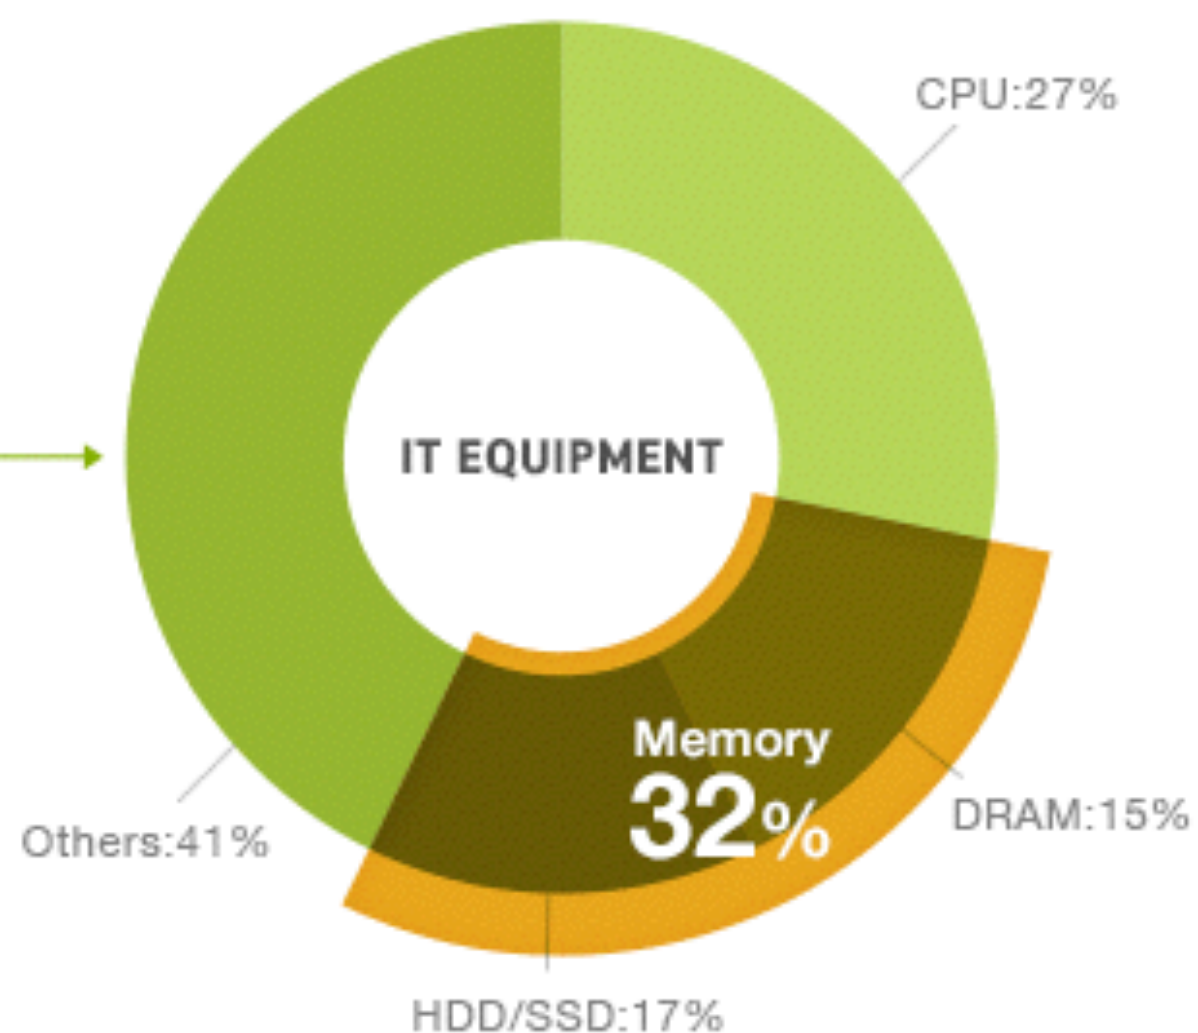


## DATACENTER POWER CONSUMPTION



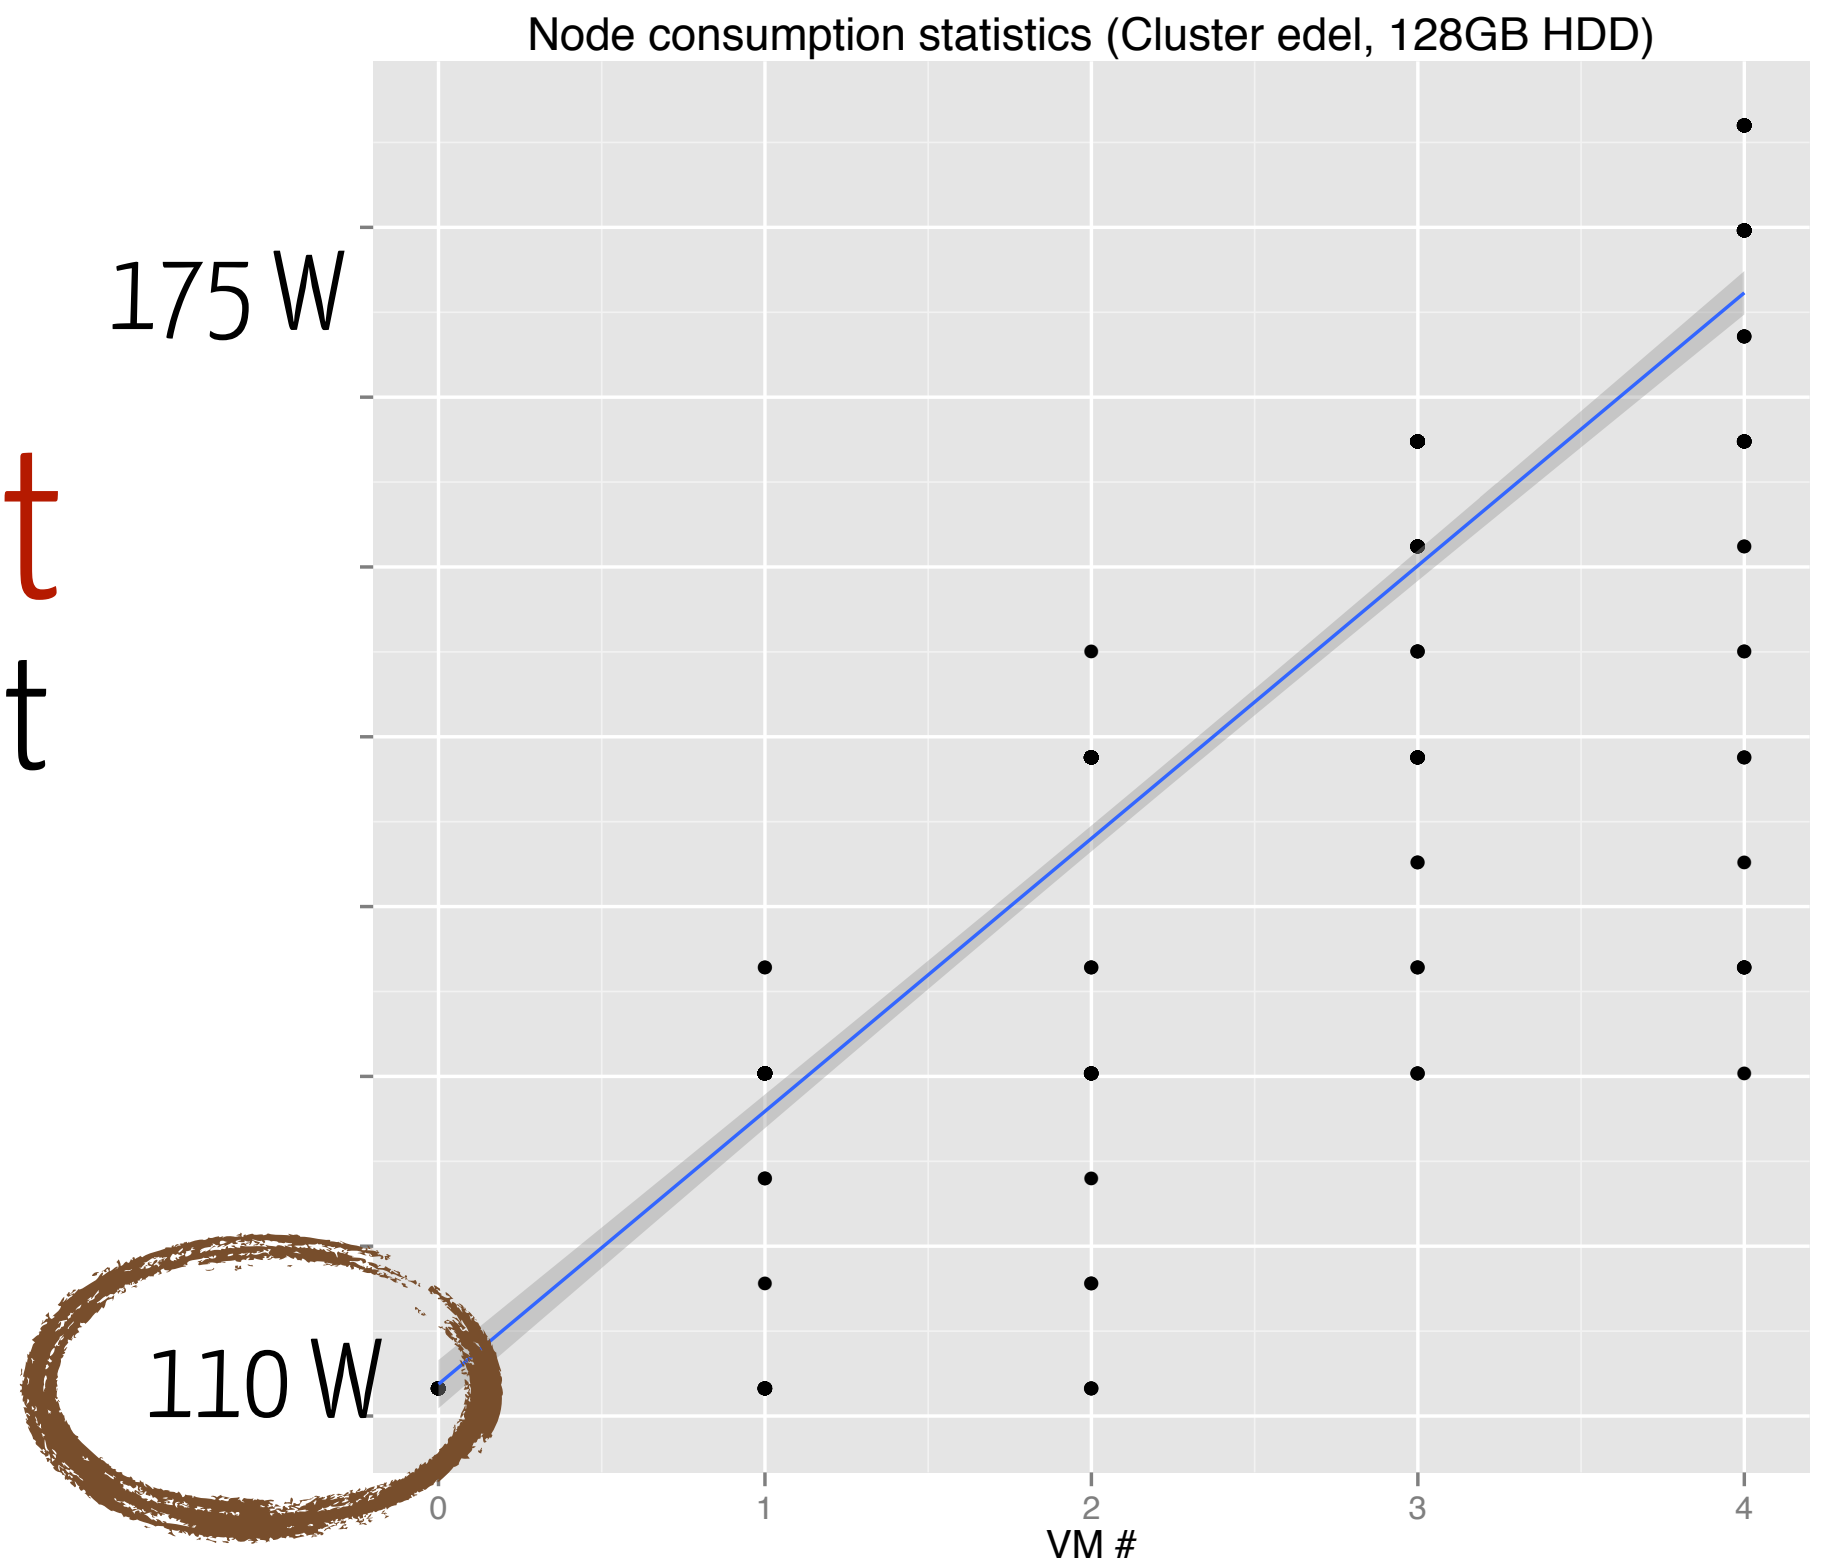
Source: Uptime Institute's 2012 Data Center Survey  
J. Koomey's Report, Aug 2011

## IT EQUIPMENT POWER DRAW



Source: Samsung, IDC, EMC

servers are **not**  
energy efficient



how to reduce consumption  
with identical nodes



# The principles

place VMs on the minimum  
number of nodes

turn off idle nodes

since/repeat for  
the dynamic version



# Best Fit Decrease (BFD)

a simple heuristic to pack VMs

**sort** VMs in desc order

for each VM

pick the **suitable node**

with the **least remaining**

space

efficiency decreases when nb. of dimensions increase

# 2-pass dynamic VM packing

to manage the running VMs

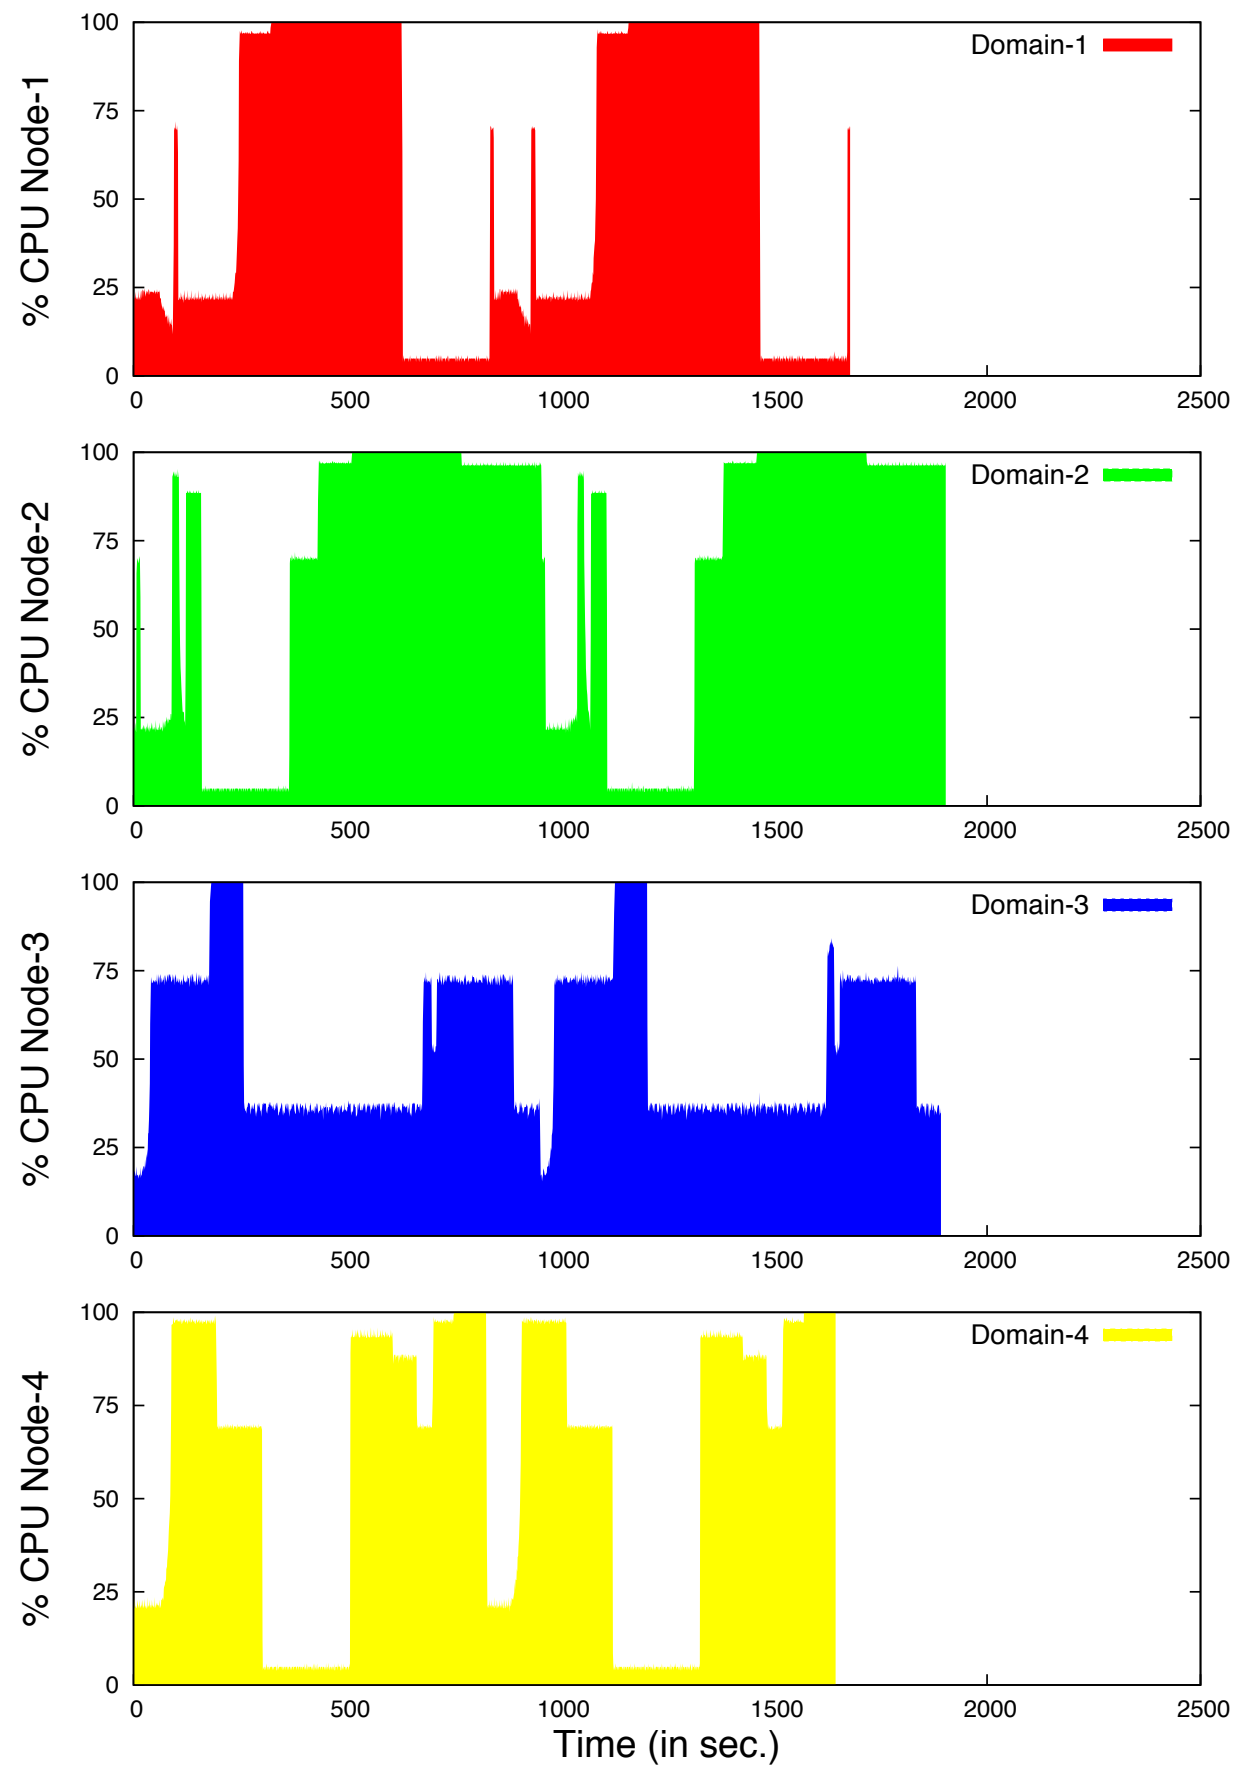
## 1) address performance issues

on saturated nodes,  
BFD to move away VMs on non-saturated nodes

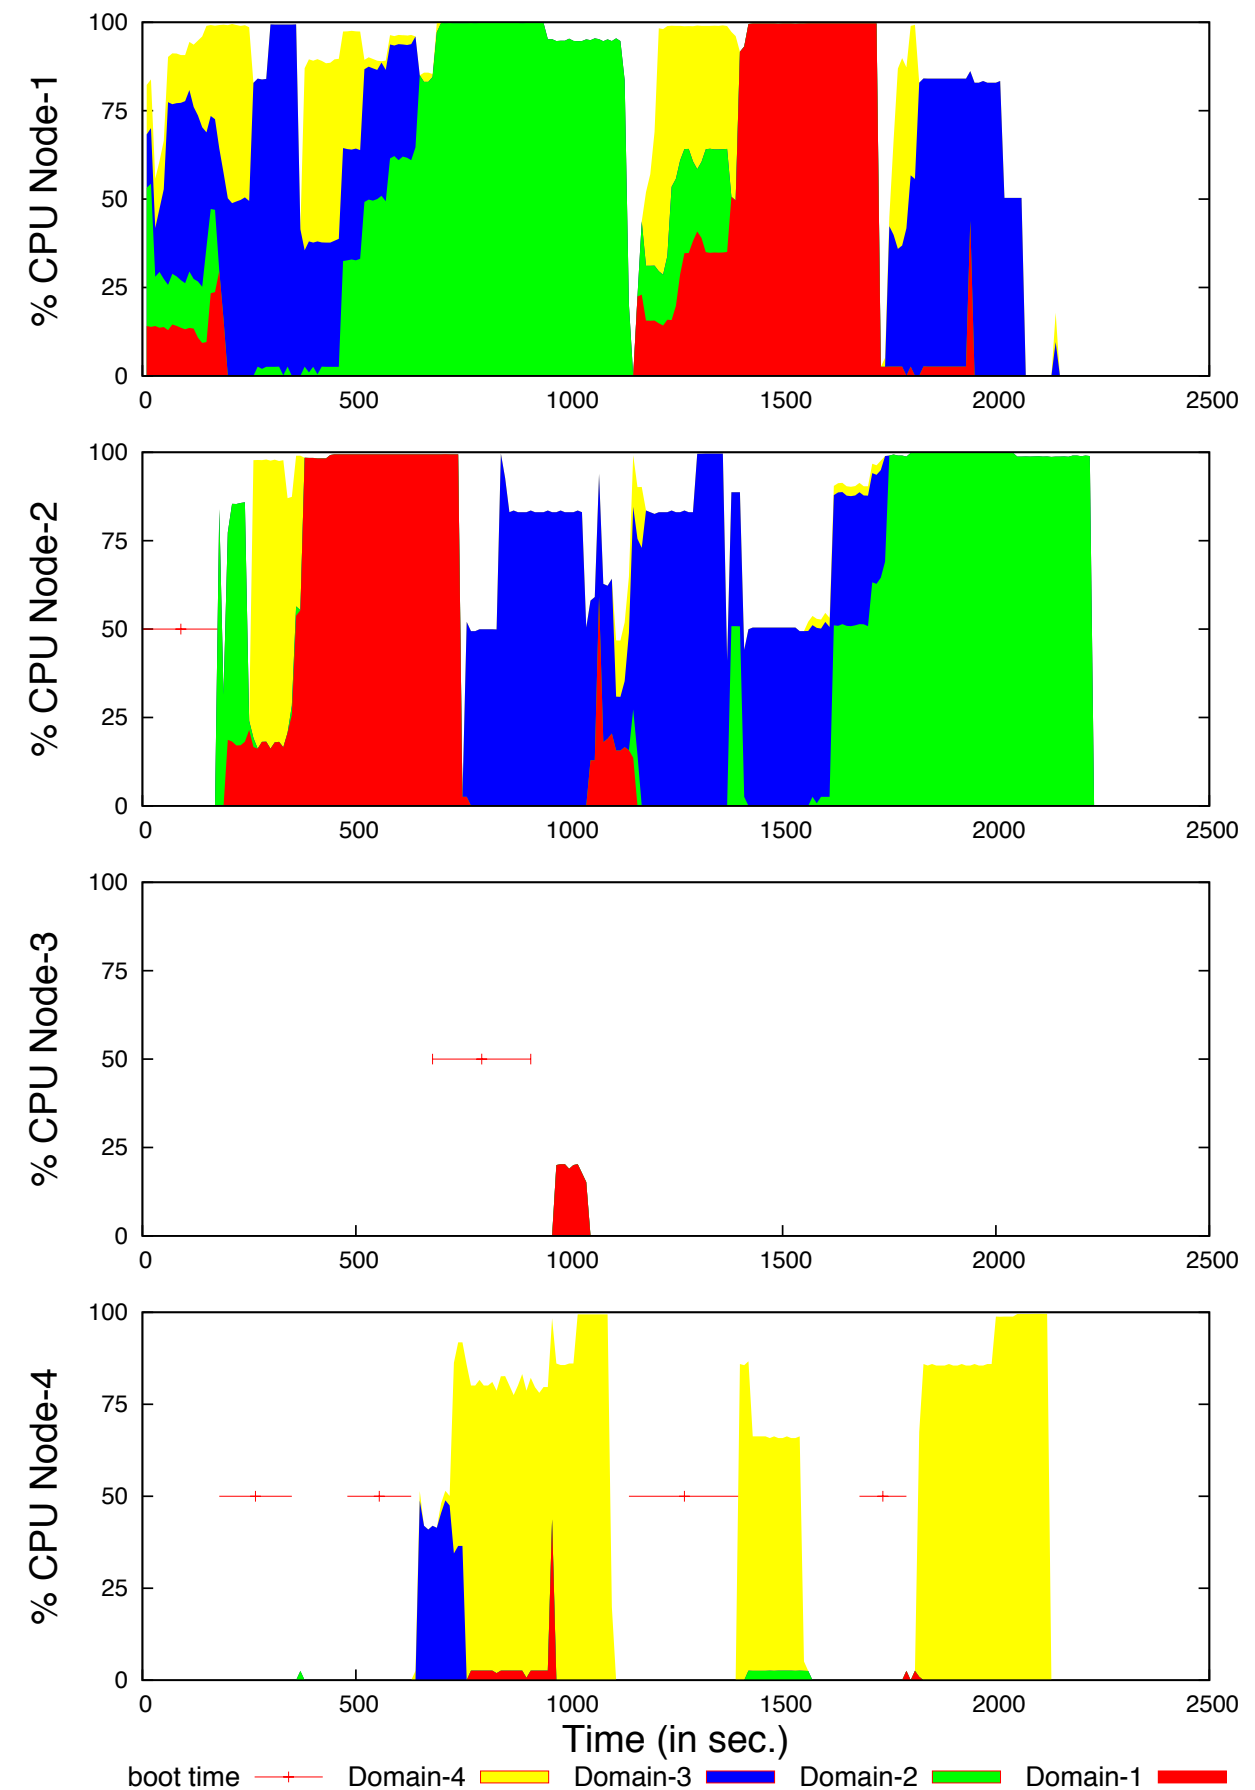
## 2) address energy efficiency issues

on low-loaded nodes,  
BFD to move away VMs on heavily loaded nodes

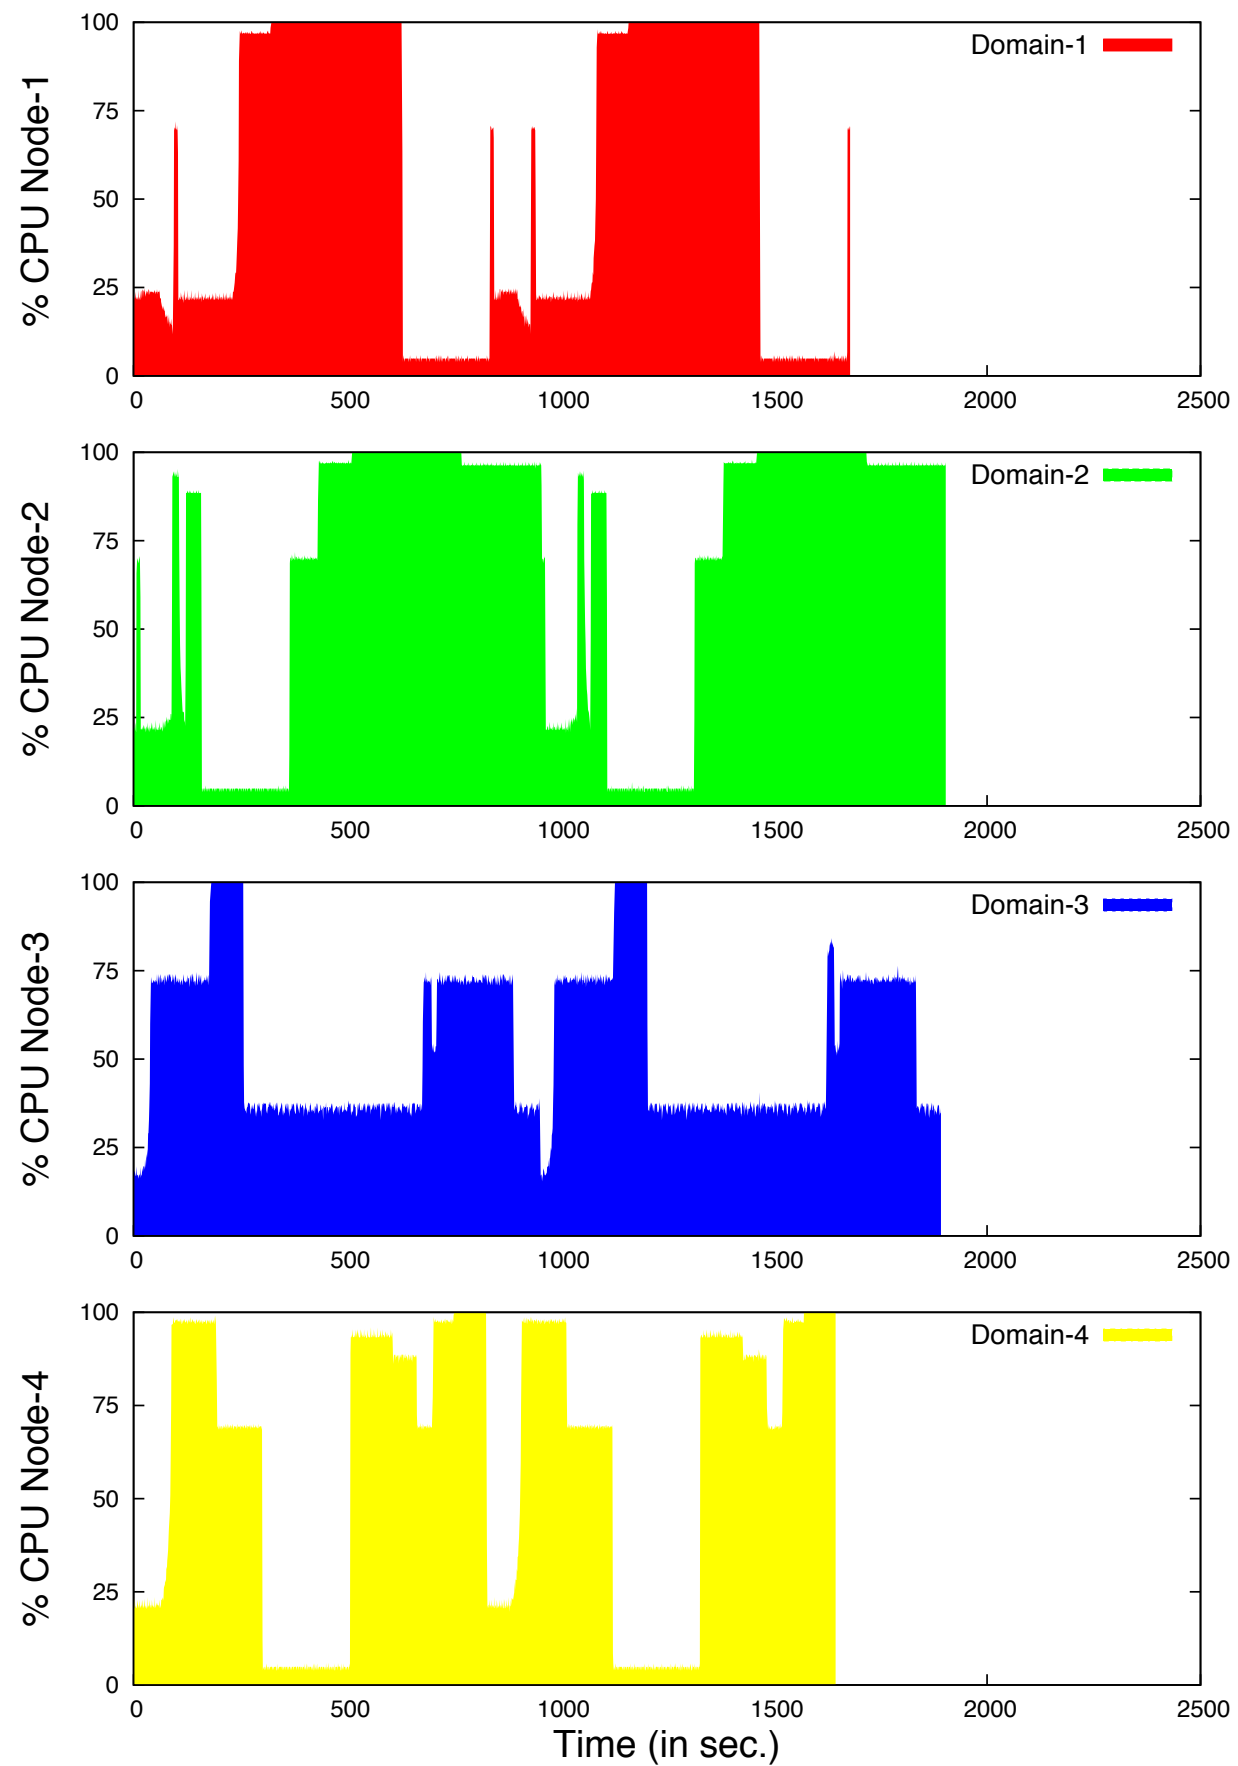
many variations using threshold based systems, load predictions ...



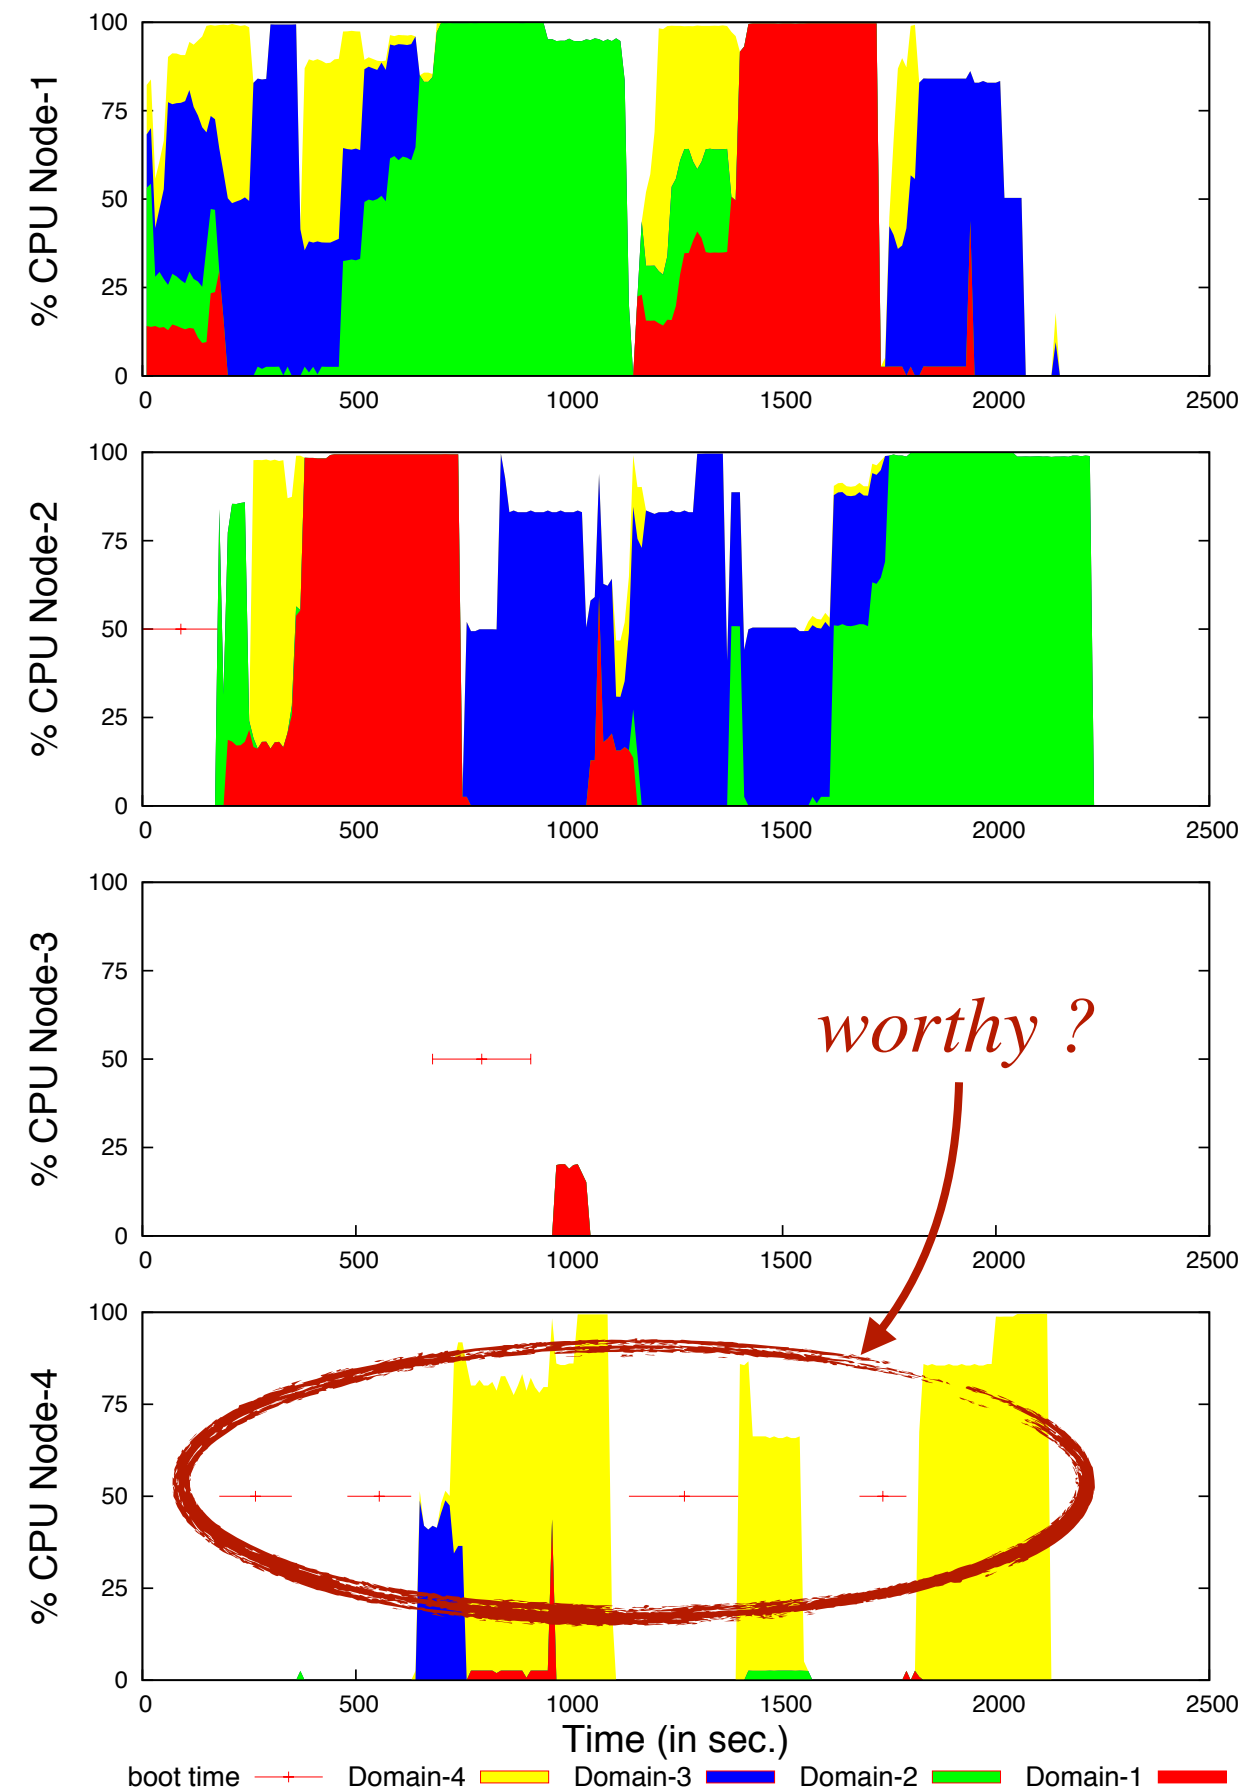
No packing



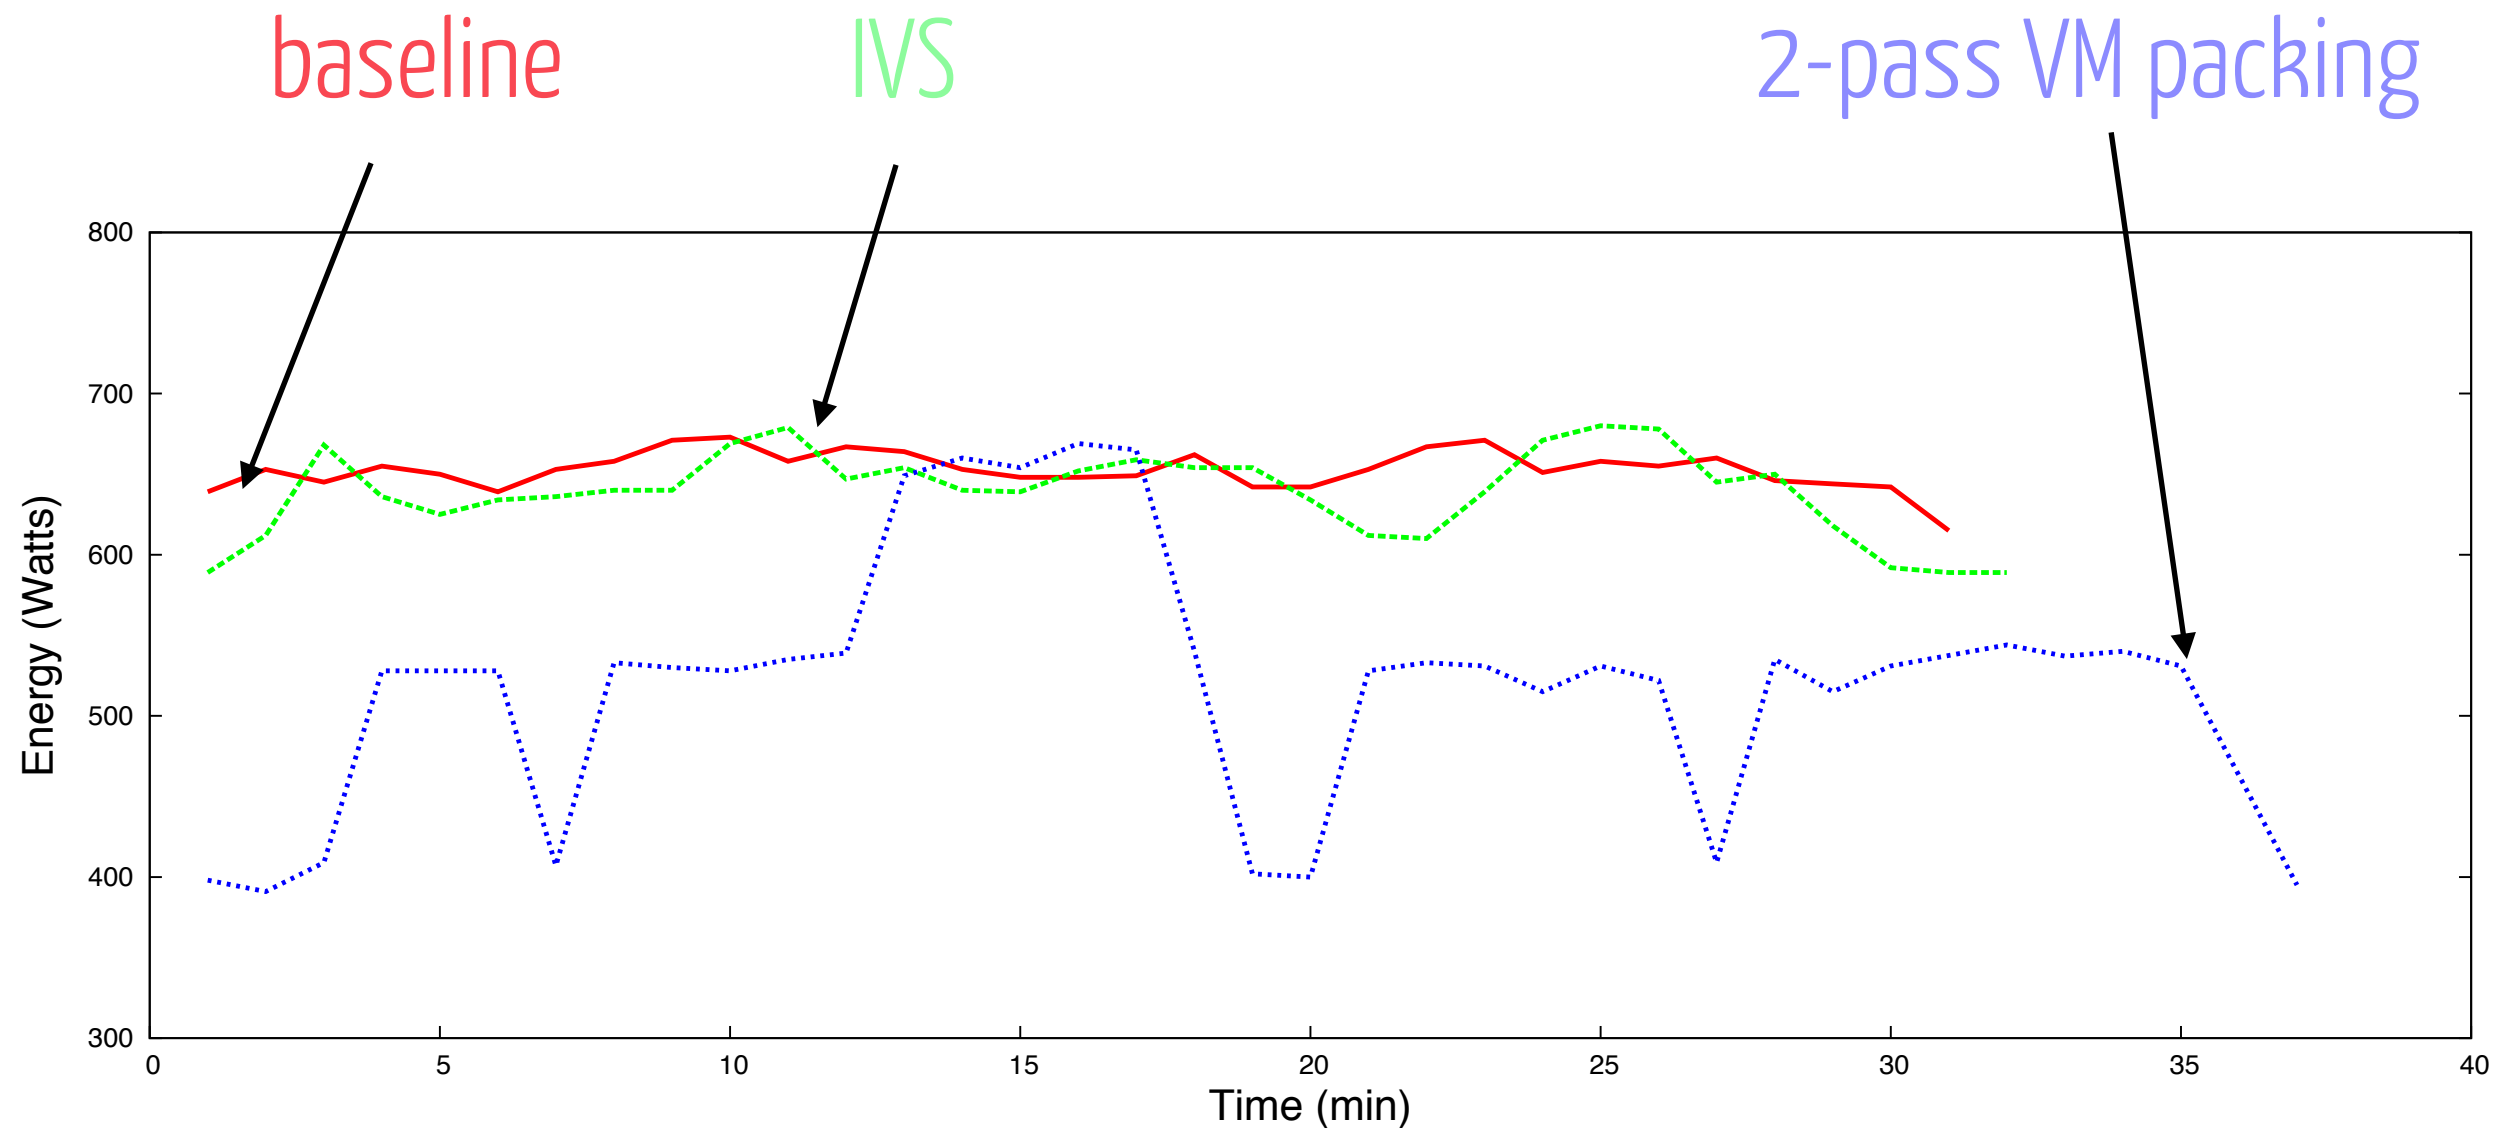
2-passes packing



No packing



2-passes packing



interesting gains,  
the node boot time was the bottleneck  
simplistic model (identical nodes)

Consequences of having  
heterogeneous nodes ?

# Consequences of having heterogeneous nodes ?

different performance  
different Power Usage Effectiveness

...

how to estimate the benefits of a migration ?

vector packing problem  
+ performance model  
+ power model  
+ cost model  
+ ...

---

a specialised model



# power models

estimate the energy consumption

model the static and the dynamic energy  
profile of the components

usually linear equations

$$W_h(\text{node}) = \alpha H_w + \beta$$

# VM performance models

a neutral performance unit  
mips, ECU, GCU, ...

a model to map  
VM performance with their host

$$\text{perf}(\text{VM}, N) = \alpha_n \% \text{cpu} + \beta$$

# hw. particularities

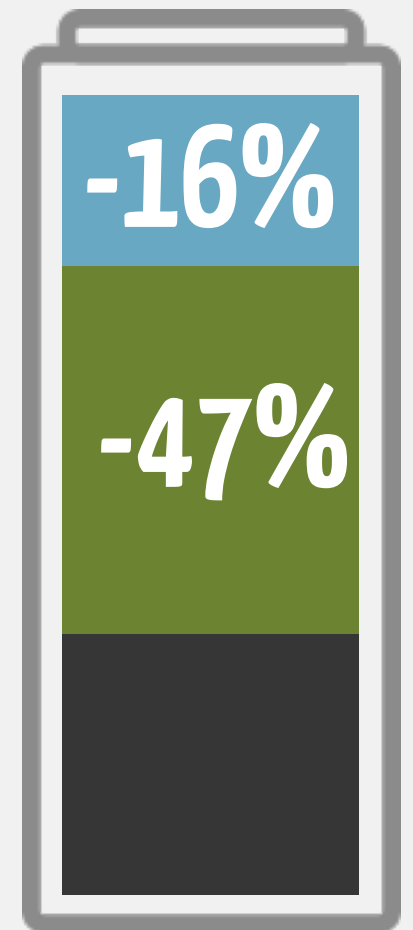
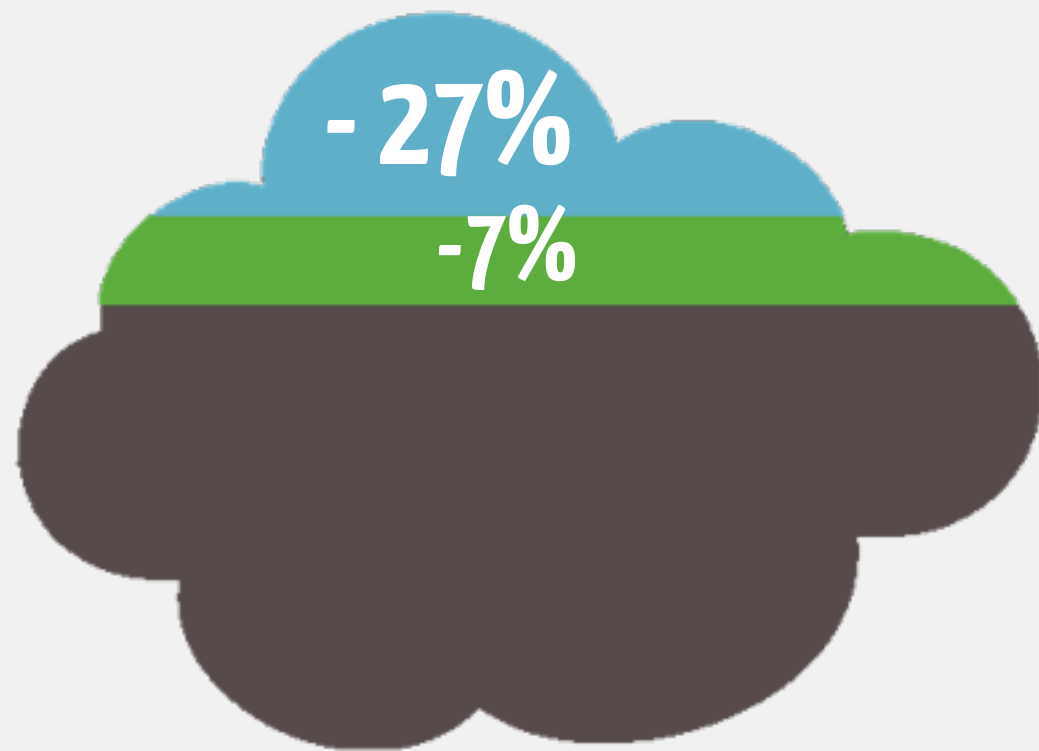
multi-core CPUs,  
DDR3 memory,  
spinning HD,  
PUE / CUE,  
boot/shutdown time



# workload particularities

VM template  
migration duration  
migration payback time

a fine-grain power model



coarse to fine  
grain optimisation

# RECAP

no holy grail

master the  
problem